

CNN (Convolutional Neural Network) Cheat Sheet

Codeium

2025-02-15

CNN (Convolutional Neural Network) Cheat Sheet

1. Basic Concepts

Convolution Operation

Takes an input image and applies a filter/kernel to produce a feature map

Example:

Input Image:		Filter:		Feature Map:
[1 2 3]		[1 0]		[8 11]
[4 5 6]	*	[2 1]	=	[17 20]
[7 8 9]				

Step-by-step calculation: - Top-left (8) = $1 \times 1 + 2 \times 0 + 4 \times 2 + 5 \times 1 = 1 + 0 + 8 + 5 = 8$ - Top-right (11) = $2 \times 1 + 3 \times 0 + 5 \times 2 + 6 \times 1 = 2 + 0 + 10 + 6 = 11$ - Bottom-left (17) = $4 \times 1 + 5 \times 0 + 7 \times 2 + 8 \times 1 = 4 + 0 + 14 + 8 = 17$ - Bottom-right (20) = $5 \times 1 + 6 \times 0 + 8 \times 2 + 9 \times 1 = 5 + 0 + 16 + 9 = 20$

Padding

Types: Valid (no padding), Same (output size = input size)

Example:

Original:		With Zero Padding:
[1 2 3]		[0 0 0 0 0]
[4 5 6]	→	[0 1 2 3 0]
[7 8 9]		[0 4 5 6 0]
		[0 7 8 9 0]
		[0 0 0 0 0]

Stride

Controls filter movement across input. Example with stride 2:

Input:	Filter:	Output:
[1 2 3 4]	[1 0]	
[5 6 7 8]	[0 1]	[12 16]
[9 10 11 12]		[28 32]
[13 14 15 16]		

2. Pooling Layers

Max Pooling (2×2, stride 2)

Input: Output:

```

[1  2  3  4]      [6   8]
[5  6  7  8]  →  [14  16]
[9 10 11 12]
[13 14 15 16]

```

Average Pooling (2×2, stride 2)

```

Input:      Output:
[1  2  3  4]  [3.5  5.5]
[5  6  7  8]  →  [11.5 13.5]
[9 10 11 12]
[13 14 15 16]

```

3. Activation Functions

ReLU (Rectified Linear Unit)

$f(x) = \max(0, x)$

Input: [-2, -1, 0, 1, 2]

Output: [0, 0, 0, 1, 2]

Softmax (for classification)

Raw scores: [2.0, 1.0, 0.1]

Softmax probabilities: [0.67, 0.24, 0.09]

4. Common Architecture Patterns

Basic CNN Structure

Input → Conv → ReLU → Pool → Conv → ReLU → Pool → Flatten → Dense → Output

Typical Layer Sizes

- Input: 224×224×3 (RGB image)
- Conv1: 32 filters of 3×3
- Pool1: 2×2 max pooling
- Conv2: 64 filters of 3×3
- Pool2: 2×2 max pooling
- Dense: 512 units
- Output: Number of classes (e.g., 10 for MNIST)

5. Training Concepts

Learning Rate

- Typical values: 0.1, 0.01, 0.001, 0.0001
- Example:

$\text{weight_new} = \text{weight_old} - \text{learning_rate} \times \text{gradient}$

If gradient = 0.5 and learning_rate = 0.01:

$\text{weight_new} = \text{weight_old} - 0.01 \times 0.5 = \text{weight_old} - 0.005$

Batch Size

- Common values: 32, 64, 128, 256
- Example with batch size 4:
Batch 1: [image1, image2, image3, image4]
Batch 2: [image5, image6, image7, image8]

6. Common Problems & Solutions

Overfitting

Solutions: - Dropout (e.g., $p=0.5$ drops 50% of neurons) - Data Augmentation - L1/L2 Regularization

Vanishing Gradients

Solutions: - Batch Normalization - ResNet Connections - Proper initialization

7. Performance Metrics

Accuracy Calculation

Accuracy = Correct Predictions / Total Predictions

Example: 90 correct out of 100 = 90%

Confusion Matrix (Binary Classification)

	Predicted	
Actual	0	1
0	TN	FP
1	FN	TP

Example:

	Predicted	
Actual	0	1
0	45	5
1	10	40

8. Tips for Training

1. Start with a simple architecture
2. Use transfer learning when possible
3. Monitor validation loss
4. Use early stopping
5. Implement learning rate scheduling

9. Popular CNN Architectures

- LeNet-5: 7 layers
- AlexNet: 8 layers
- VGG-16: 16 layers
- ResNet-50: 50 layers
- Inception/GoogleNet: 22 layers

1. Different Types of Convolutions

- Standard Convolution

- Input \rightarrow Depthwise Conv \rightarrow Pointwise Conv (1×1)
Reduces parameters by $\sim 9x$

- ```

Regular: Dilated (rate=2):
1 2 3 1 _ 2 _ 3
4 5 6 → _ _ _ _
7 8 9 4 _ 5 _ 6
 _ _ _ _
 7 _ 8 _ 9

```

- ## 2. Advanced Architectures

- Input  $\rightarrow$  Split  $\rightarrow$   $\begin{matrix} [1\times1 \text{ Conv}] \\ [3\times3 \text{ Conv}] \\ [5\times5 \text{ Conv}] \\ [\text{MaxPool}] \end{matrix} \rightarrow$  Concatenate

- Input  
↓  
→ Conv → ReLU → Conv → Add → ReLU → Output  
↓

- Cross-Entropy Loss:

where:

$y_i = \text{true label}$

$p_i$  = predicted probability

- Focal Loss (for imbalanced datasets)
- Contrastive Loss (Siamese networks)
- Triplet Loss (face recognition)

- Learning Rate Scheduling:

Step Decay:

$$lr = lr\_initial \times 0.1^{(epoch/N)}$$

### Cosine Annealing:

$$\text{lr} = \text{lr\_initial} \times (1 + \cos(\times \text{epoch}/\text{total\_epochs}))/2$$

- 4

- Knowledge Distillation

## 5. Model Optimization

- Quantization:

32-bit  $\rightarrow$  8-bit

Size  $\downarrow$  ~4x

Speed  $\uparrow$  ~3x

- Pruning
- Model Distillation
- Weight Sharing

## 6. Visualization Techniques

- Gradient-weighted Class Activation Mapping (Grad-CAM)
- Filter Visualization
- t-SNE for Feature Space
- Attention Maps

## 7. Advanced Applications

- Object Detection
  - YOLO
  - SSD
  - Faster R-CNN
- Semantic Segmentation
  - U-Net
  - DeepLab
  - FCN
- Instance Segmentation
  - Mask R-CNN
- Style Transfer
  - Content/Style Loss
  - AdaIN

## 8. Recent Developments

- Vision Transformers (ViT)
- MobileNet Architectures
- EfficientNet Family
- Self-Attention in CNNs

## 11. Troubleshooting Guide

### Common Issues

1. Checkerboard Artifacts
  - Solution: Use proper stride in transposed convolutions
2. Feature Map Collapse
  - Solution: Check initialization, learning rate
3. Mode Collapse in GANs
  - Solution: Use conditional inputs, proper architecture

## **Debugging Tips**

1. Feature Map Visualization
2. Gradient Flow Analysis
3. Learning Rate Range Test
4. BatchNorm Statistics Monitoring

## **12. Best Practices**

### **Architecture Design**

1. Start with proven architectures
2. Use skip connections for deep networks
3. Consider computational budget
4. Balance model capacity

### **Training Strategy**

1. Progressive learning
2. Curriculum learning
3. Multi-task learning
4. Domain adaptation

### **Deployment Considerations**

1. Model quantization
2. Platform-specific optimization
3. Latency vs accuracy tradeoff
4. Memory constraints

Remember: - Start small and scale up - Monitor training/validation curves - Use GPU acceleration when possible - Save checkpoints regularly