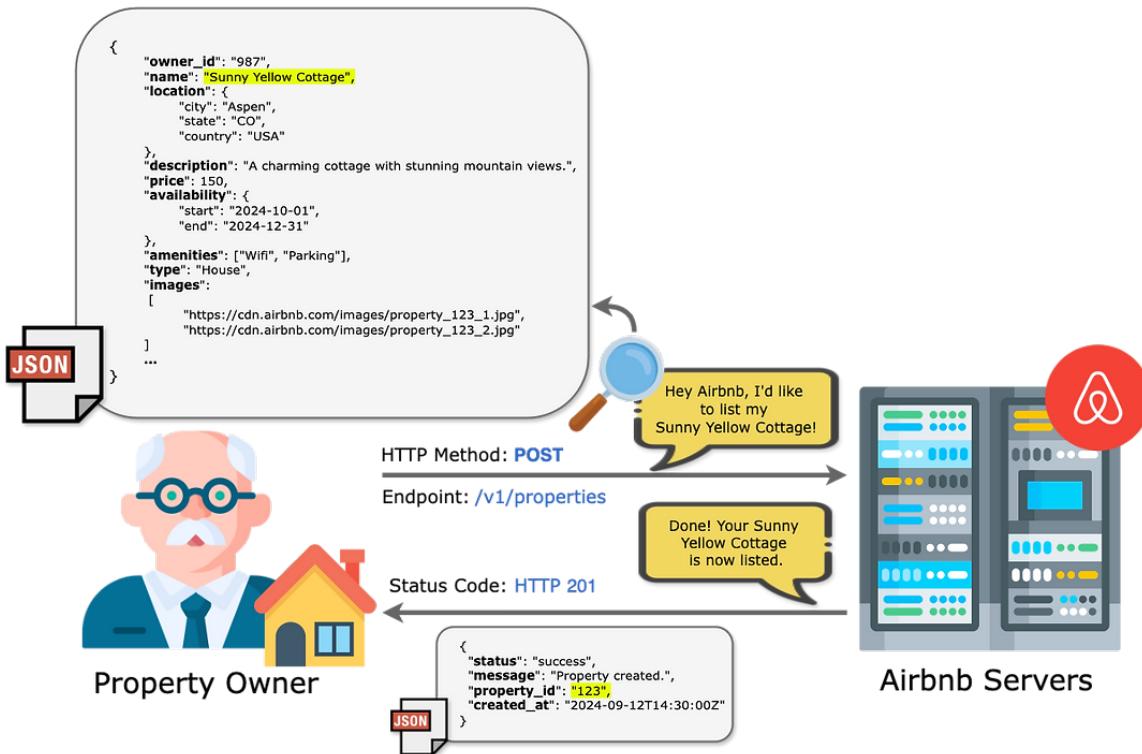
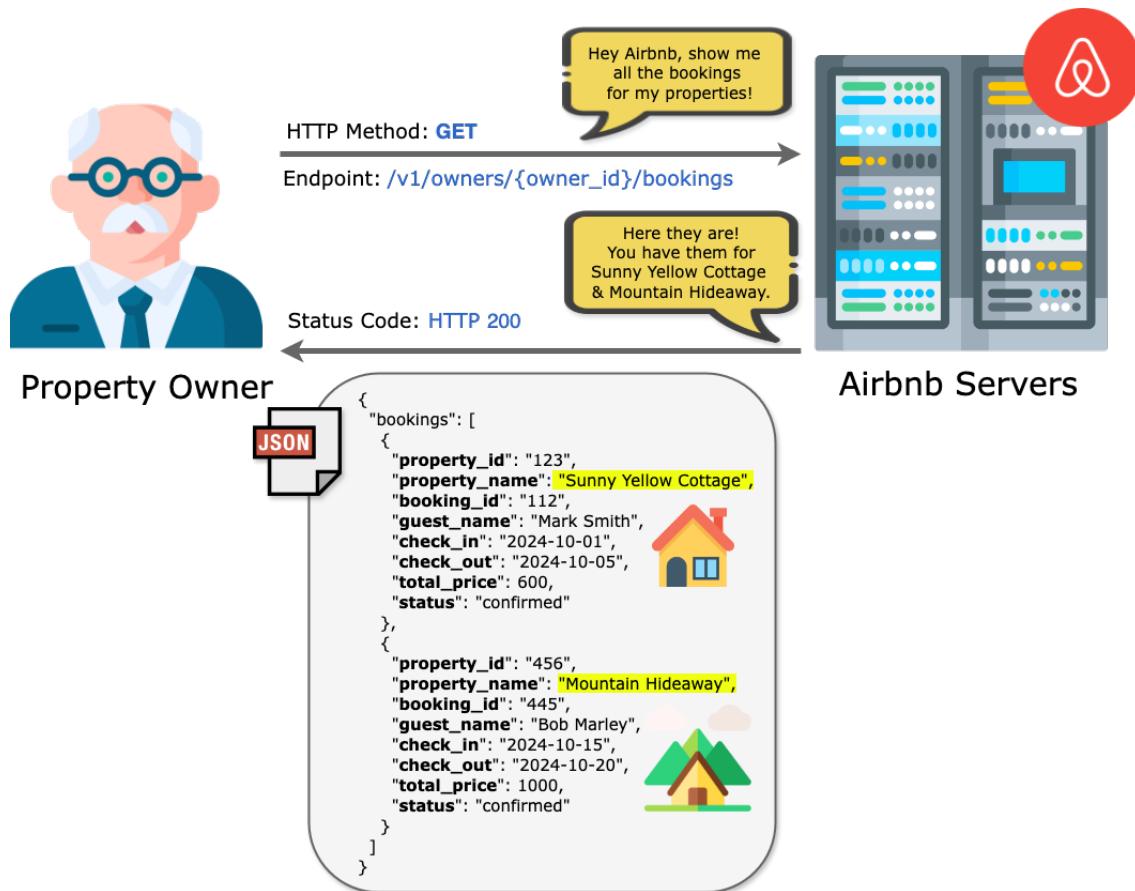


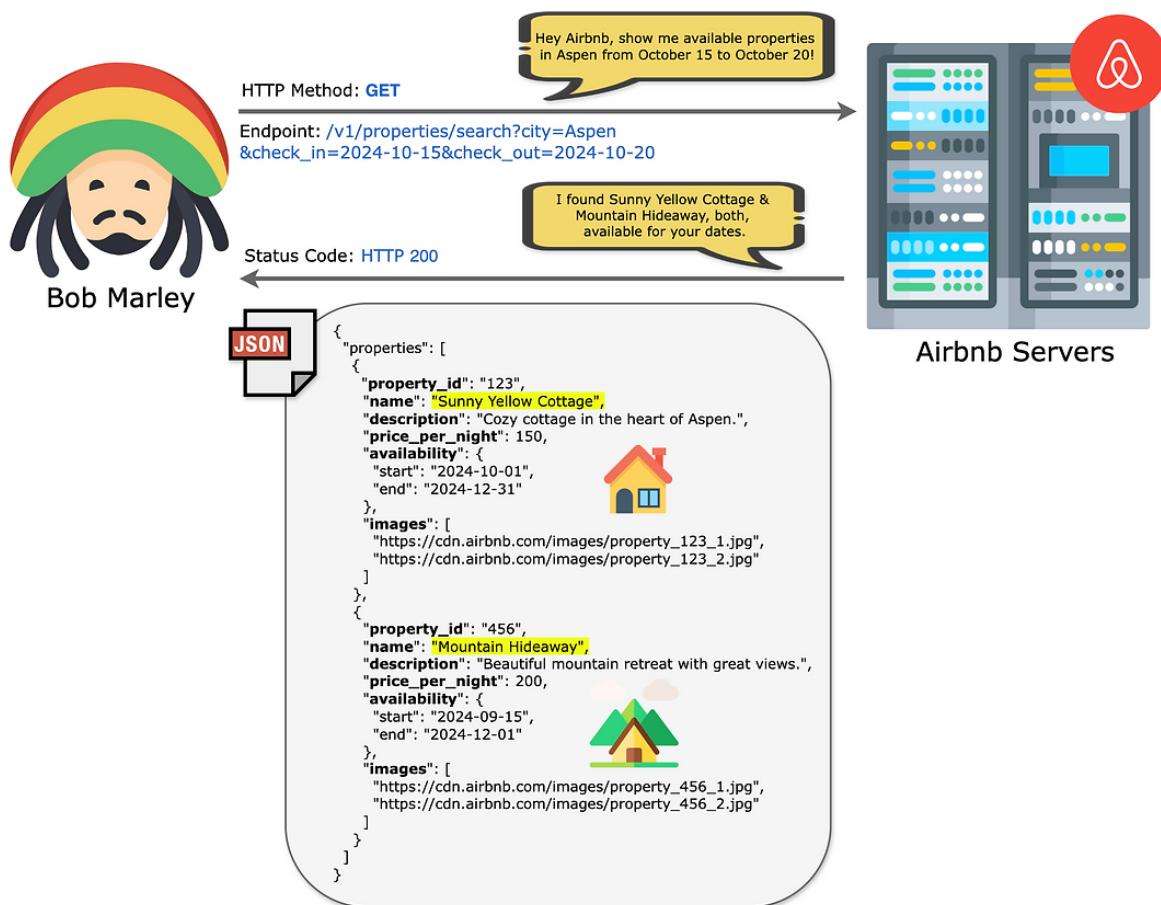
API Design-Add Property (Property Owner)



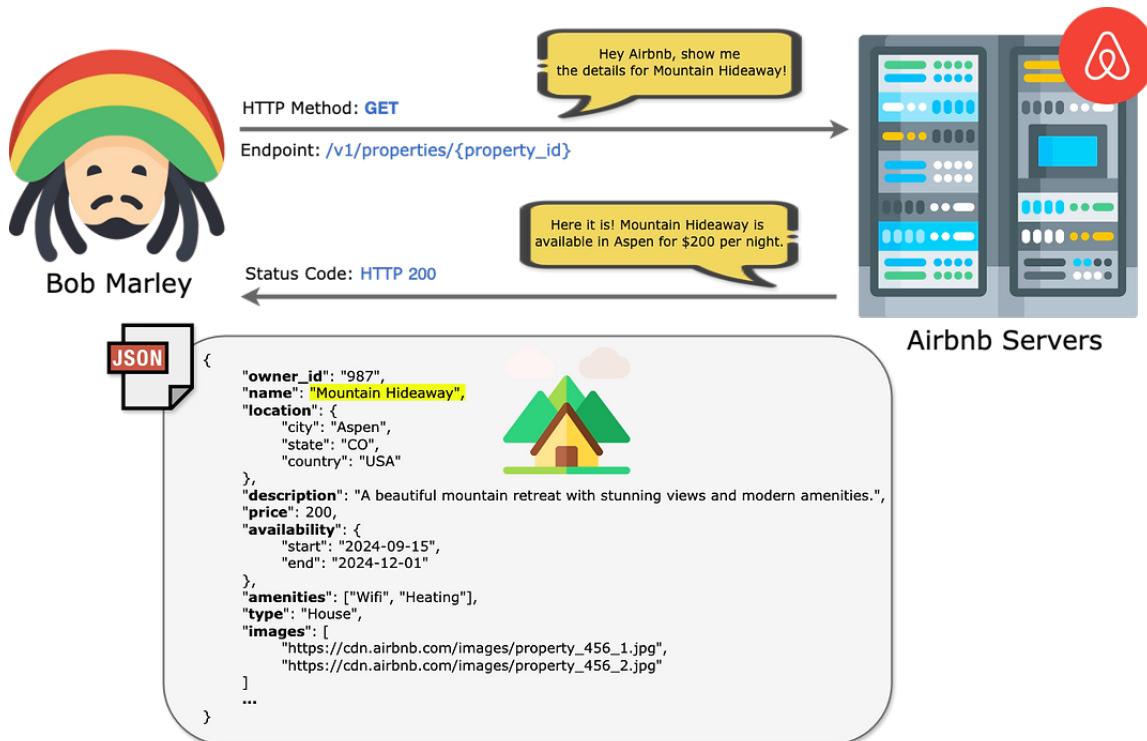
API Design-View Bookings (Property Owner)



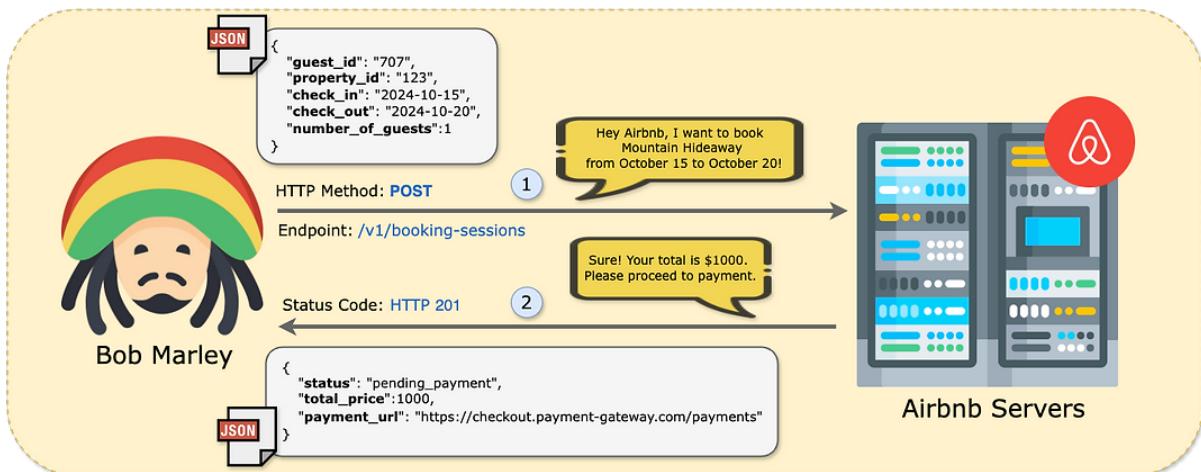
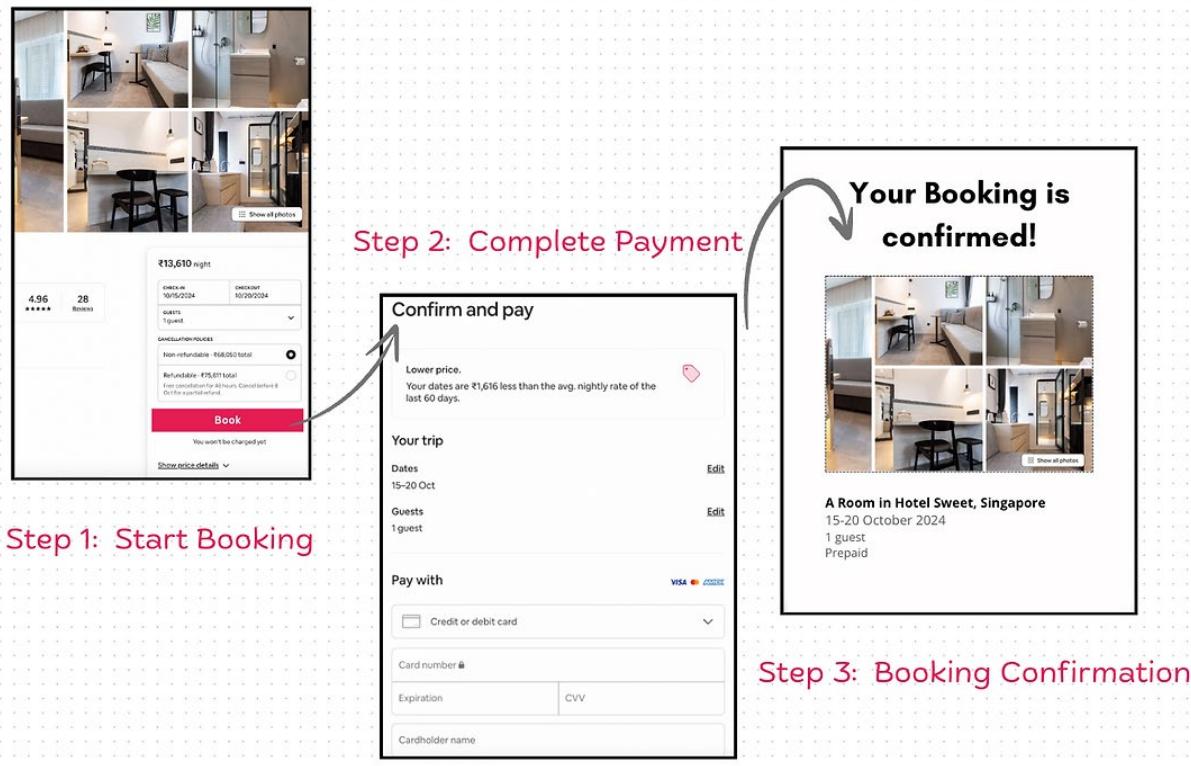
API Design-Search Properties (Guest)

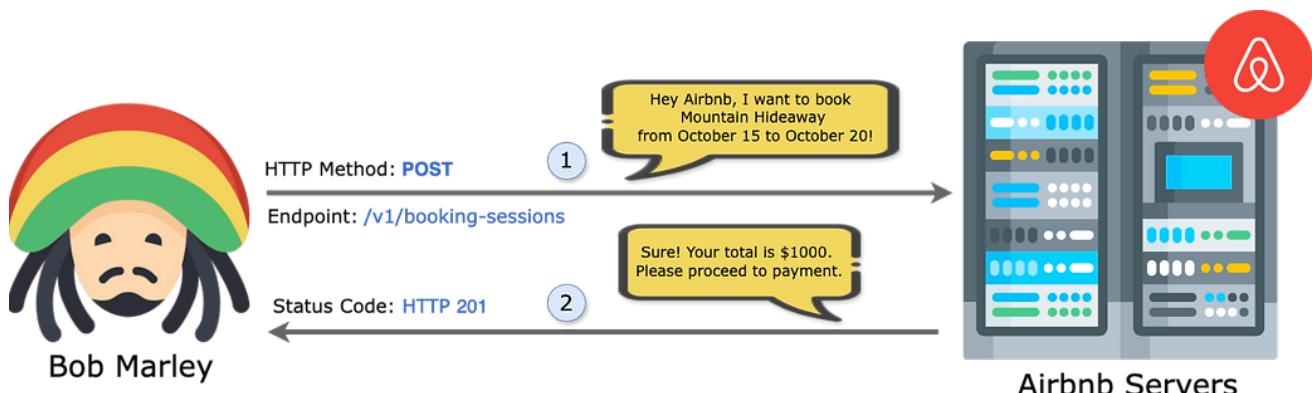


API Design-View Property (Guest)



API Design-Book Property (Guest)



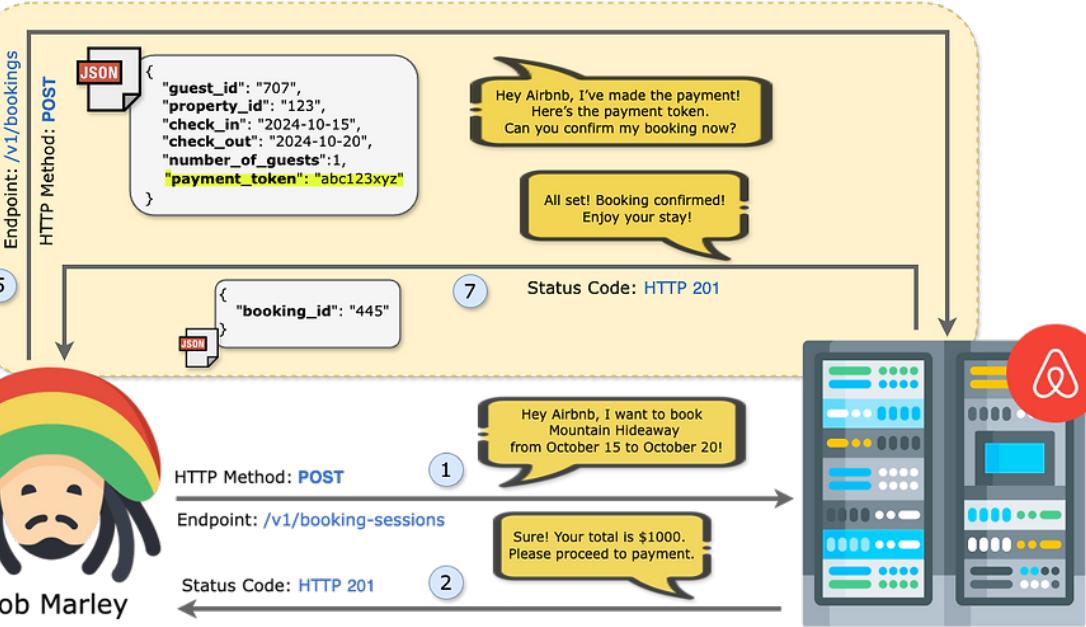
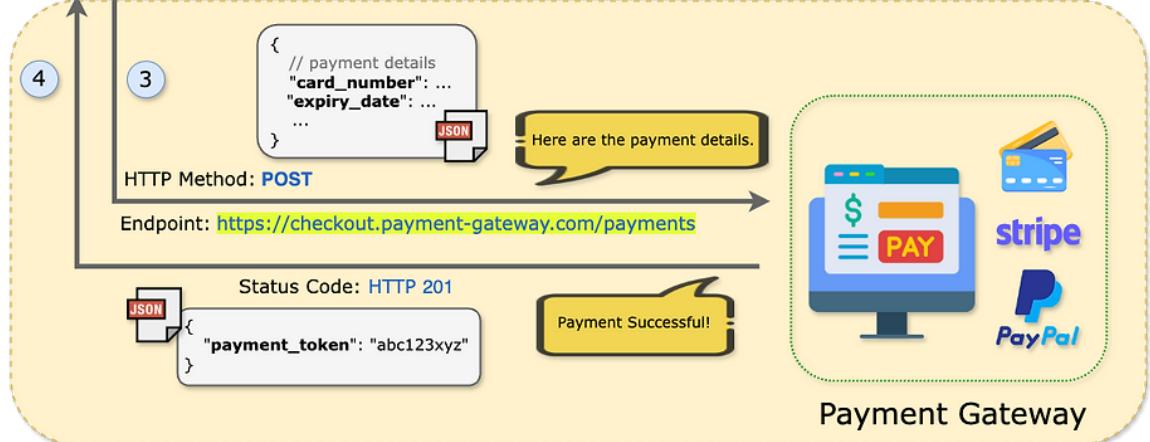


1 Hey Airbnb, I want to book Mountain Hideaway from October 15 to October 20!

2 Sure! Your total is \$1000. Please proceed to payment.

Bob Marley

Airbnb Servers



Endpoint: /v1/bookings
HTTP Method: POST

JSON

{
 "guest_id": "707",
 "property_id": "123",
 "check_in": "2024-10-15",
 "check_out": "2024-10-20",
 "number_of_guests": 1,
 "payment_token": "abc123xyz"
}

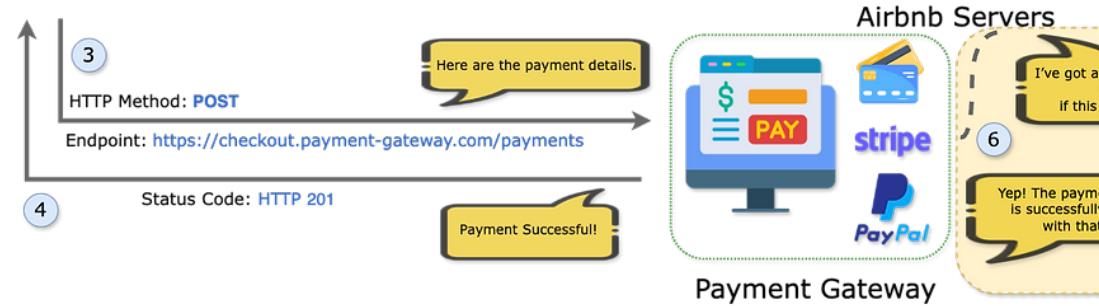
Hey Airbnb, I've made the payment!
Here's the payment token.
Can you confirm my booking now?

All set! Booking confirmed!
Enjoy your stay!

Bob Marley



Airbnb Servers



Endpoint: /v1/booking-sessions
HTTP Method: POST

JSON

{
 "guest_id": "707",
 "property_id": "123",
 "check_in": "2024-10-15",
 "check_out": "2024-10-20",
 "number_of_guests": 1,
 "payment_token": "abc123xyz"
}

Hey Airbnb, I want to book
Mountain Hideaway
from October 15 to October 20!

Sure! Your total is \$1000.
Please proceed to payment.

Bob Marley



Airbnb Servers



Endpoint: https://checkout.payment-gateway.com/payments
HTTP Method: POST

JSON

{
 "card_number": "...",
 "expiry_date": "...",
 ...
}

Here are the payment details.

Payment Successful!

Bob Marley



Airbnb Servers



Endpoint: https://checkout.payment-gateway.com/payments
HTTP Method: POST

JSON

{
 "card_number": "...",
 "expiry_date": "...",
 ...
}

Here are the payment details.

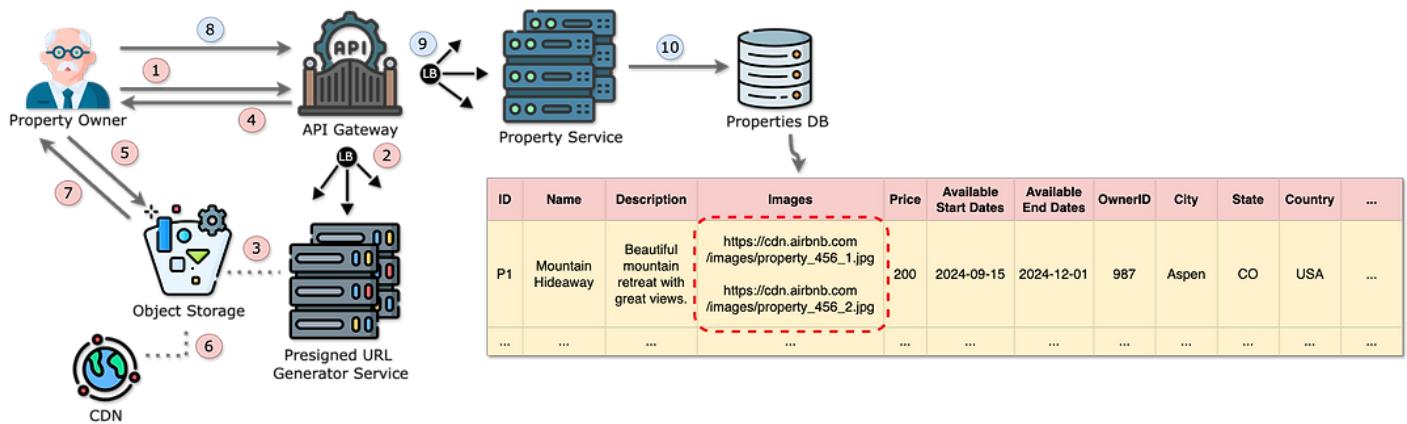
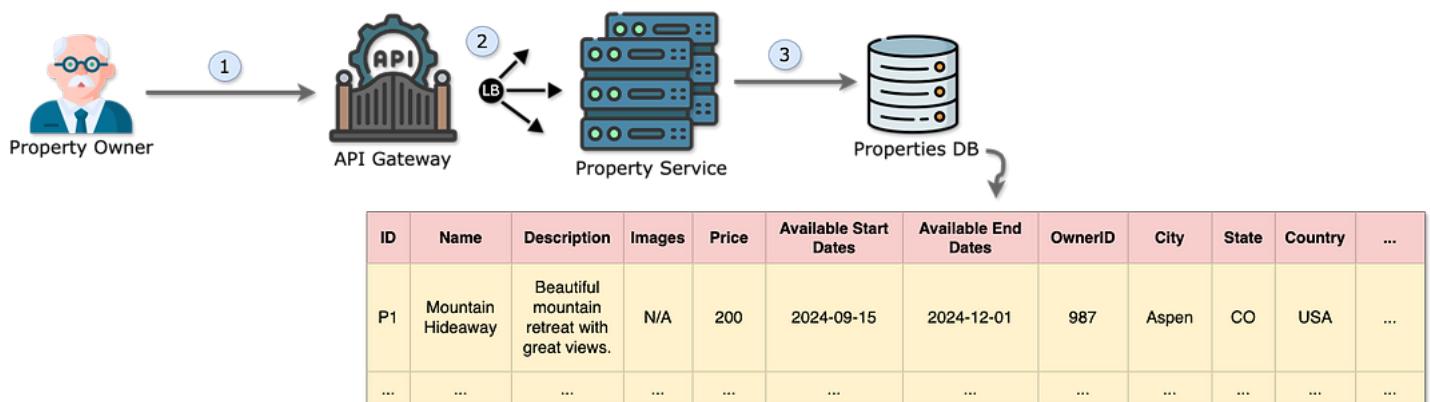
Payment Successful!

Bob Marley

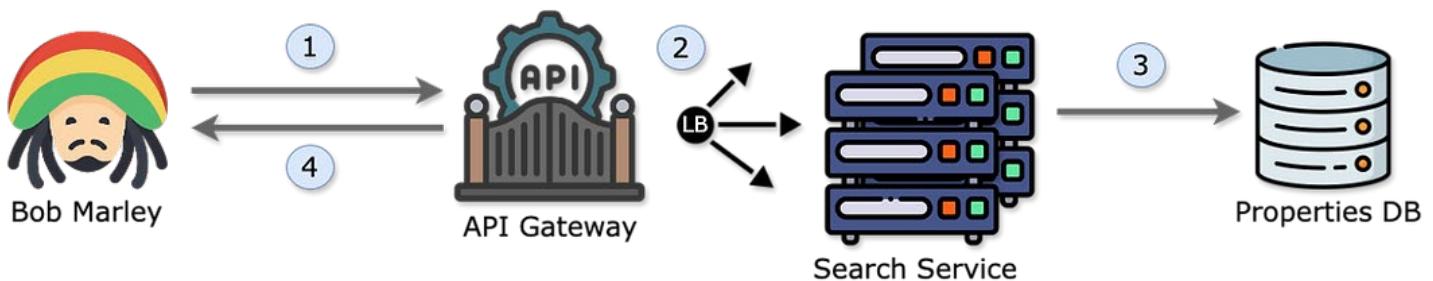


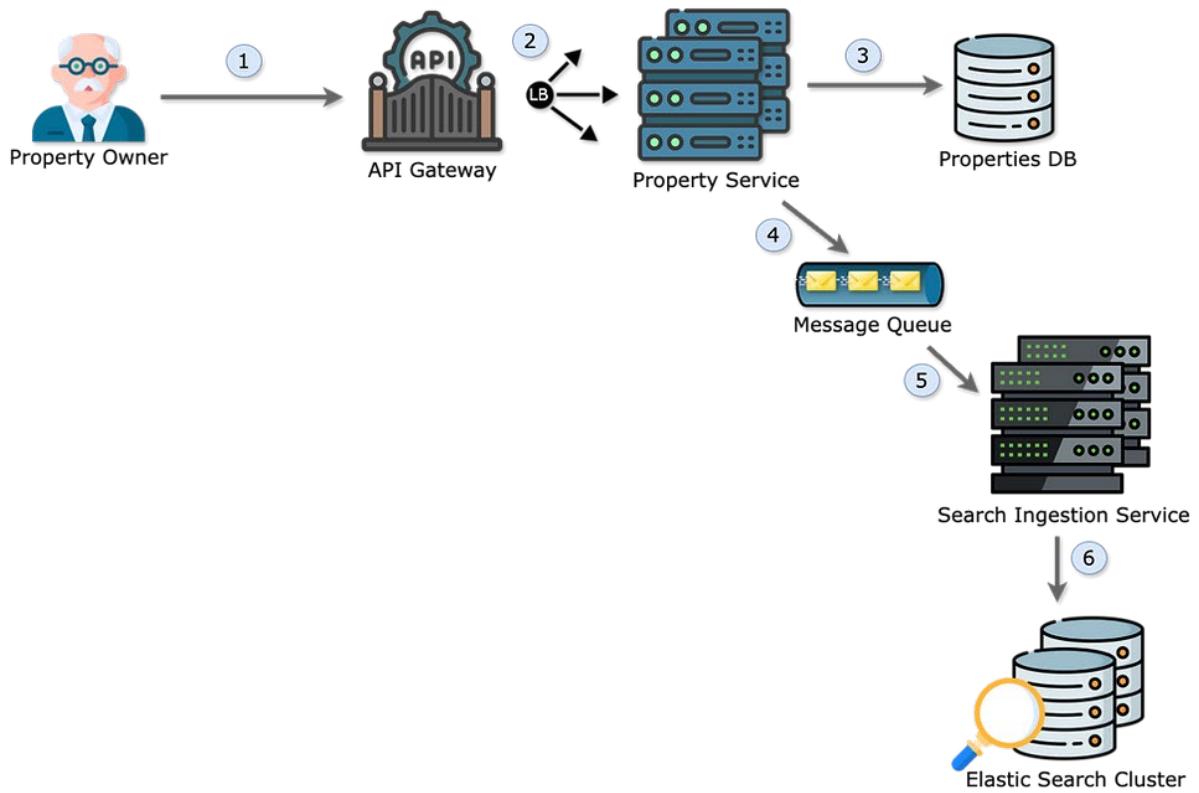
Airbnb Servers

High Level Design-Add/Update property

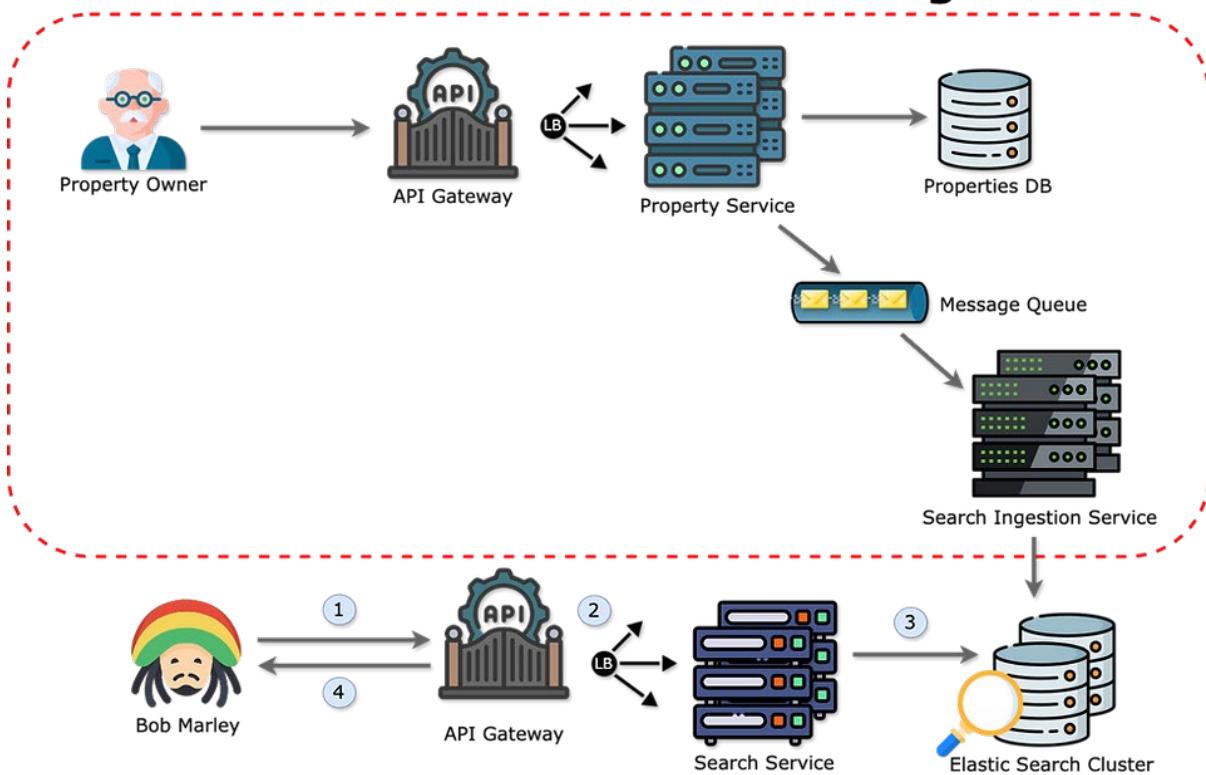


High Level Design- Search Property (Guest)

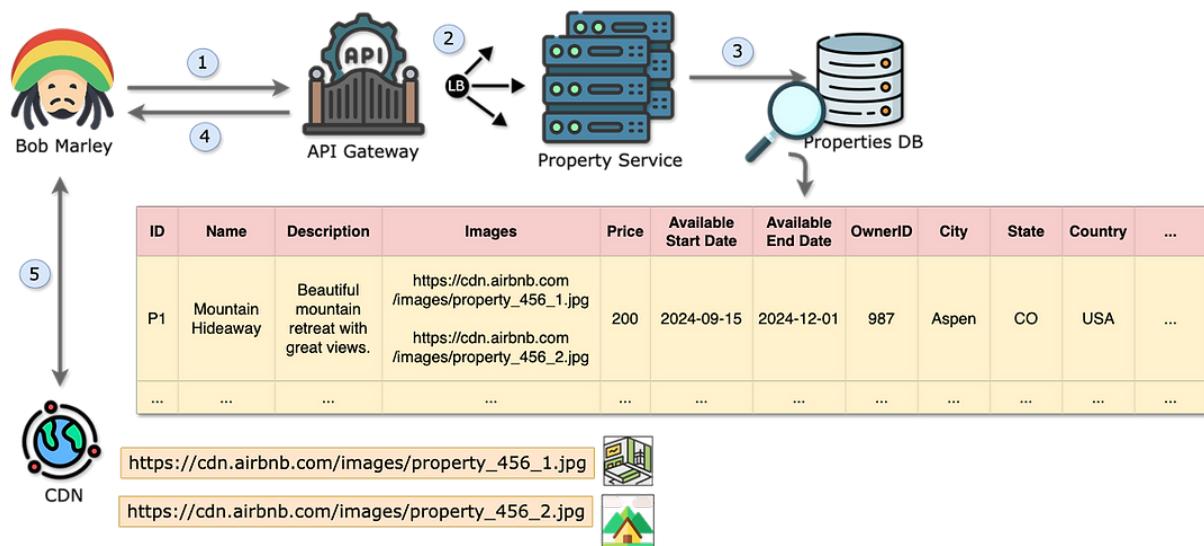




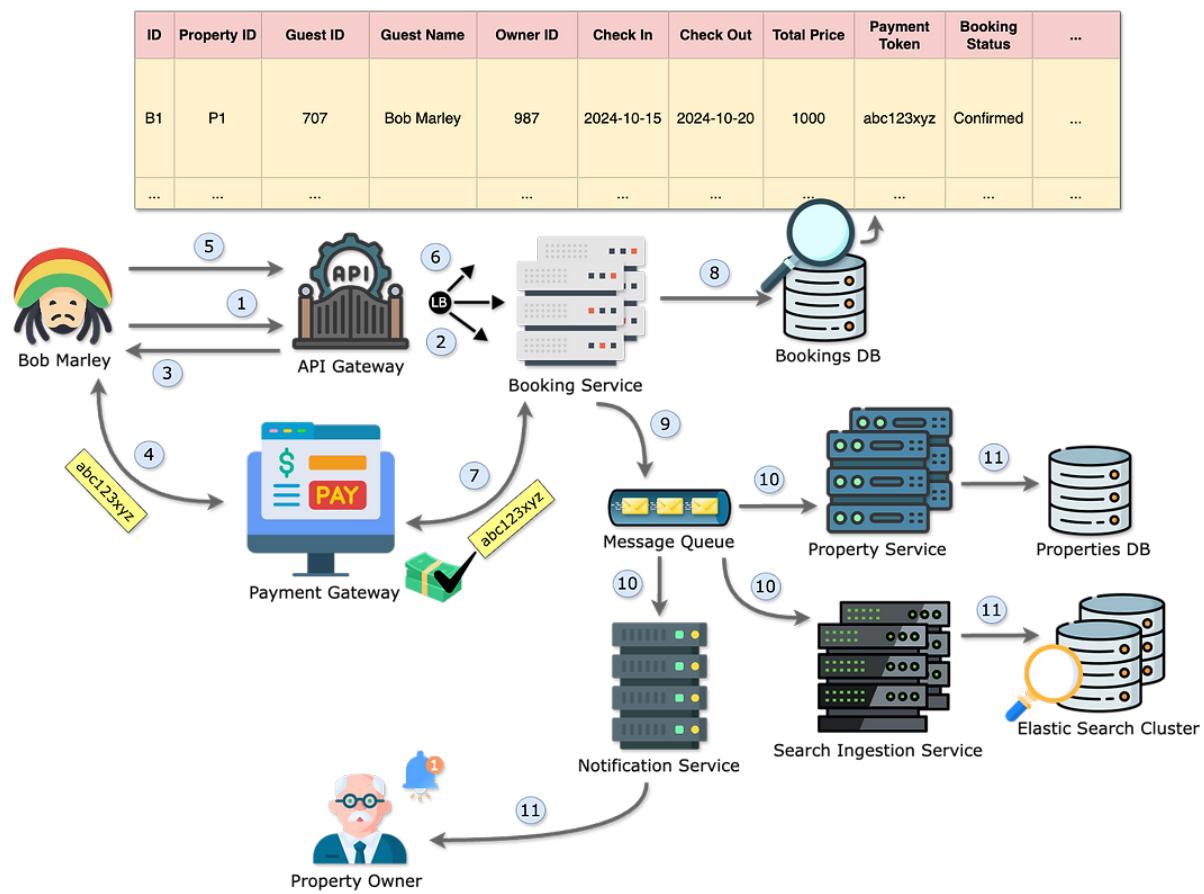
Search Data 'Pre'-Processing



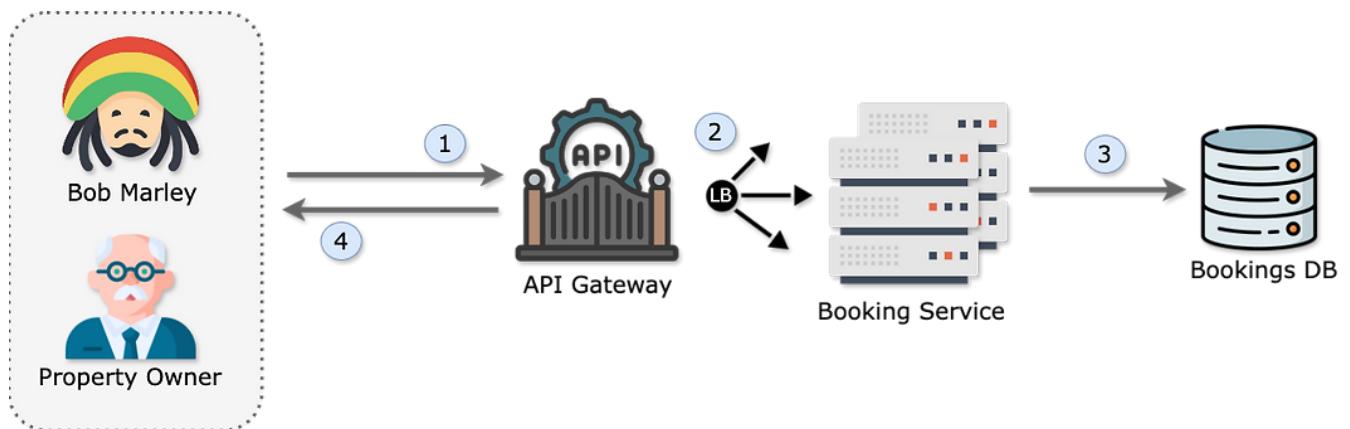
High Level Design-View Property (Guest)



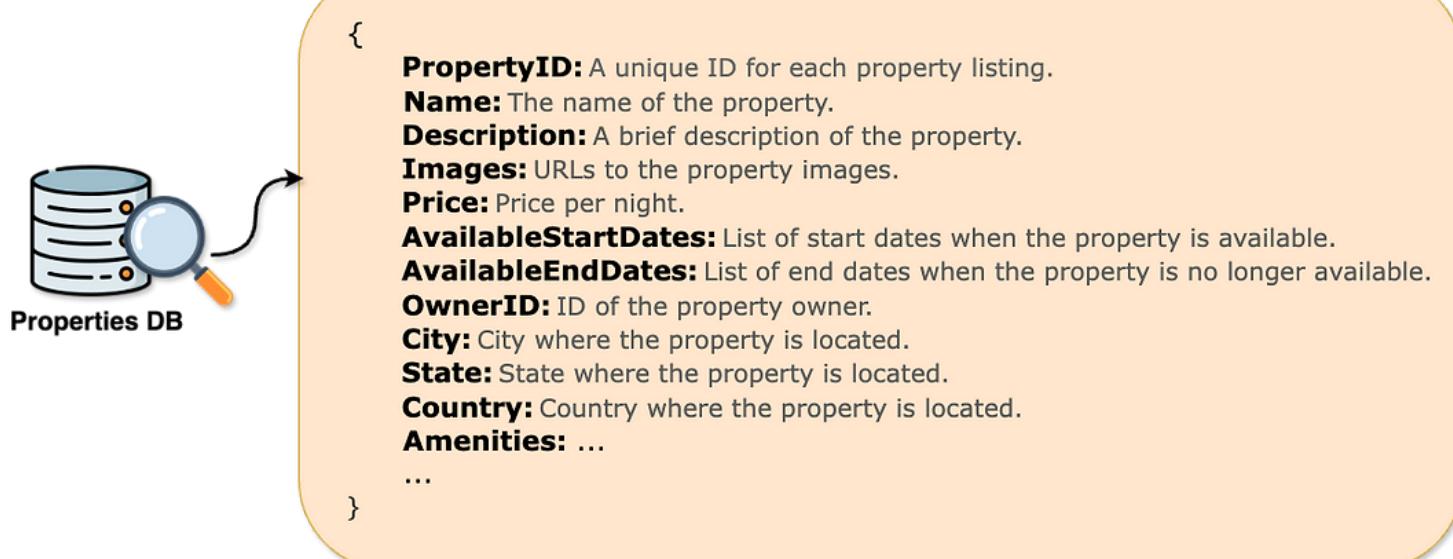
High Level Design-Book Property (Guest)



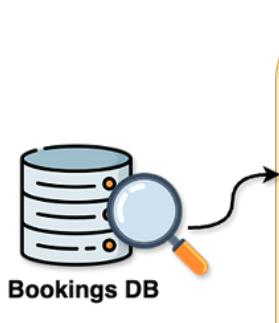
High Level Design-View Bookings (Property Owner/Guest)



Properties DB Schema

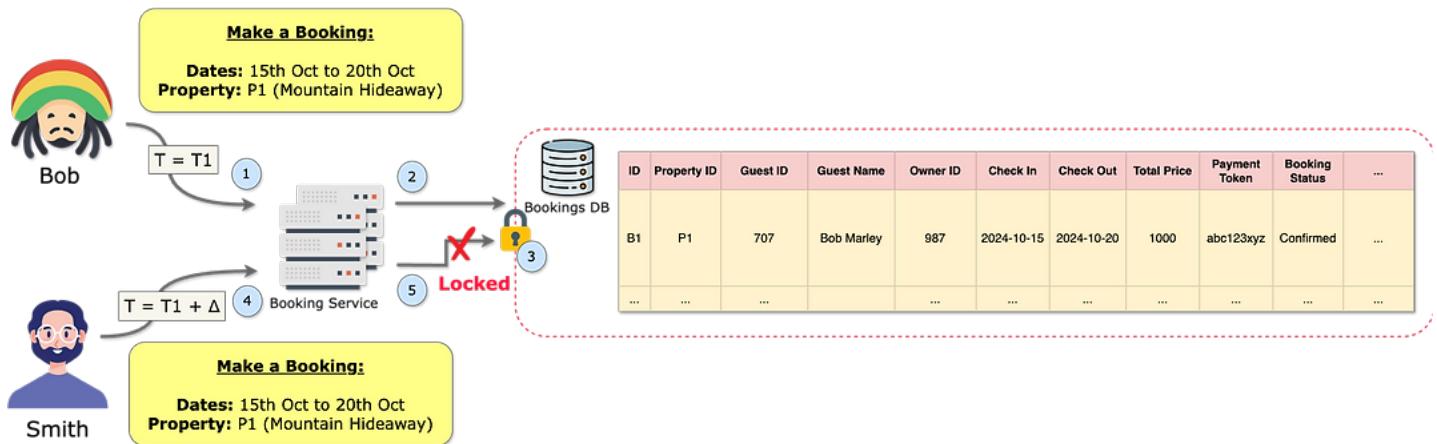


Bookings DB Schema



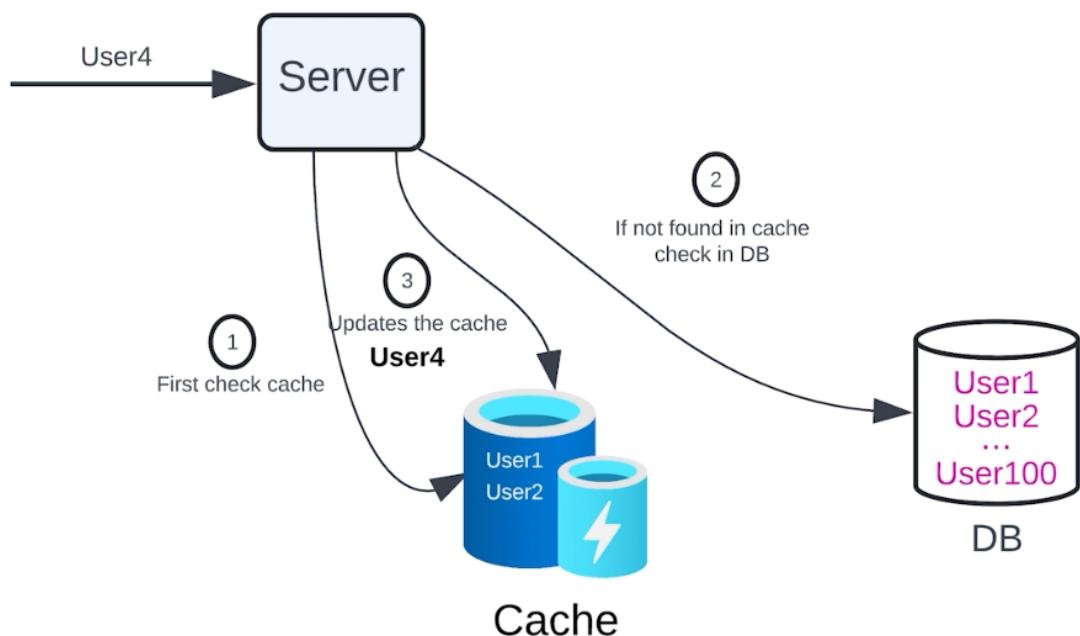
- {
 - BookingID:** A unique identifier for each group.
 - PropertyID:** The ID of the property being booked.
 - GuestID:** The ID of the guest making the booking.
 - GuestName:** Name of the guest.
 - OwnerID:** The ID of the property owner.
 - CheckInDate:** Date when the guest will check in.
 - CheckOutDate:** Date when the guest will check out.
 - TotalPrice:** Total price for the stay.
 - PaymentToken:** Token received after successful payment.
 - BookingStatus:** Status of the booking (e.g., "Confirmed", "Completed", "Cancelled")
- ...

Handling Concurrent Bookings

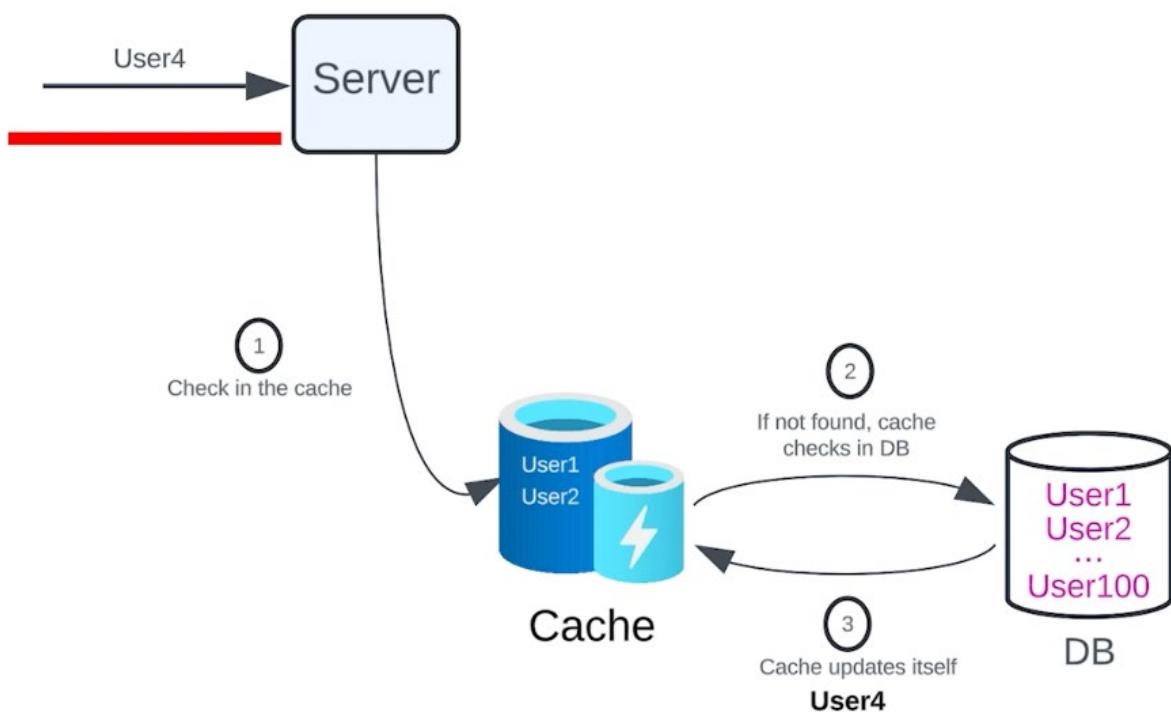


Caching Strategies (Read)

Cache Aside Strategy

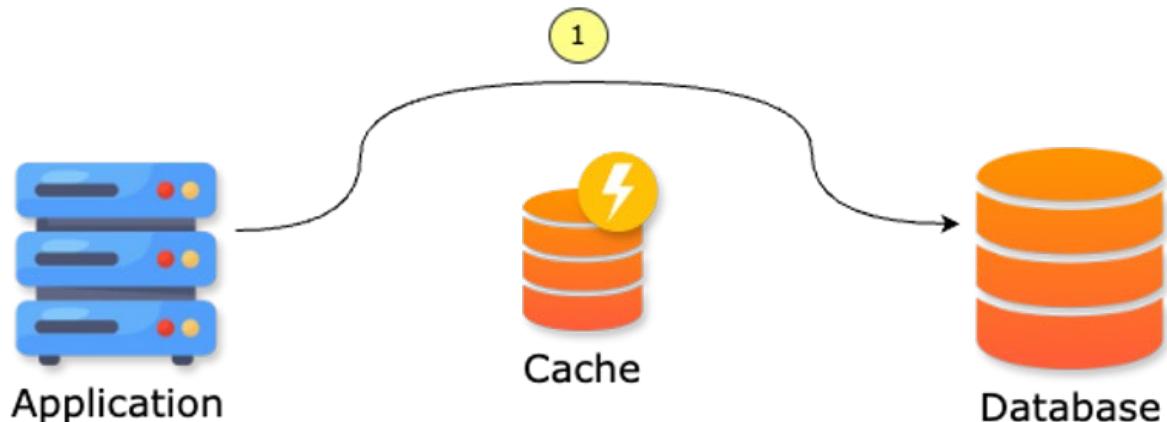


Read Through Strategy

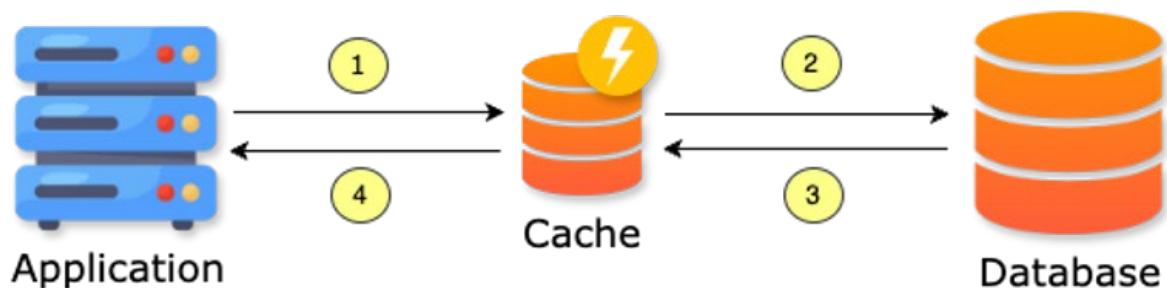


Caching Strategies (Write)

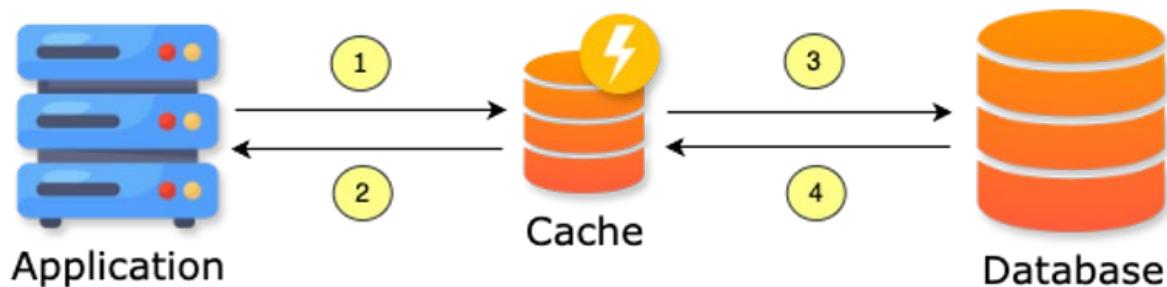
Write Aside Strategy



Write Through Strategy

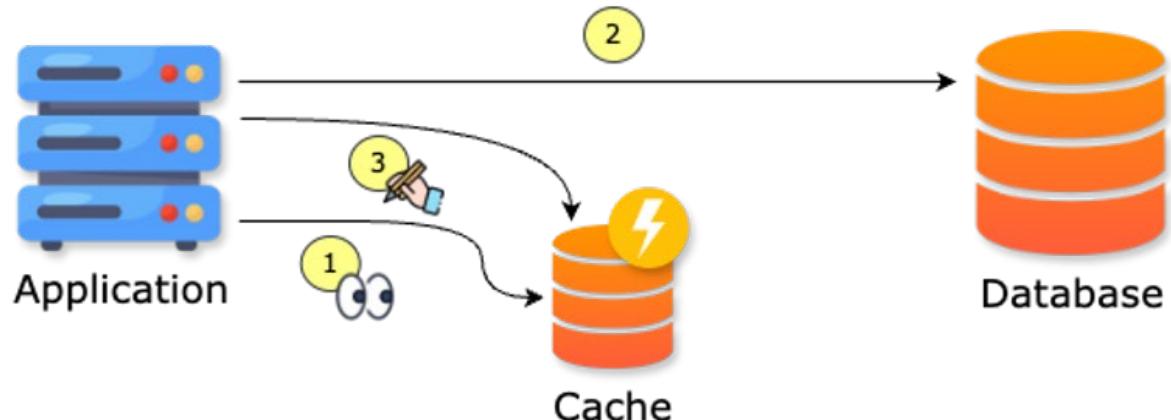


Write Behind Strategy

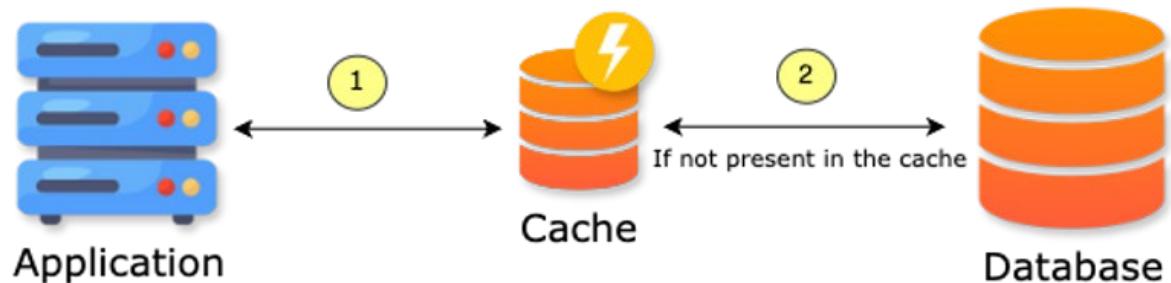


Read Caching Strategies (Pros & Cons)

Cache Aside Strategy

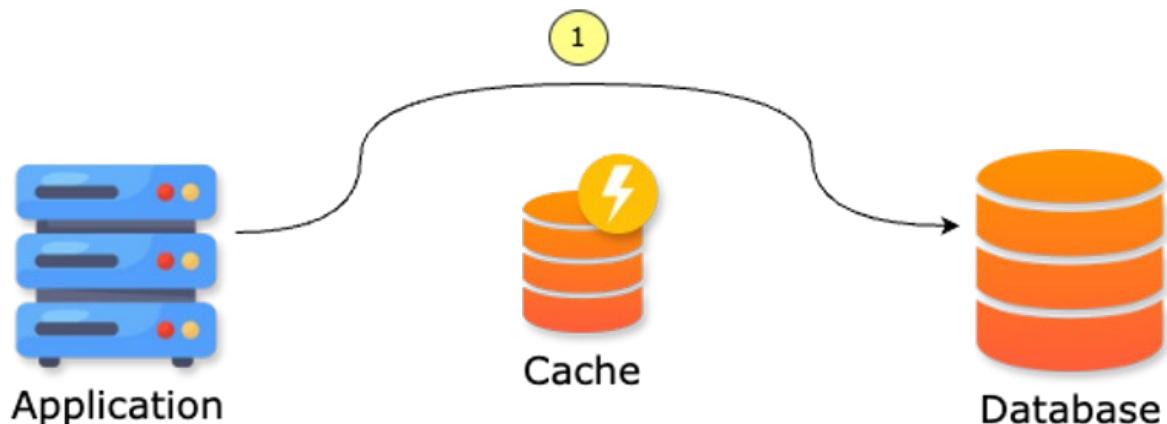


Read Through Strategy

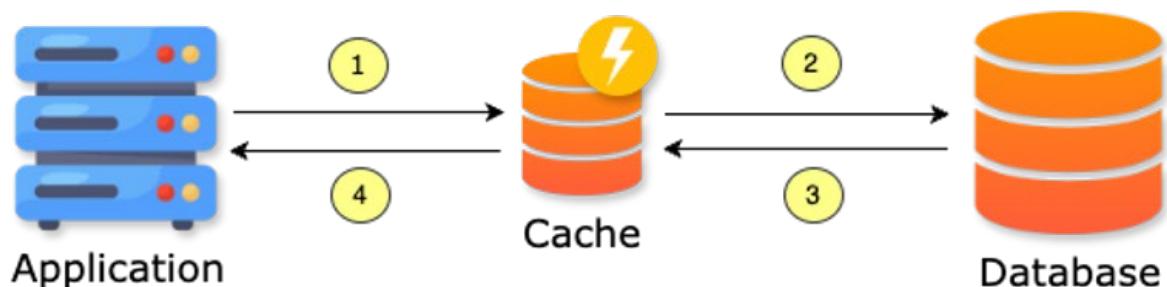


Read Caching Strategies (Pros & Cons)

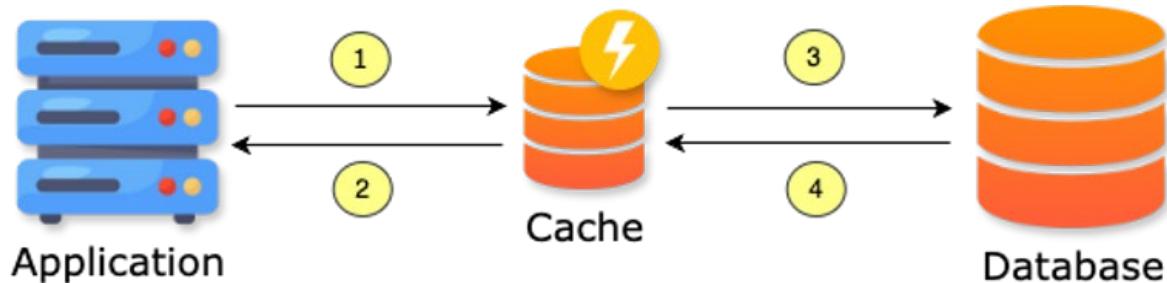
Write Aside Strategy



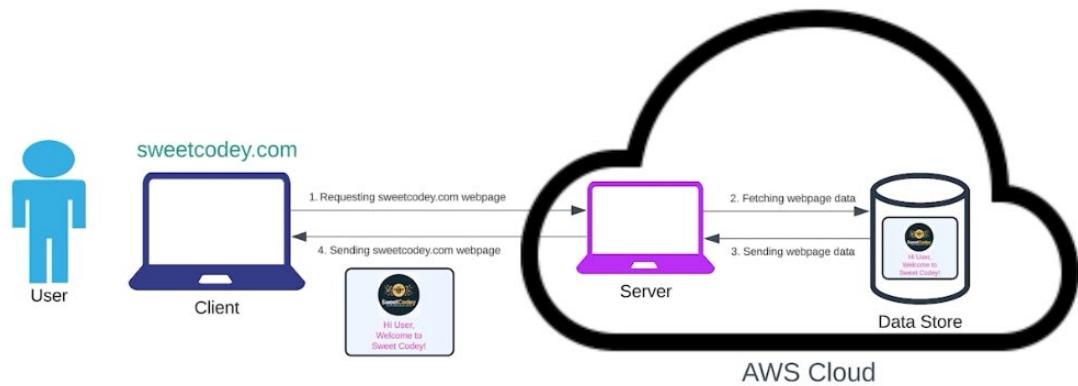
Write Through Strategy



Write Behind Strategy



Cloud Computing



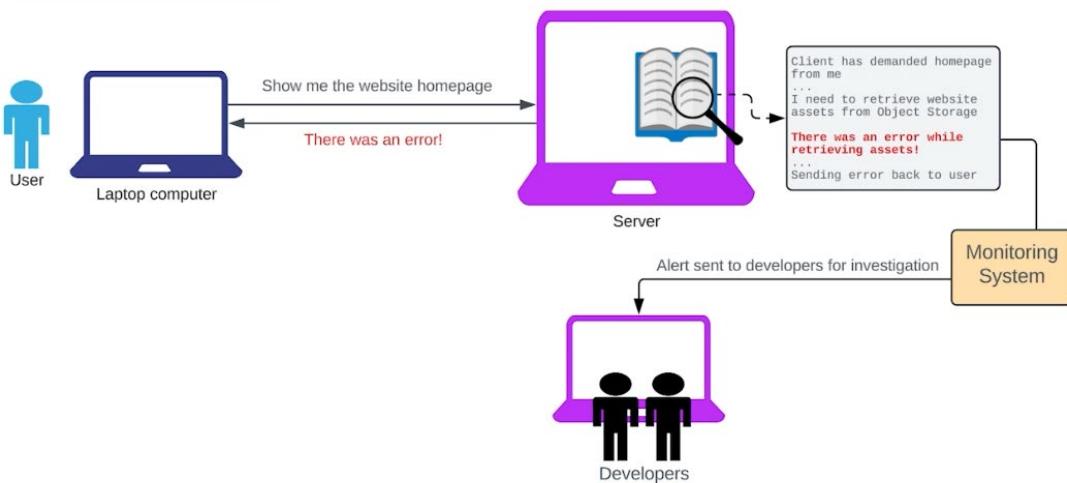
Logging & Monitoring

Logging

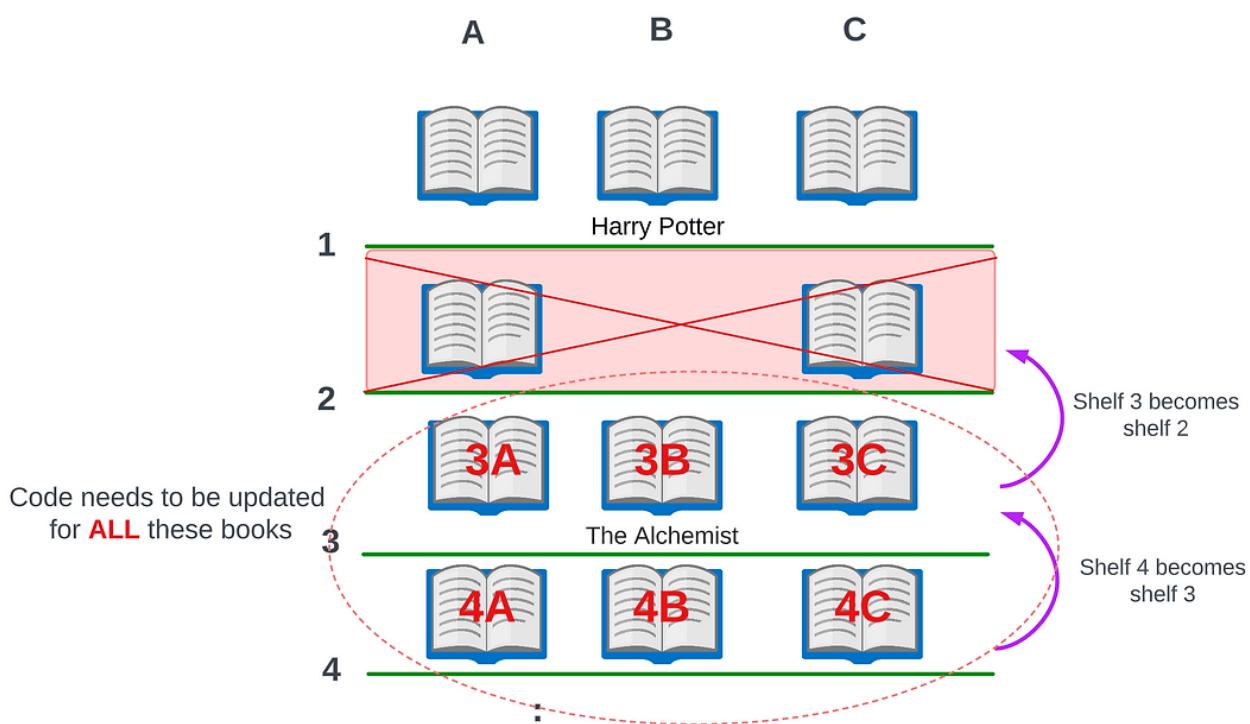
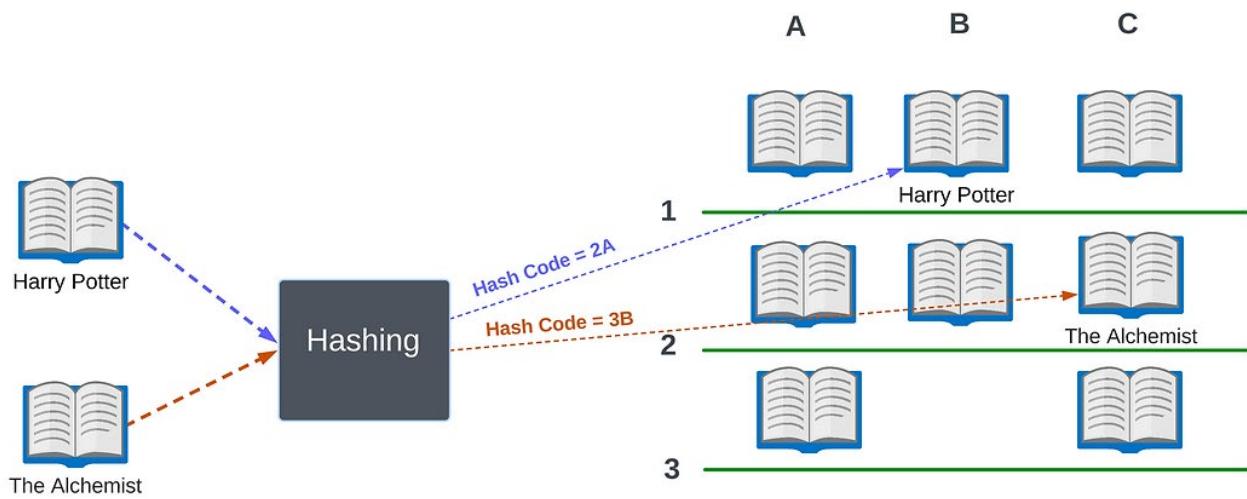
Server Writing Diary == Logging



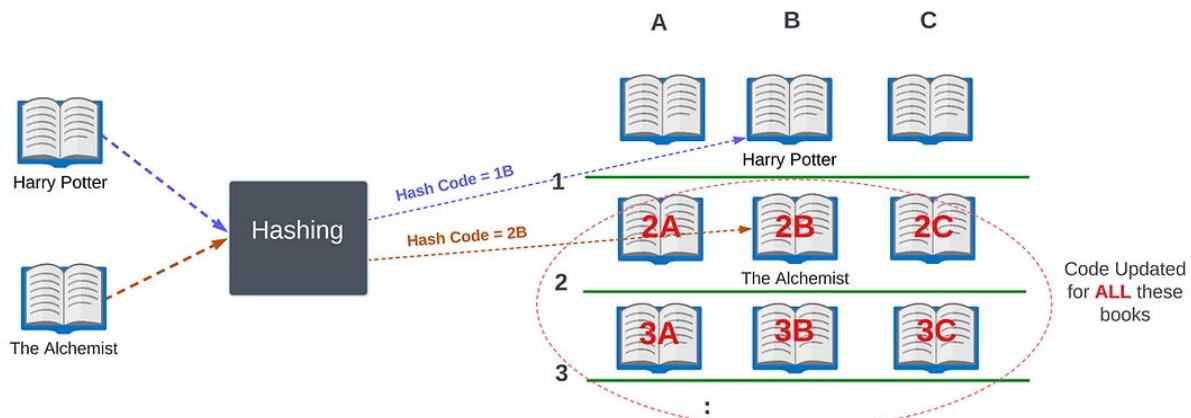
Monitoring



Hashing

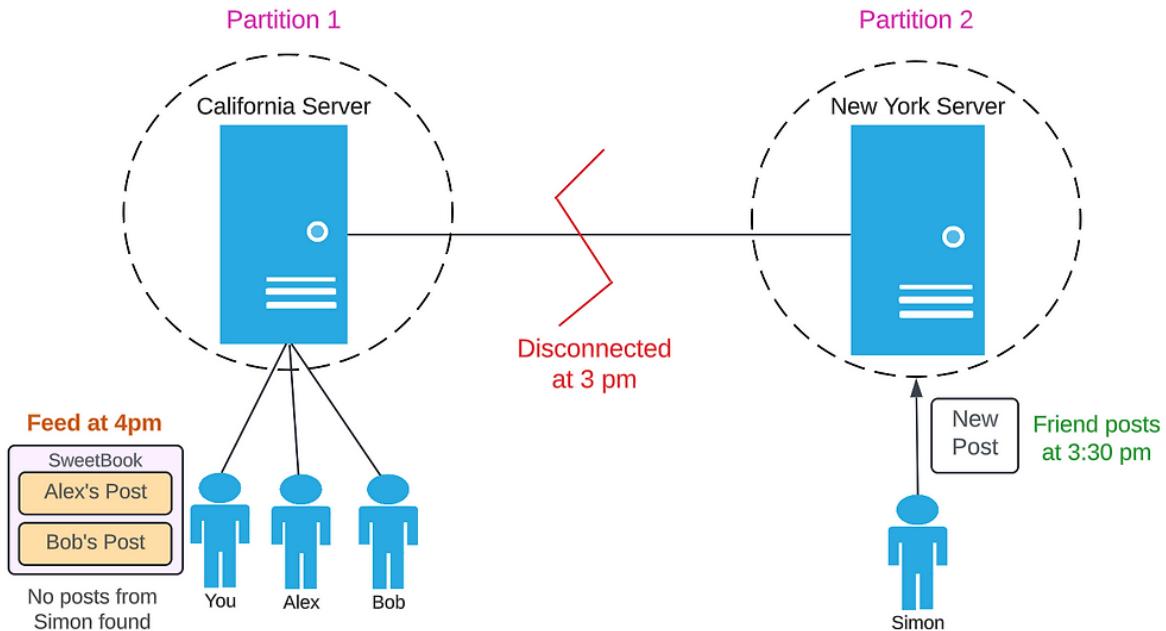


Consistent Hashing

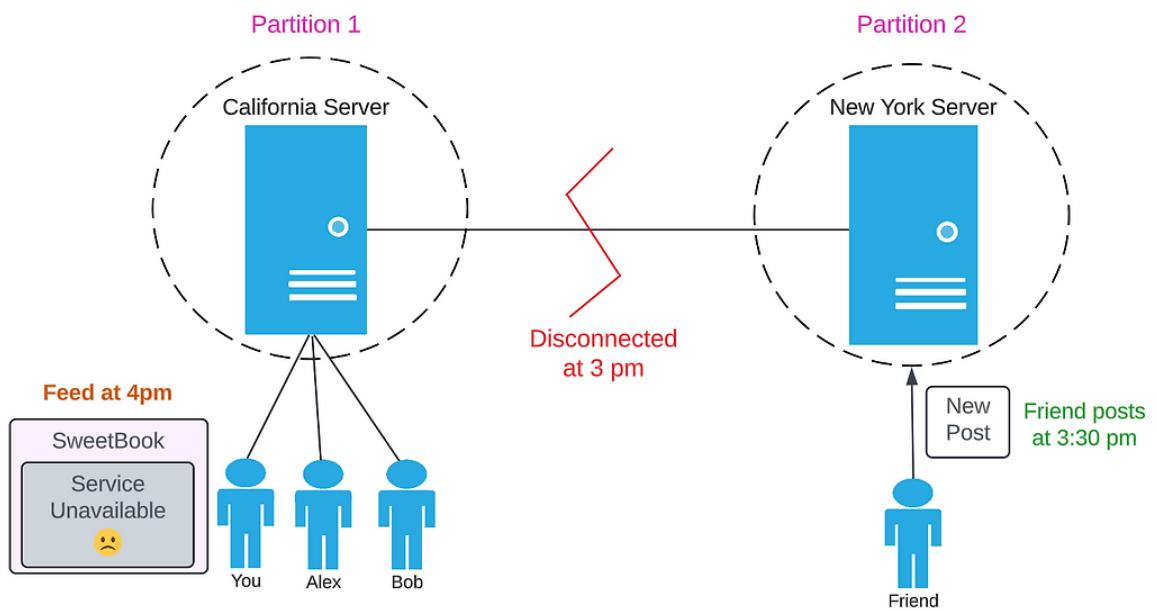


CAP THEOREM

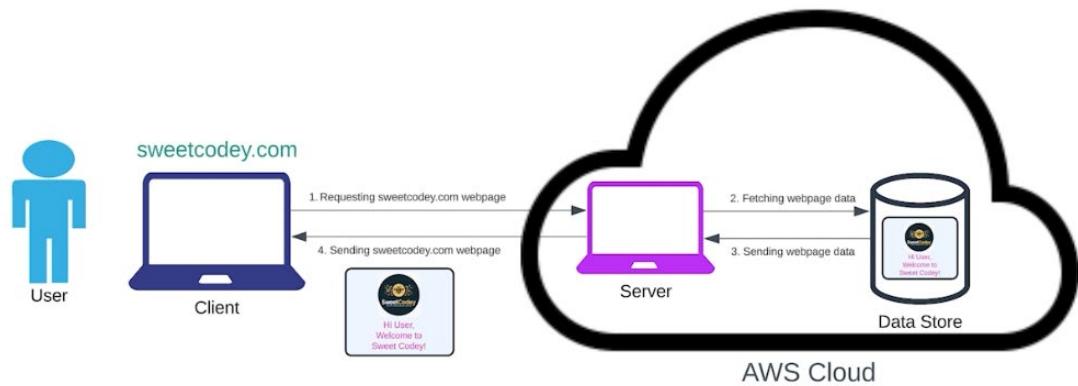
Prioritizing Availability



Prioritizing Consistency



Cloud Computing



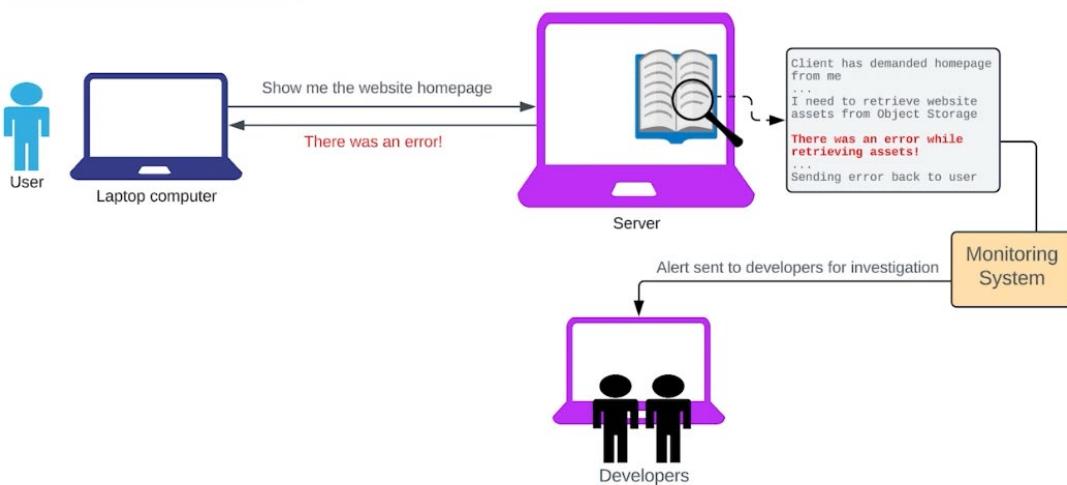
Logging & Monitoring

Logging

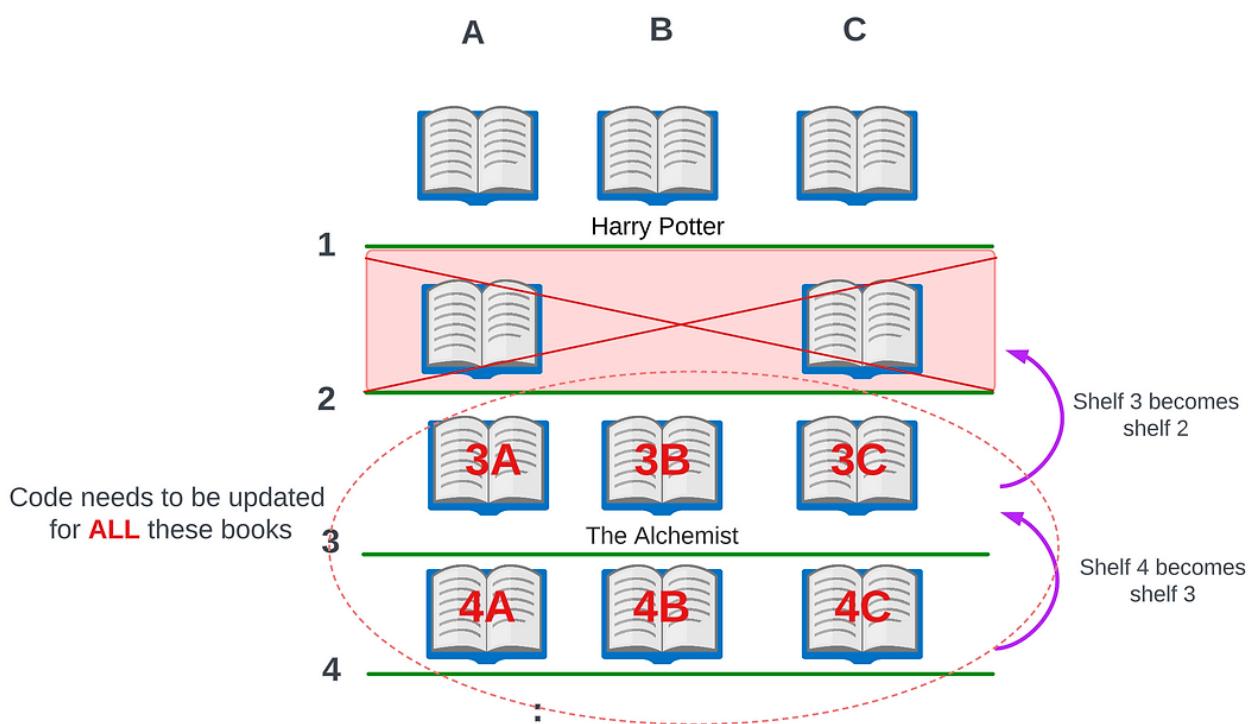
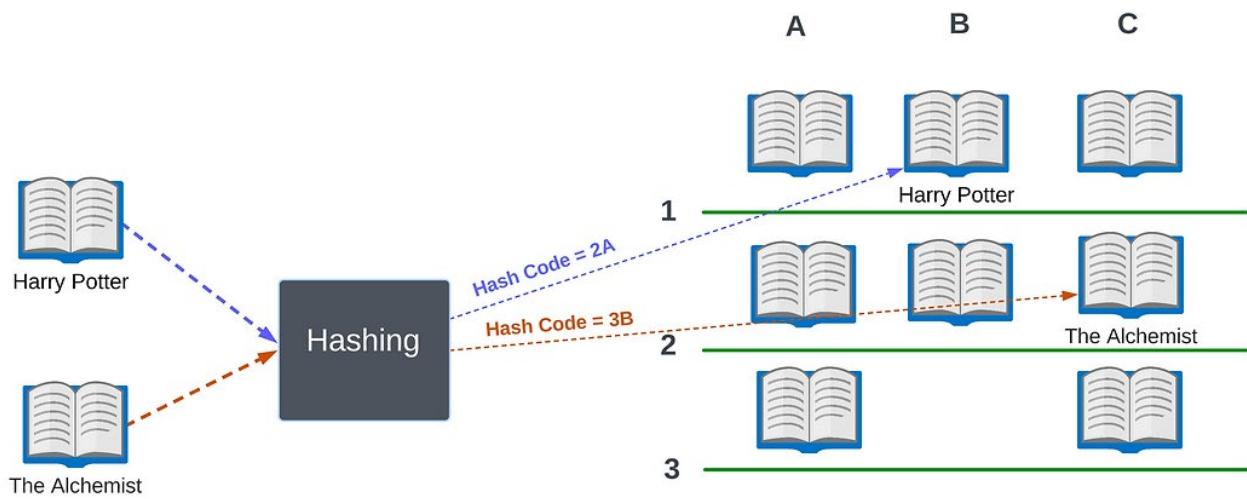
Server Writing Diary == Logging



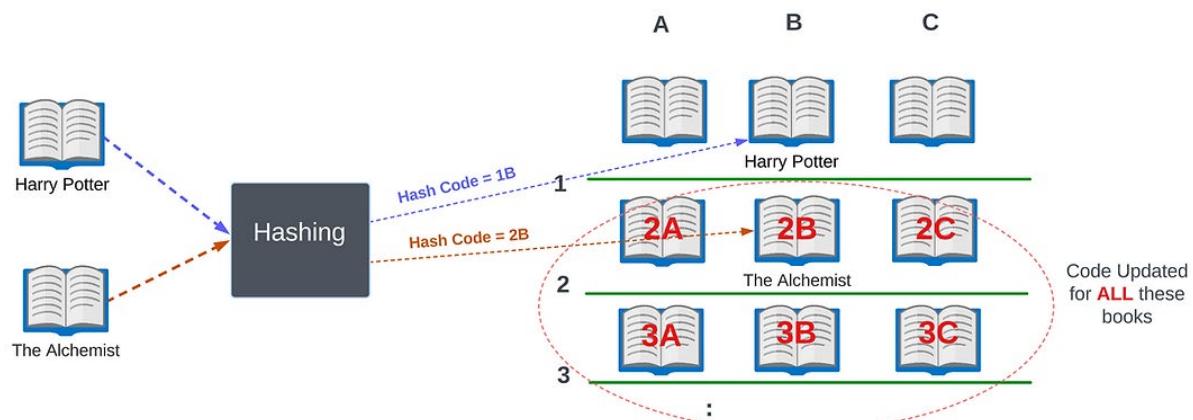
Monitoring



Hashing

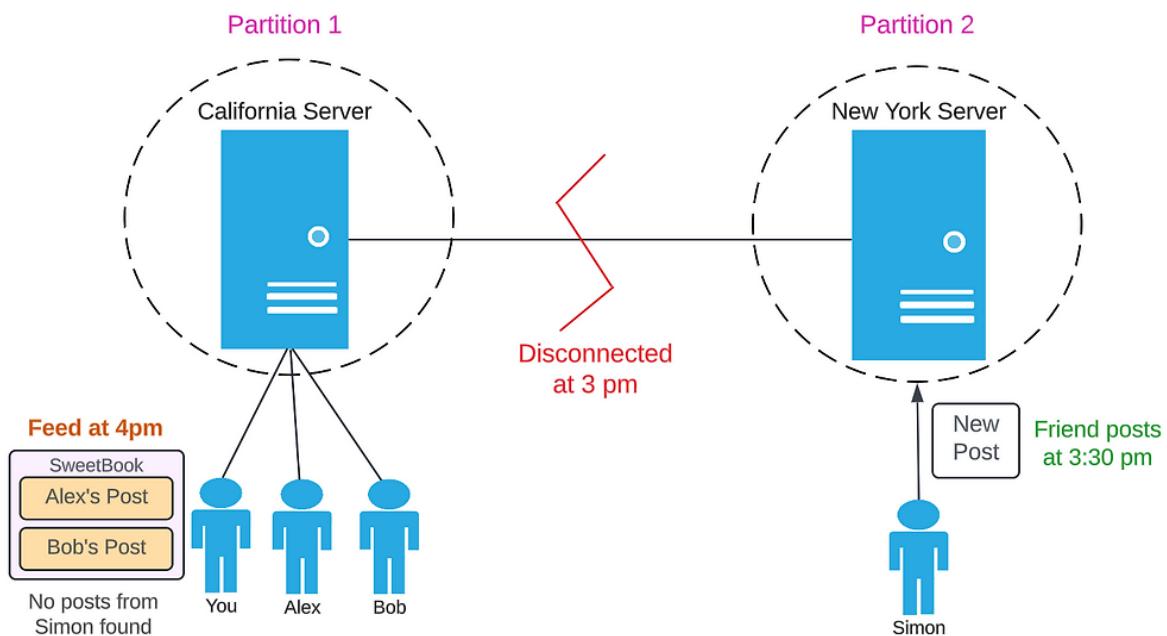


Consistent Hashing

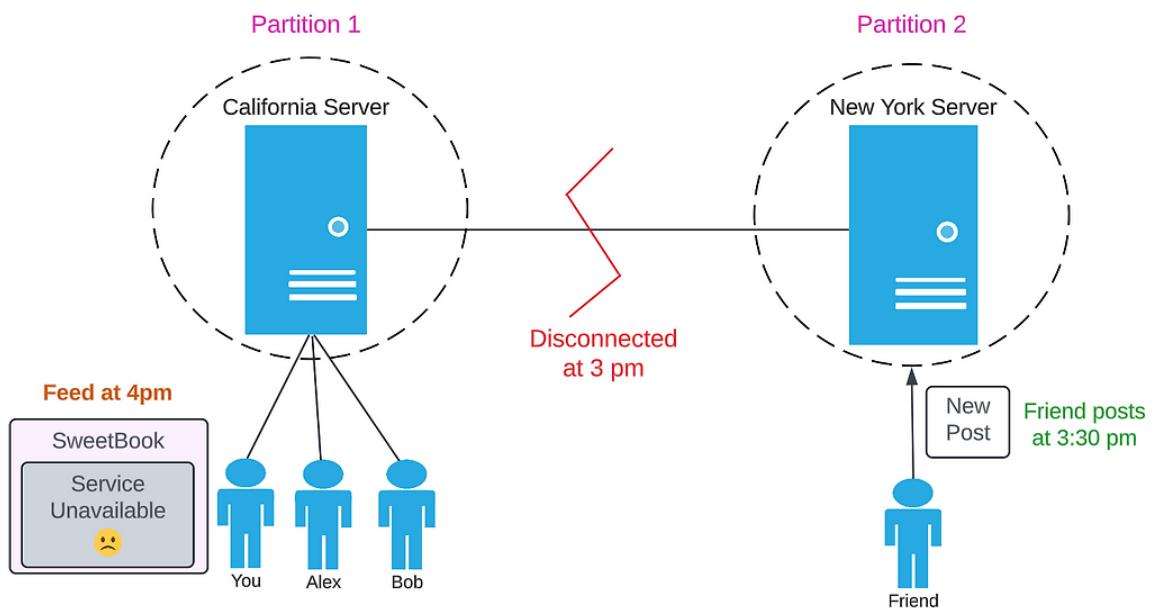


CAP THEOREM

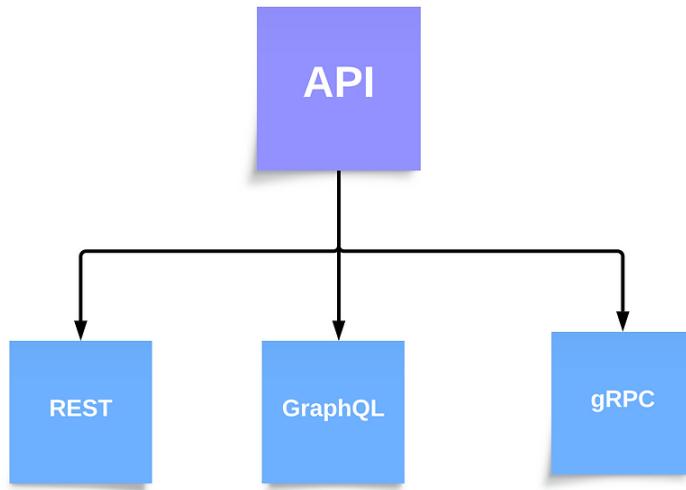
Prioritizing Availability



Prioritizing Consistency



API

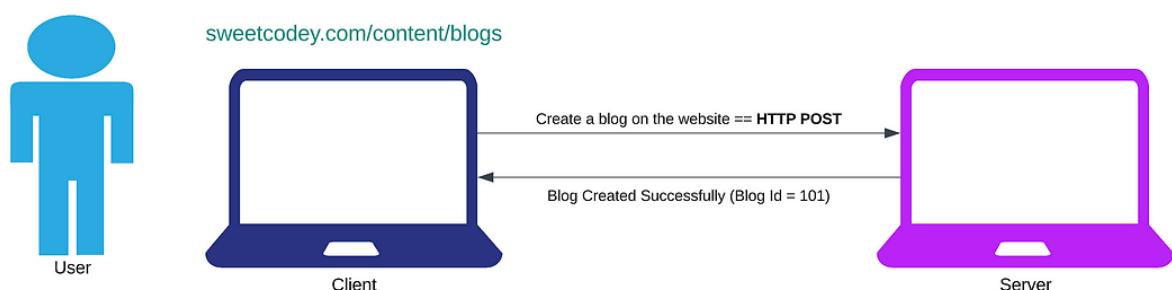


REST API

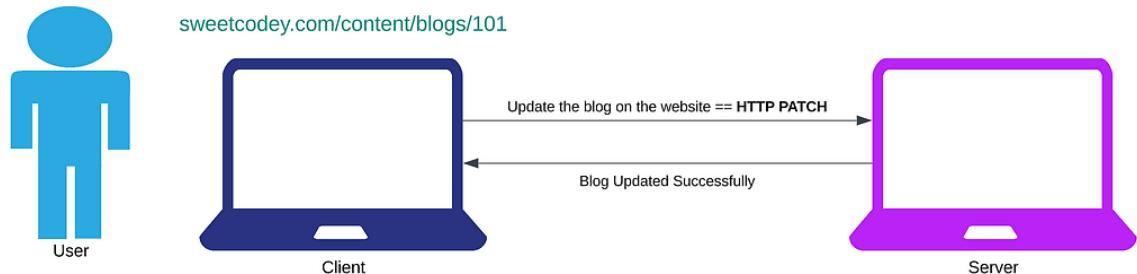
HTTP GET



HTTP POST



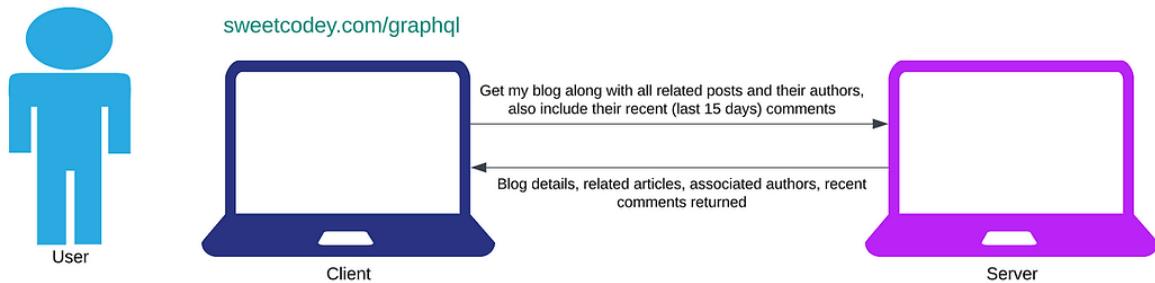
HTTP PATCH



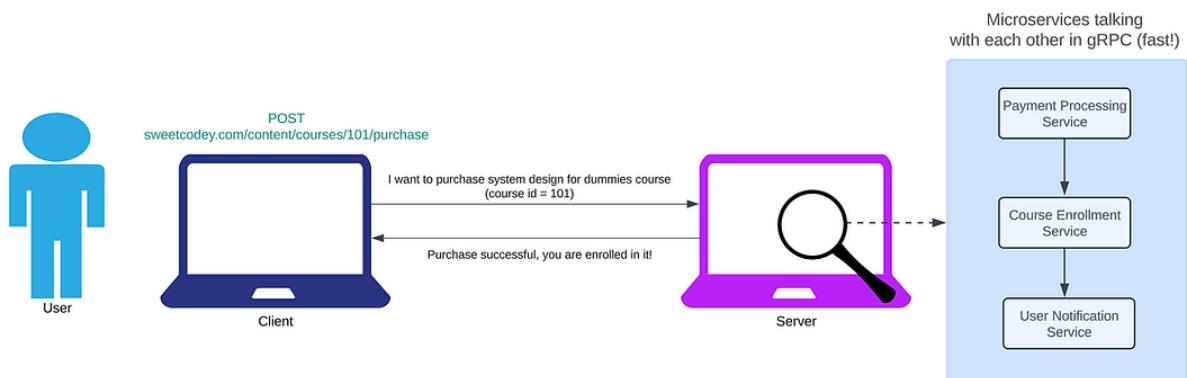
HTTP DELETE



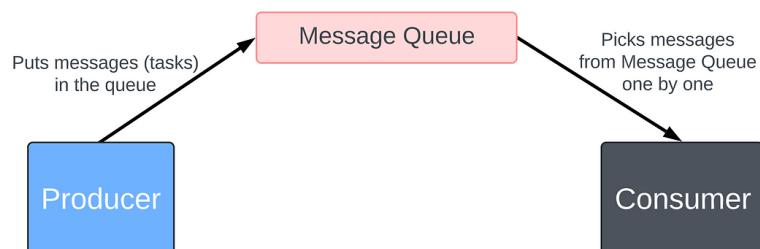
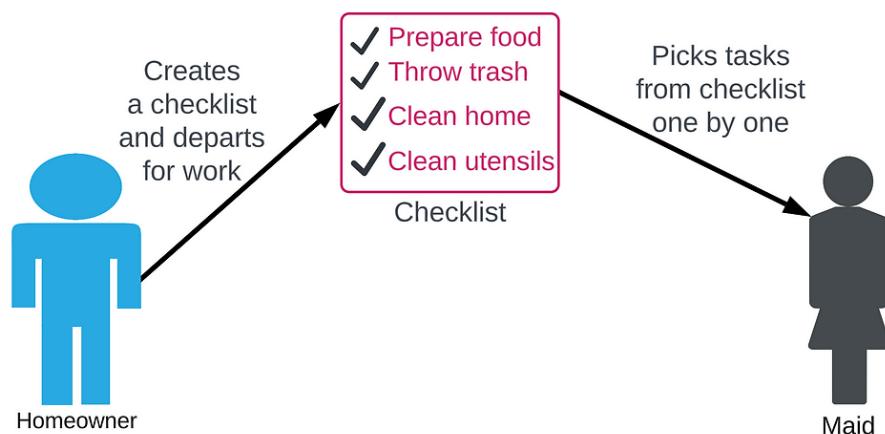
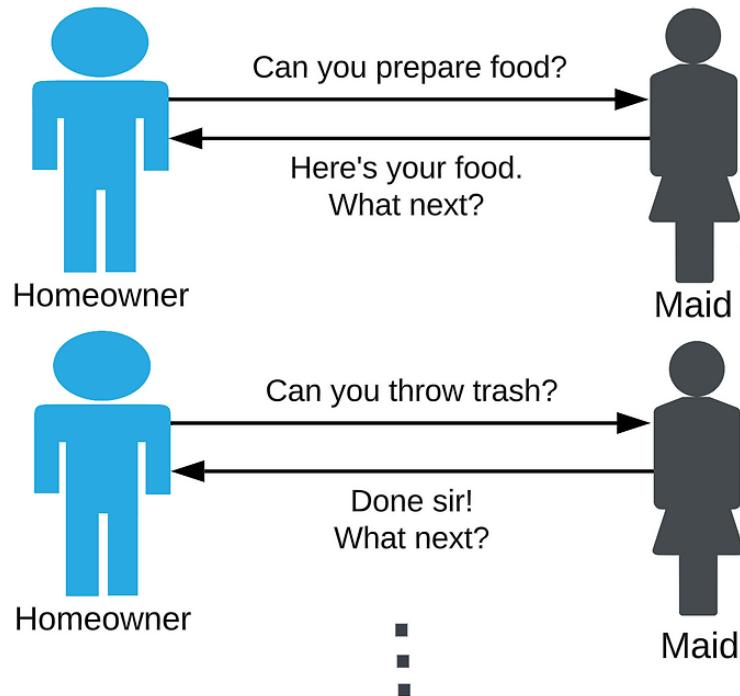
GRAPHQL



GRPC



Message Queue



Relational Database (SQL)

User Table

1	Bob	Male	28	Single
2	Mark	Male	34	Married
3	Alice	Female	50	Married

User data has a defined structure

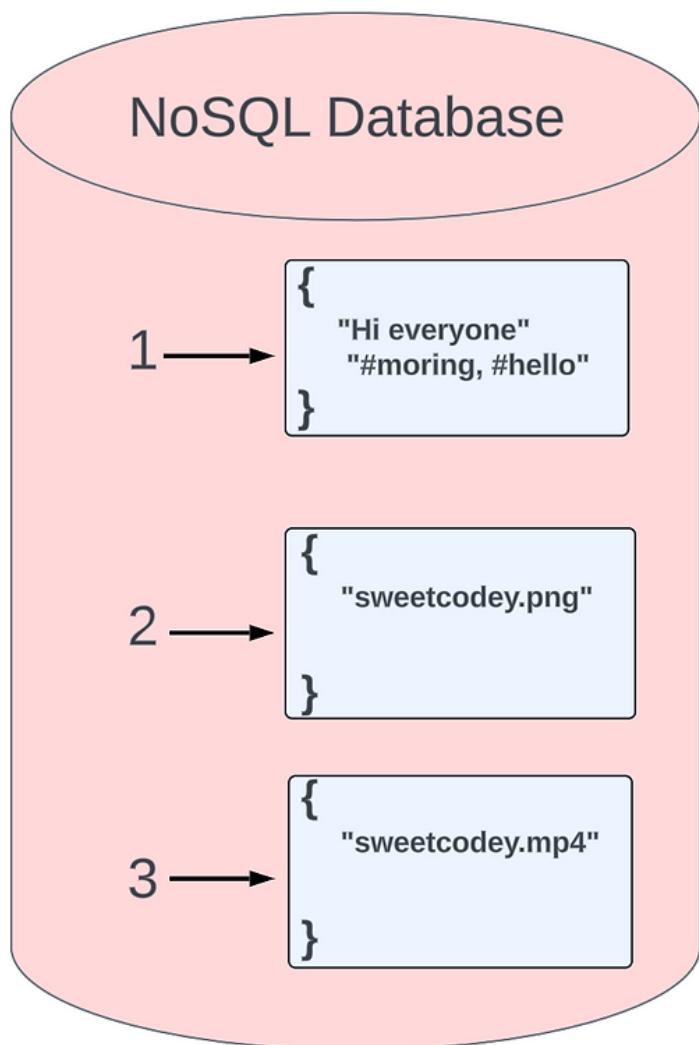
Non-Relational Database (NOSQL)

Social Media Post Table

1	"Hi everyone"			#morning, #hello
2		 sweetcodey.png		
3			 sweetcodey.mp4	

Lot of Empty Entries == Table Space Wasted!

Social Media Post data doesn't have a defined structure

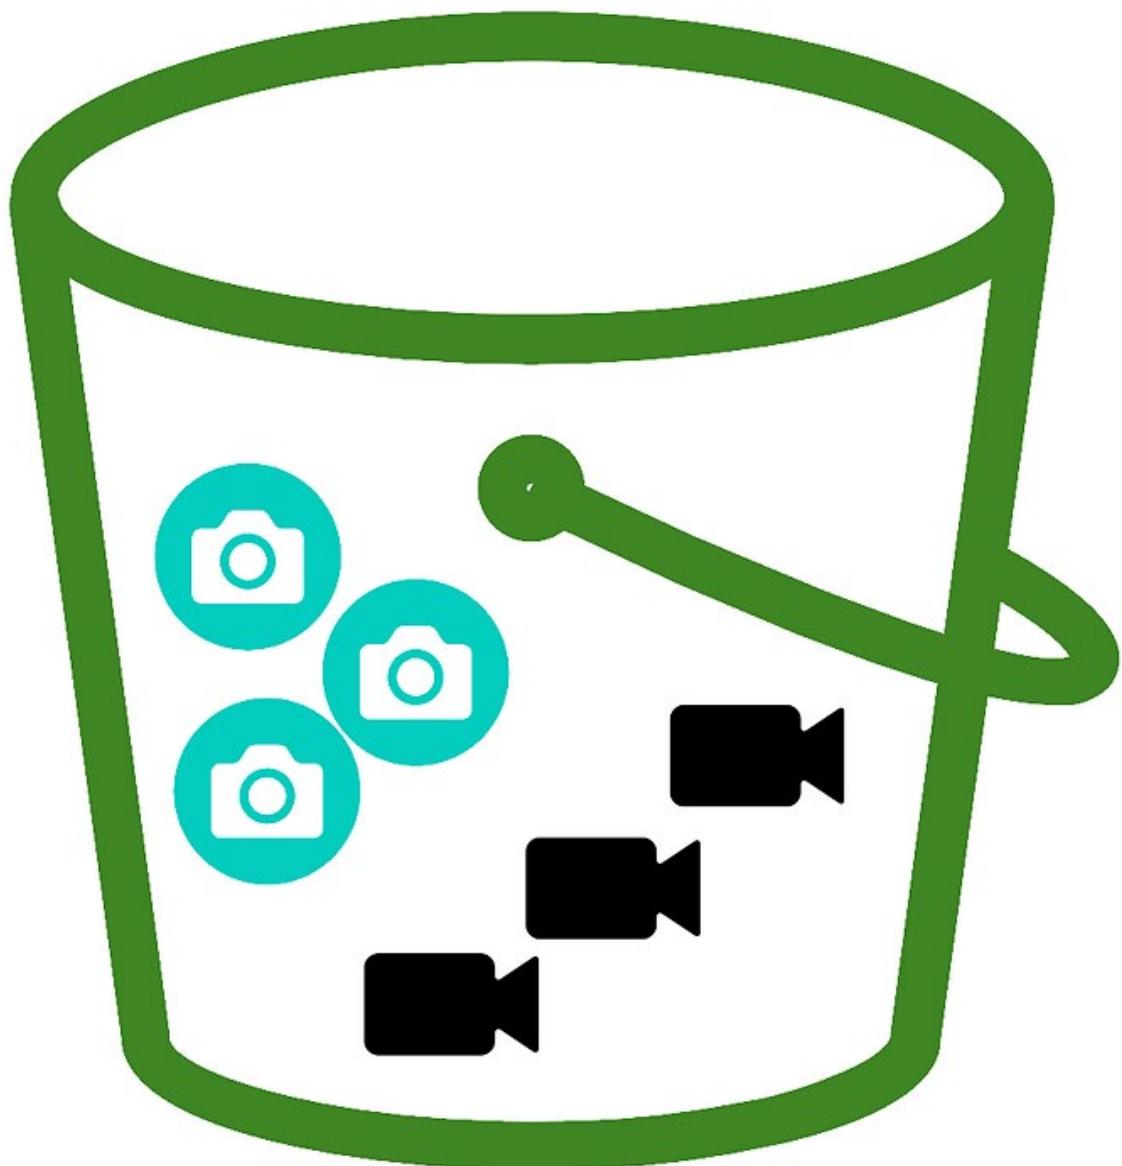


Each post (being unstructured) is stored in a JSON document.

Can be accessed by Key (PostId)

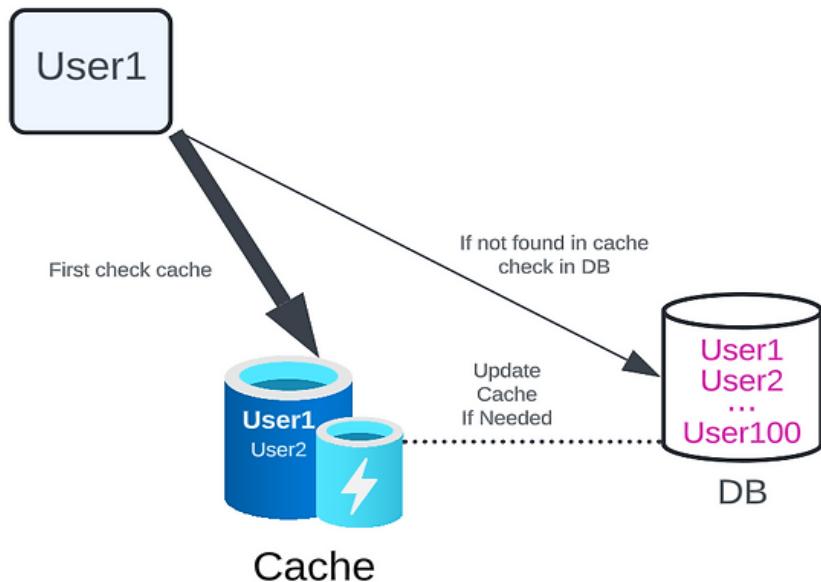
Object Storage

Object Storage



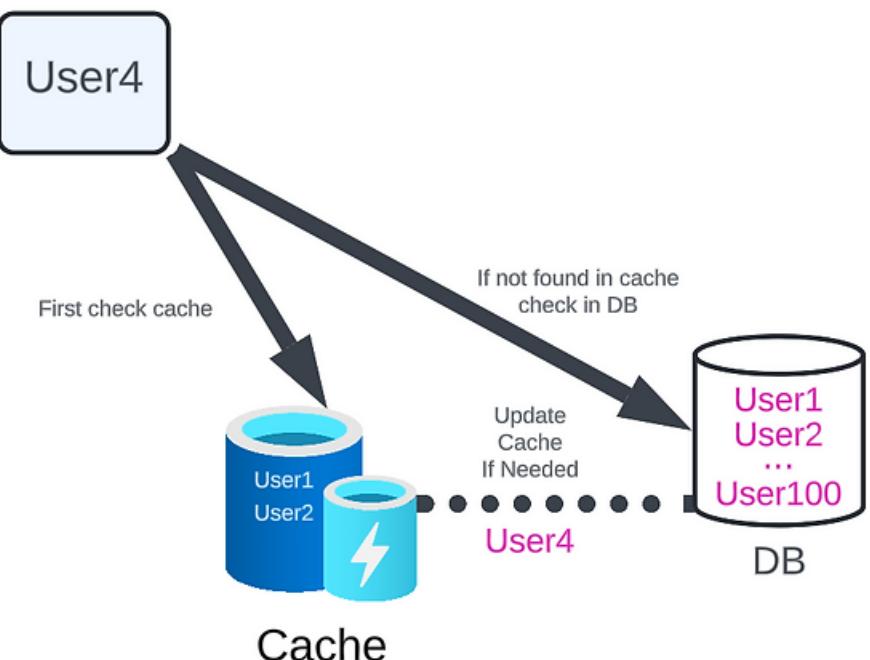
Cache

Cache Hit



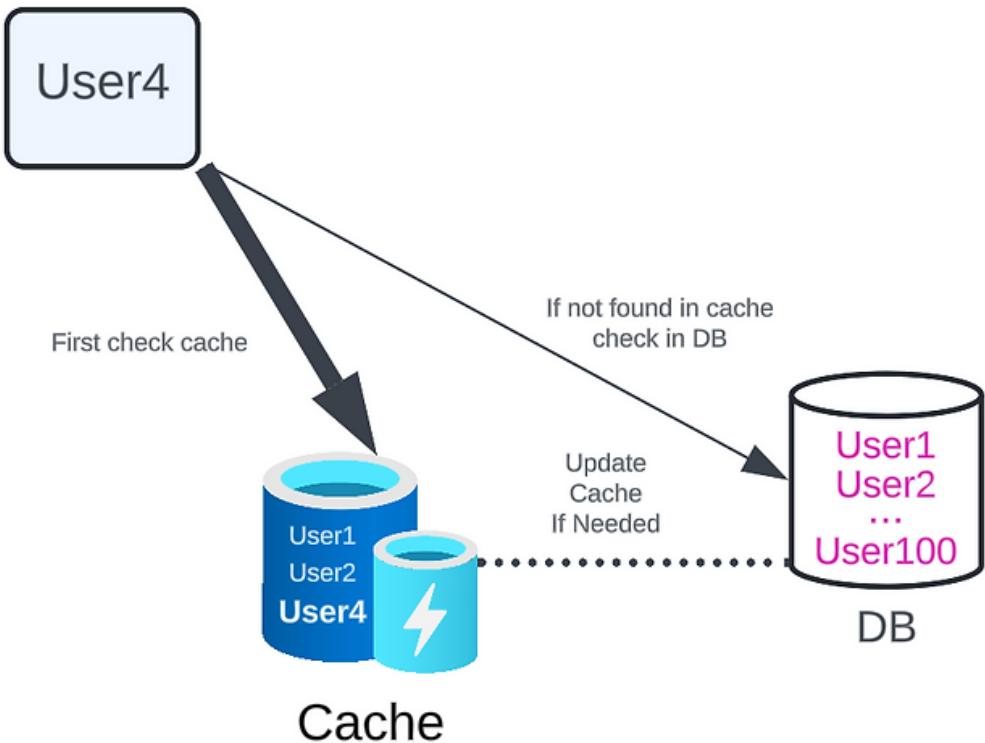
User1 already in Cache

Cache Miss



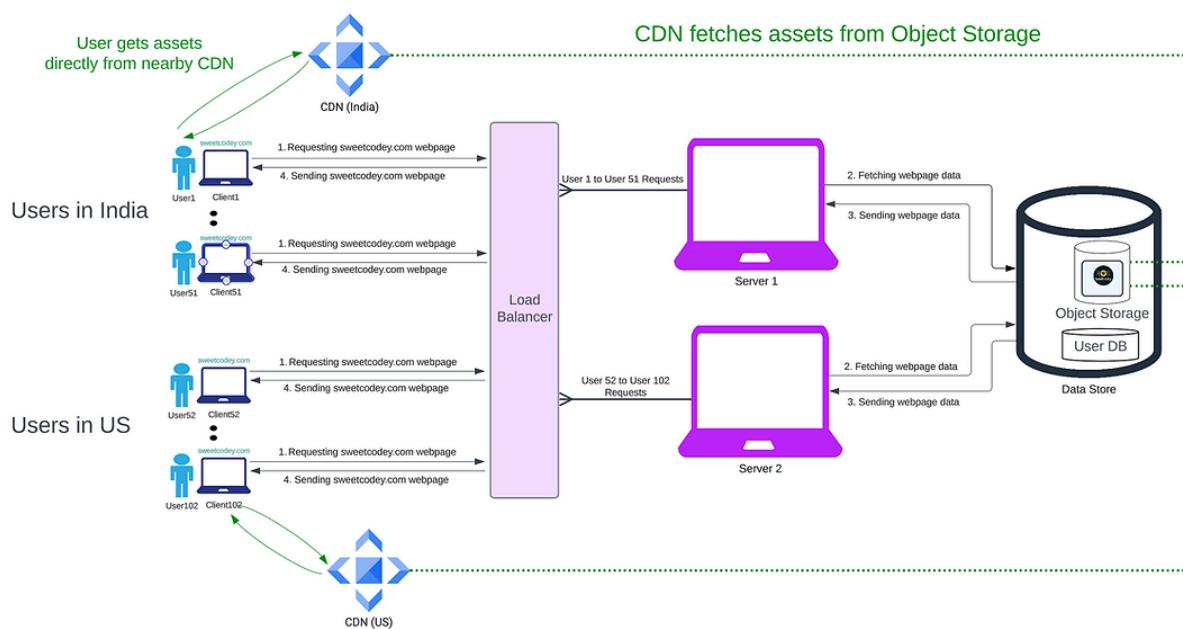
User4 NOT in Cache
Cache Updated with User4

Cache Hit

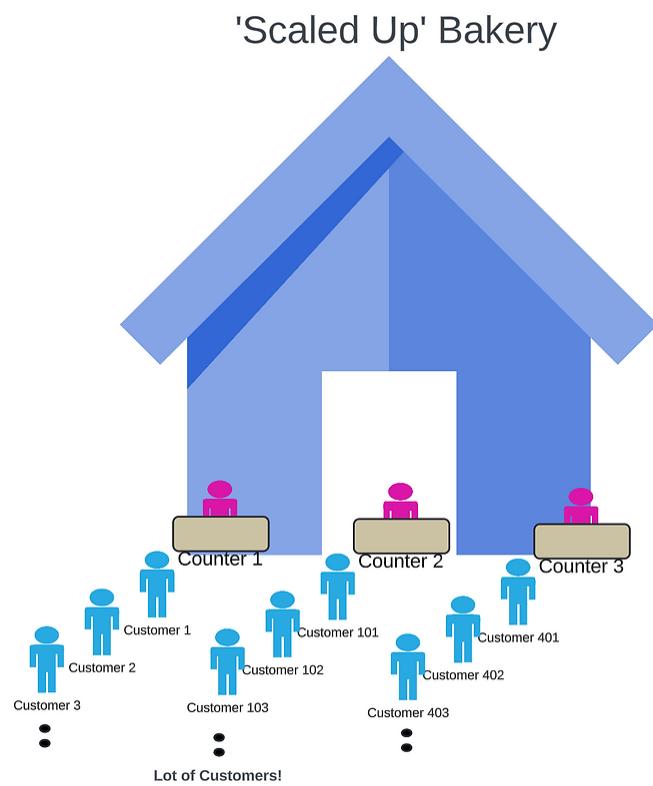
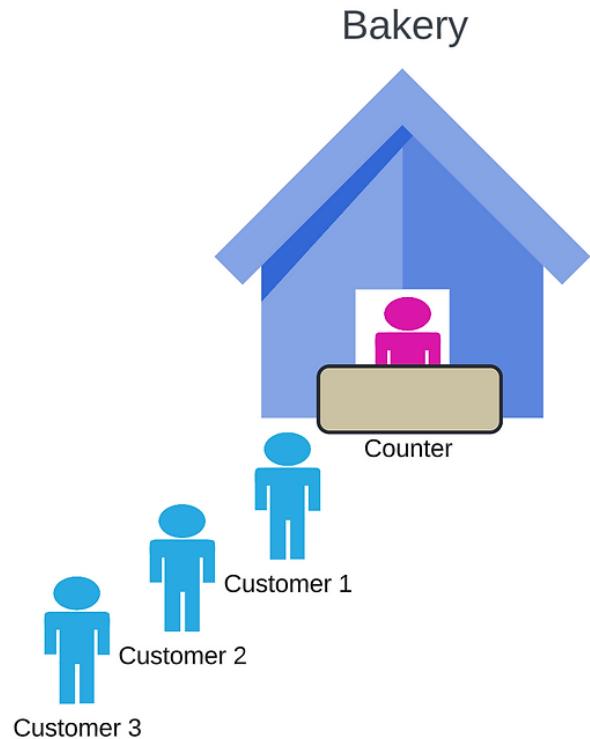


User4 is in Cache

CDN

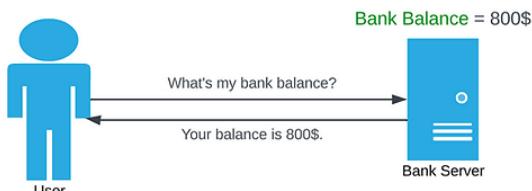
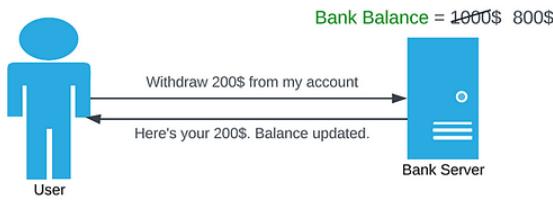


Scalability

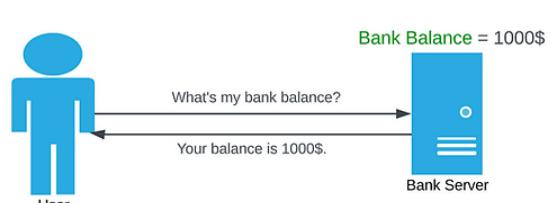
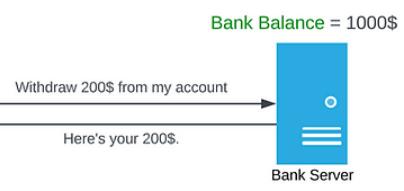


Consistency

Consistent

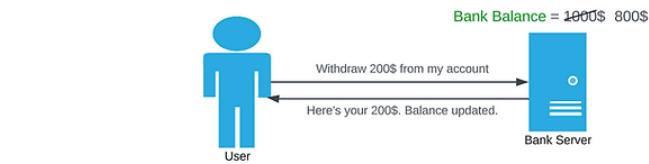


Inconsistent

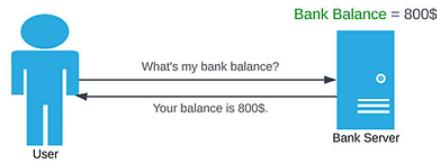


Strong/Eventual Consistency

Strong Consistency

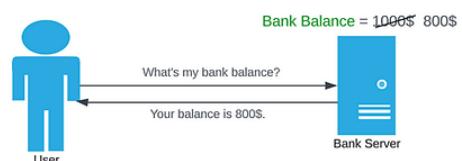
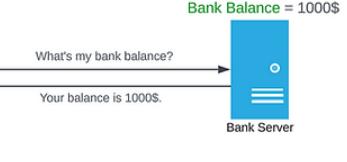
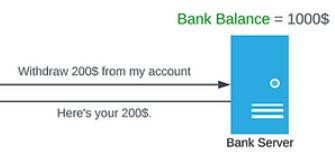


After 1 second

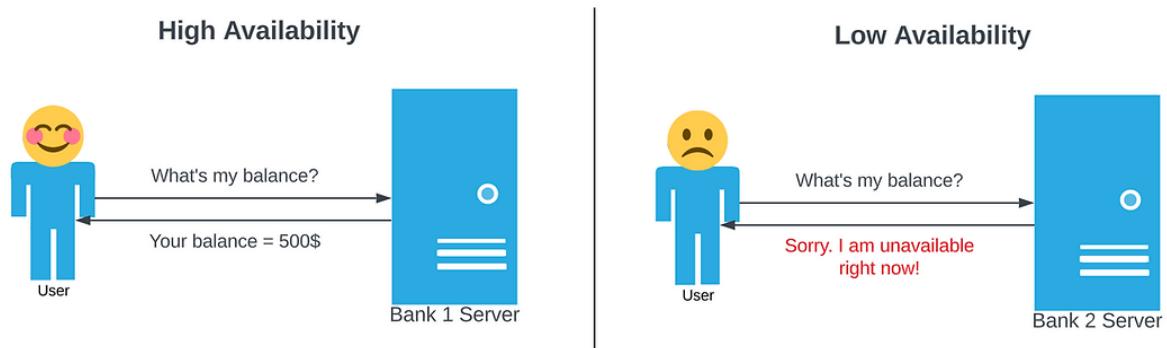


After 1 hour

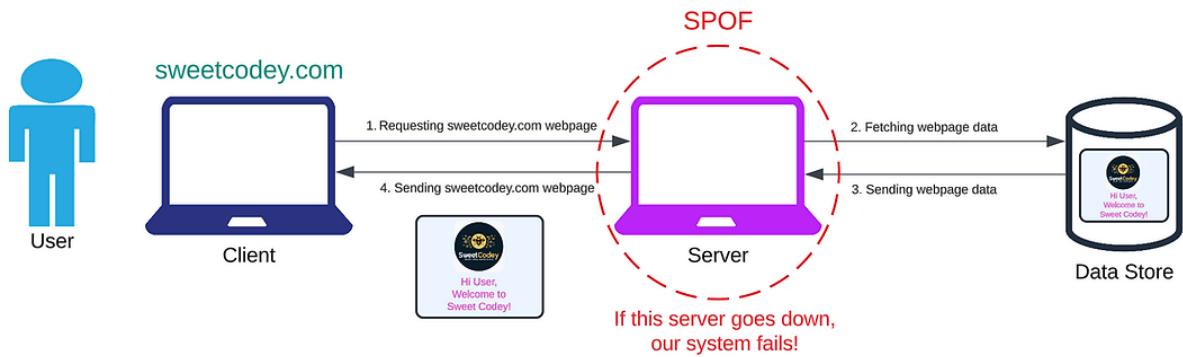
Eventual consistency



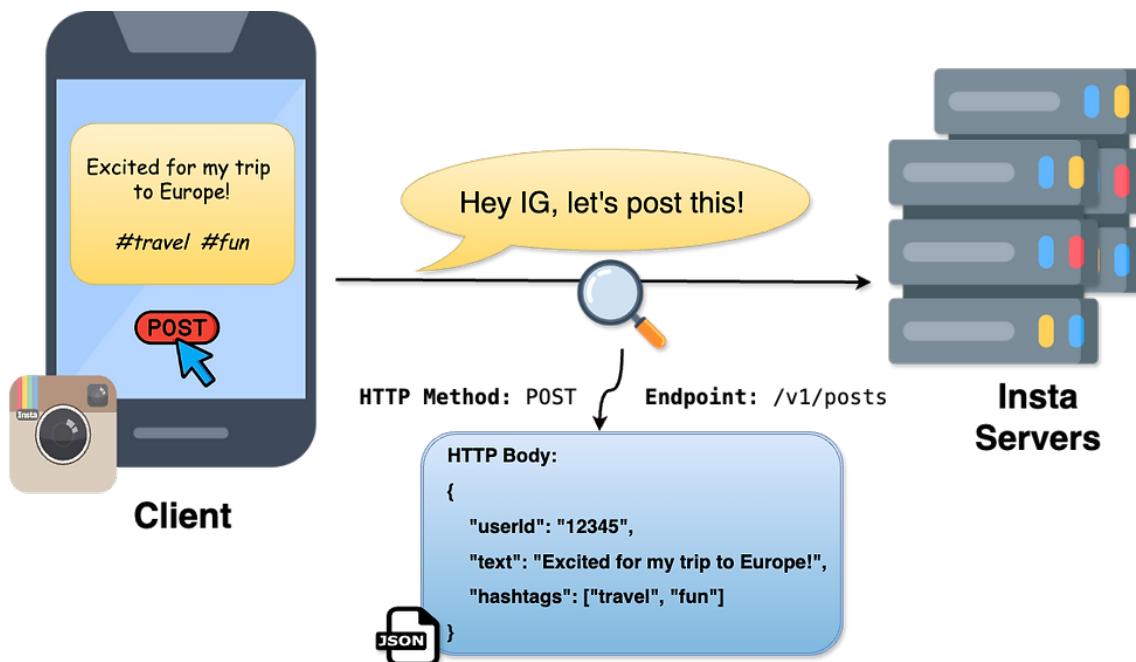
Availability



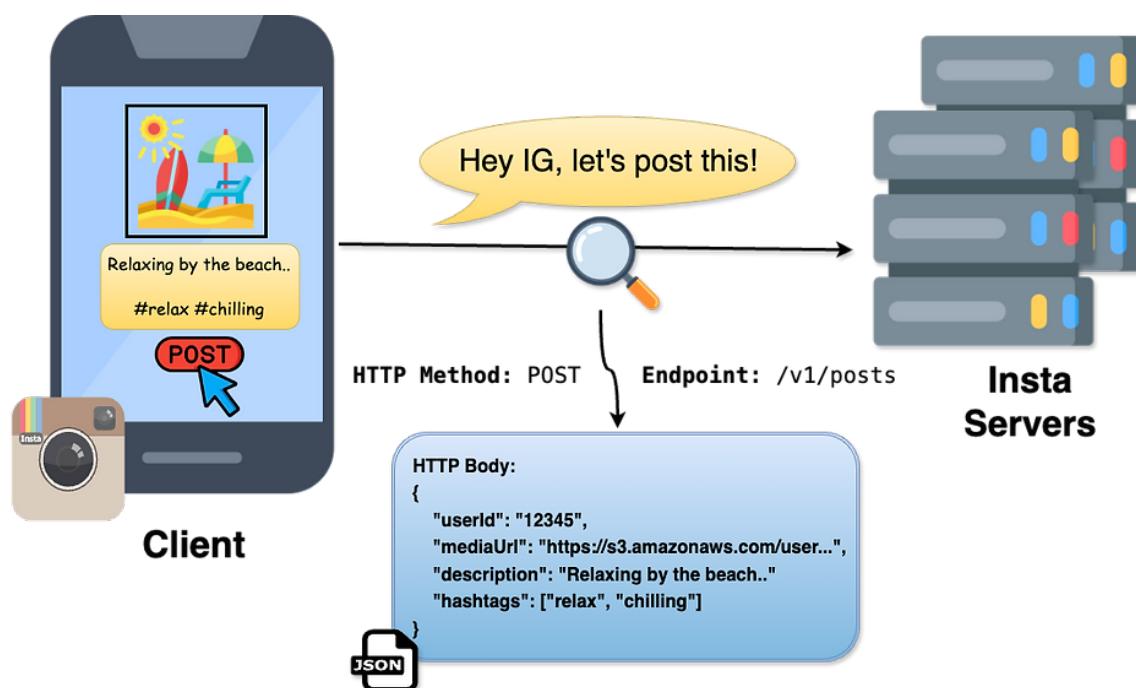
Single Point of Failure



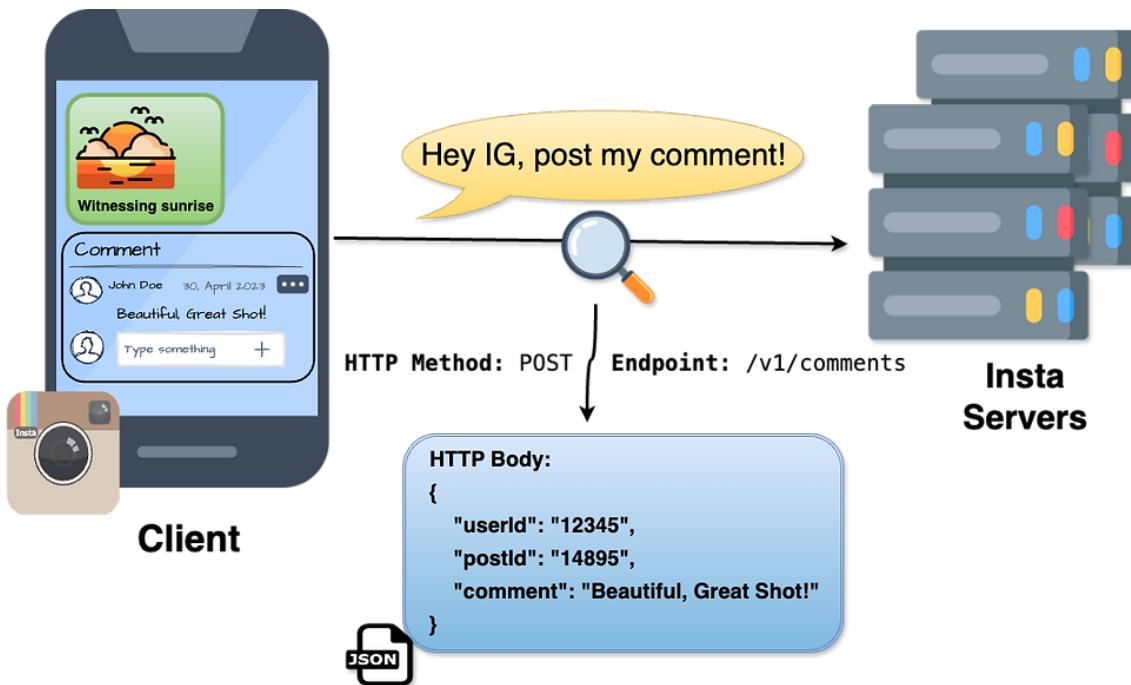
API Design-Create a Text Post



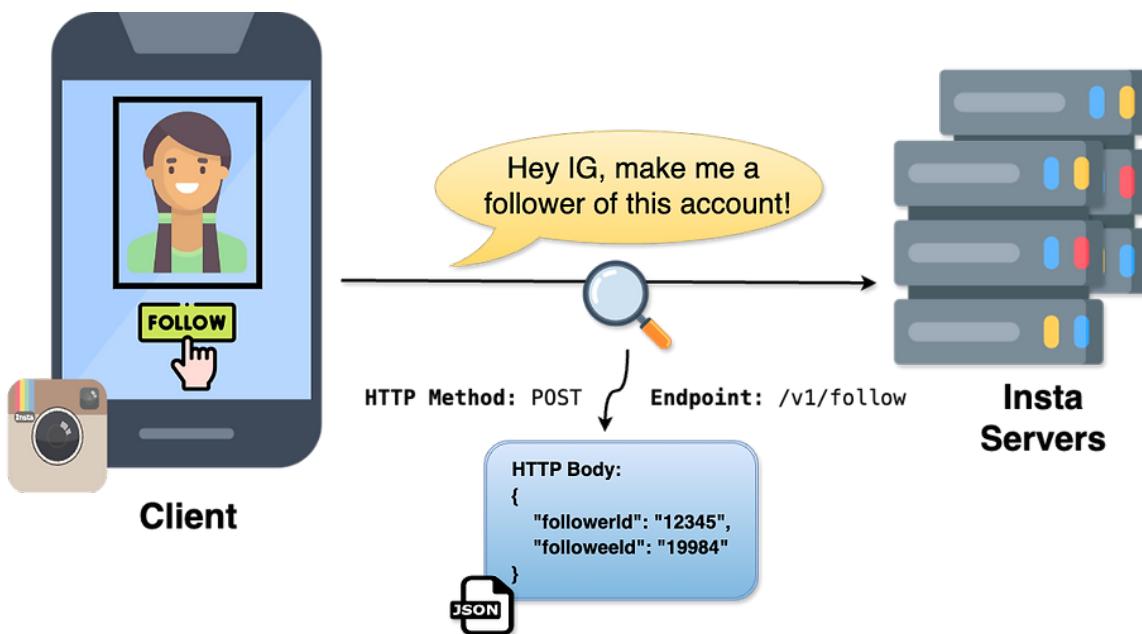
API Design-Create Image/ Video Post



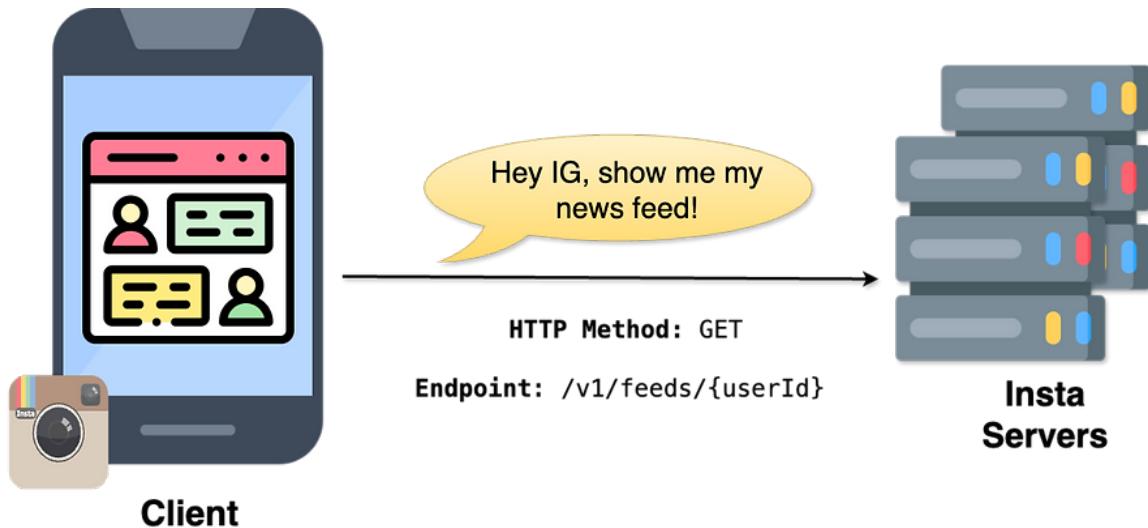
API Design-Like /Comment a Post



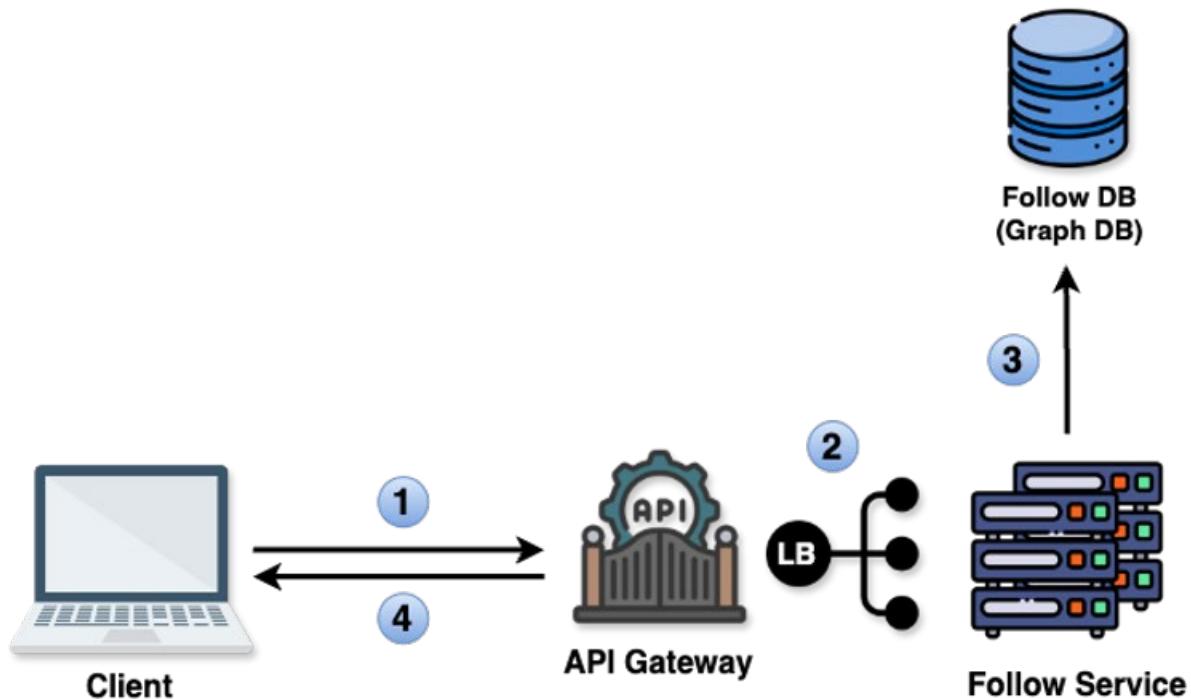
API Design-Follow / Unfollow another User



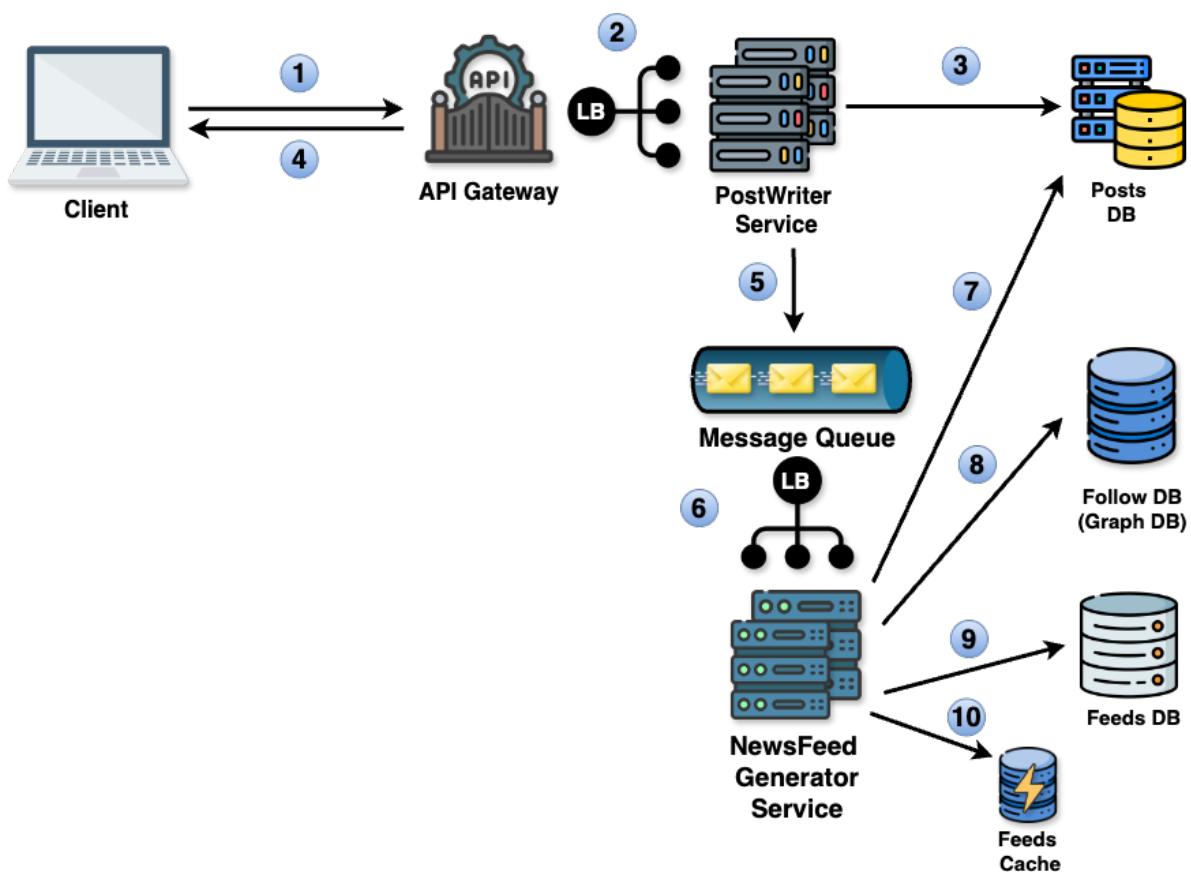
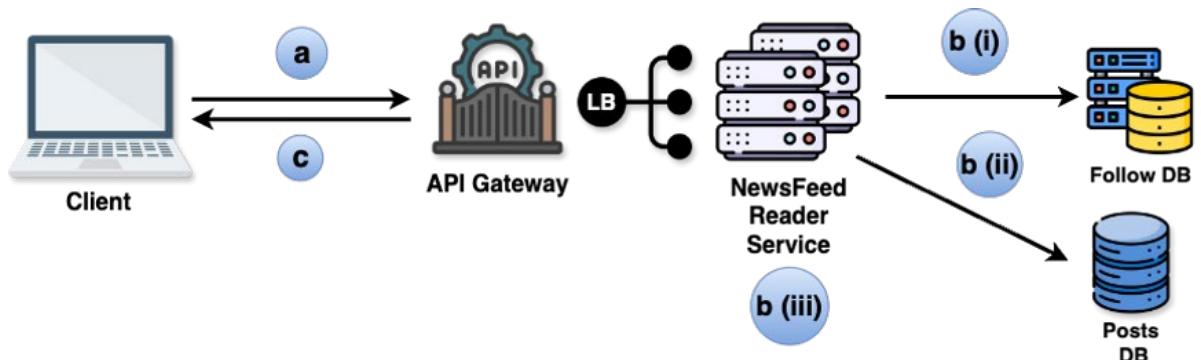
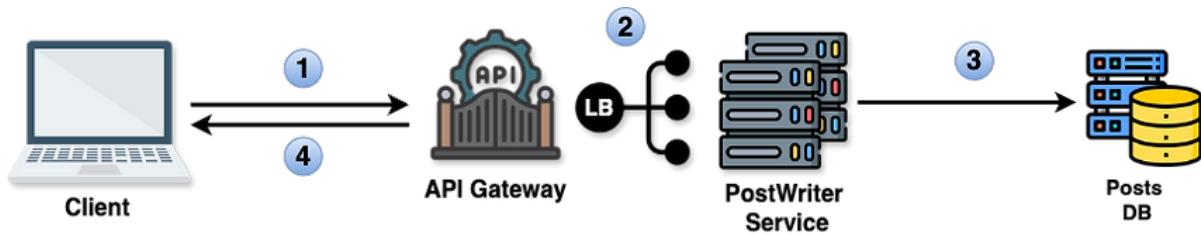
API Design-Read the Newsfeed (aka Timeline)



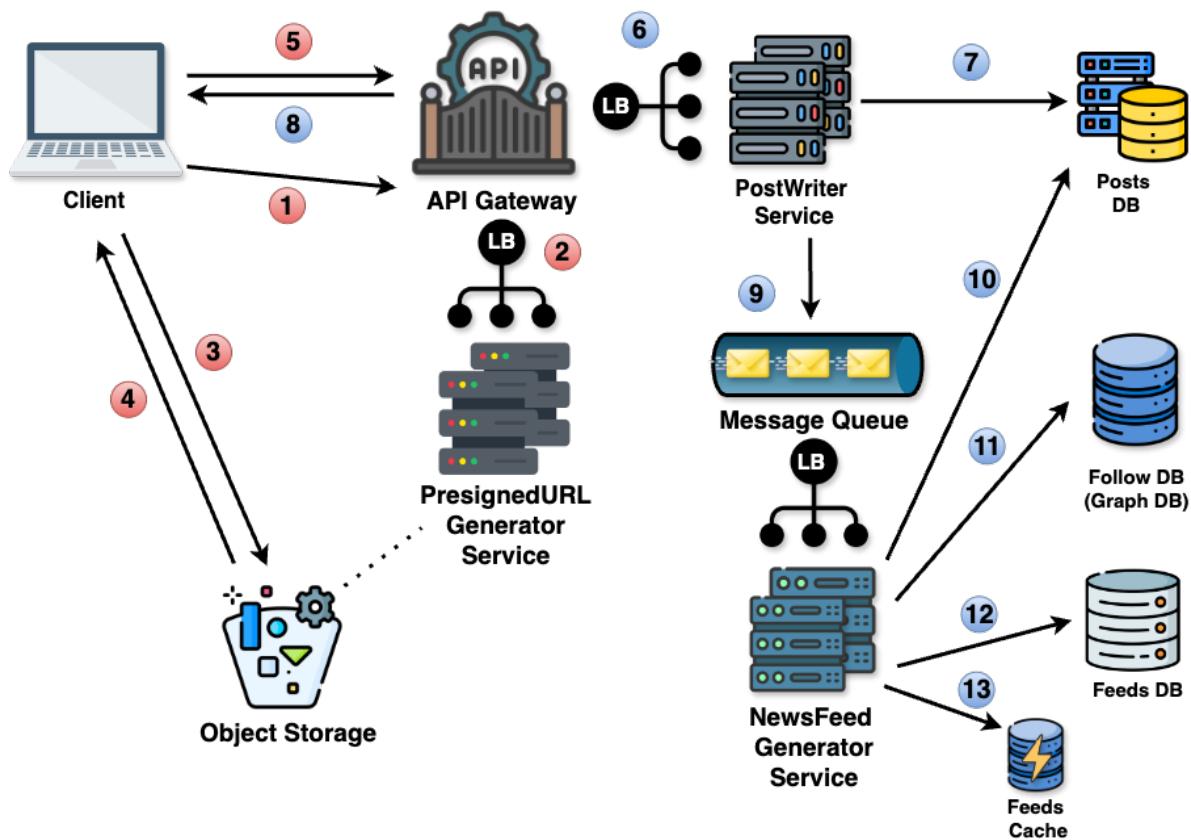
High level Design- Follow / Unfollow another User



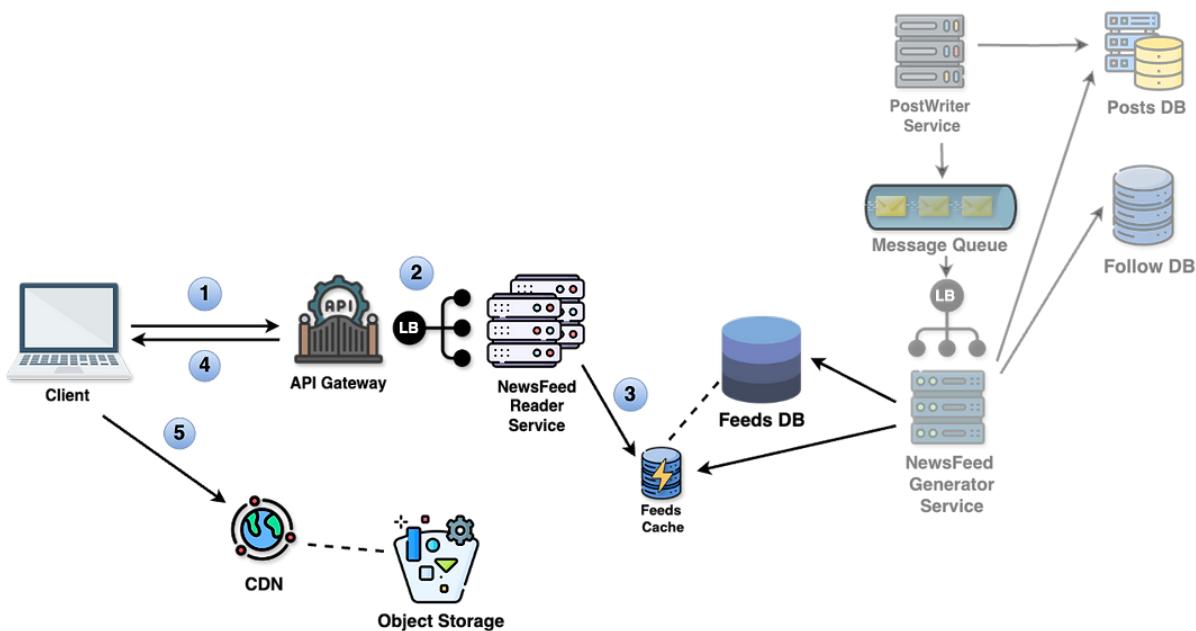
High level Design- Create Text Post



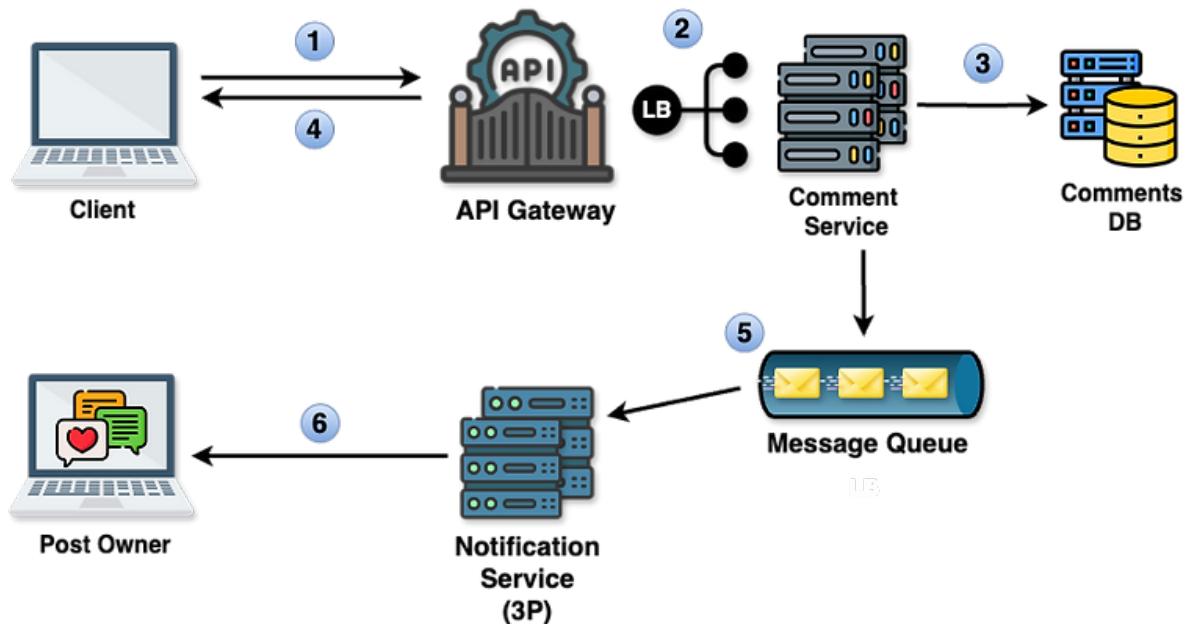
High level Design- Create Image/Video Post



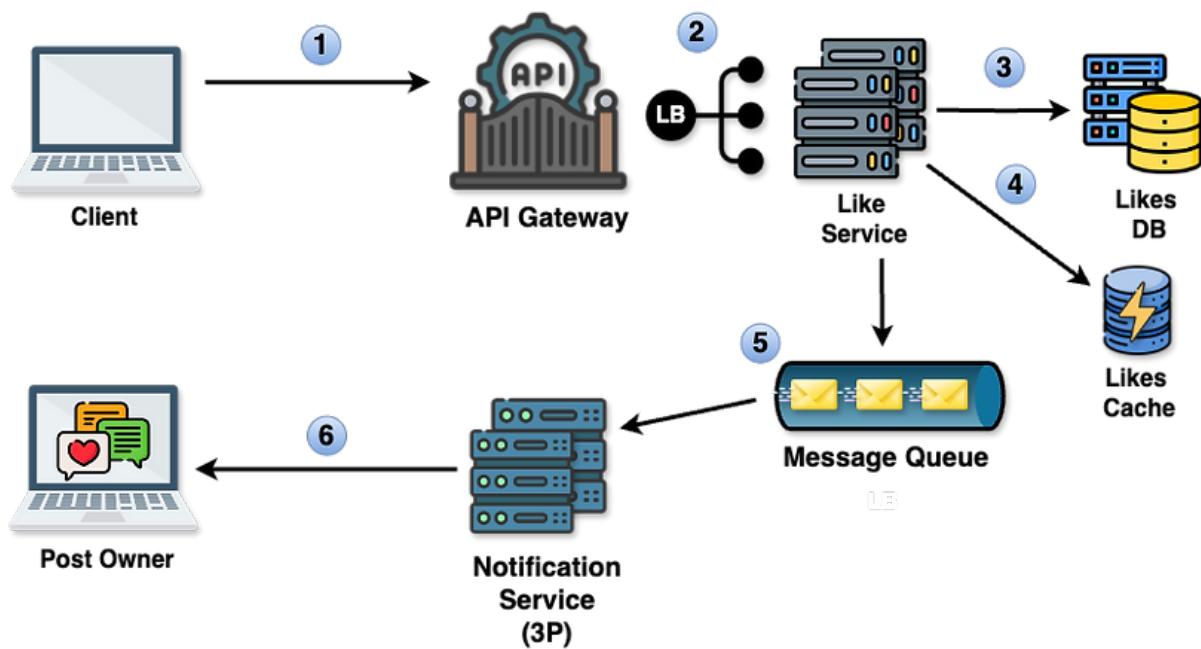
High level Design- Read the Newsfeed (aka timeline)

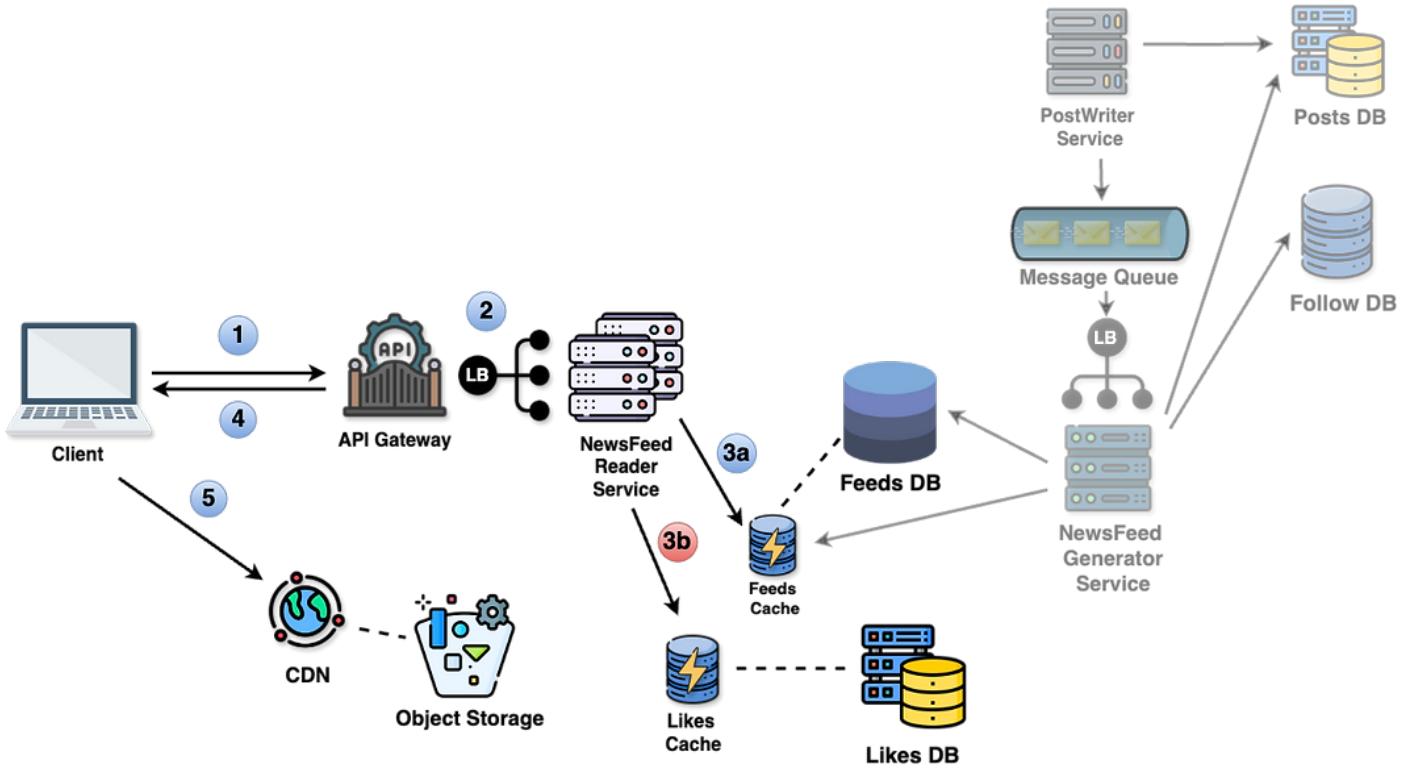


High level Design-Comment on a Post

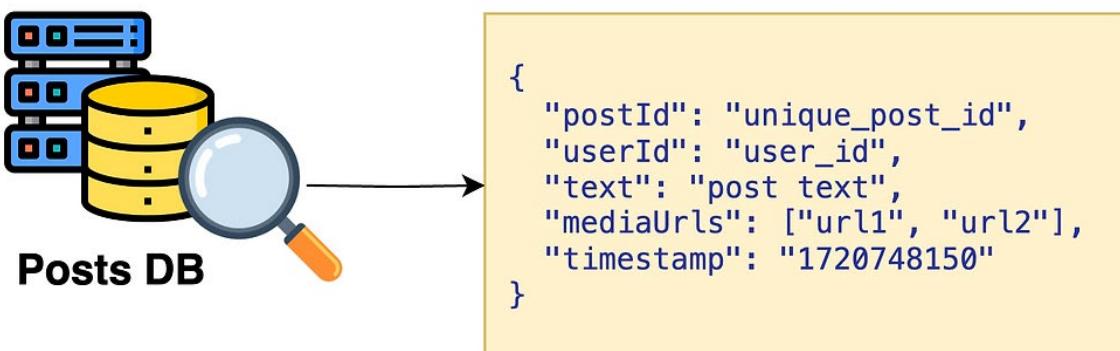


High level Design-Like on a Post

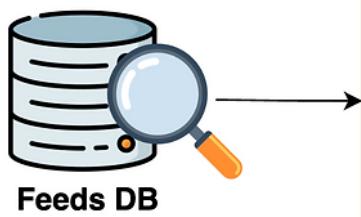




Posts Schema

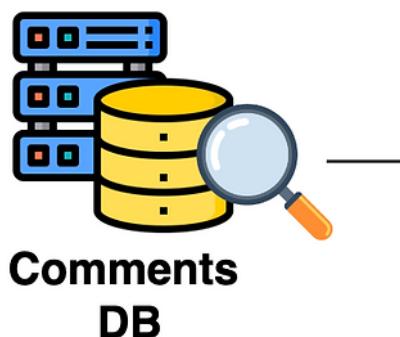


Feeds Schema



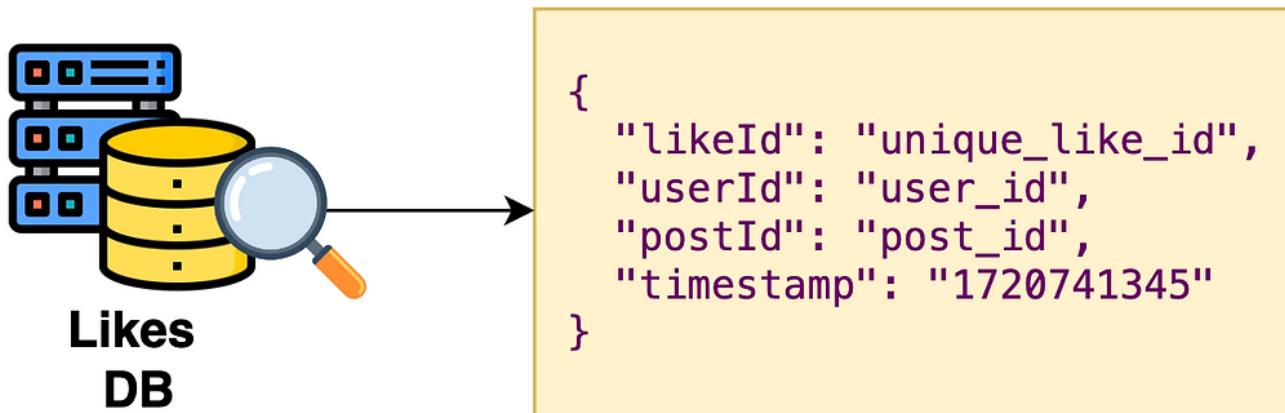
```
{  
  "userId": "unique_user_id",  
  "feedItems": [  
    {  
      "postId": "13452",  
      "userId": "84231",  
      "text": "sample post text 1",  
      "timestamp": "1720748111"  
    },  
    {  
      "postId": "86475",  
      "userId": "12394",  
      "text": "sample post text 2",  
      "mediaUrls": ["url1", "url2", ...]  
      "timestamp": "1720713459"  
    }  
    // More posts...  
  ]  
}
```

Comments Schema

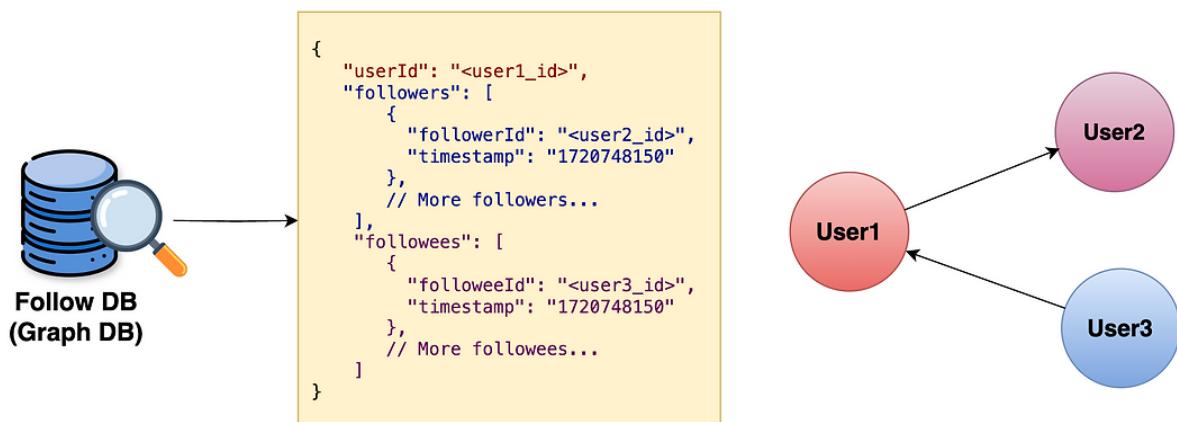


```
{  
  "commentId": "unique_comment_id",  
  "userId": "user_id",  
  "postId": "post_id",  
  "comment": "comment text",  
  "timestamp": "1720748150"  
}
```

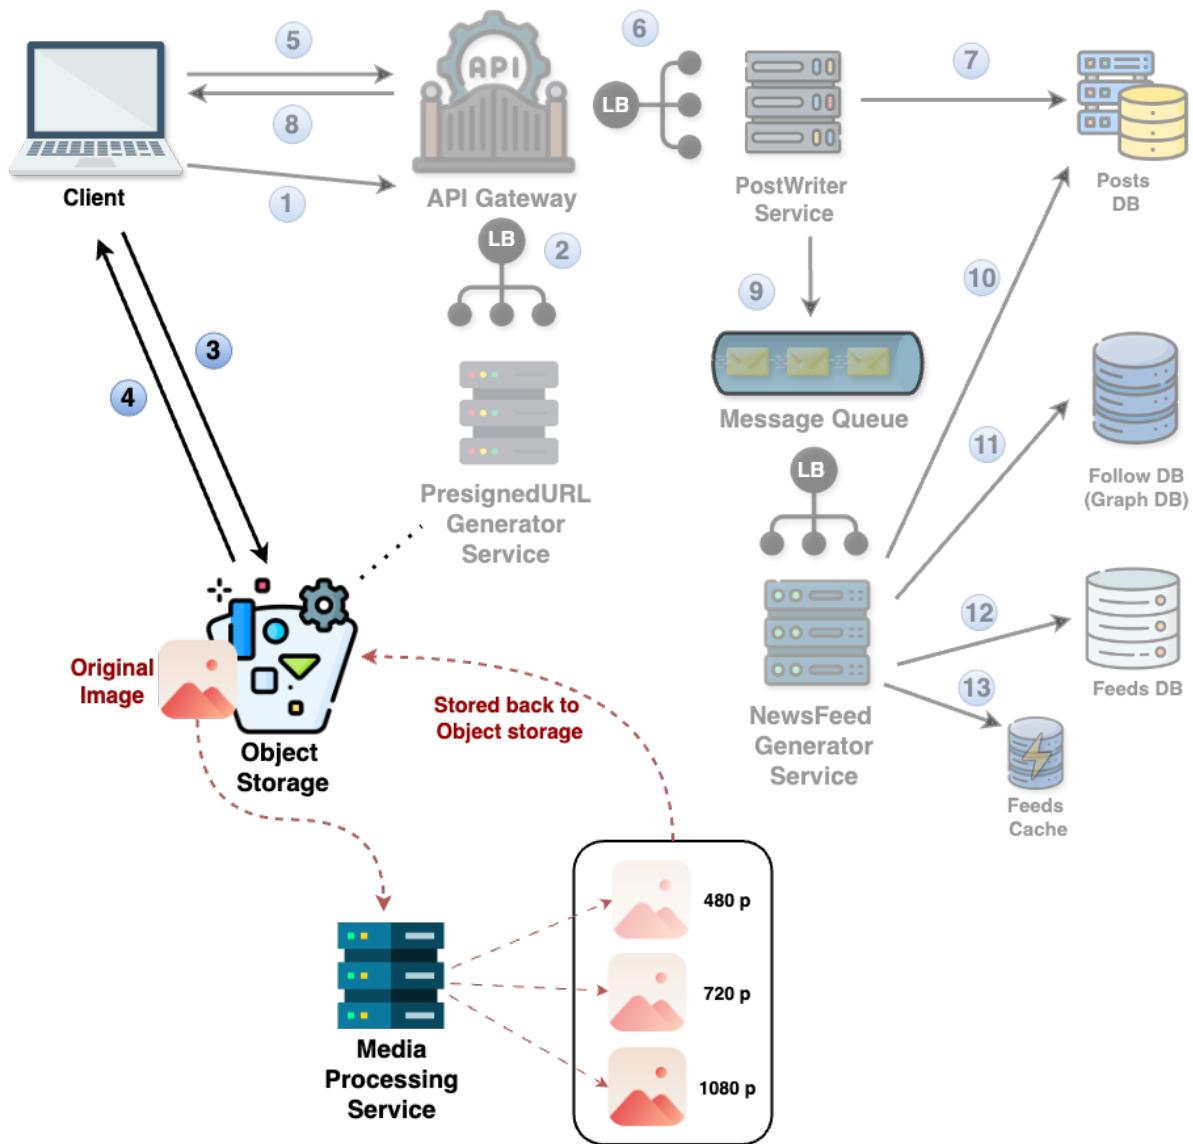
Likes Schema



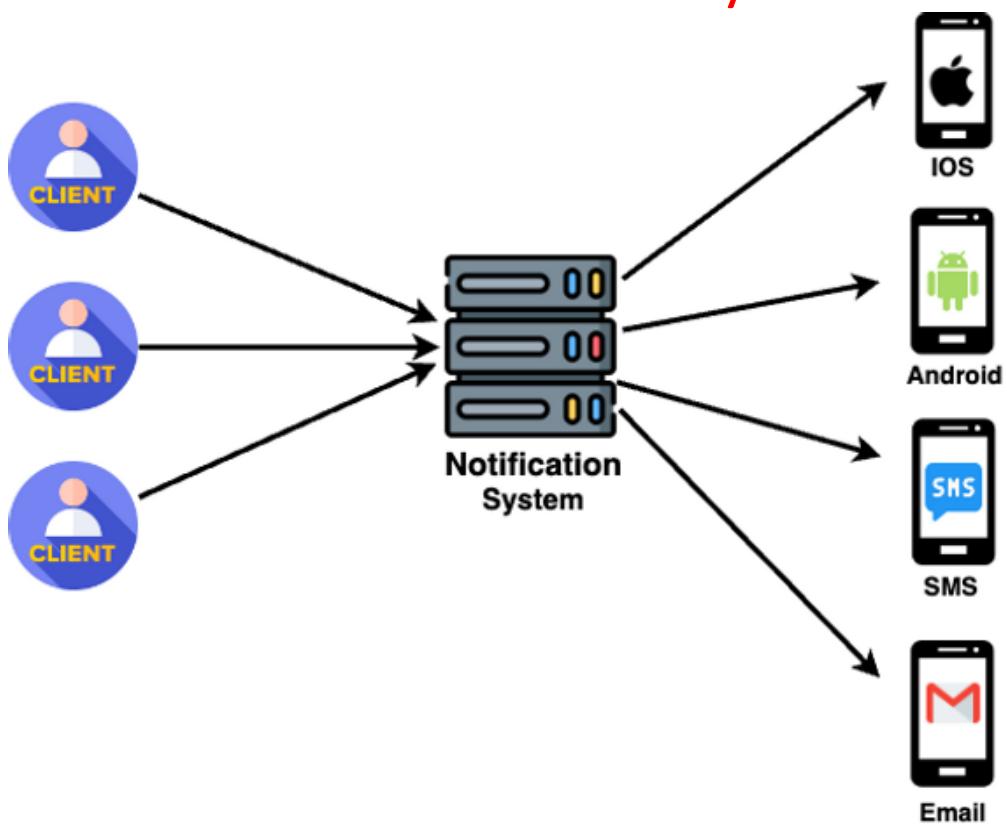
Follow Schema



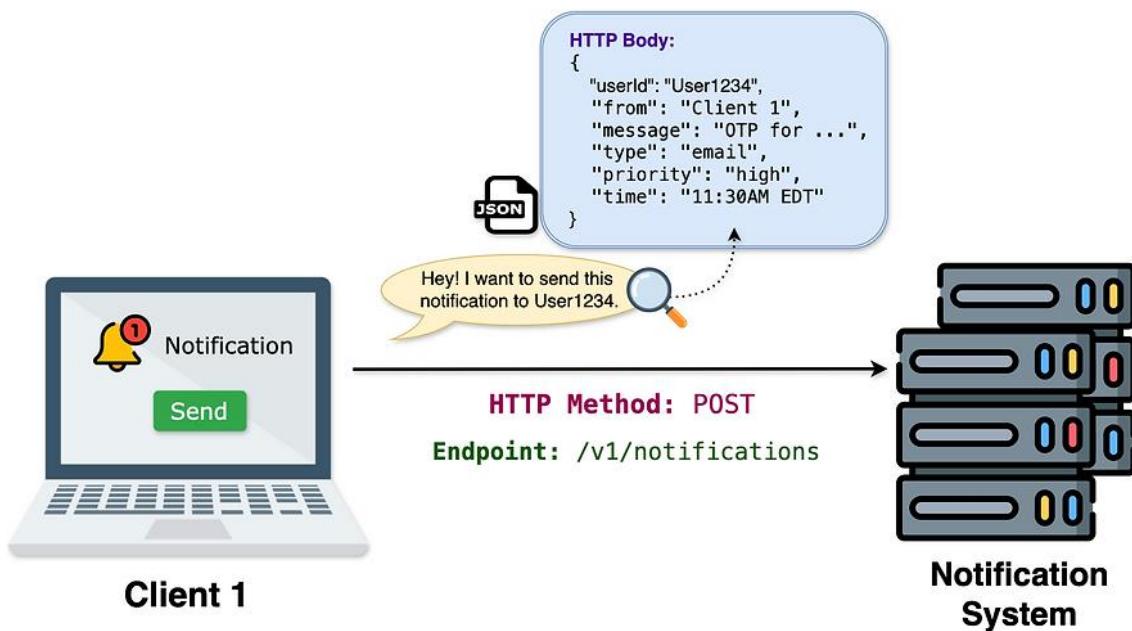
Media Processing



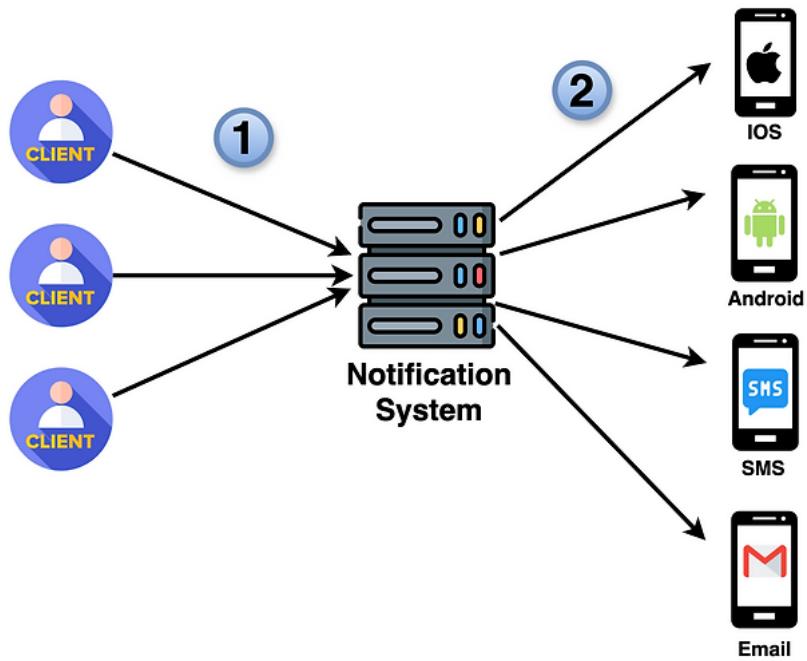
What is Notification System?



API Design- Sending a Notification



High Level Design- Basic Flow

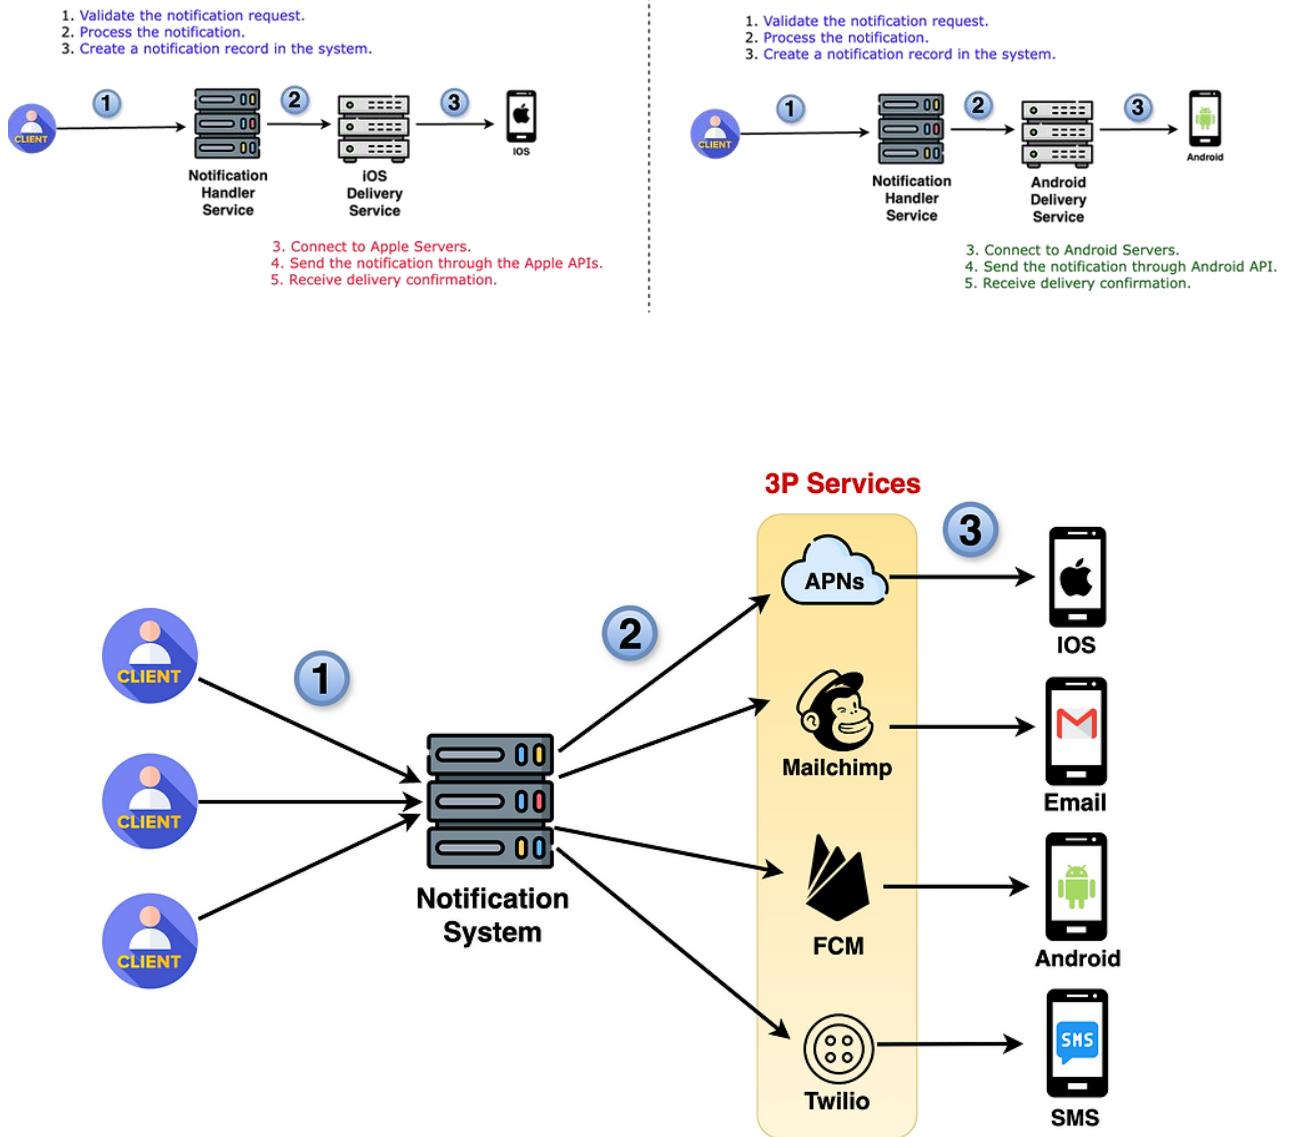


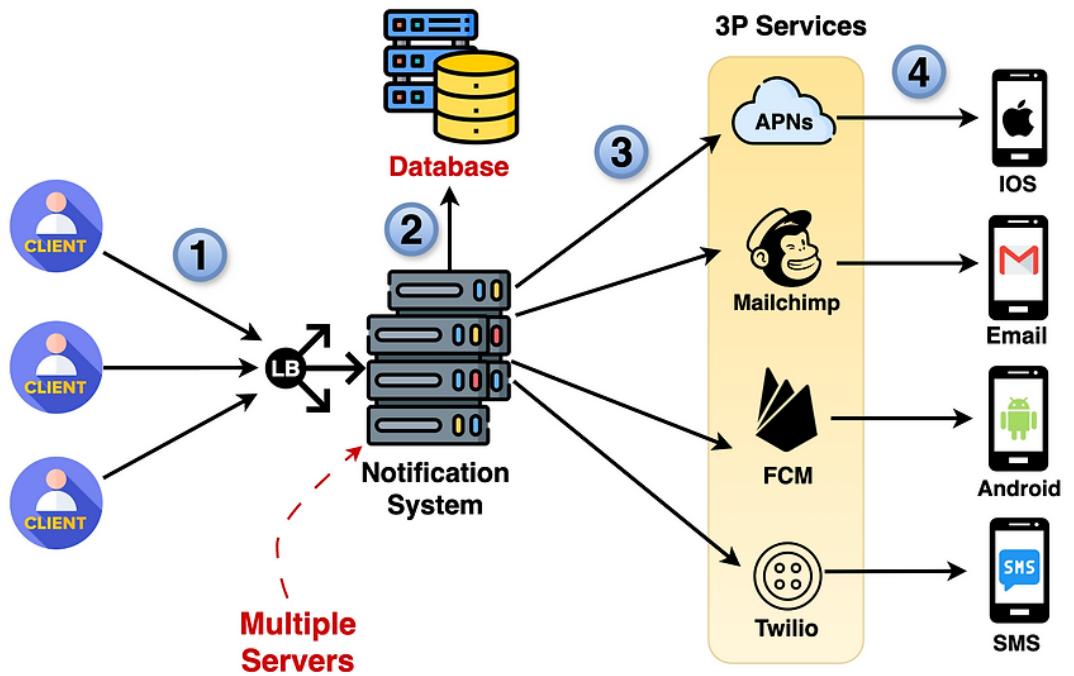
High Level Design- Problems with Basic High Level Design

1. Validate the notification request.
2. Process the notification.
3. Create a notification record in the system.
4. Connect to Apple Servers.
5. Send the notification through the Apple APIs.
6. Receive delivery confirmation.

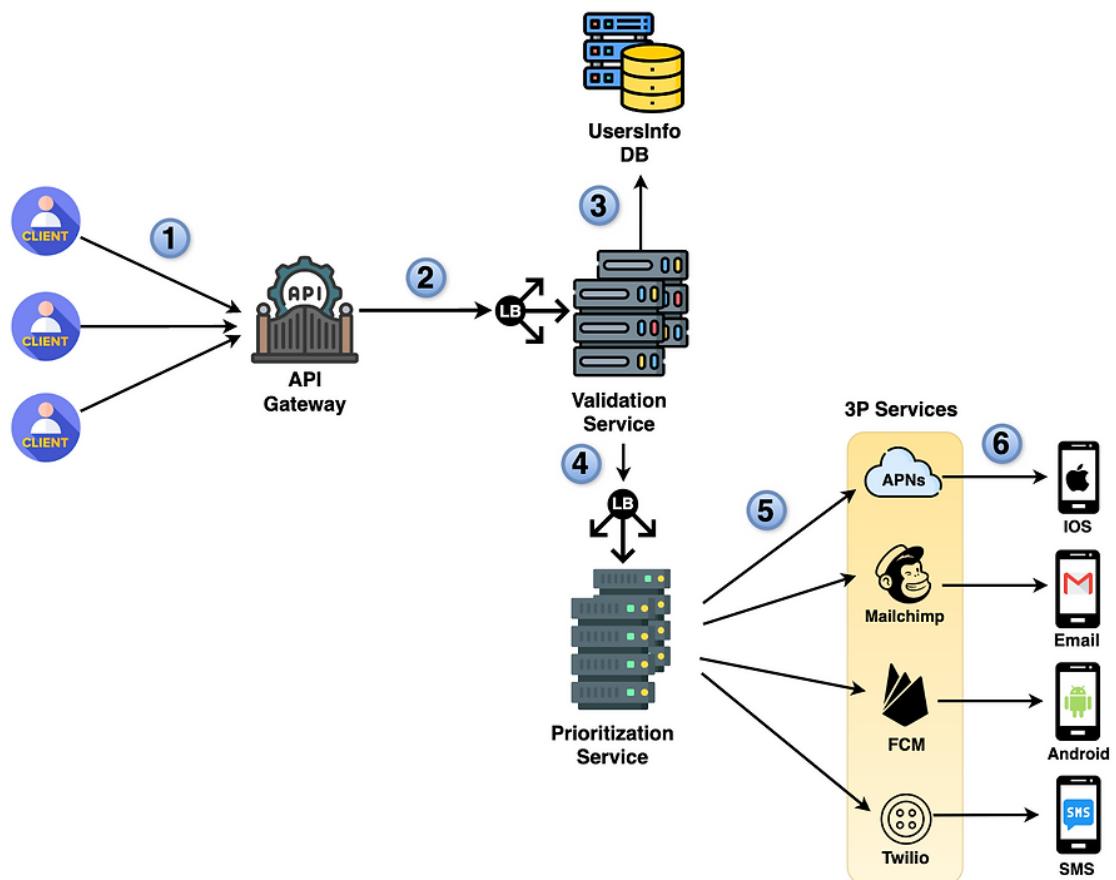
1. Validate the notification request.
2. Process the notification.
3. Create a notification record in the system.
4. Connect to Android Servers.
5. Send the notification through Android APIs.
6. Receive delivery confirmation.



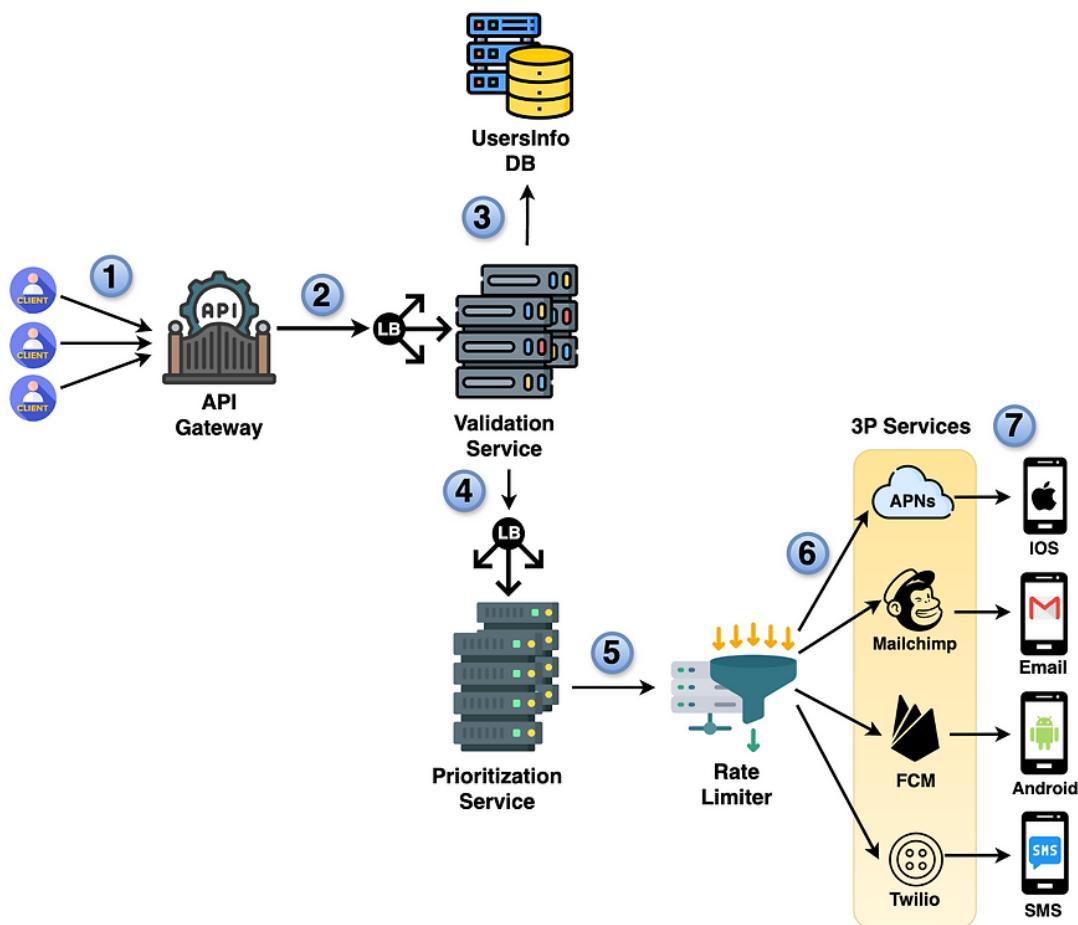




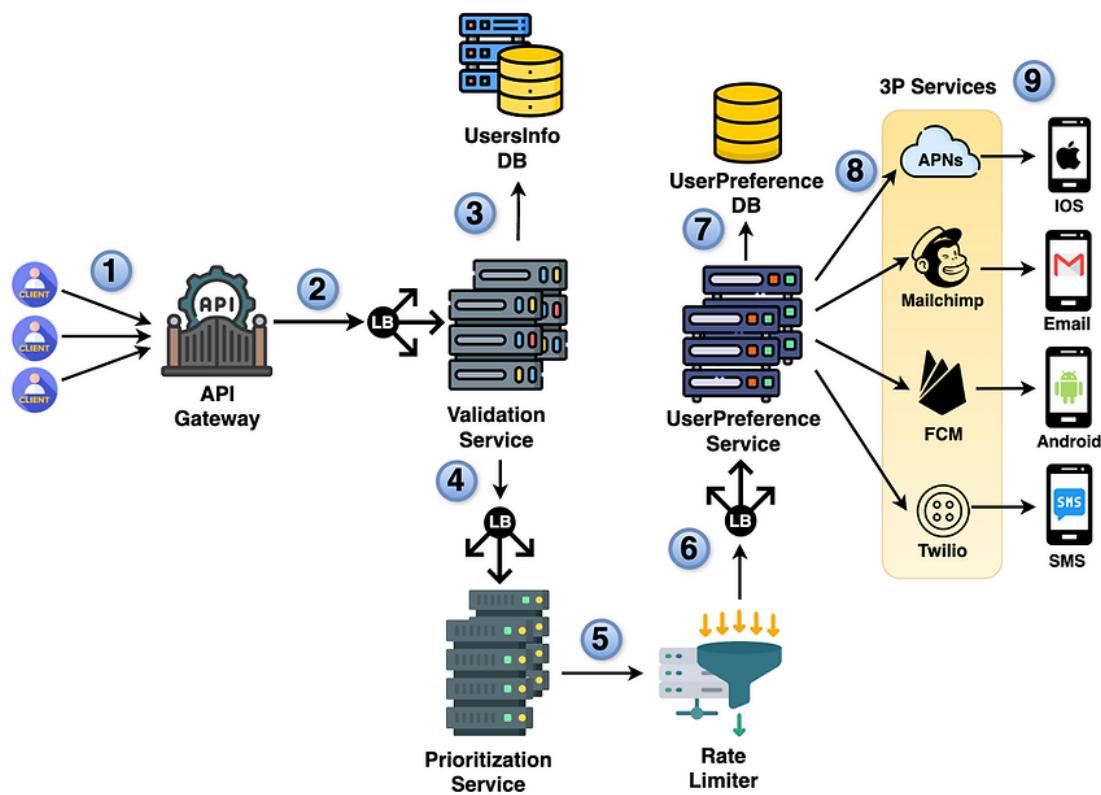
High Level Design-Validation and Prioritization



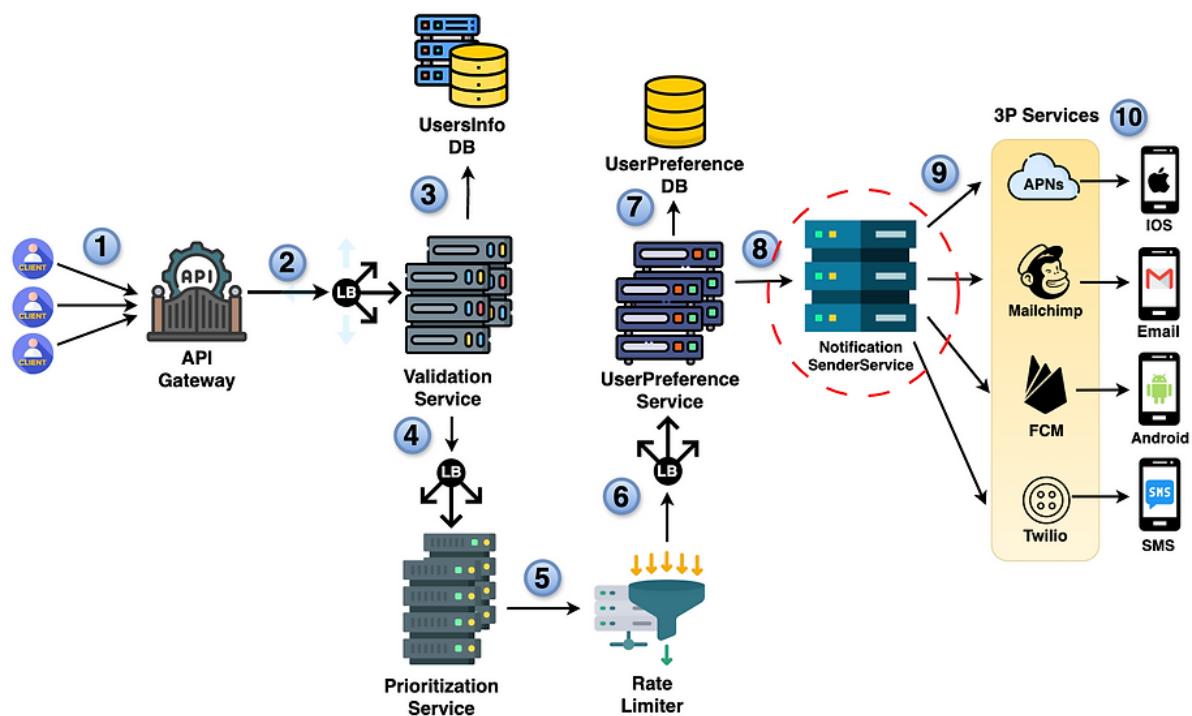
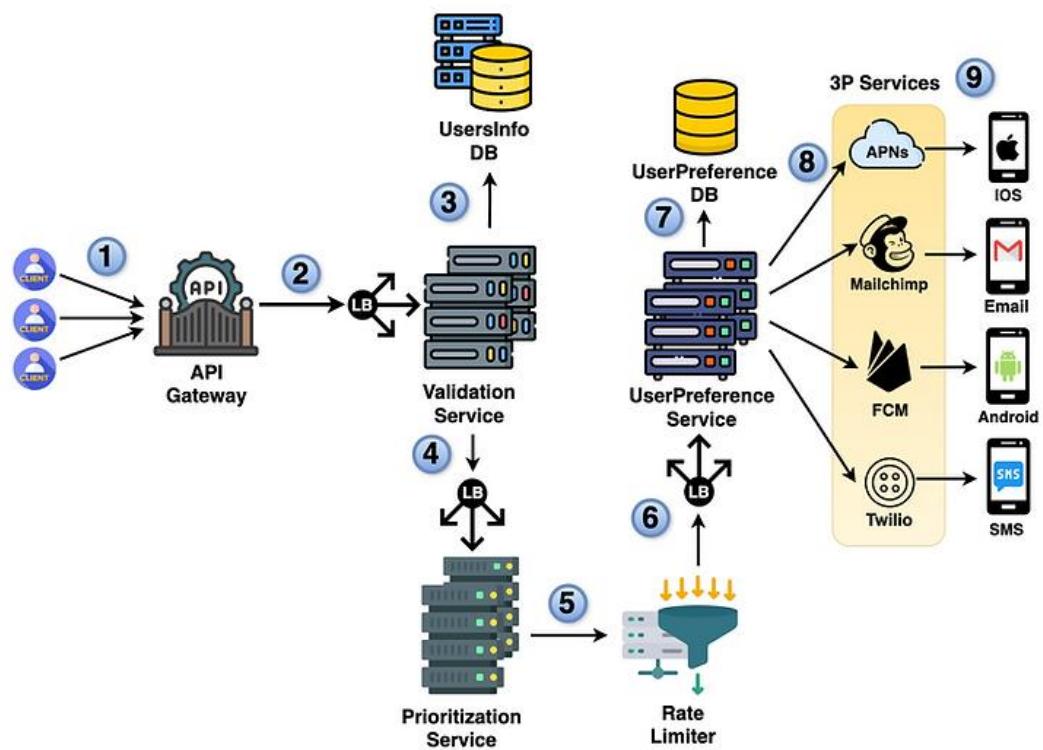
High Level Design-Rate Limiting Notifications

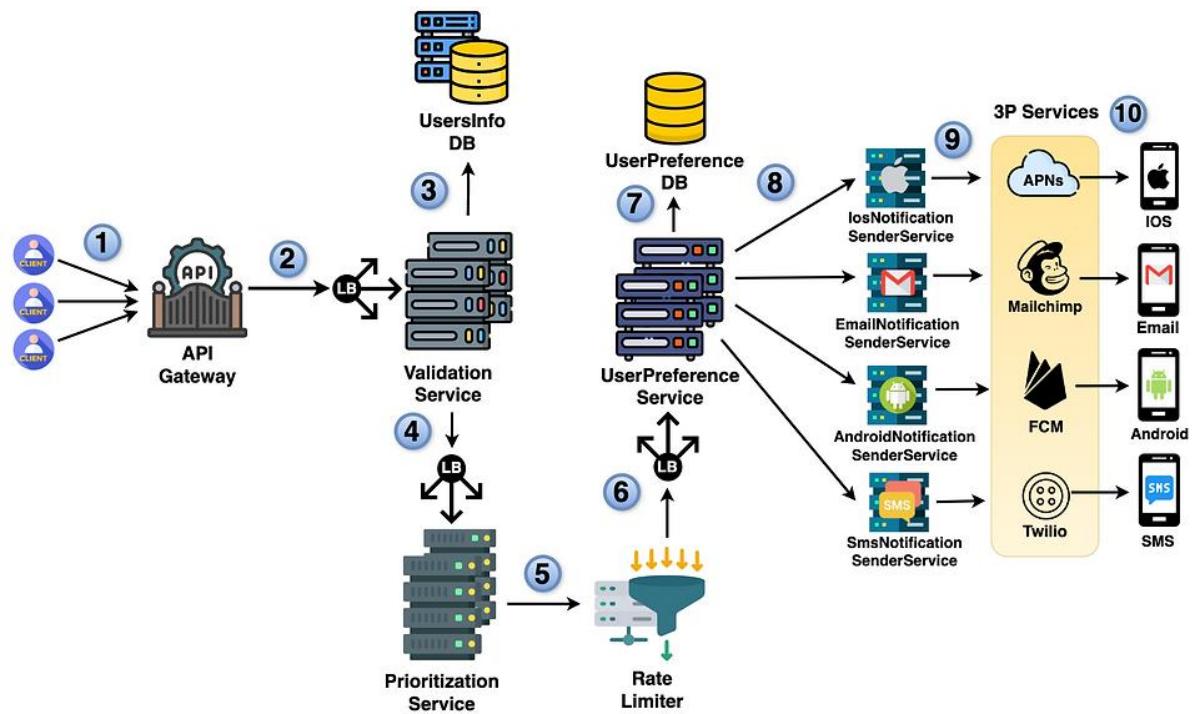


High Level Design-Adjusting to User Preferences

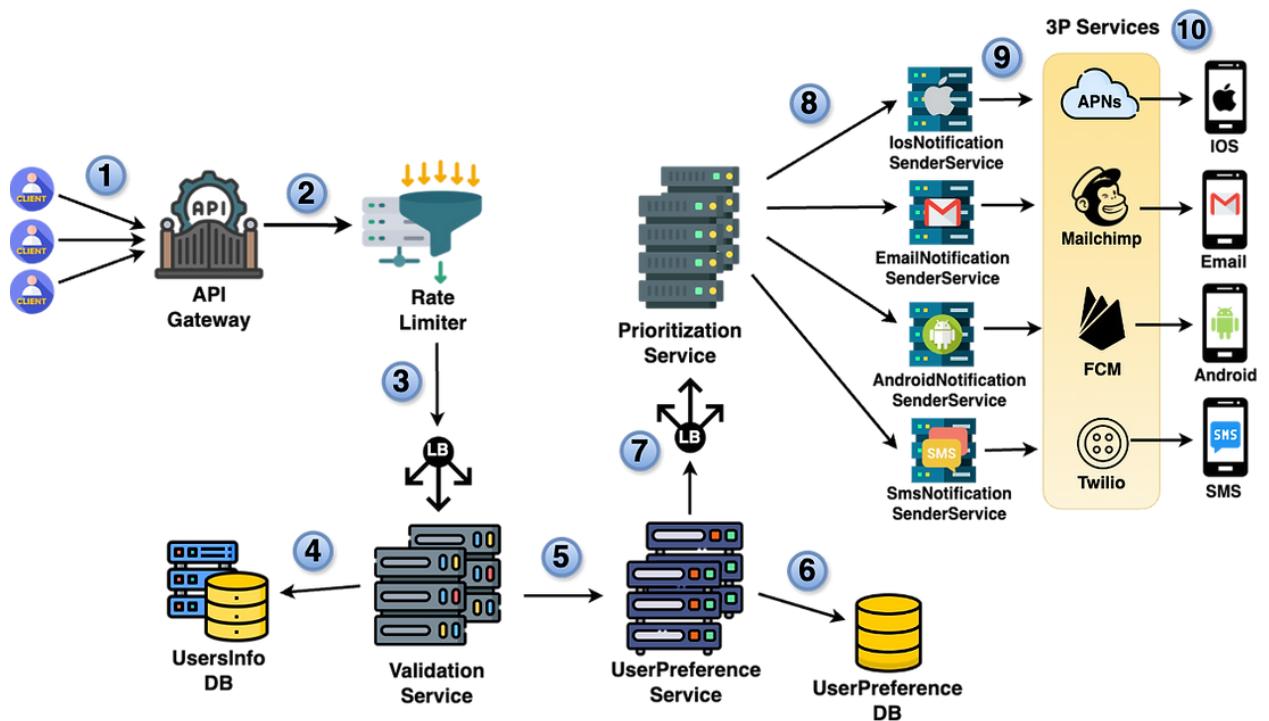


High Level Design-Handling different types of Notifications

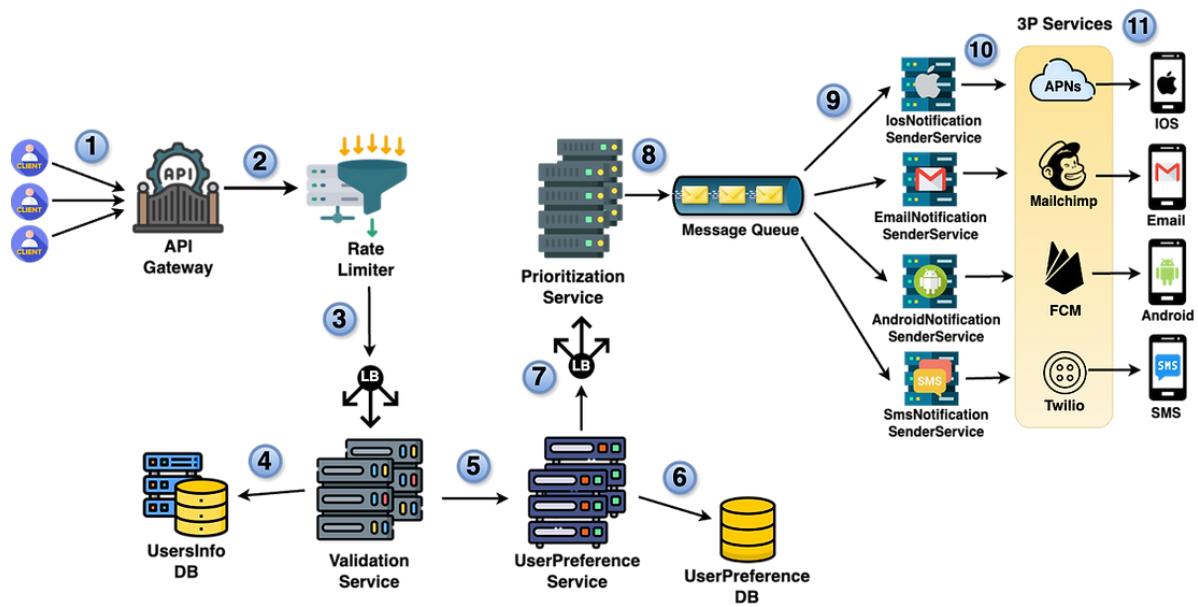




High Level Design-Reorganizing Services



High Level Design-Decoupling & Optimizing Design



Deep Dive Insights- Database Modeling (User Info DB Schema)

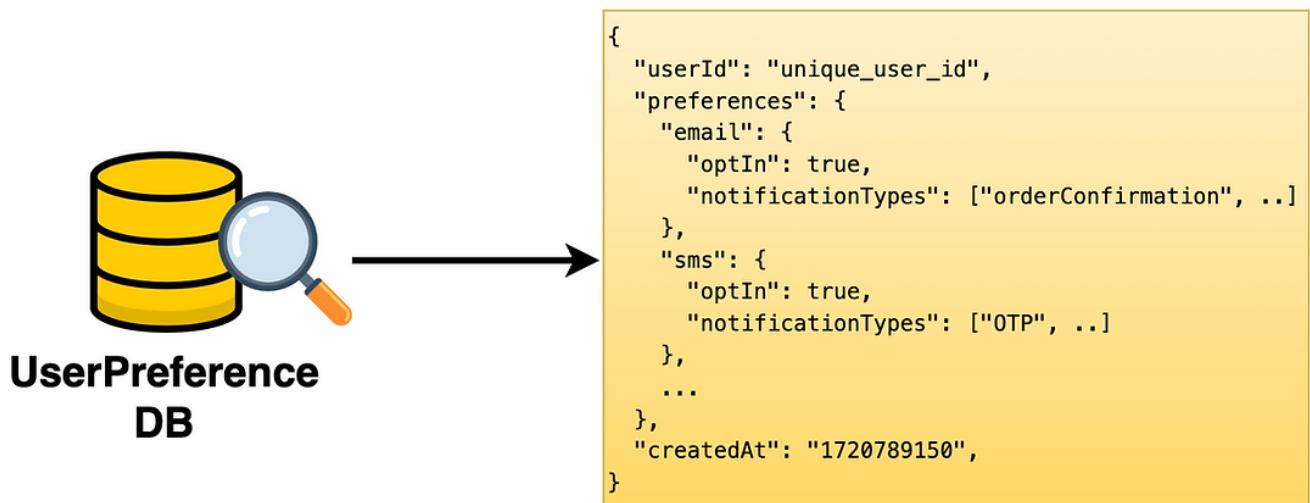


**UserInfo
DB**

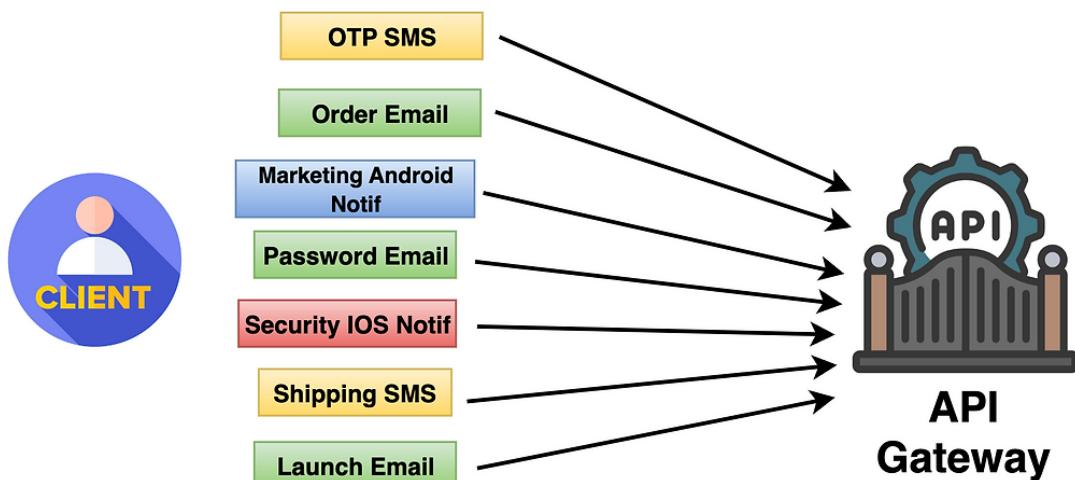


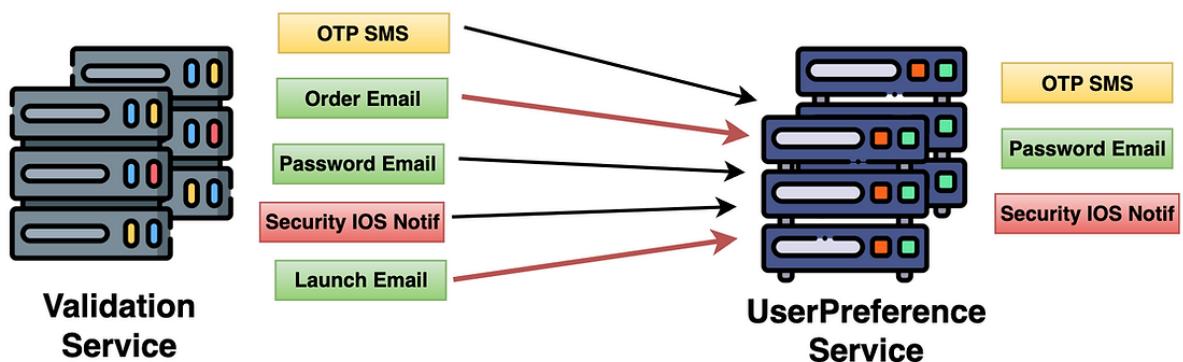
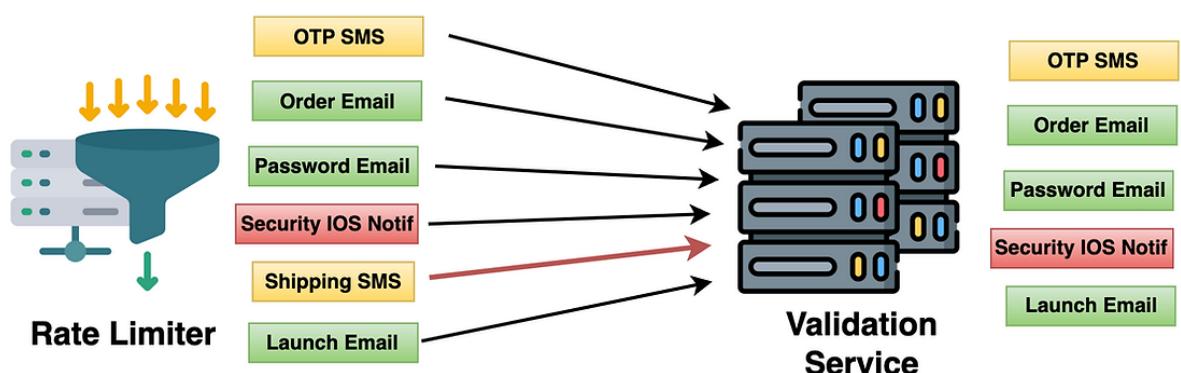
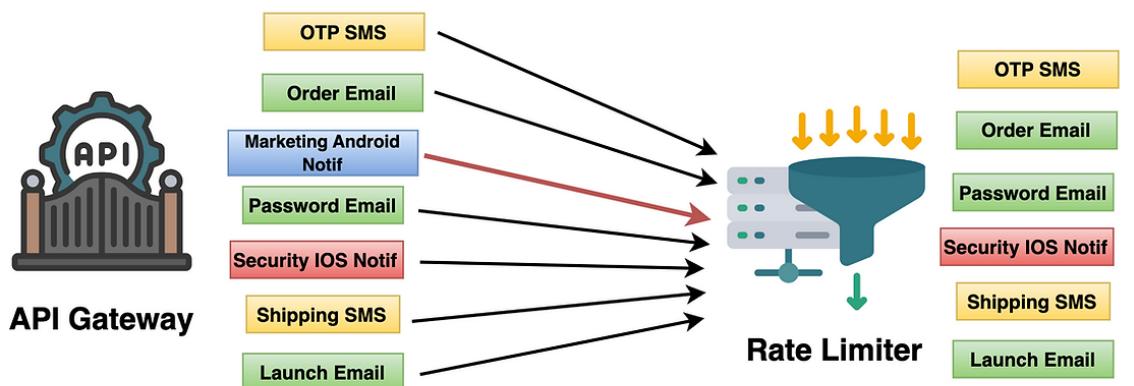
Column Name	Data Type	Description
userId	VARCHAR	The unique identifier for each user (PK).
name	VARCHAR	The user's full name.
email	VARCHAR	The user's email address (unique).
phoneNumber	VARCHAR	The user's phone number (unique).
isActive	BOOLEAN	Indicates if the user's account is active.
createdAt	TIMESTAMP	When the user was created.
updatedAt	TIMESTAMP	When the user's information was last updated.

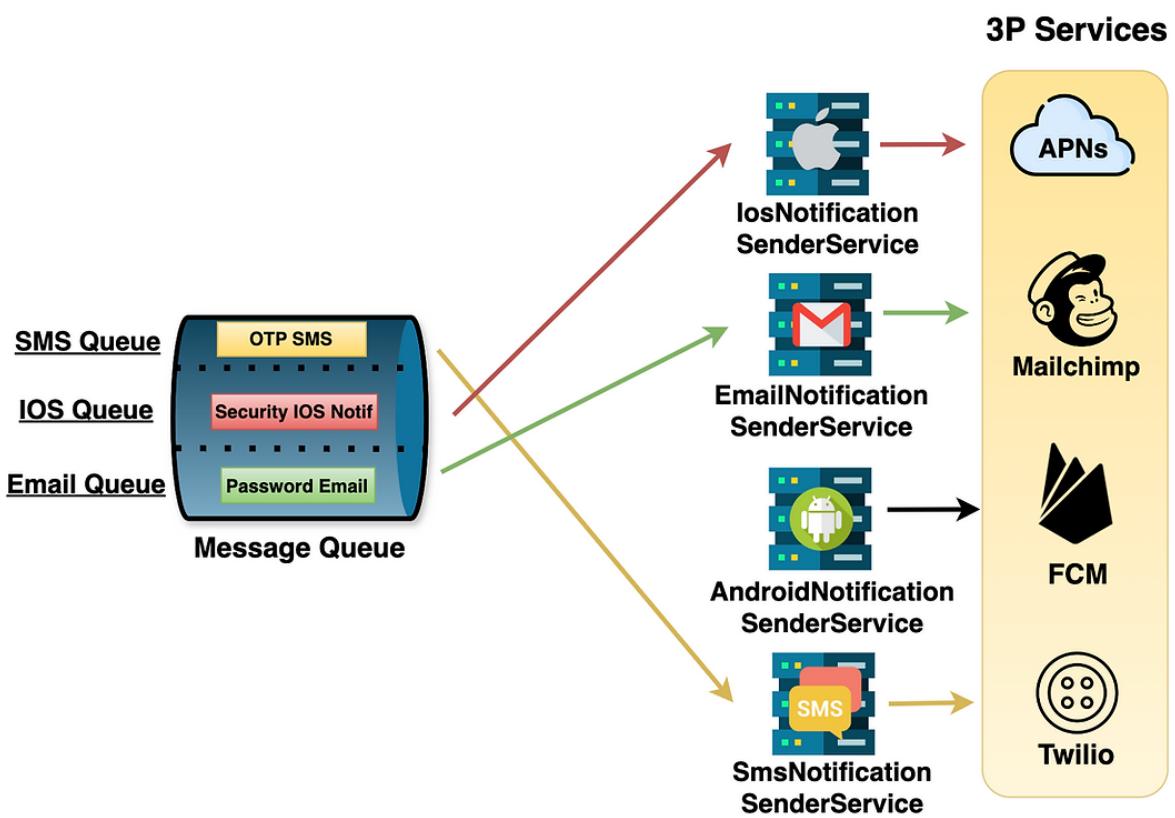
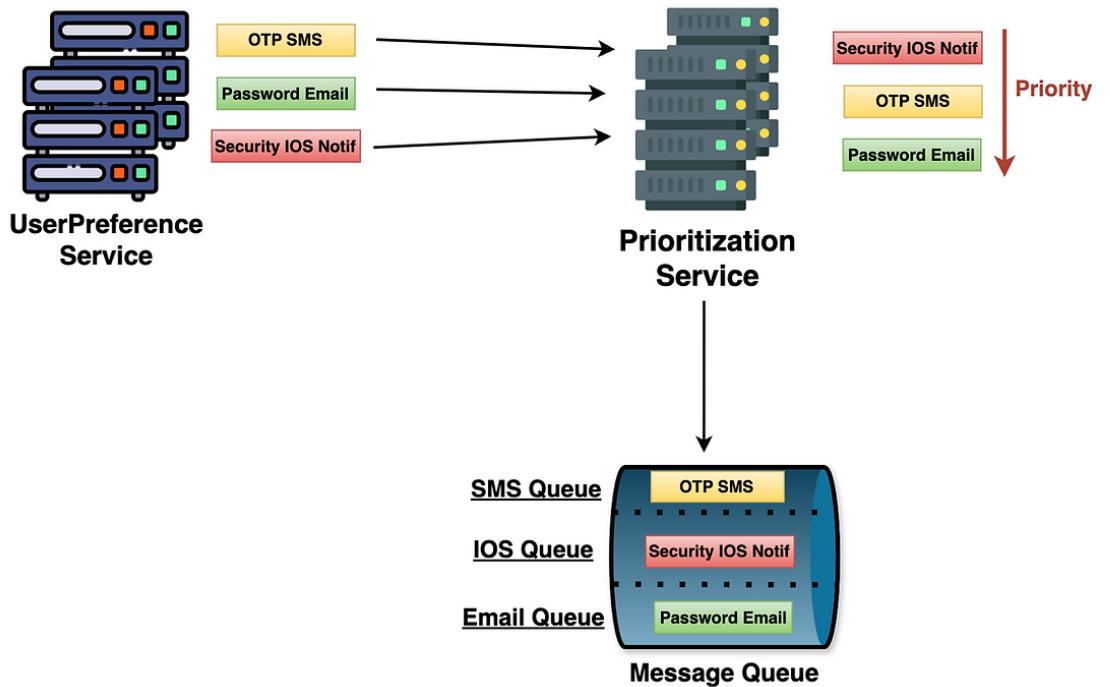
Deep Dive Insights- Database Modeling (User Preference DB Schema)



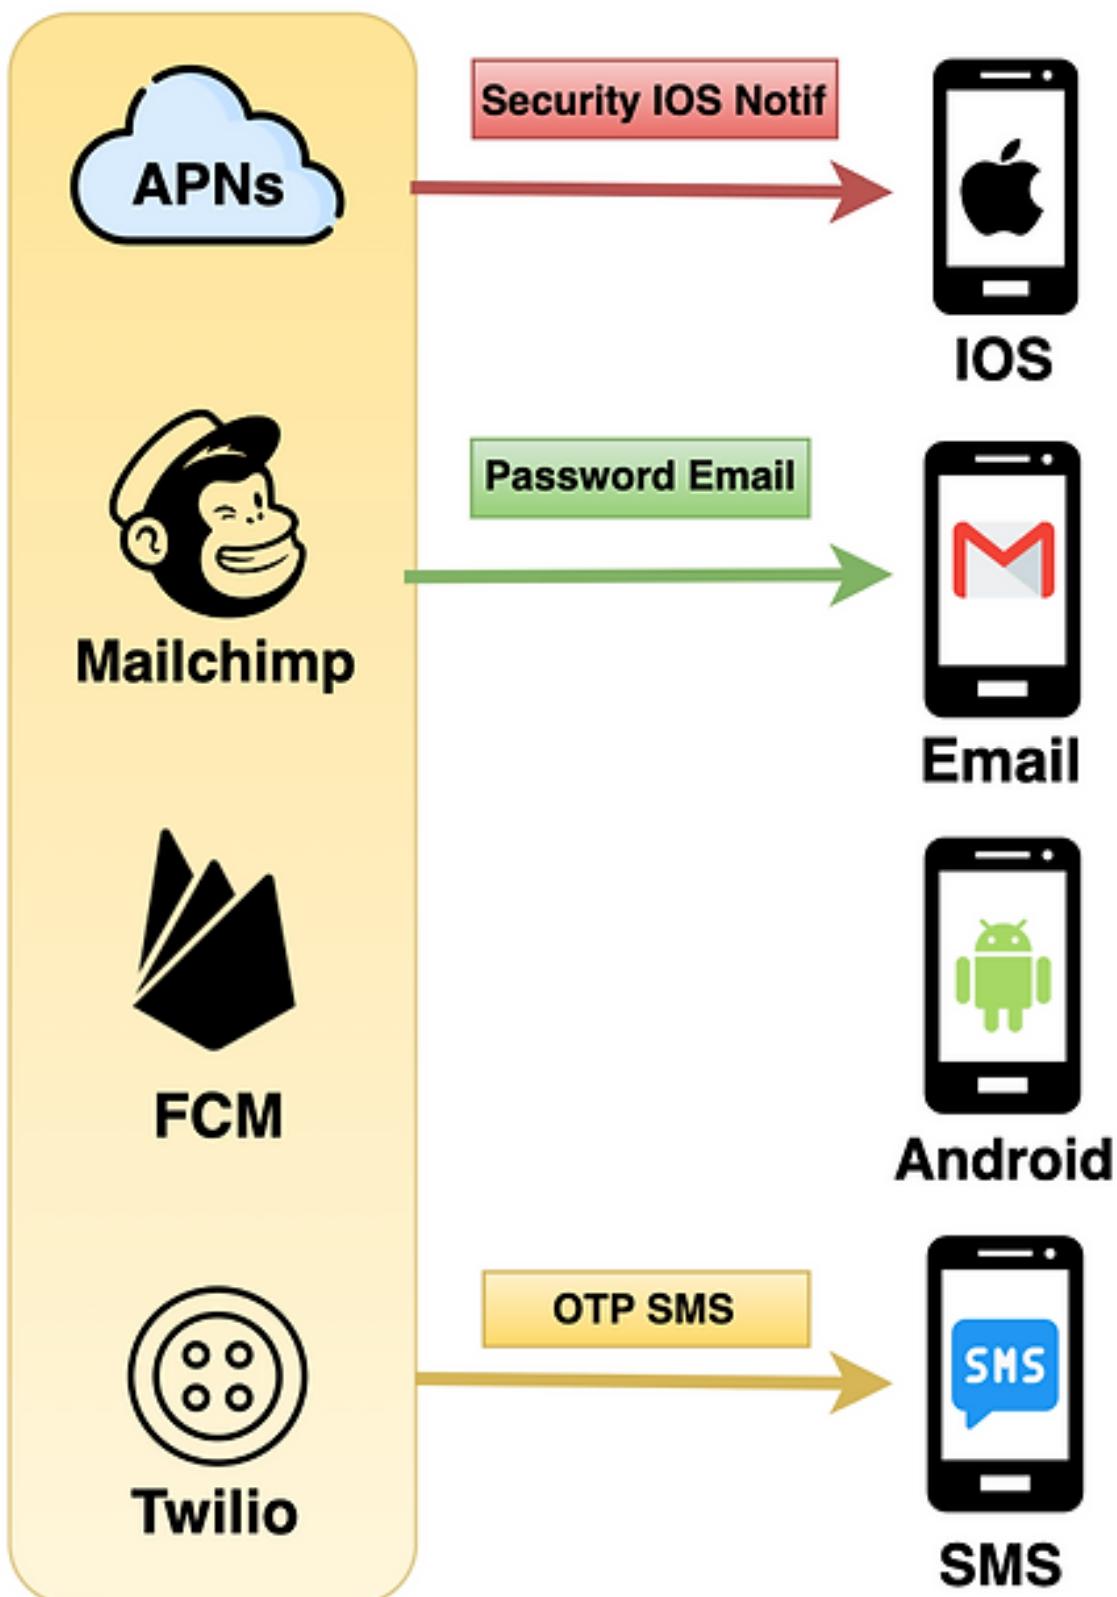
Deep Dive Insights- Notifications Overall Flow



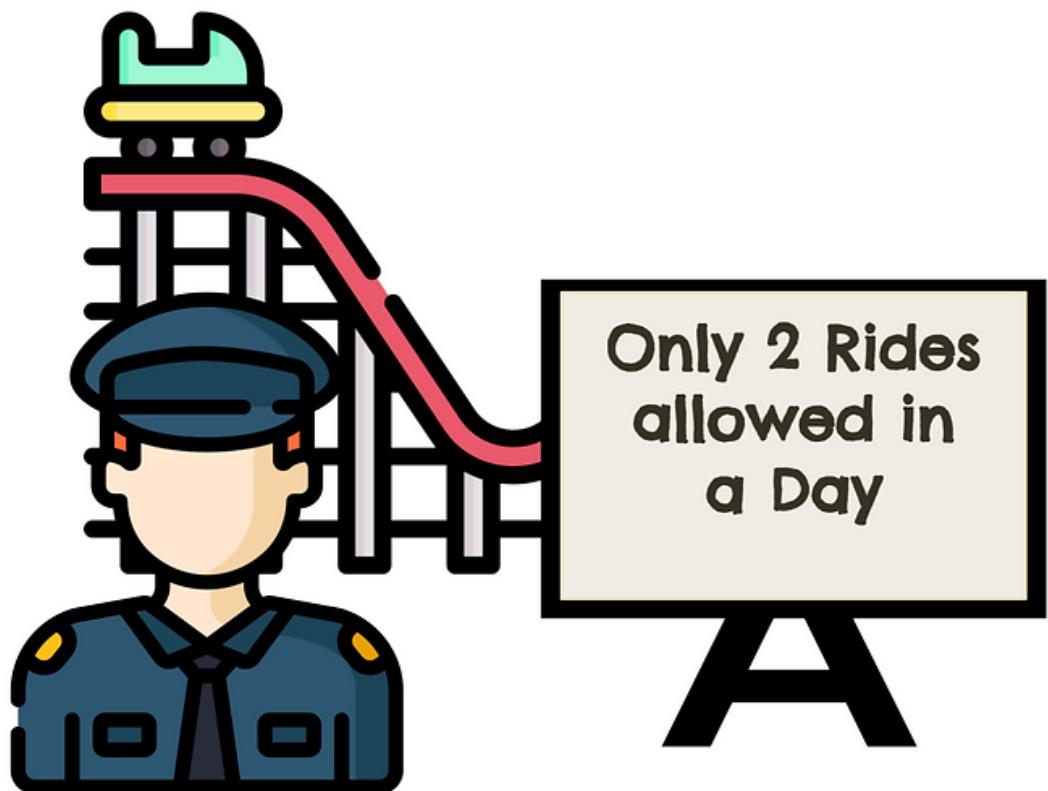




3P Services



Analogy

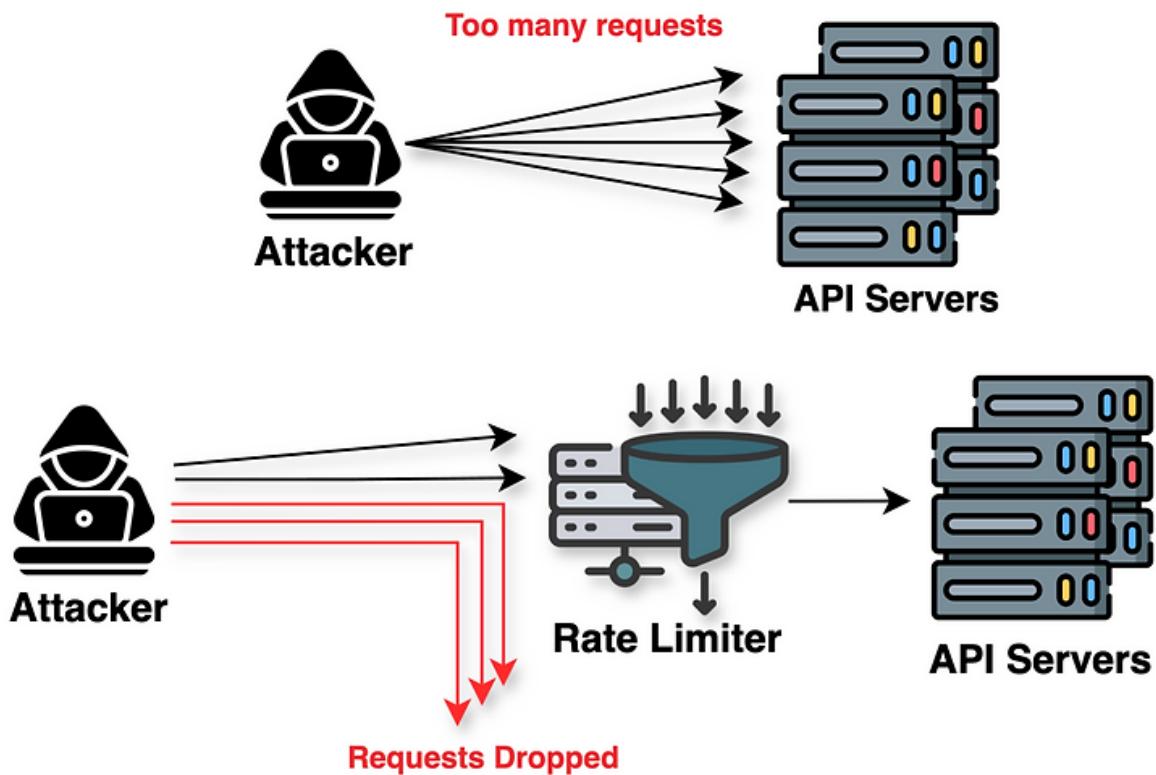


What is a Rate Limiter?

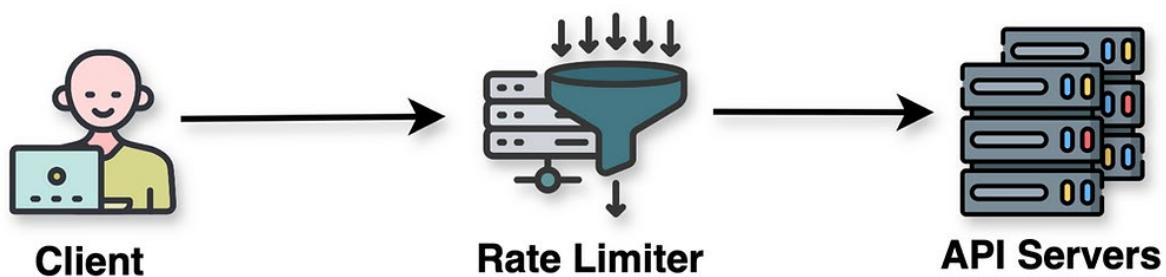


Rate Limiter

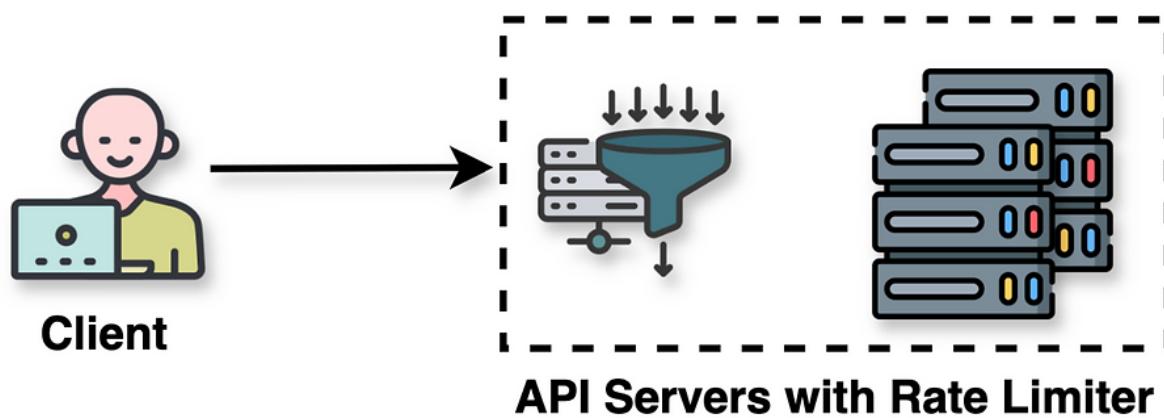
Why do we need a Rate Limiter?



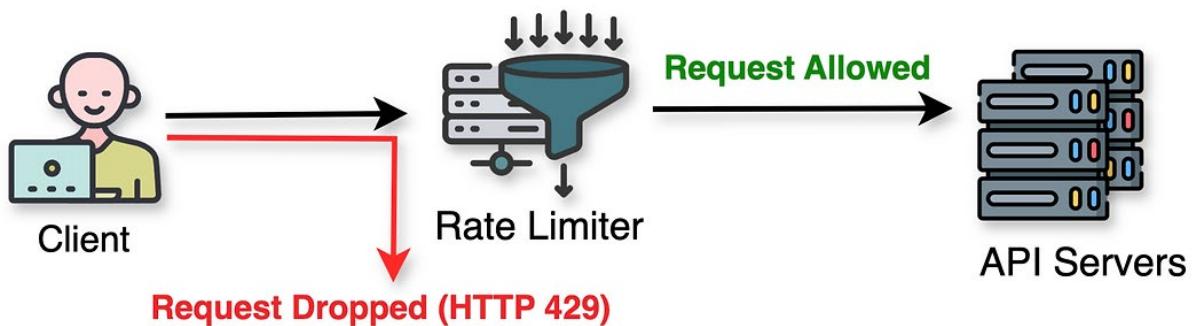
Rate Limiter comes before API Servers



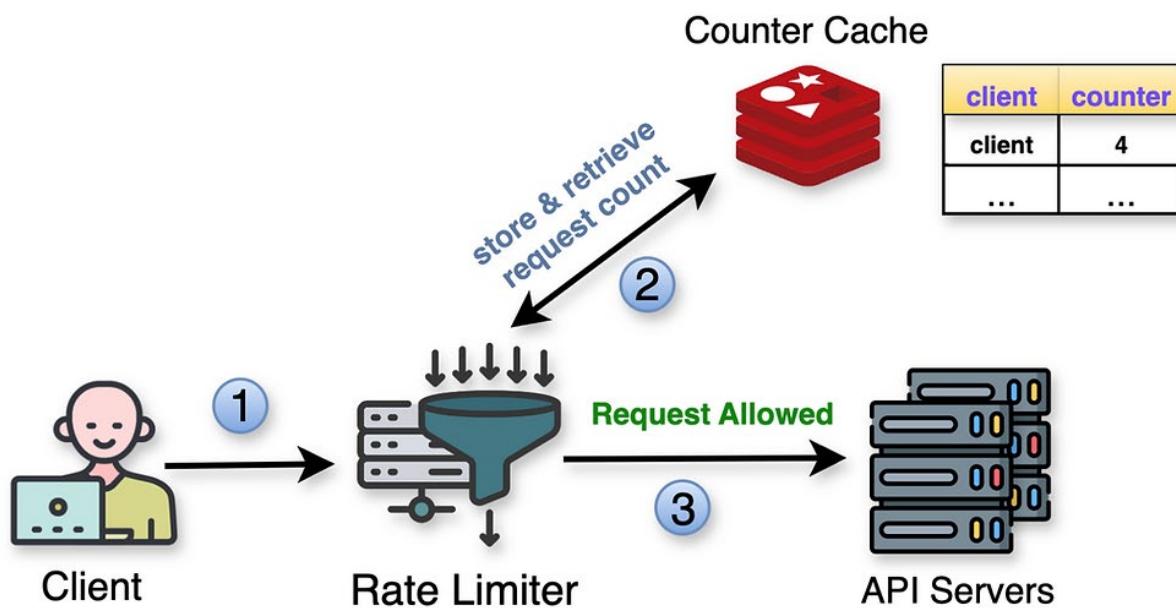
Rate Limiter comes with API Servers



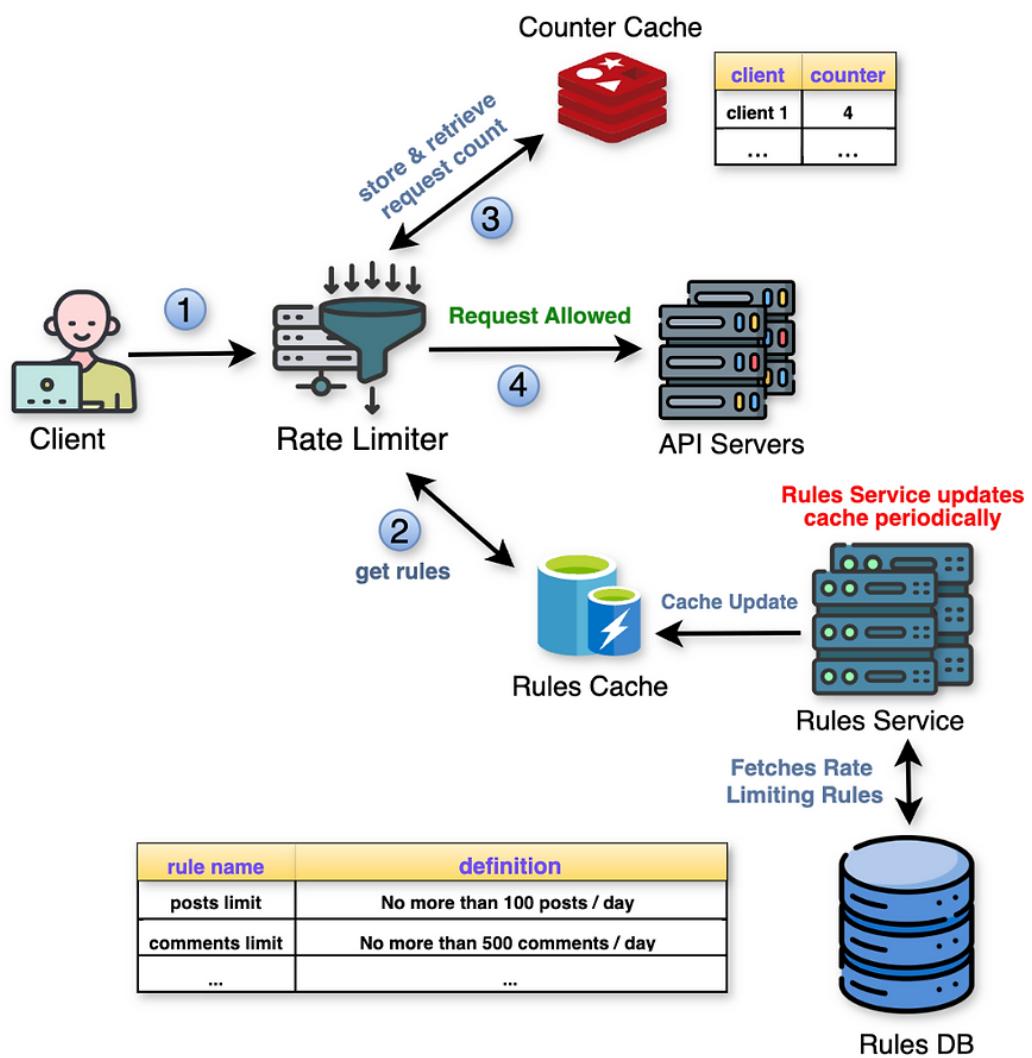
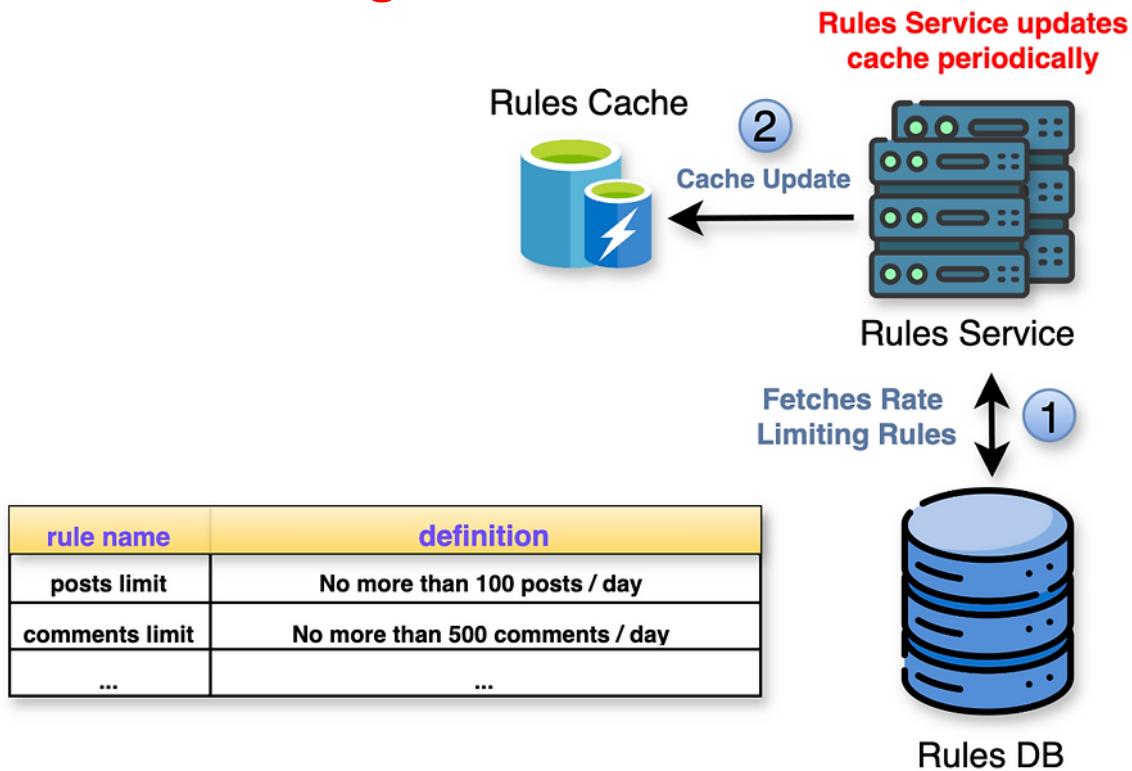
Basic High-Level Flow



Storing Request Counters

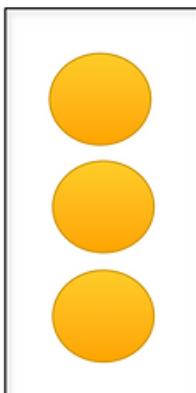


Storing Rate Limiter Rules



Token Bucket Algorithm

Bucket Refiller

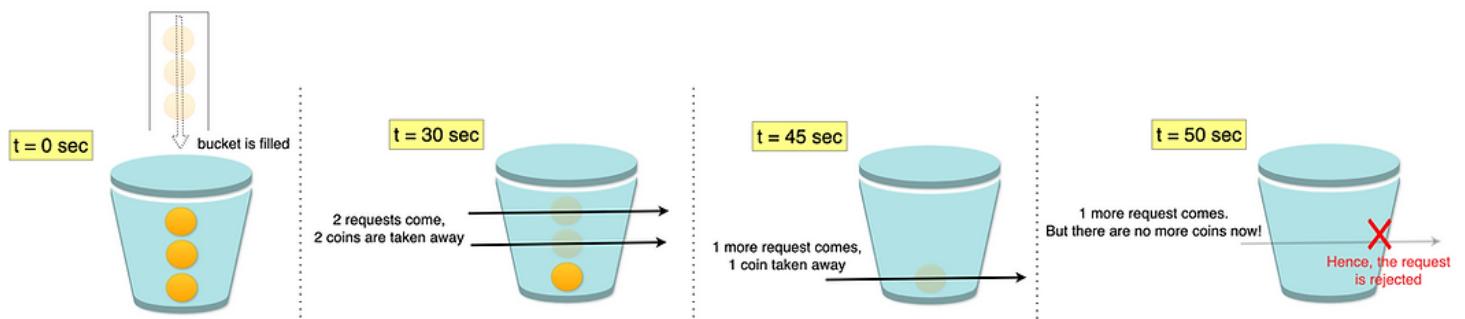


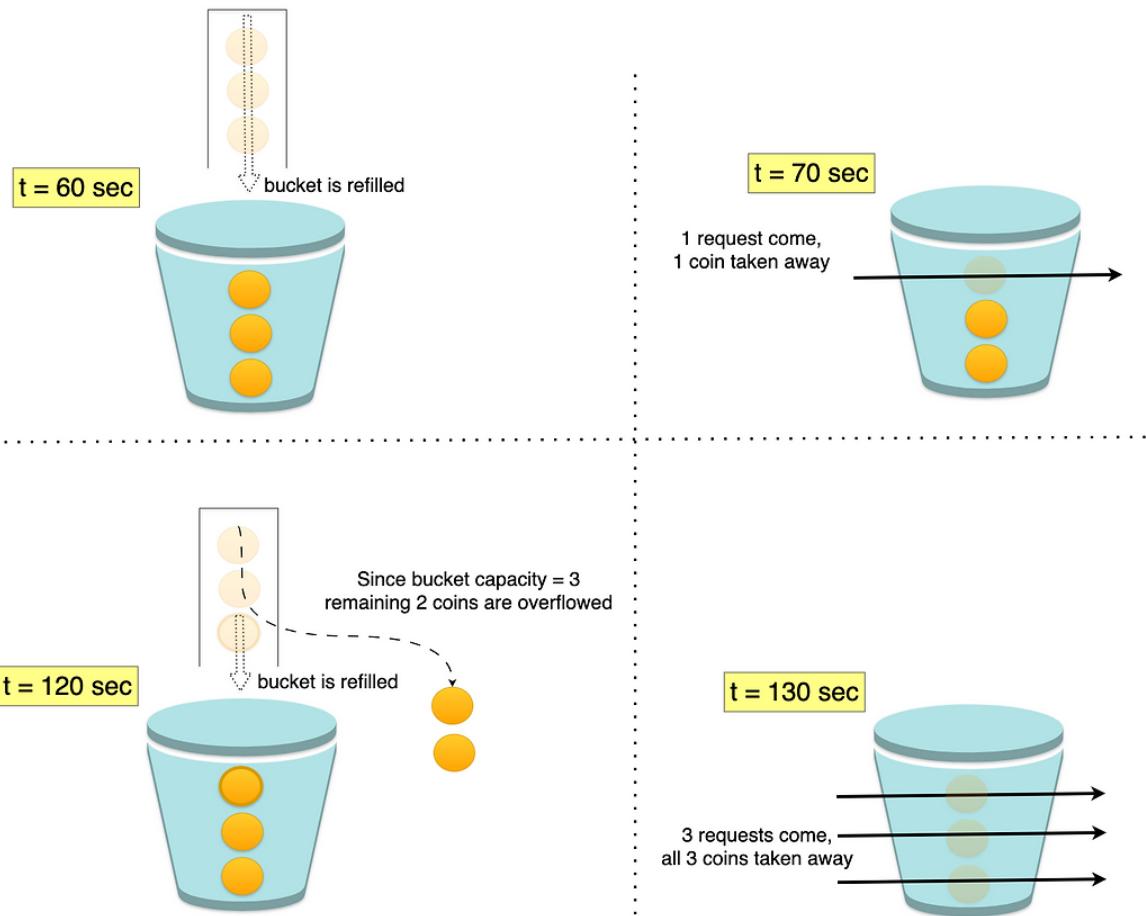
Bucket Refiller's job is to refill the bucket every 60 sec.



Bucket Capacity = 3

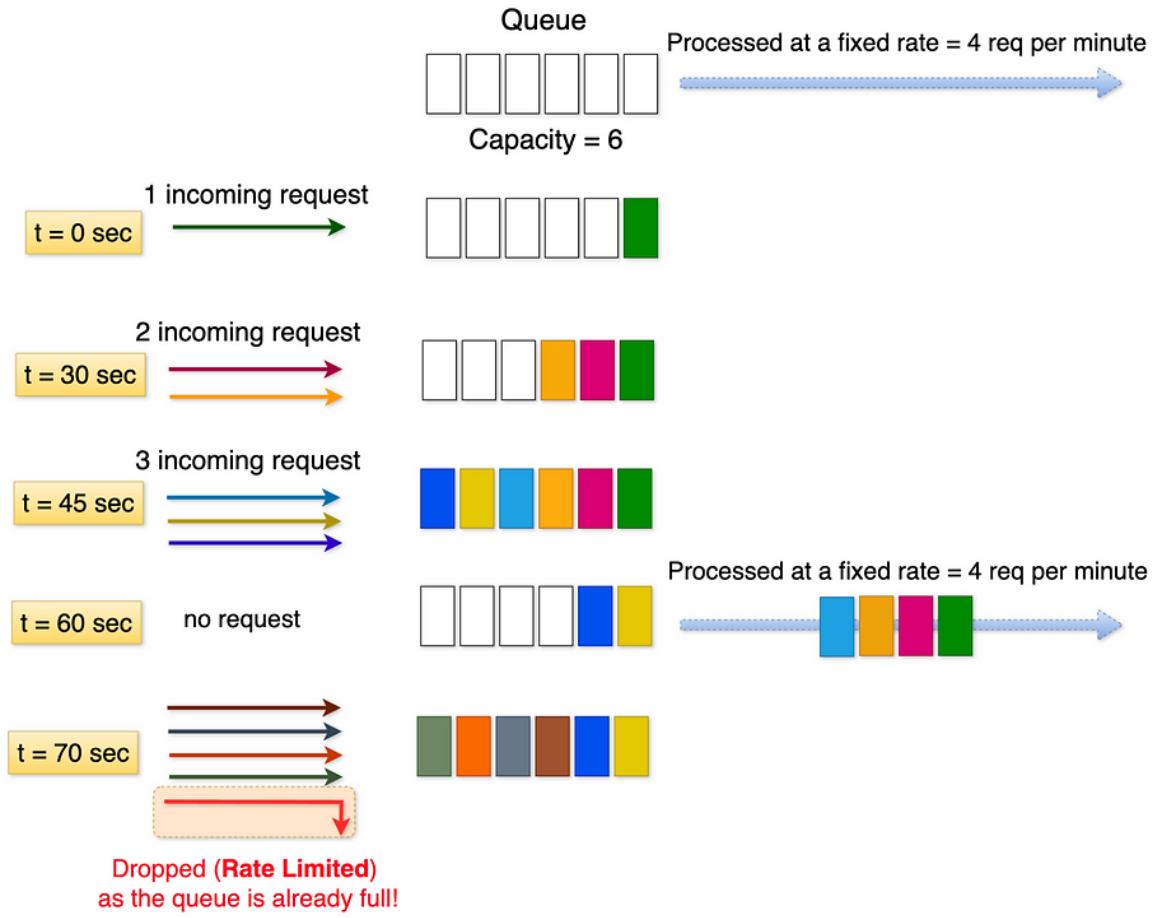
Bucket





Leaky Bucket Algorithm



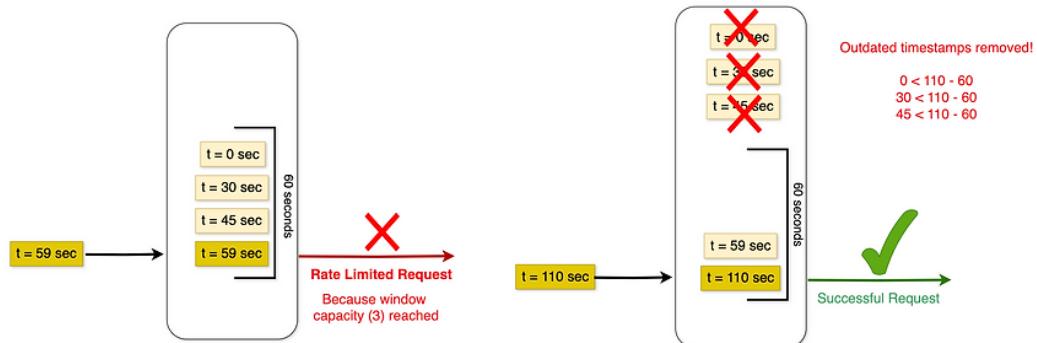
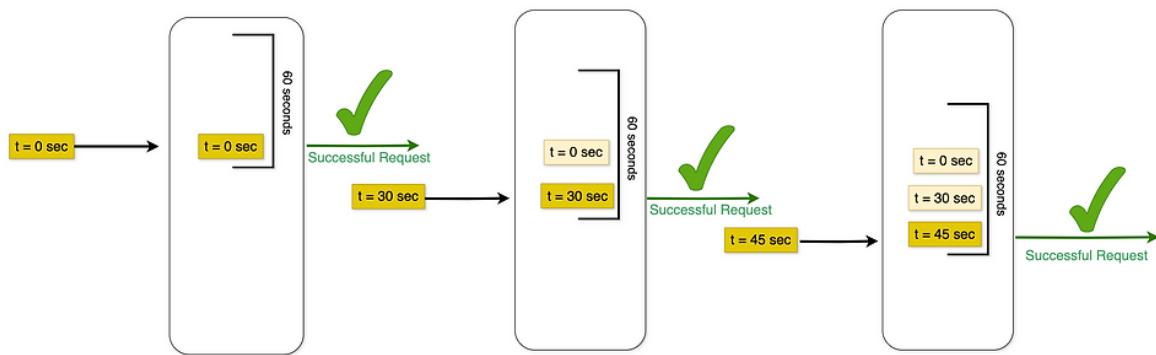
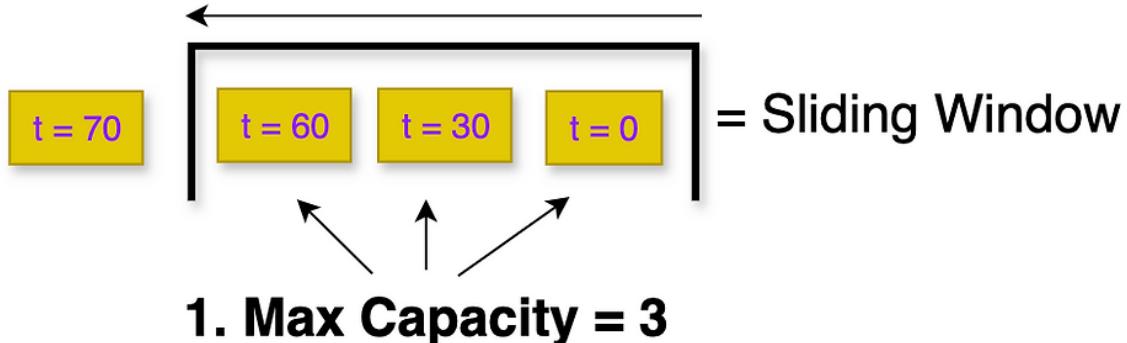


Fixed Window Counter Algorithm

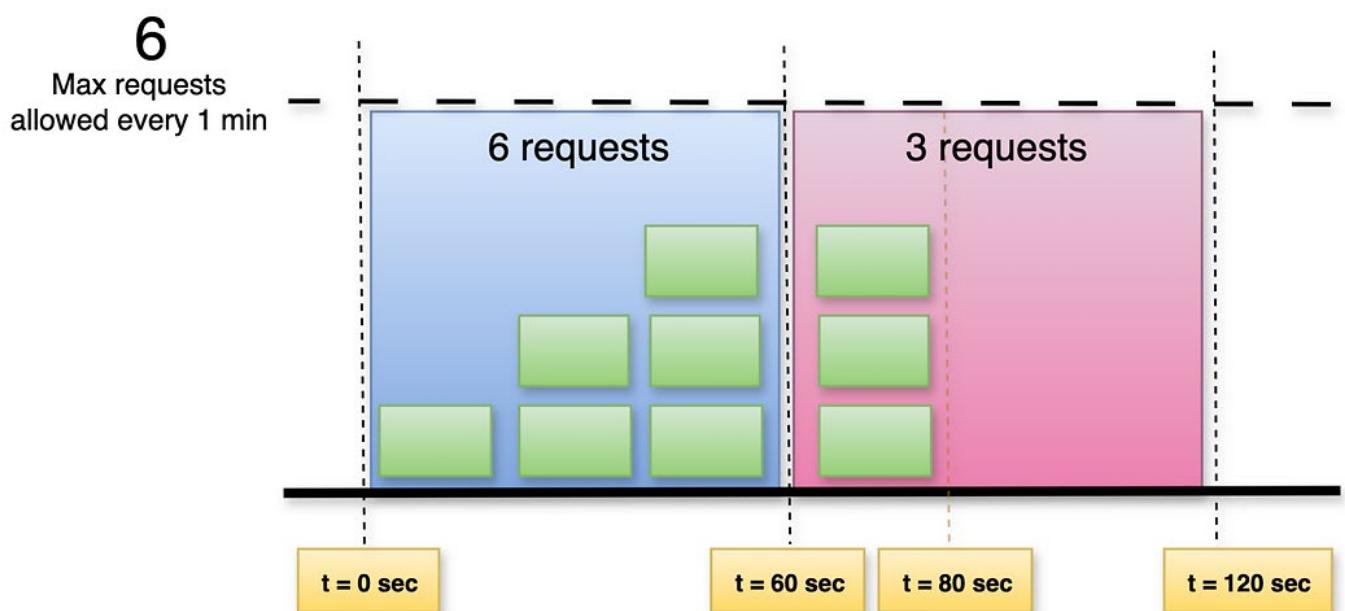


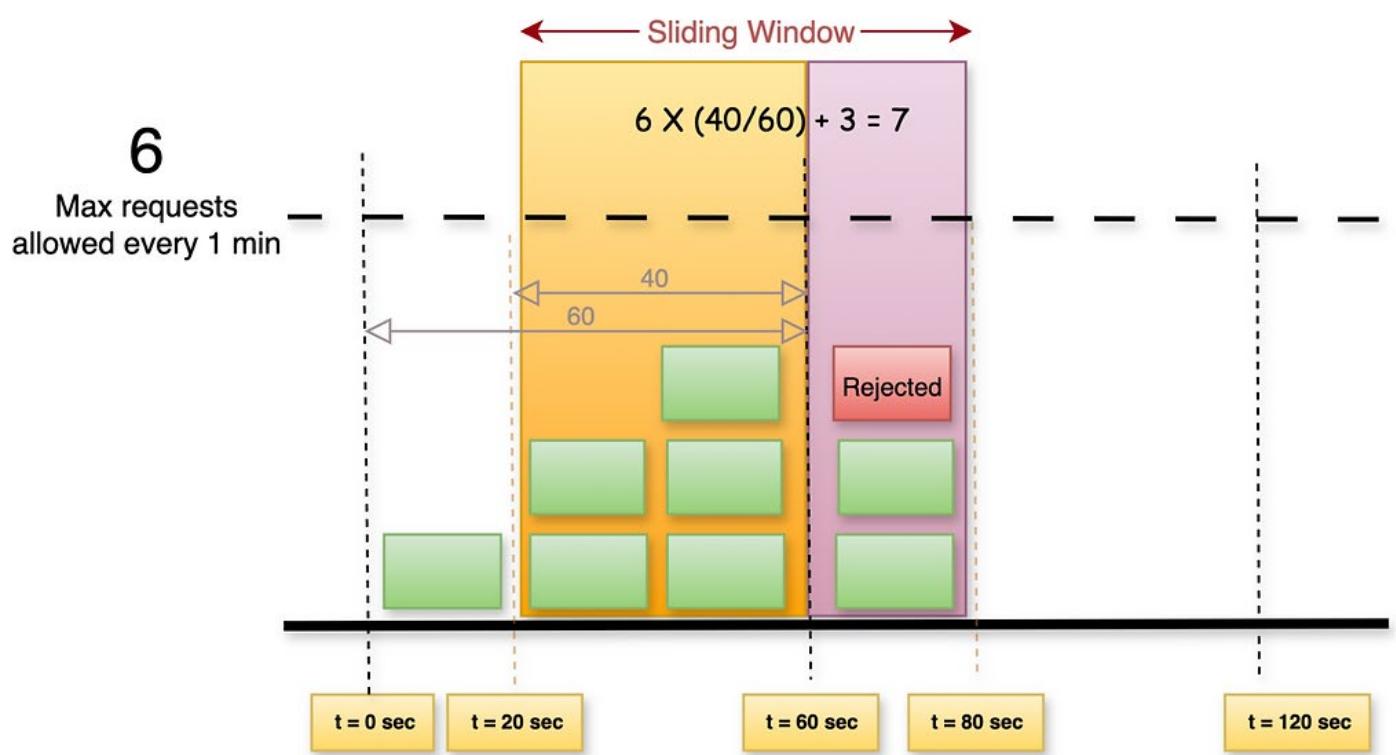
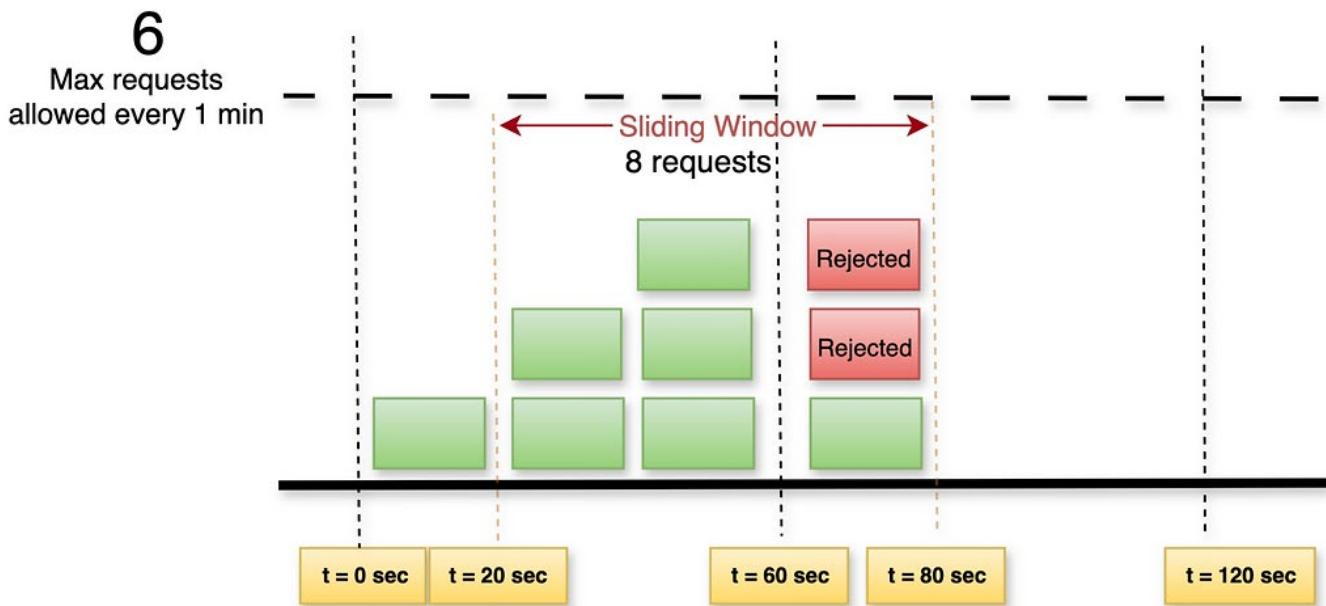
Sliding Window Log Algorithm

2. Lookback = 60 sec

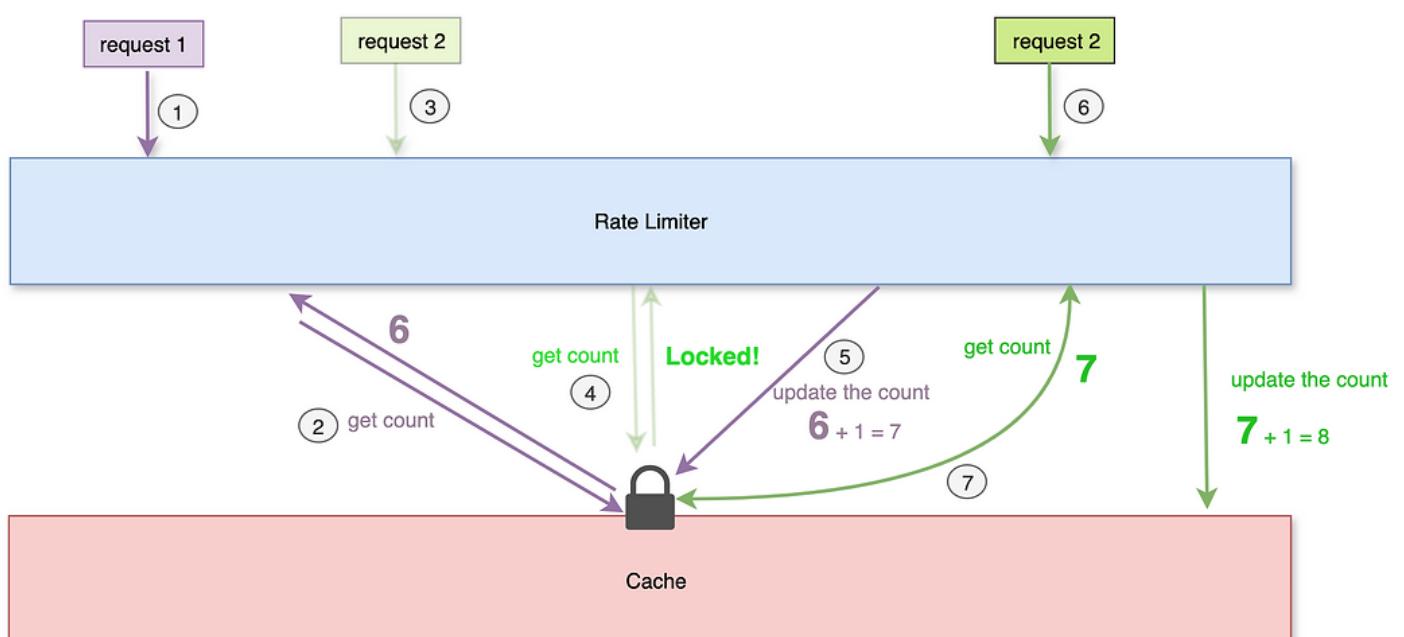
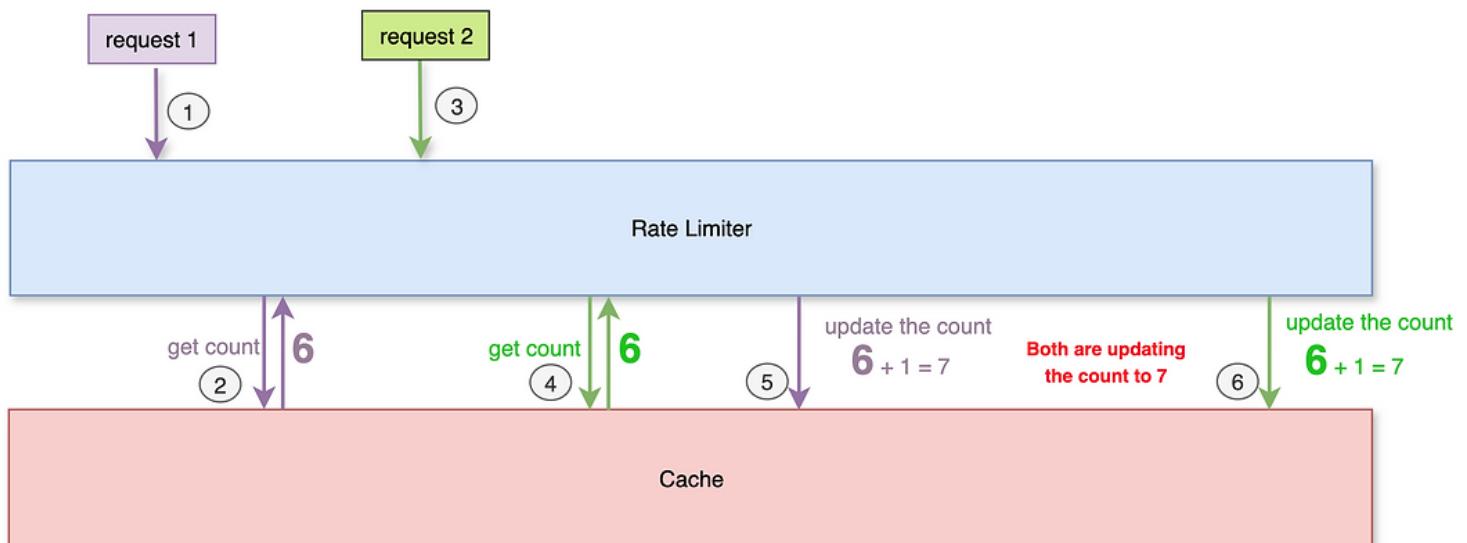


Sliding Window Counter Algorithm

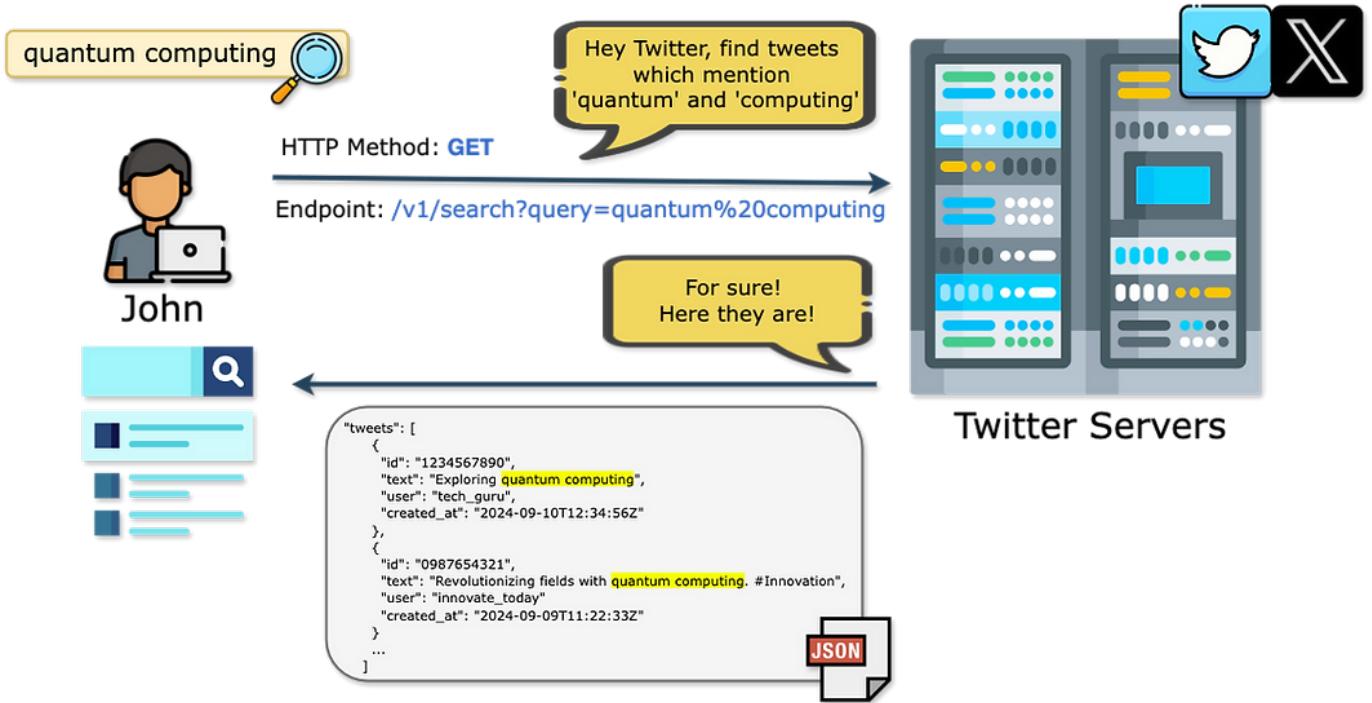




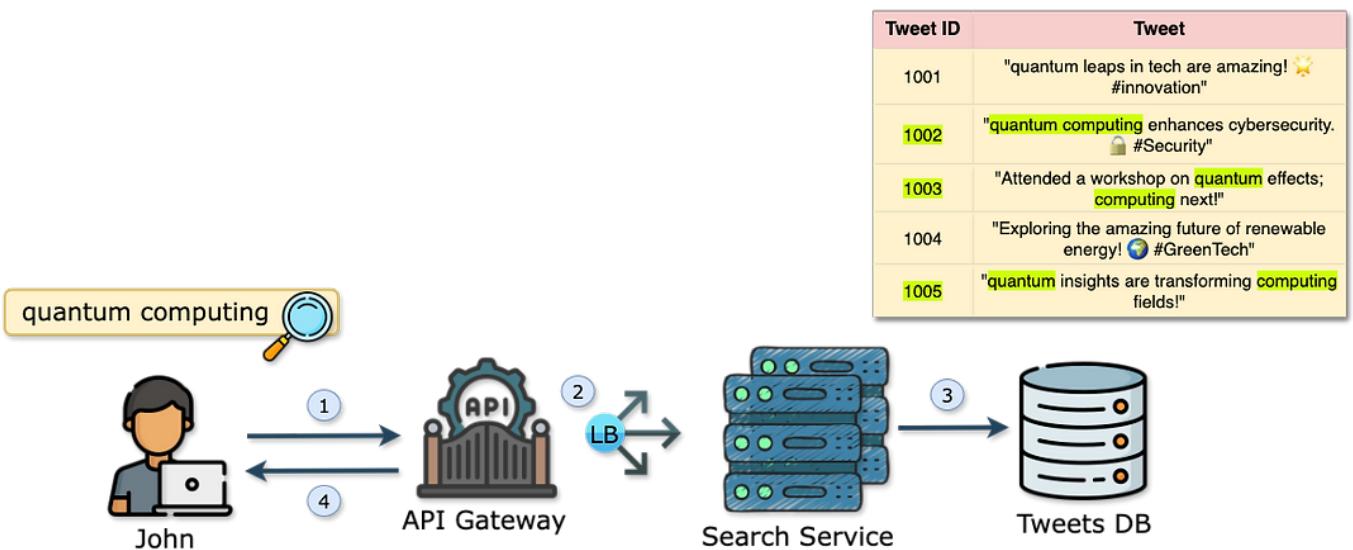
Handling Race Condition



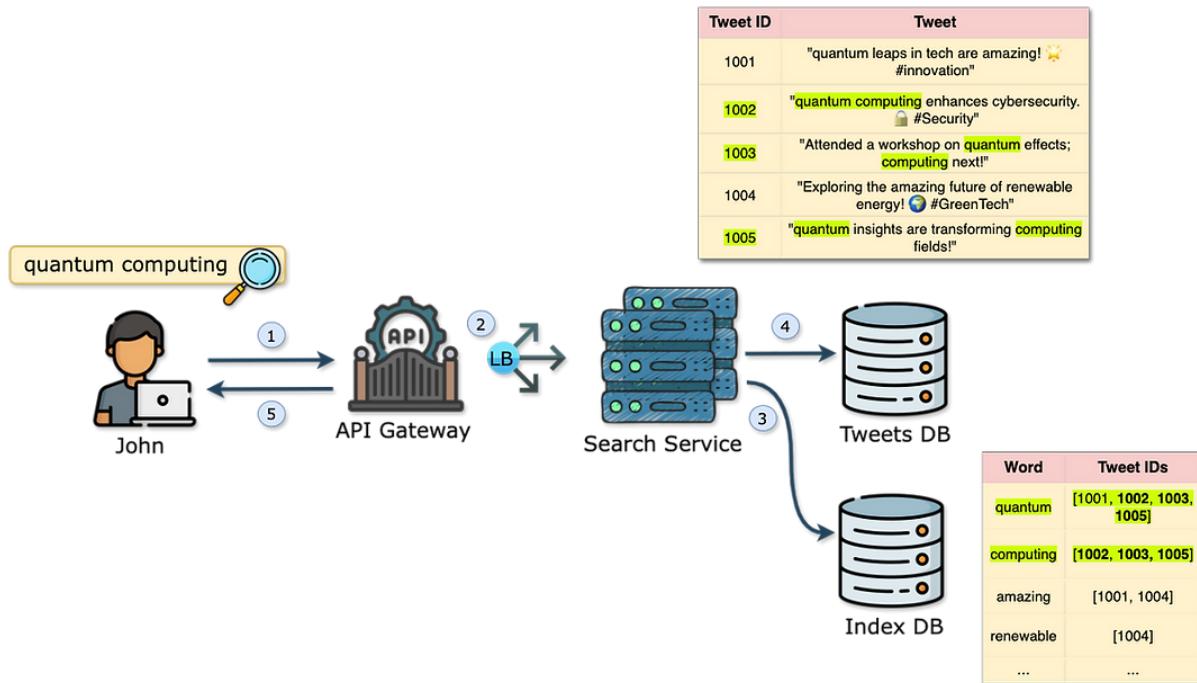
API Design- Twitter Search



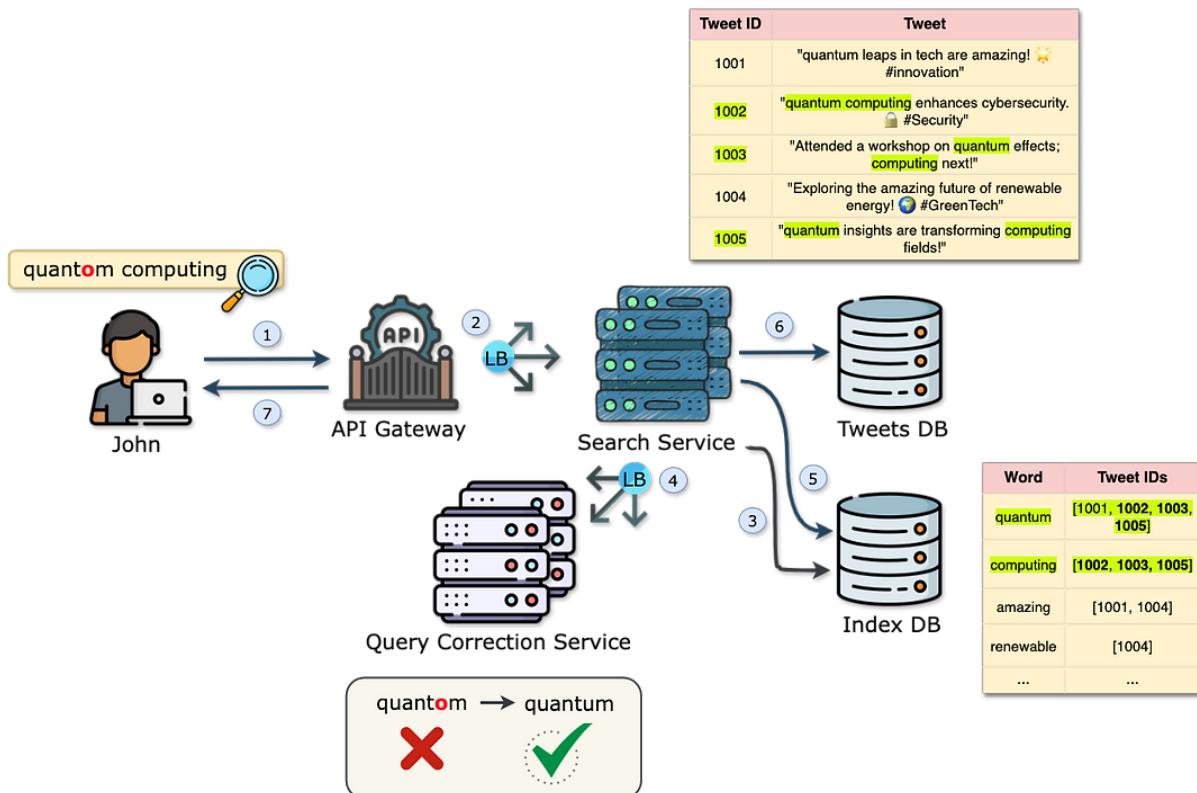
High Level Design- Twitter Search



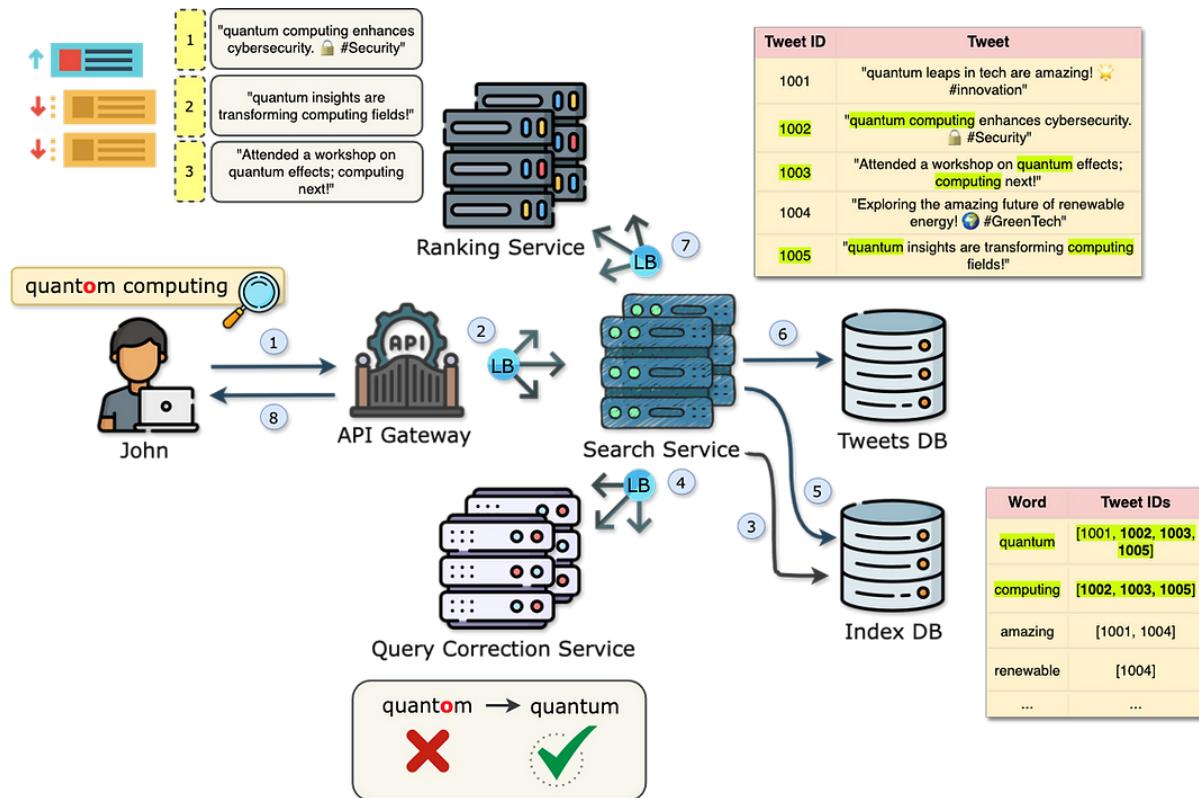
High Level Design- Twitter Search (Indexing)



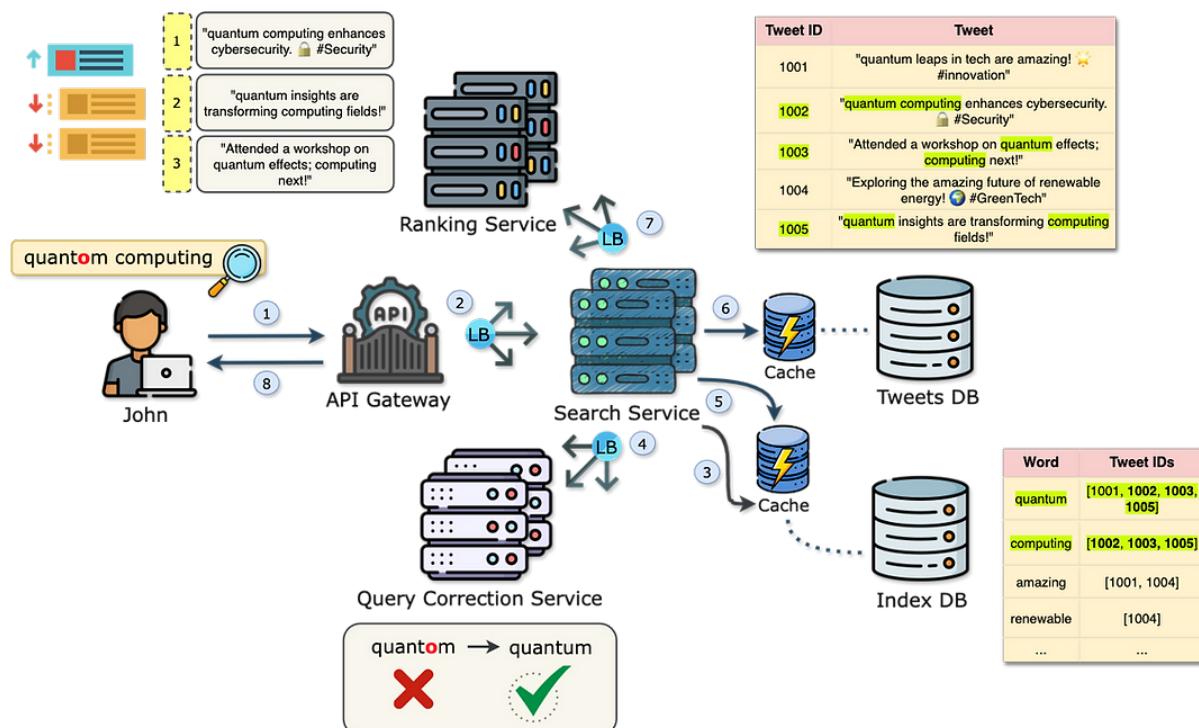
High Level Design- Twitter Search (Fuzzy Search)



High Level Design- Twitter Search (Ranking)

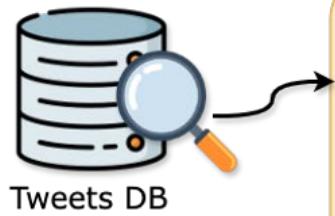


High Level Design- Twitter Search (Final Design)



DATABSE SCHEMAS

Tweets DB Schema



```
{  
  "TweetId": "Unique identifier for each tweet",  
  "Text": "Content of the tweet",  
  "UserId": "Identifier for the user who posted the tweet",  
  "CreatedAt": "Timestamp when the tweet was created"  
}
```

Index DB Schema

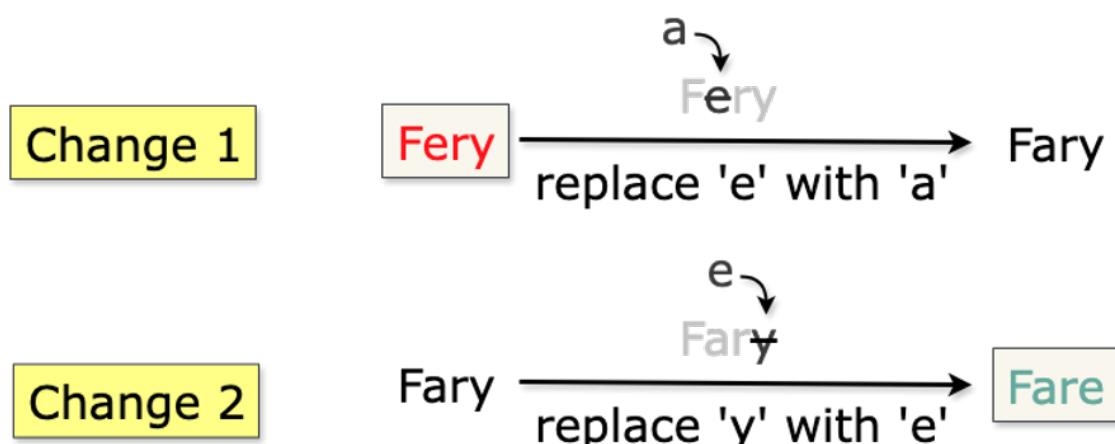


```
{  
  "Keyword": "A specific word or hashtag",  
  "TweetIds": ["List of tweet Ids having this keyword in the tweet"]  
}
```

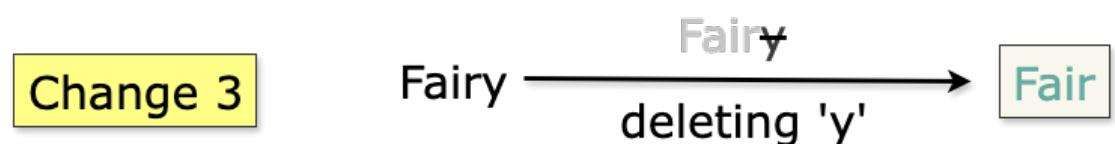
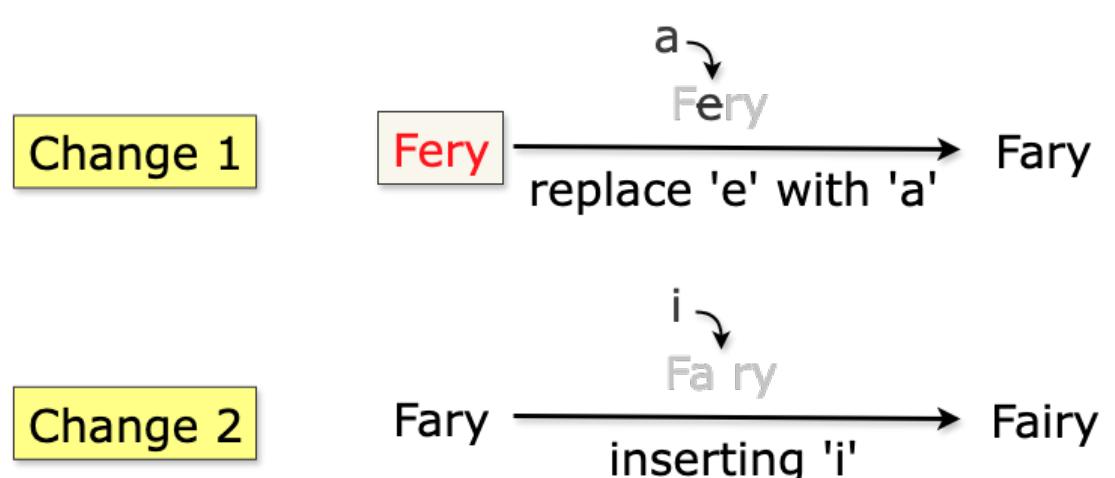
Query Correction Service (Fery to Ferry)



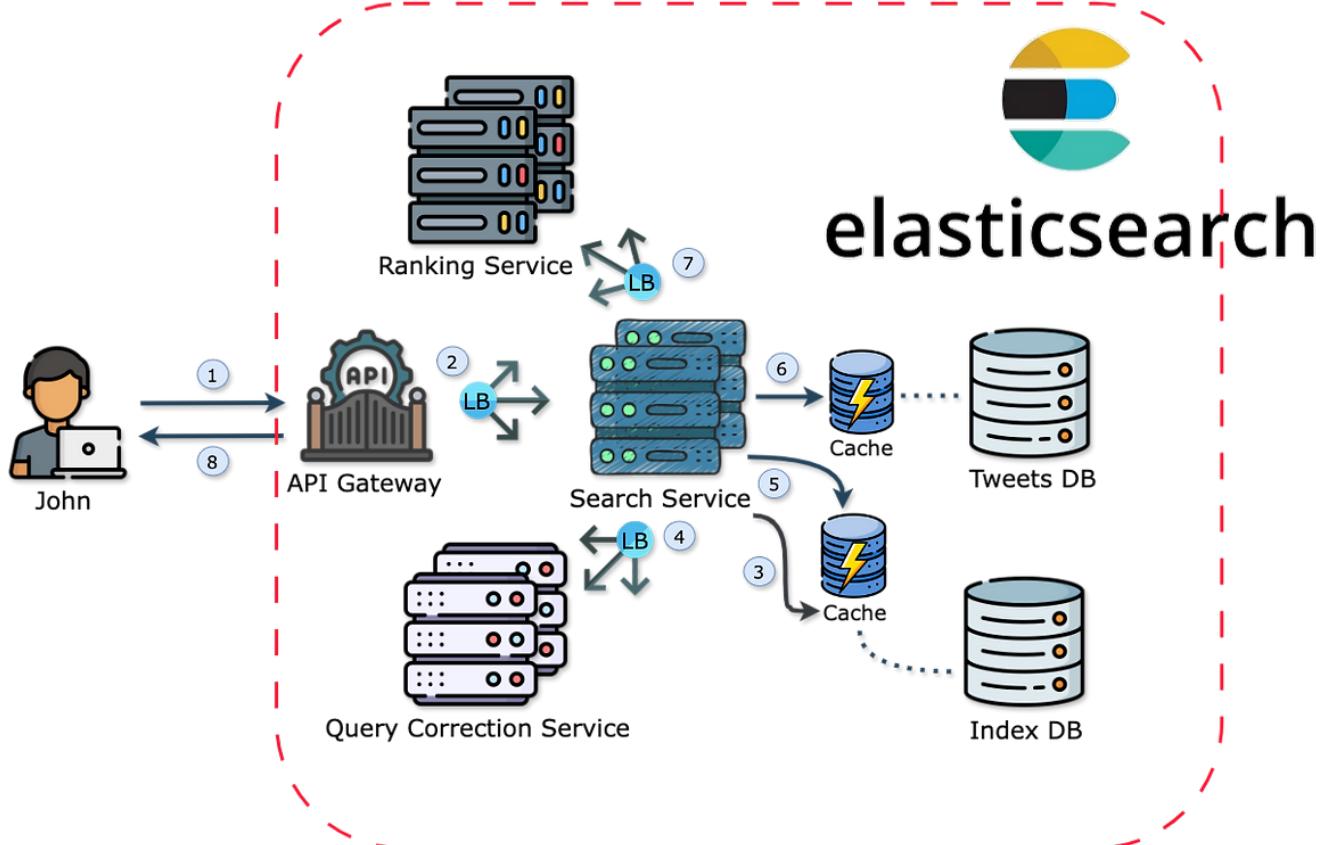
Query Correction Service (Fery to Fare)



Query Correction Service (Fery to Fair)



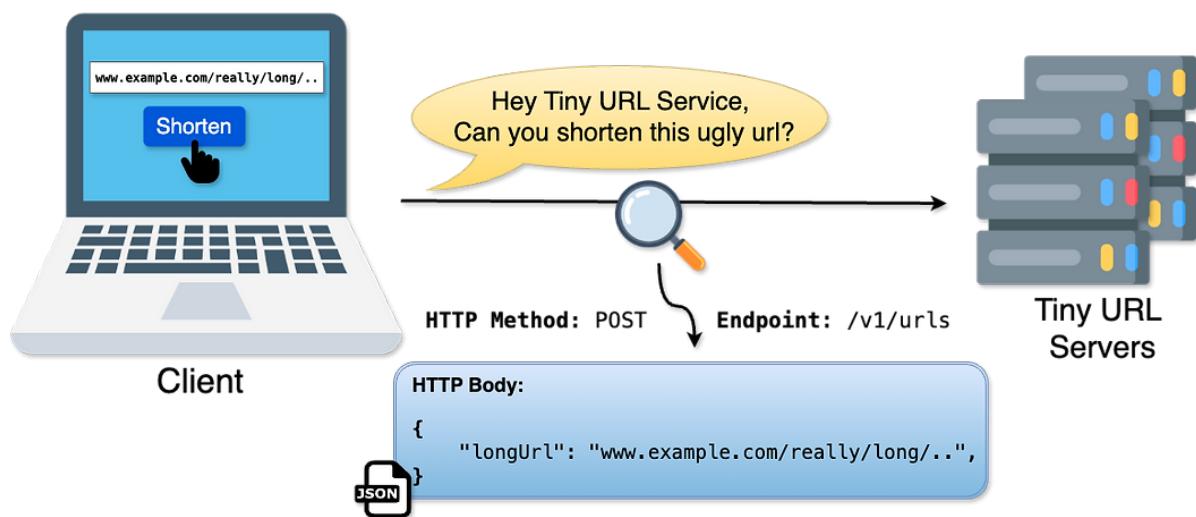
Elastic Search



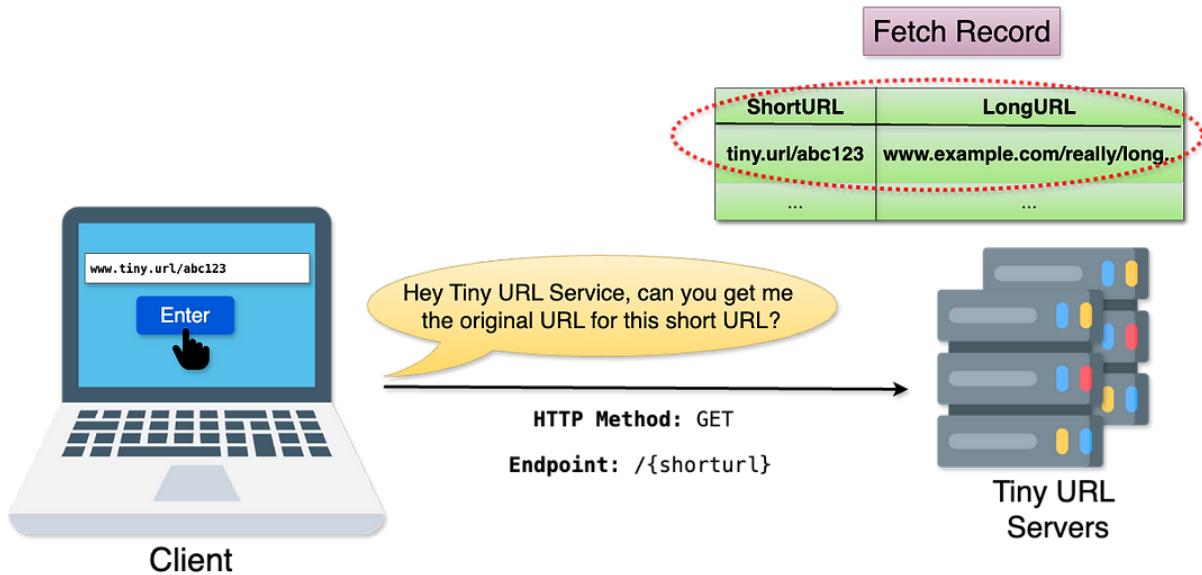
What is a Tiny URL?



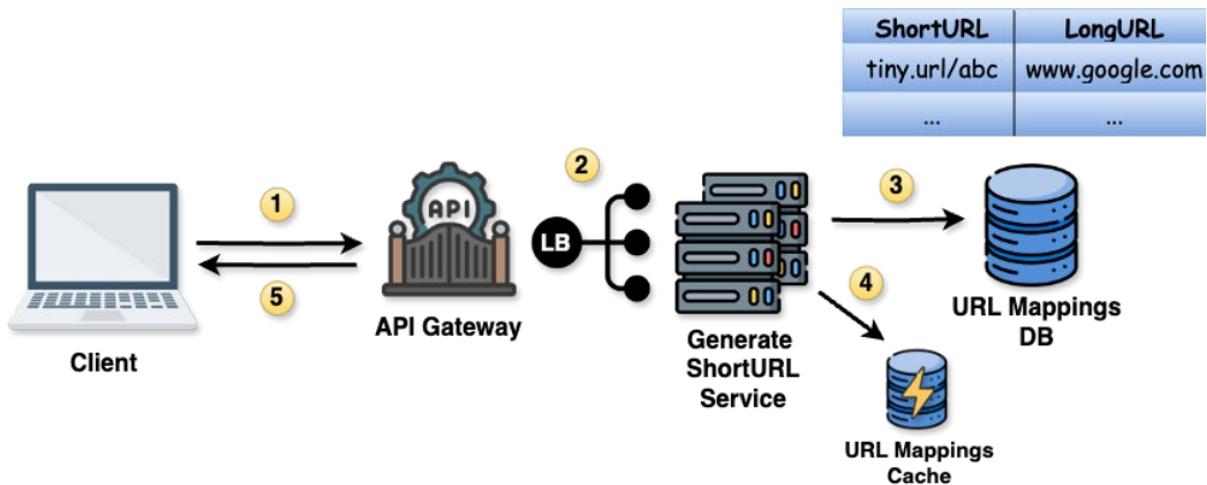
API Design- Generate Short URL



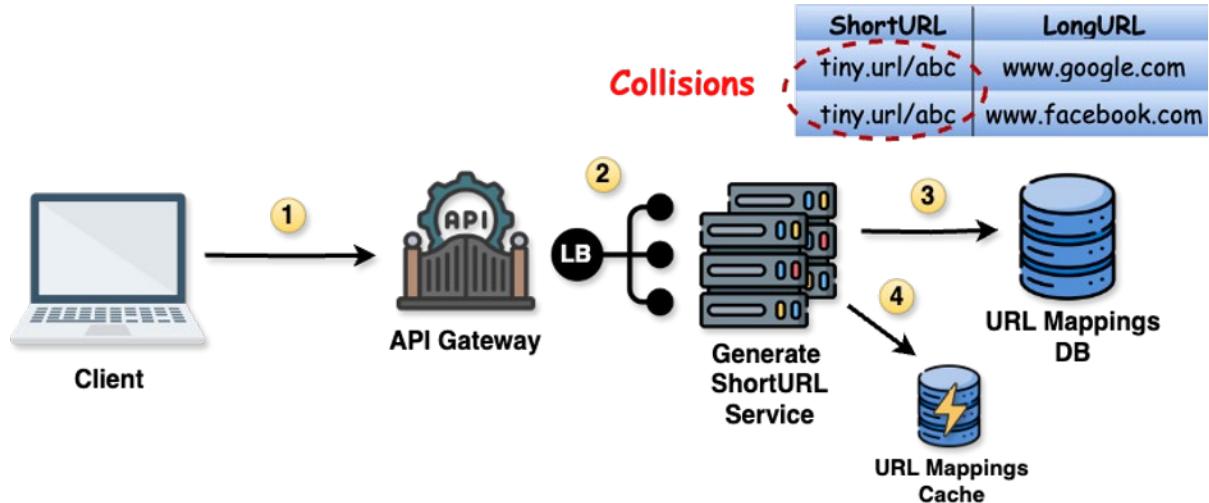
API Design- Get Long URL



High Level Design-Generate Short URL

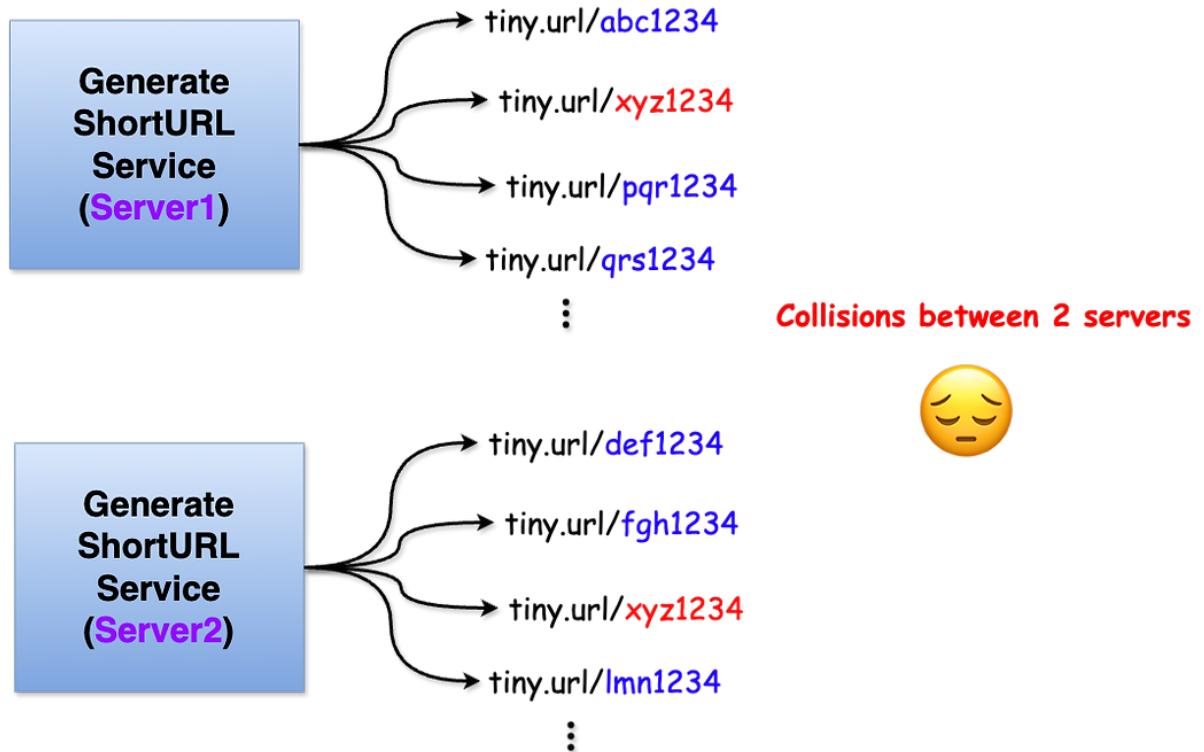


High Level Design-Collision's

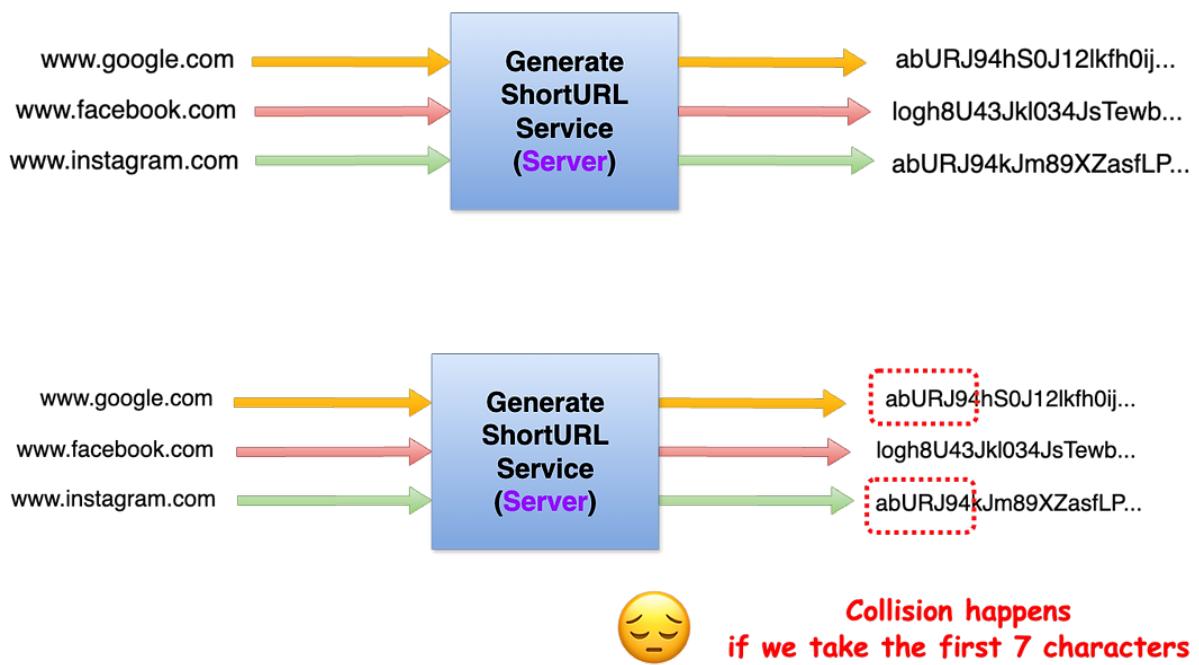


High Level Design-Approach1(Random String Generation)

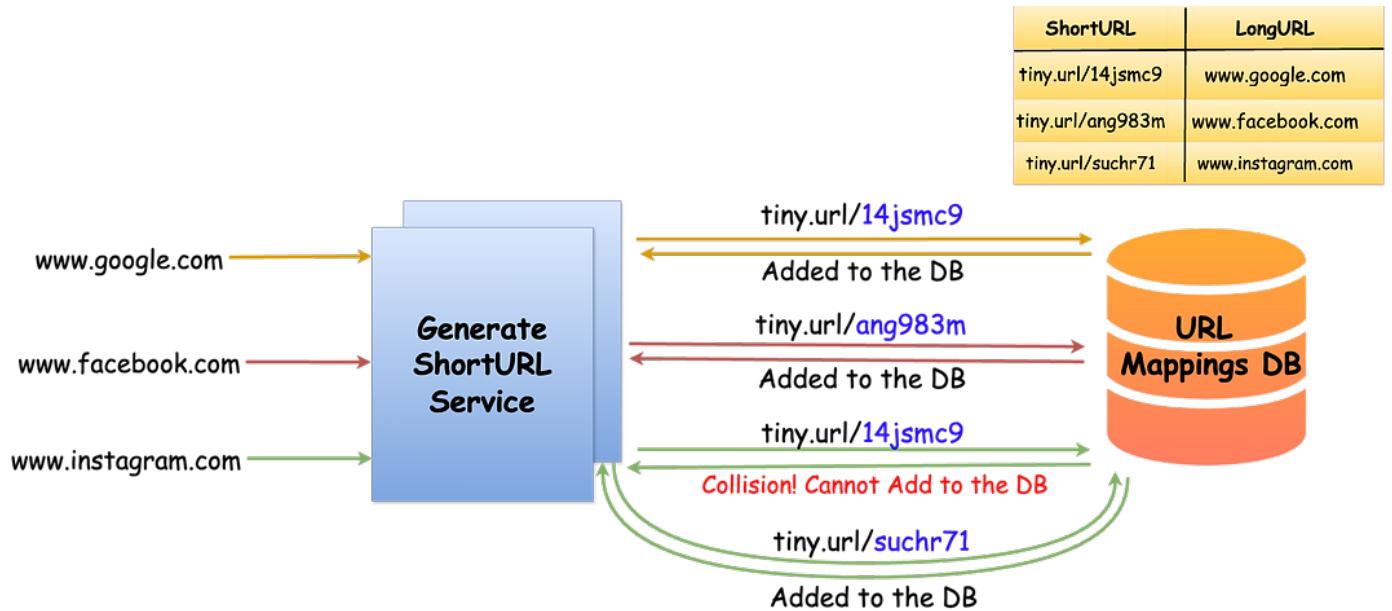




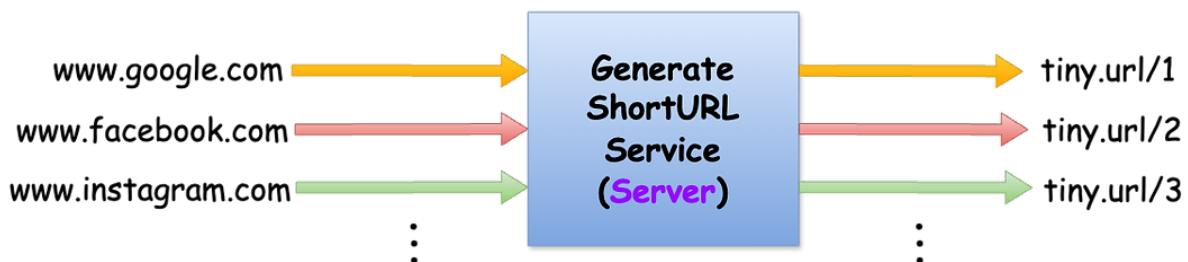
High Level Design- Approach 2: Random String Generation using Long URL

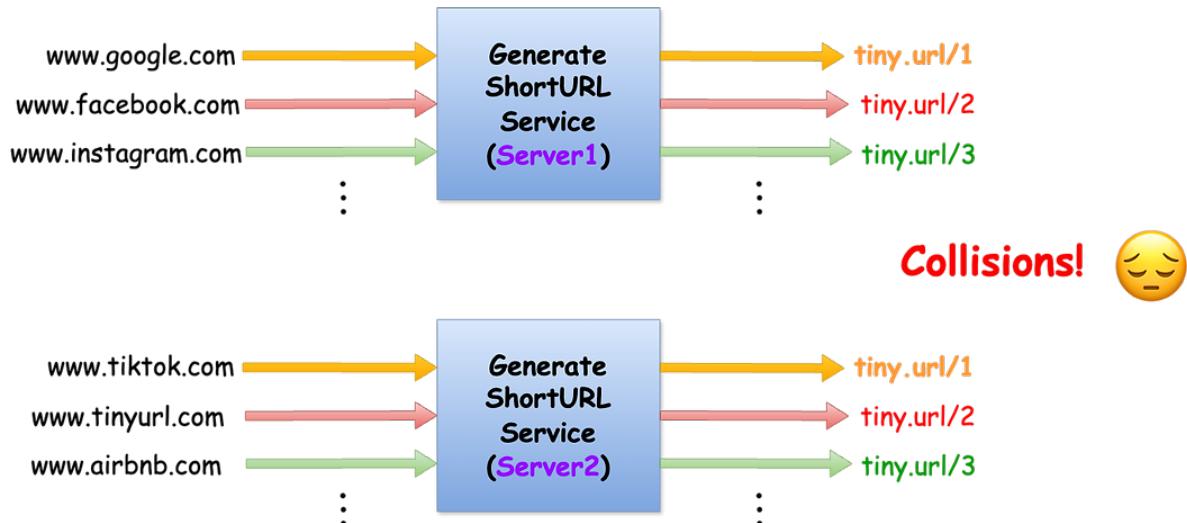


High Level Design- Approach 3: Check DB for Collisions

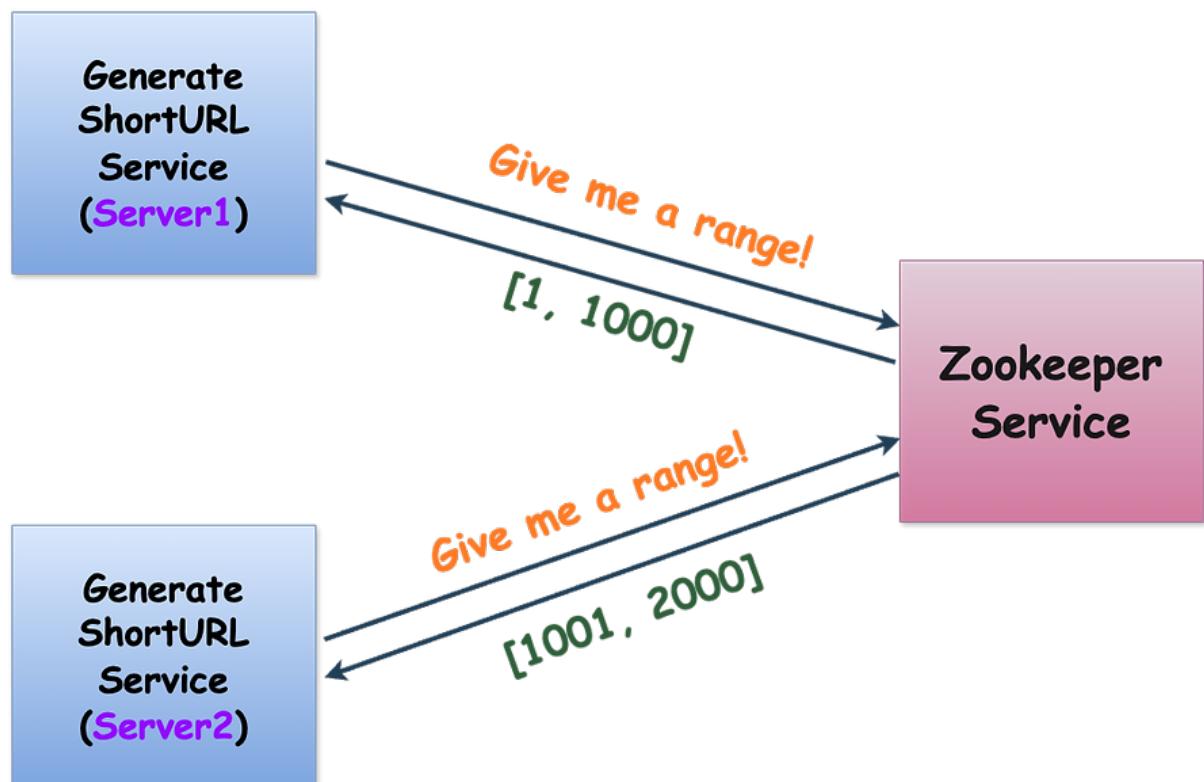


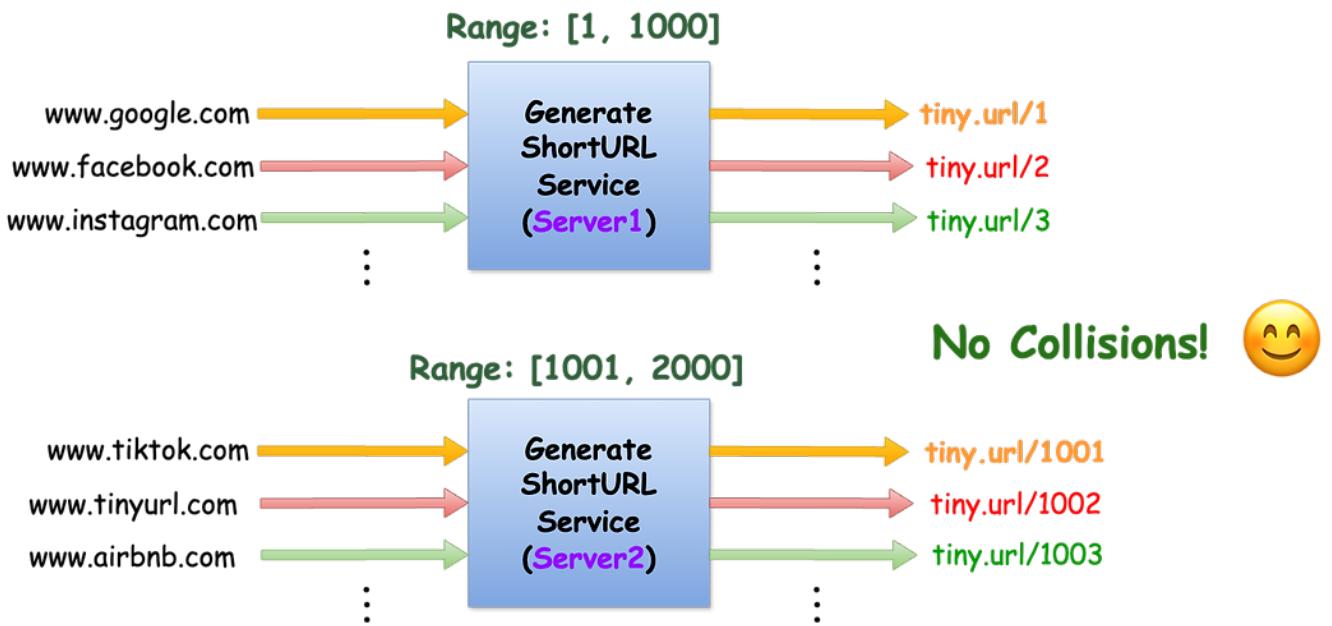
High Level Design- Approach 4: Let's Keep Counters





High Level Design- Approach 4: Zookeeper

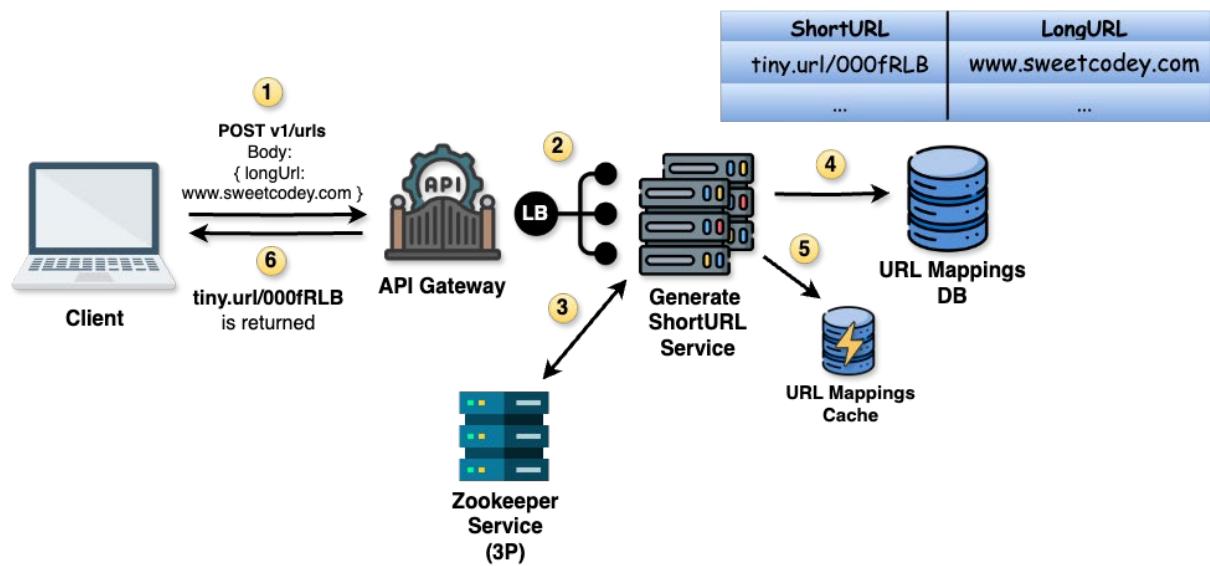




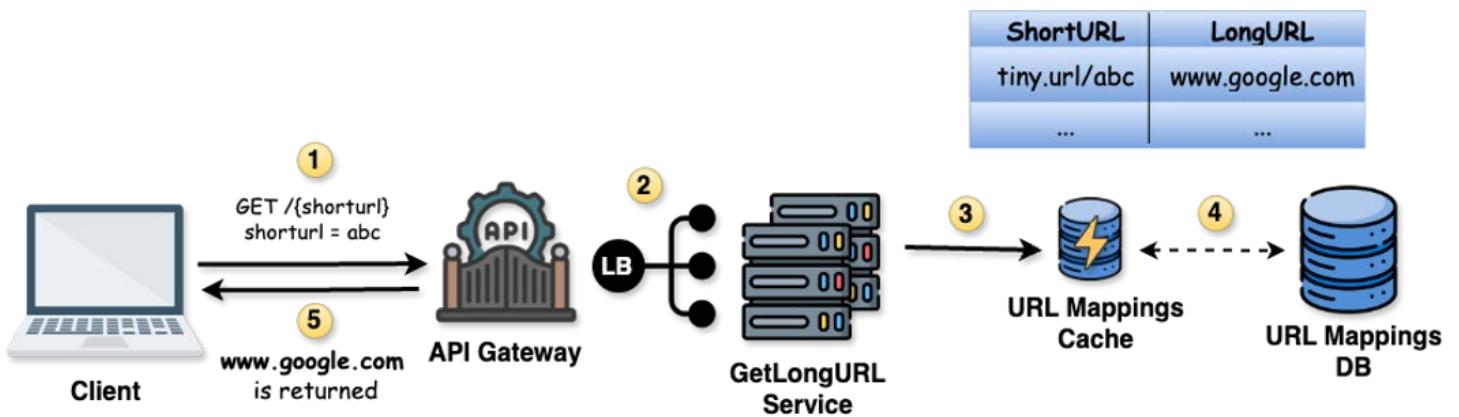
High Level Design- Base-62 Table

Decimal	Binary	Base62									
0	000000	0	16	010000	G	32	100000	W	48	110000	m
1	000001	1	17	010001	H	33	100001	X	49	110001	n
2	000010	2	18	010010	I	34	100010	Y	50	110010	o
3	000011	3	19	010011	J	35	100011	Z	51	110011	p
4	000100	4	20	010100	K	36	100100	a	52	110100	q
5	000101	5	21	010101	L	37	100101	b	53	110101	r
6	000110	6	22	010110	M	38	100110	c	54	110110	s
7	000111	7	23	010111	N	39	100111	d	55	110111	t
8	001000	8	24	011000	O	40	101000	e	56	111000	u
9	001001	9	25	011001	P	41	101001	f	57	111001	v
10	001010	A	26	011010	Q	42	101010	g	58	111010	w
11	001011	B	27	011011	R	43	101011	h	59	111011	x
12	001100	C	28	011100	S	44	101100	i	60	111100	y
13	001101	D	29	011101	T	45	101101	j	61	111101	z
14	001110	E	30	011110	U	46	101110	k			
15	001111	F	31	011111	V	47	101111	l			

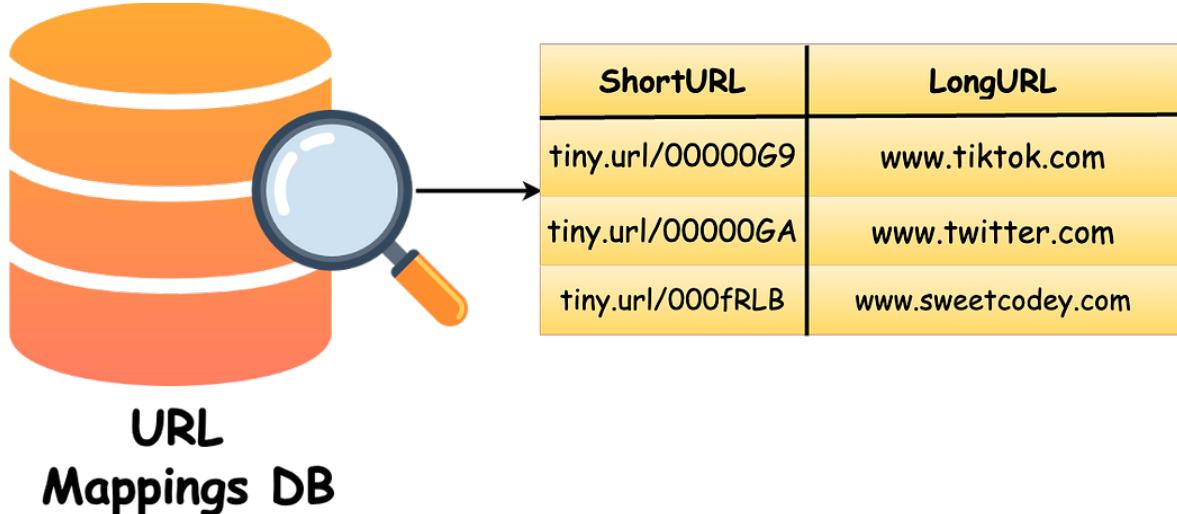
High Level Design-Approach:4-Final Design Diagram



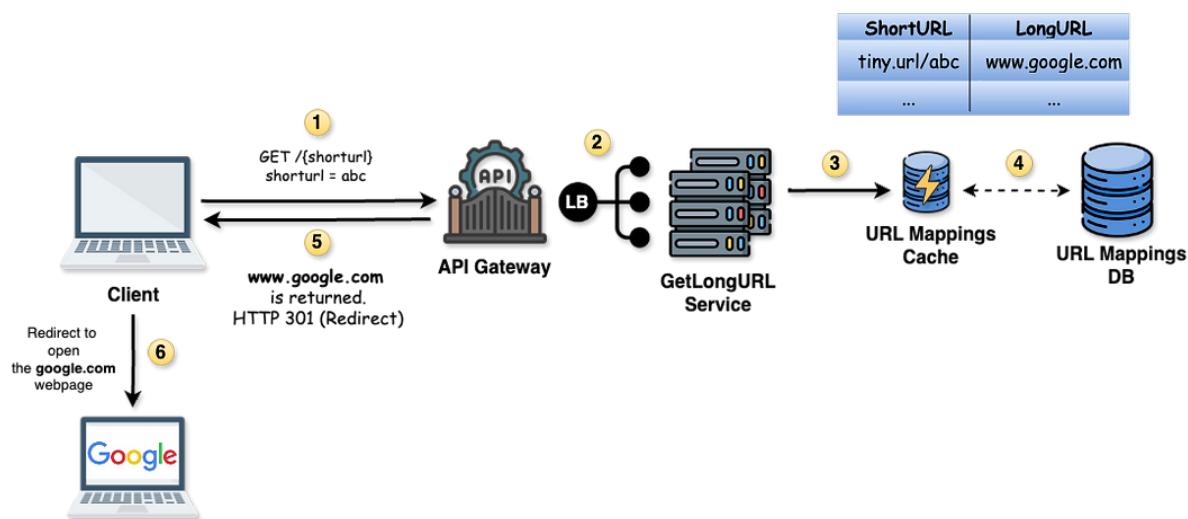
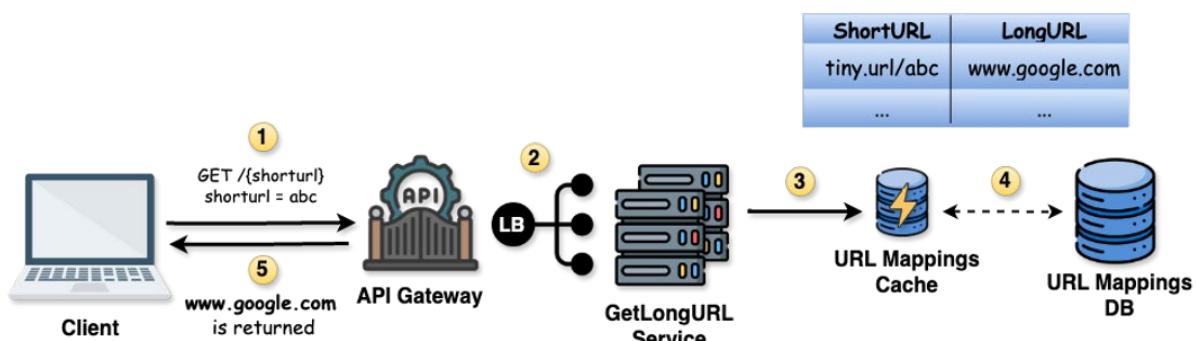
High Level Design- Get Long URL



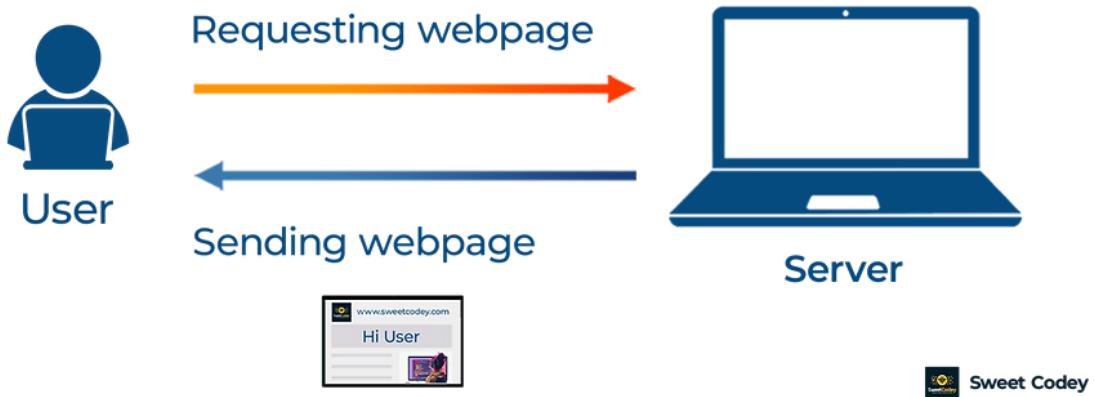
URL Mappings Schema



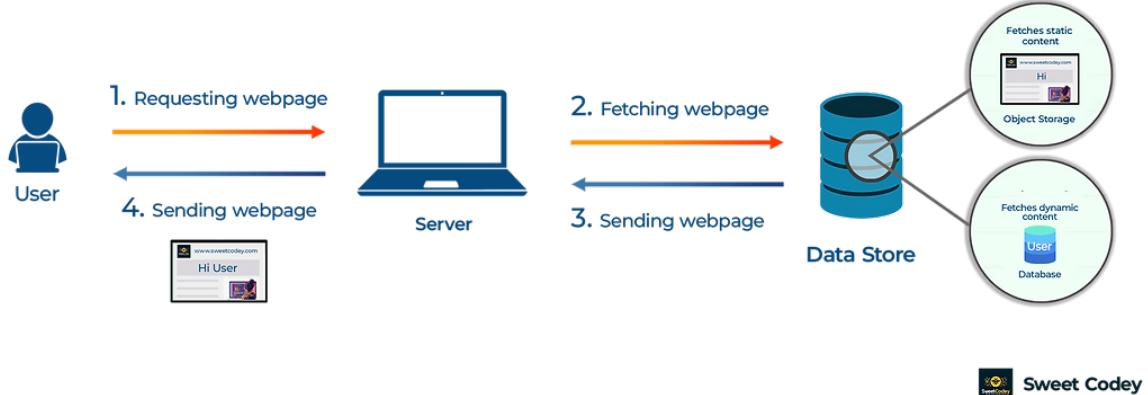
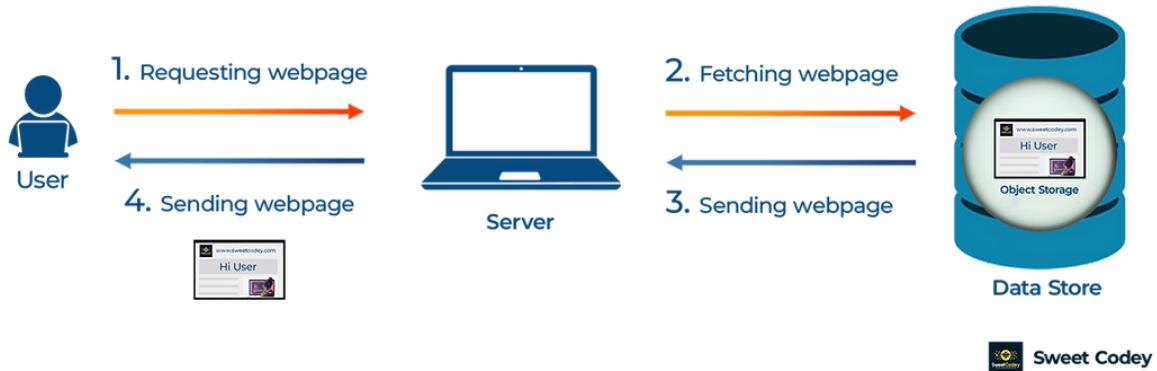
Redirection from Short URL



Step 1: Client & Server



Step 2: Database

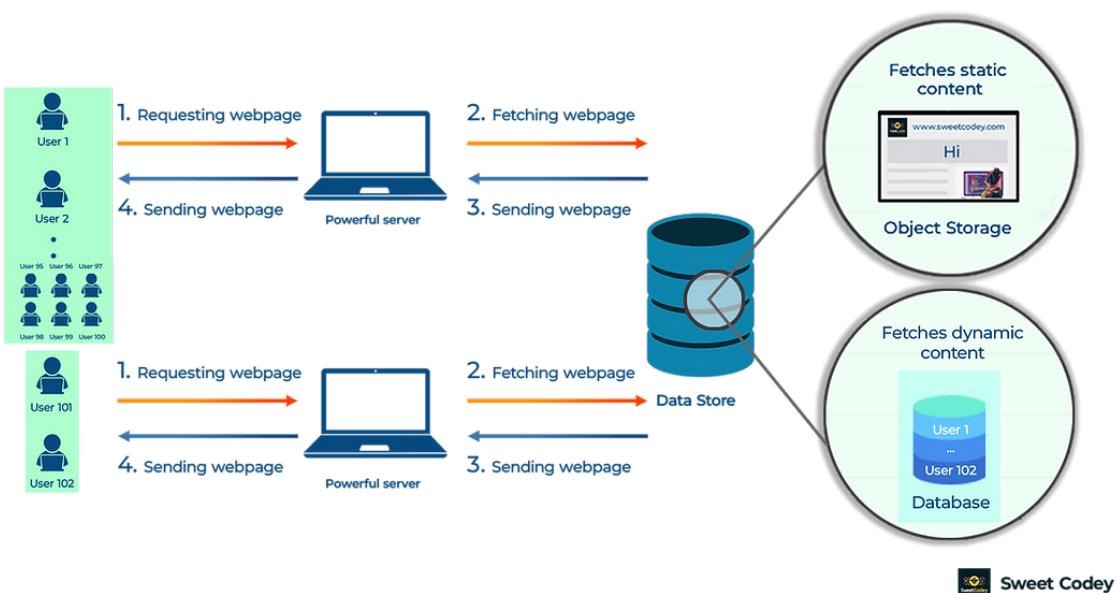


Step 3: Vertical vs Horizontal Scaling

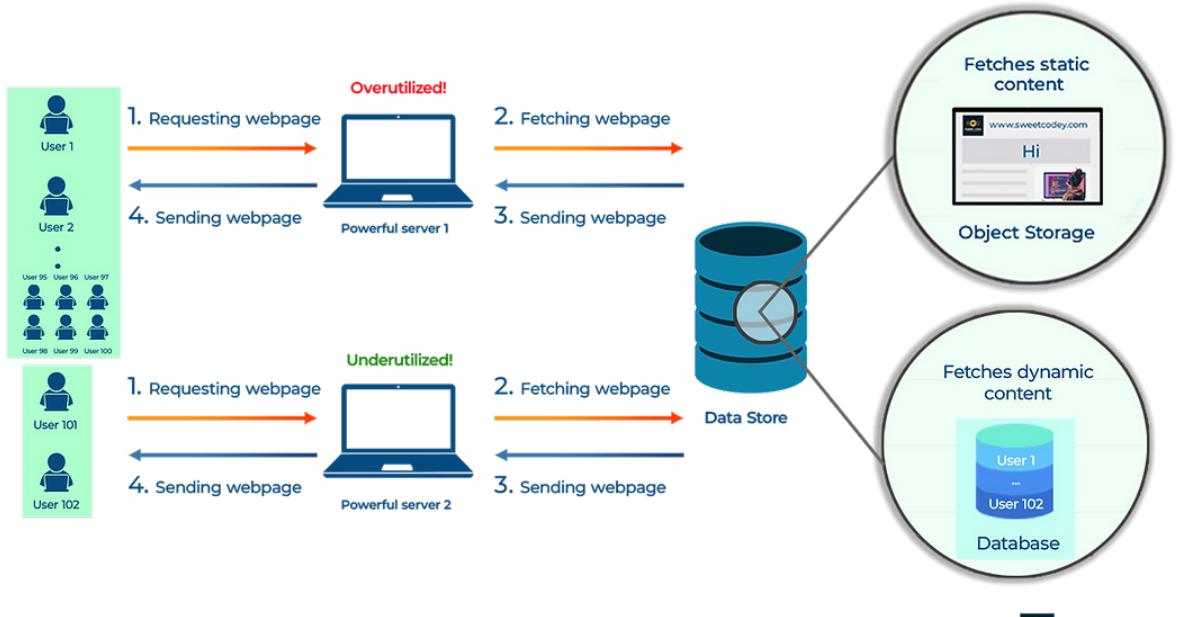
Vertical Scaling



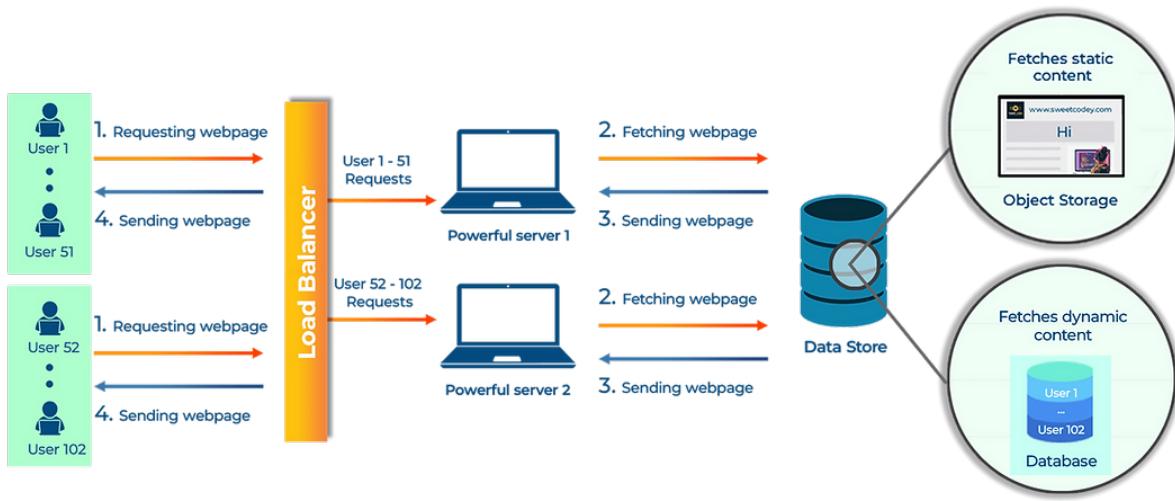
Horizontal Scaling



Step 4: Load Balancer

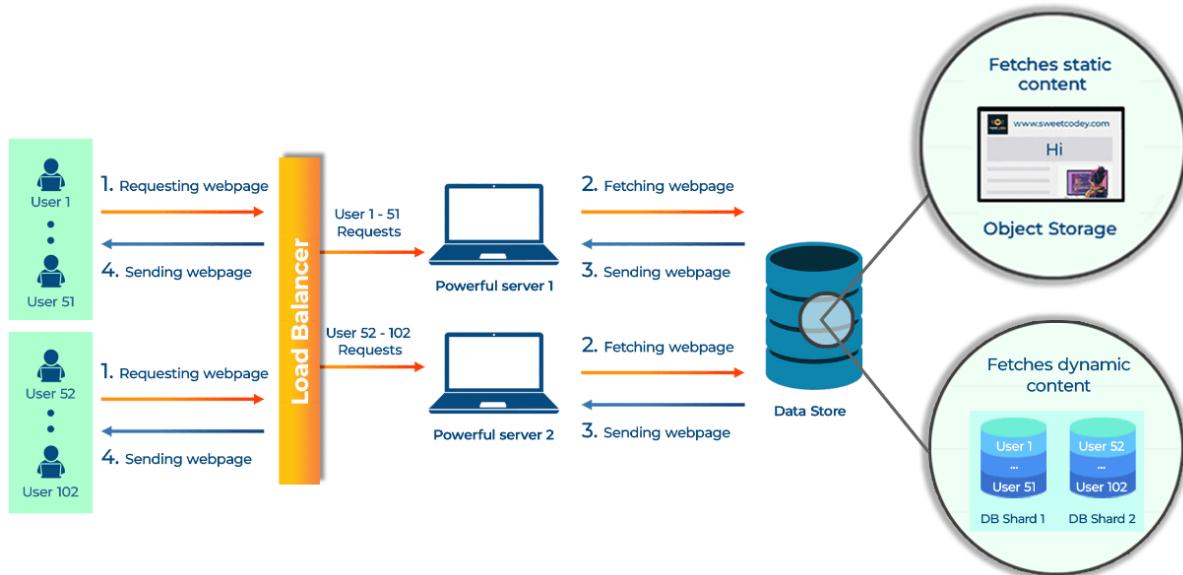


Sweet Codex

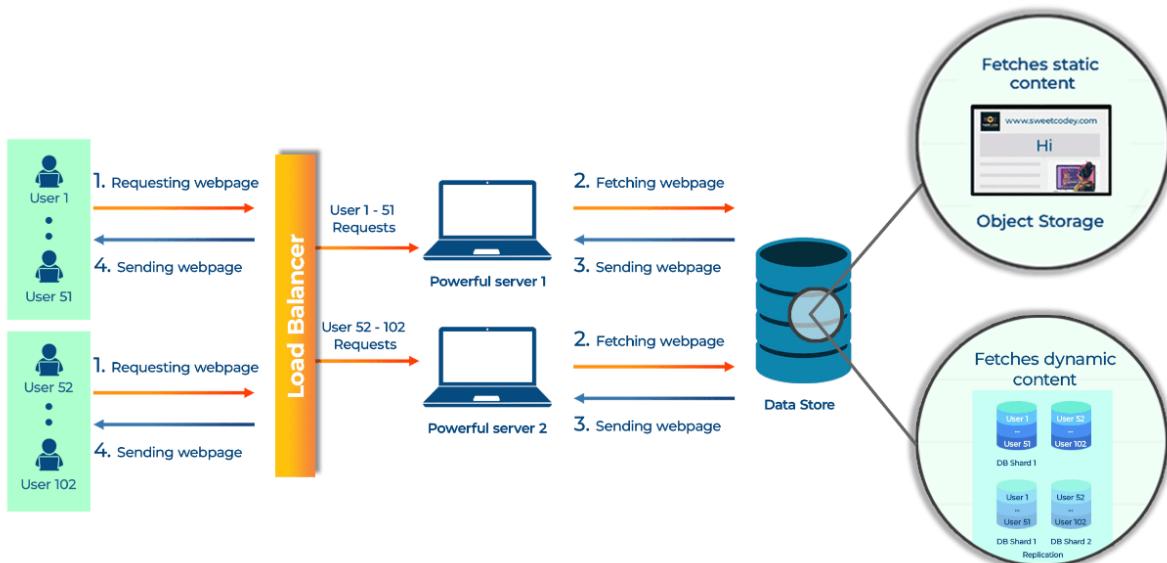


Sweet Codex

Step 5a: Database Sharding



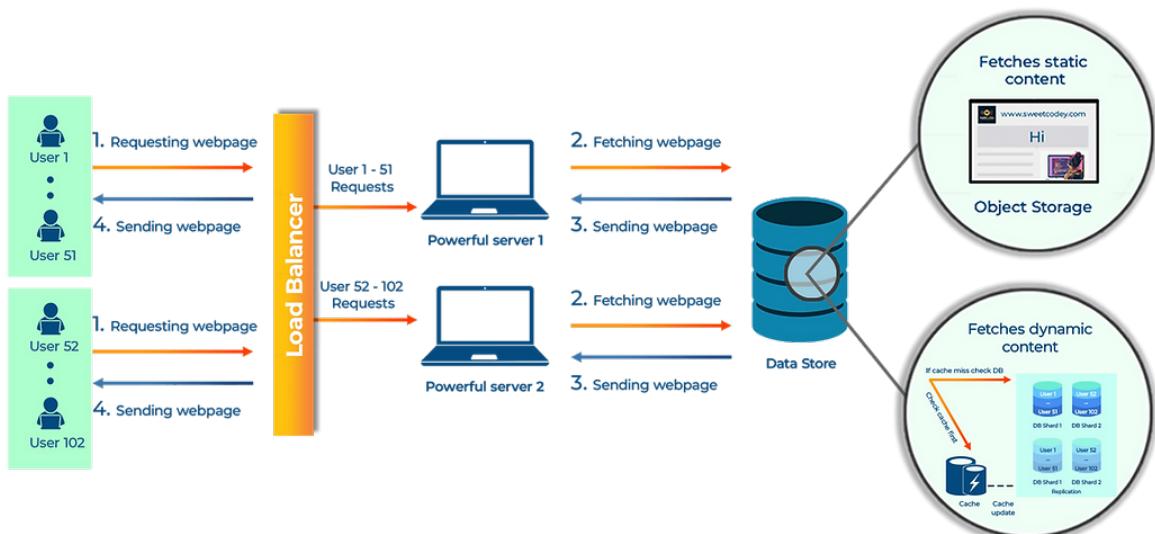
Step 5b: Database Replication



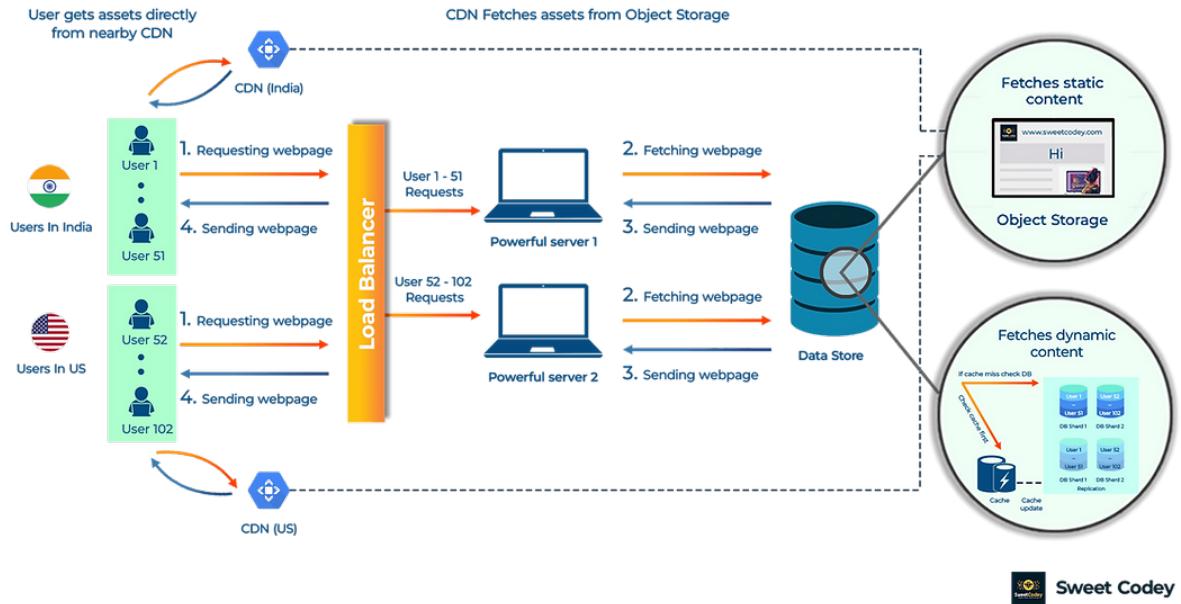
Step 6: Cache



Data Store

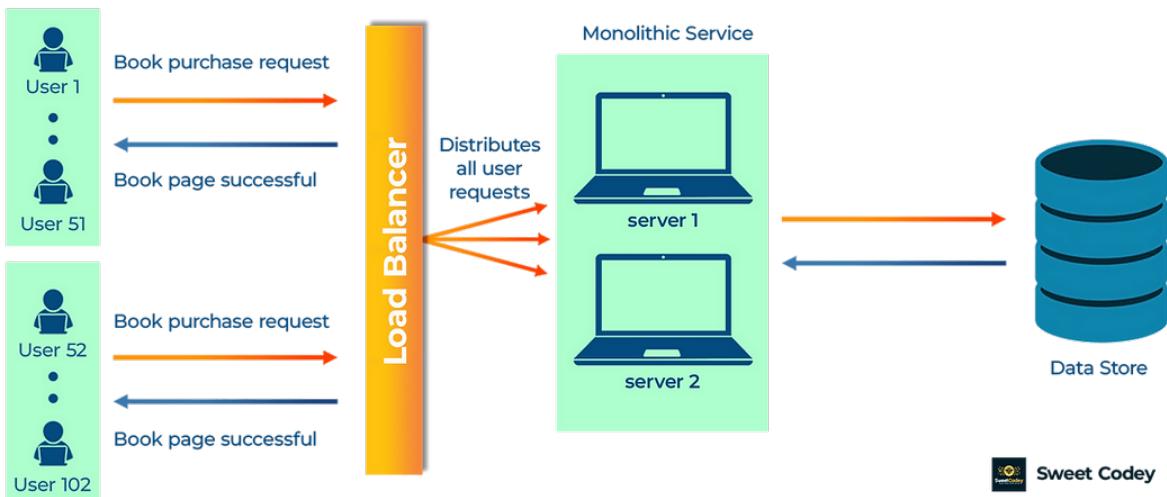
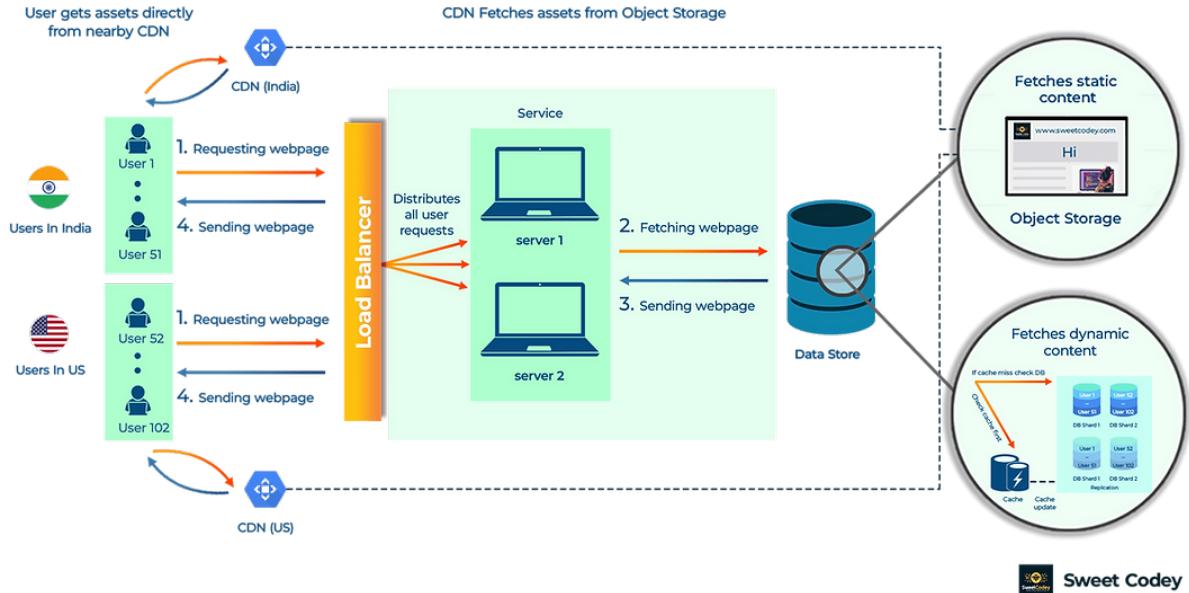


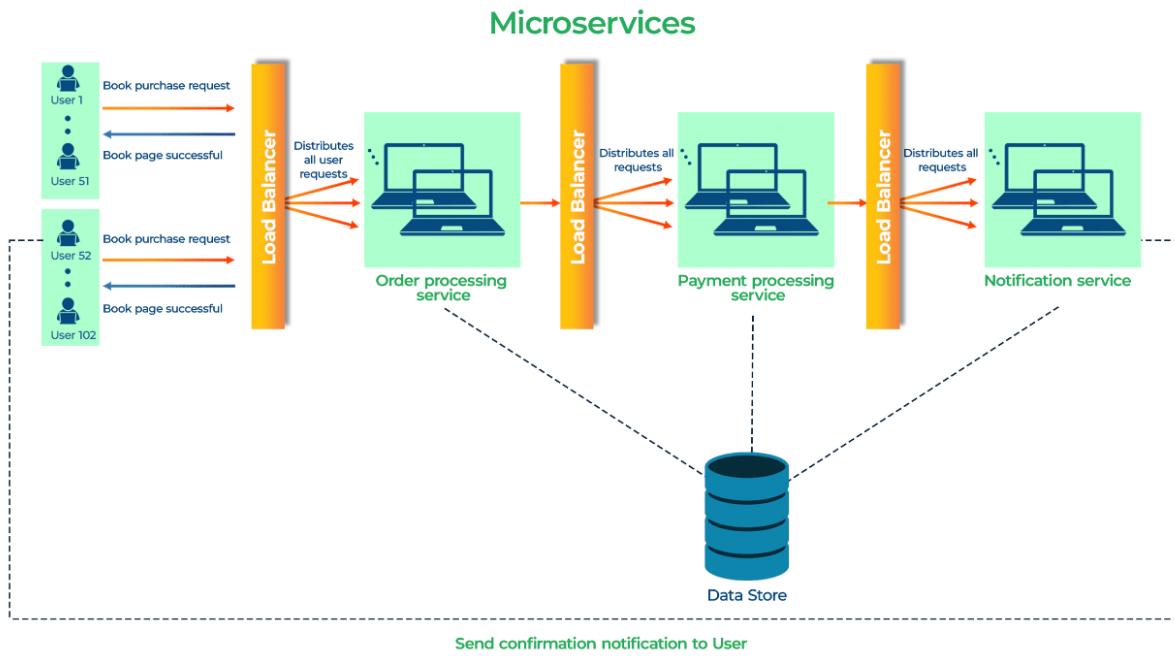
Step 7: Content Delivery Network (CDN)



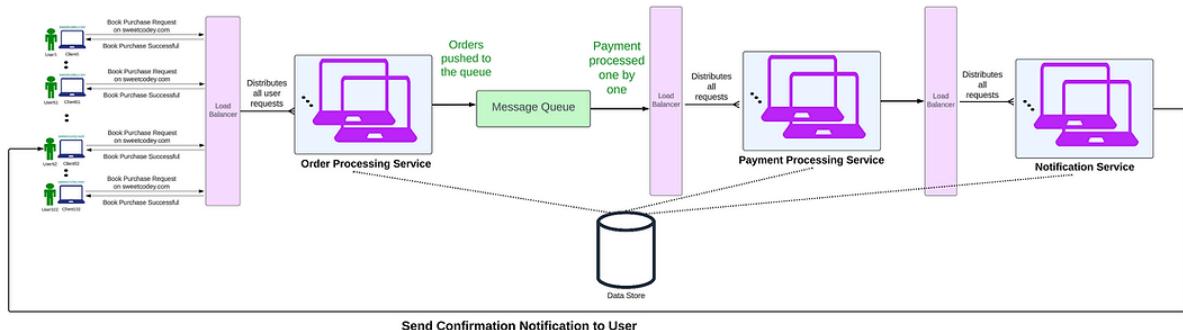
Step 8: Monolith & Microservice

Monolith

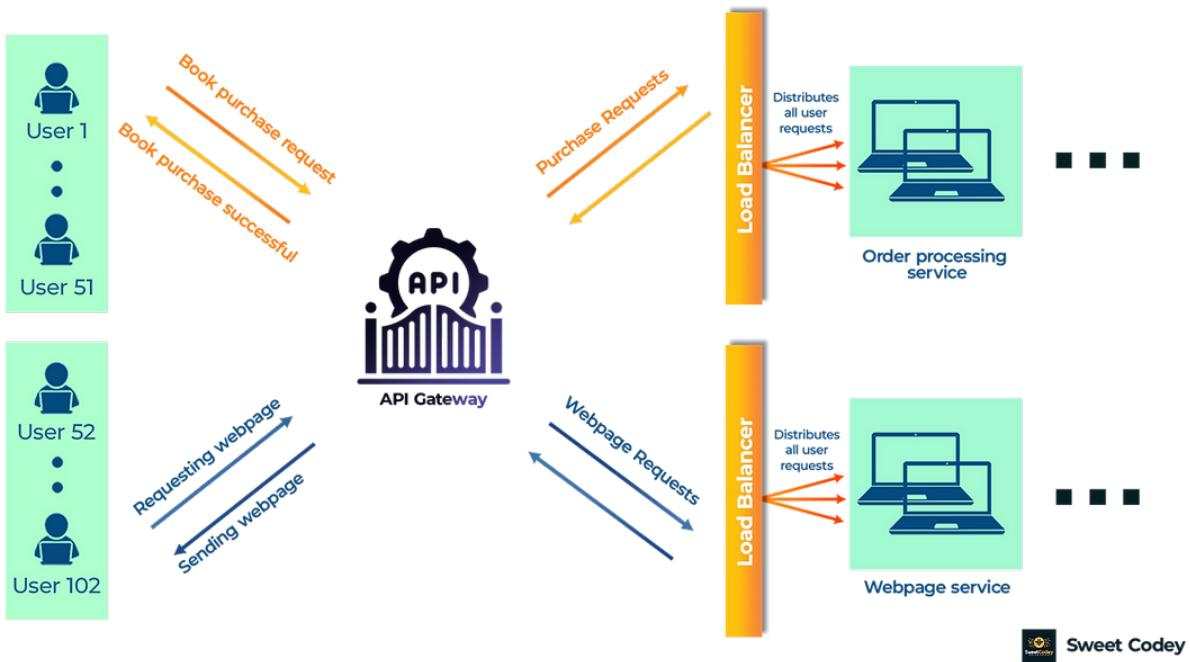




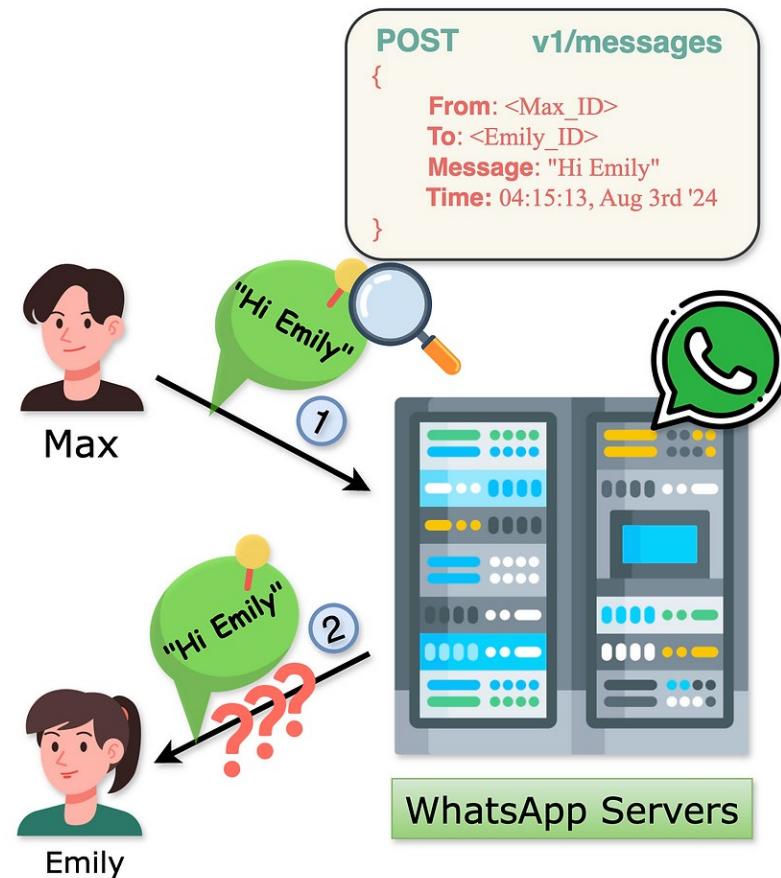
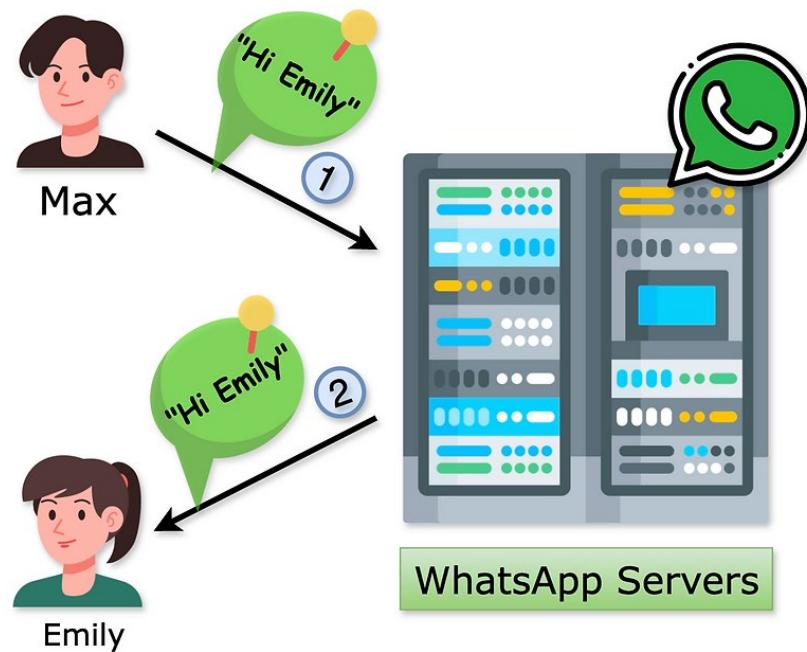
Step 9: Message Queue

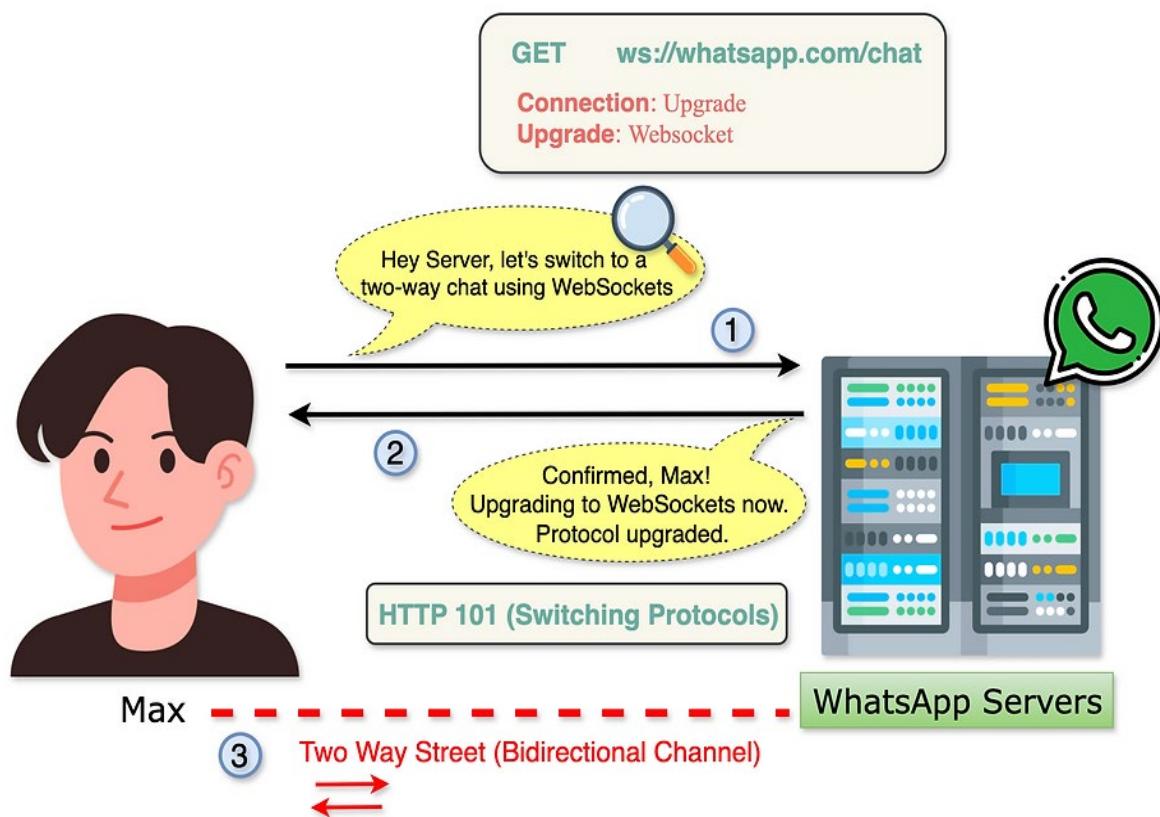


Step 10: API Gateway

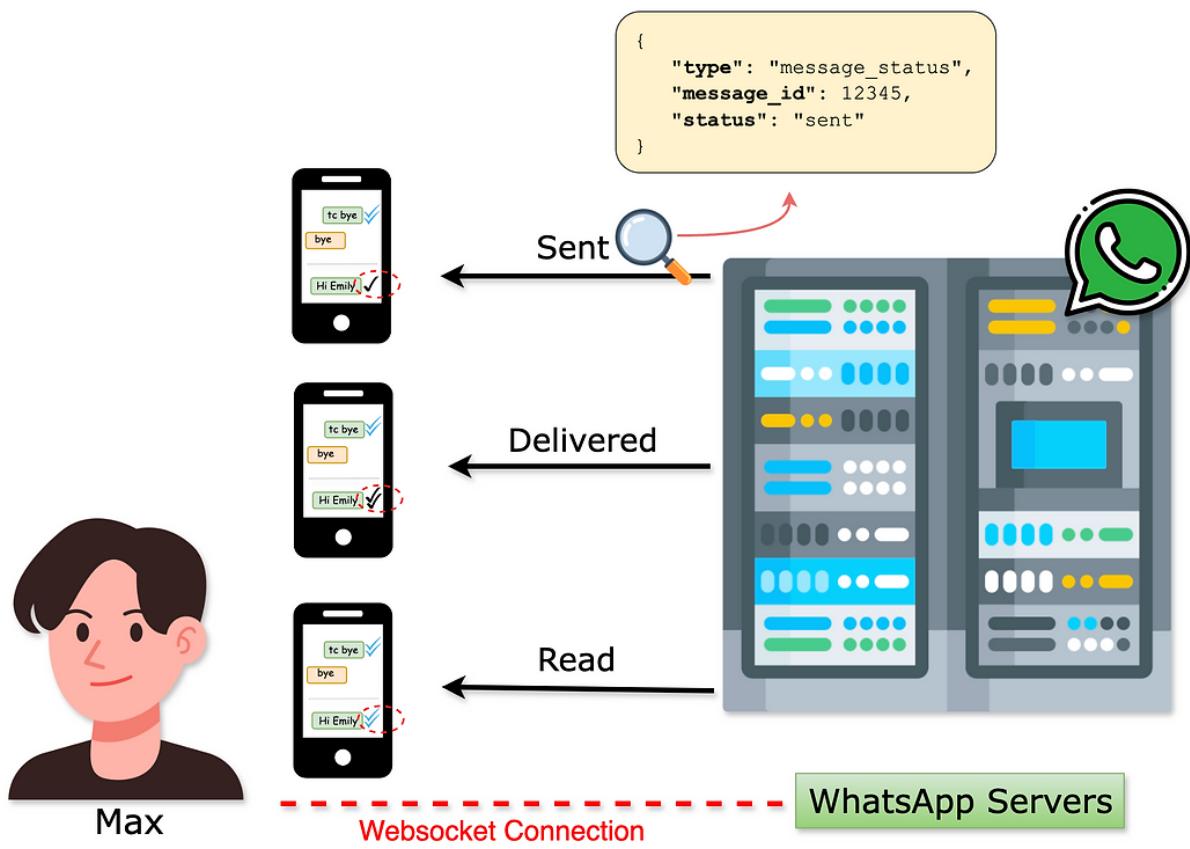


API Design: One-One Messaging

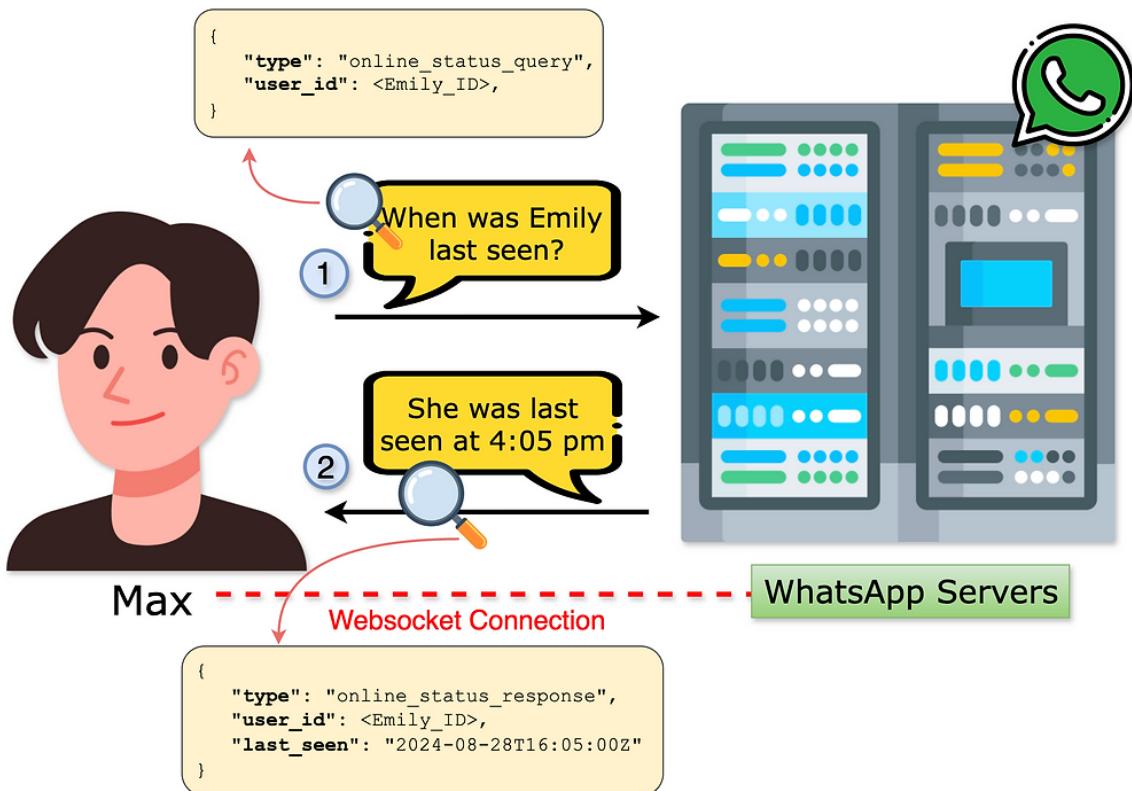
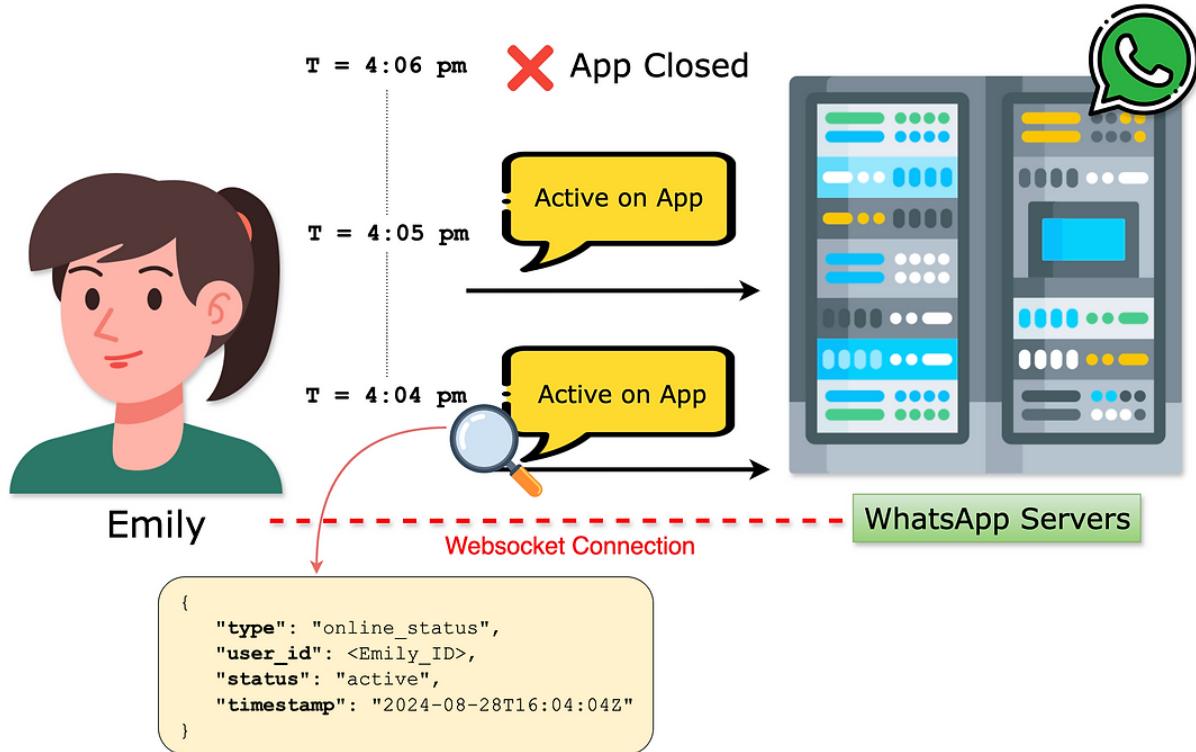




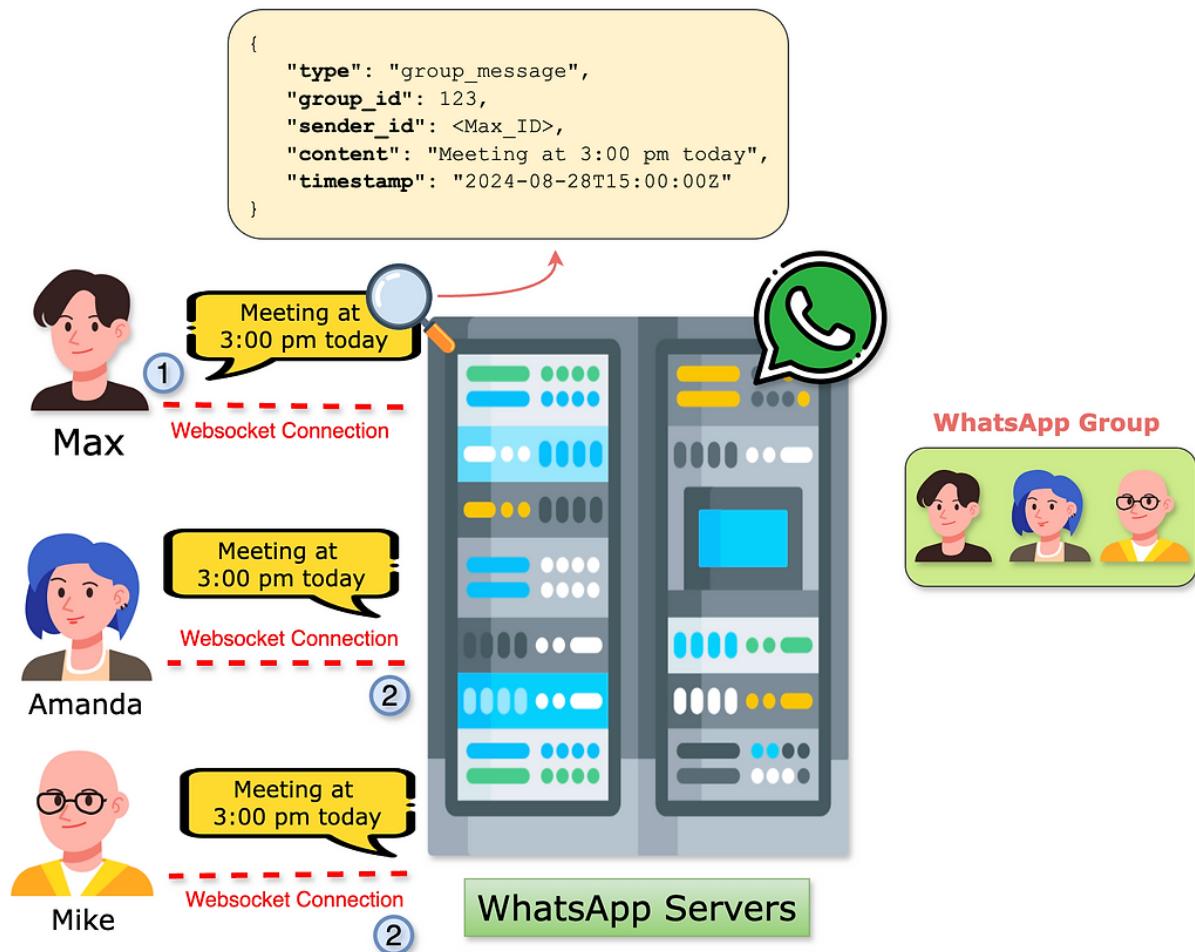
API Design: Message Status



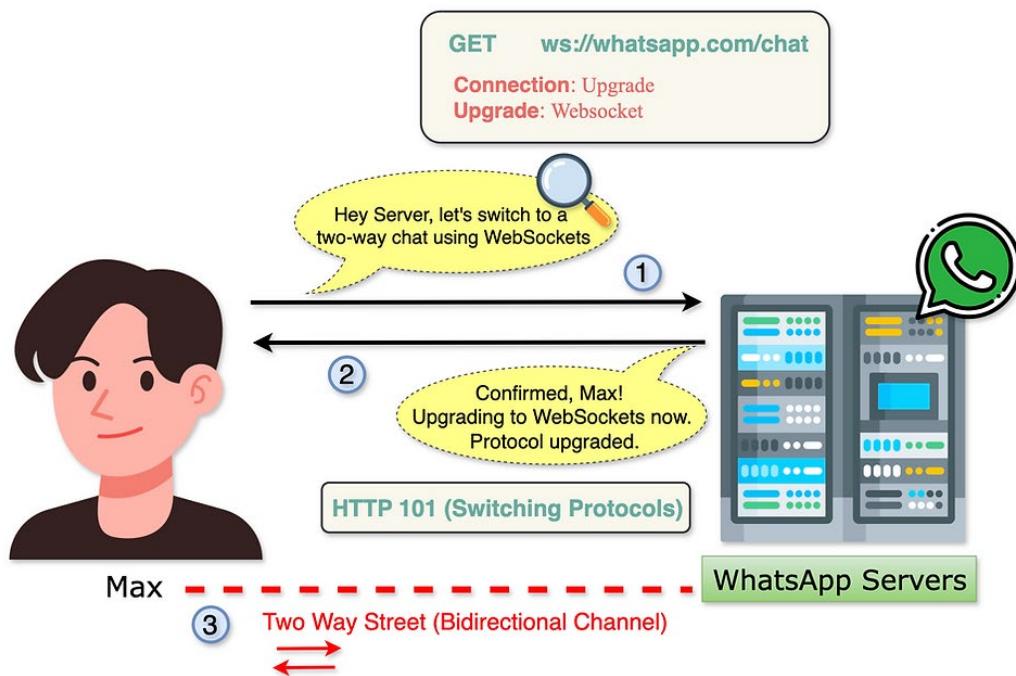
API Design: Online Status

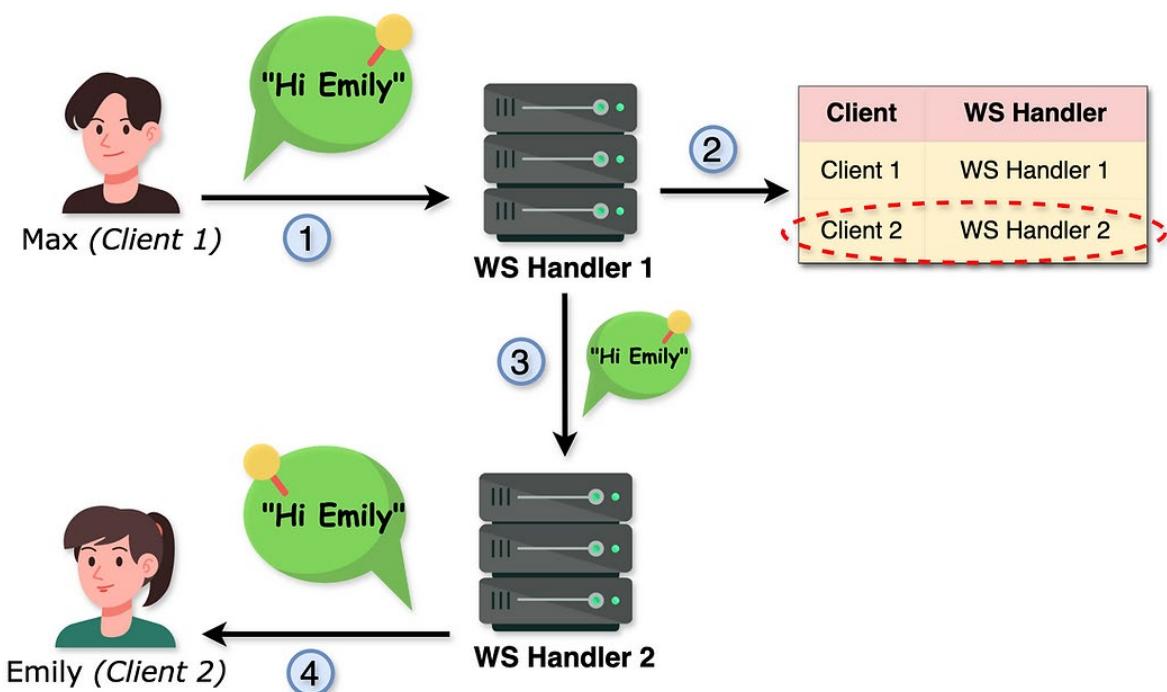
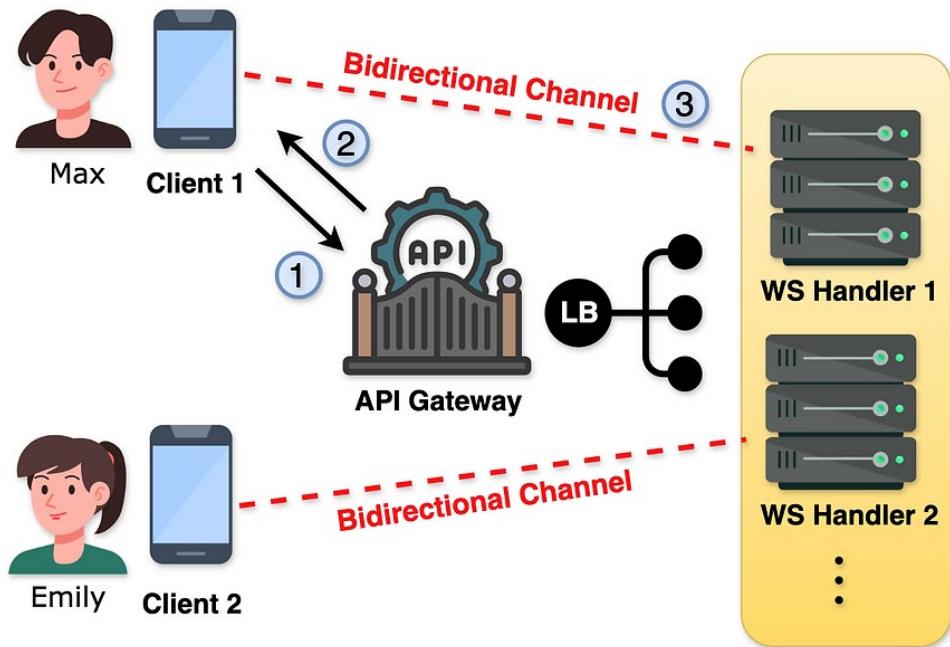


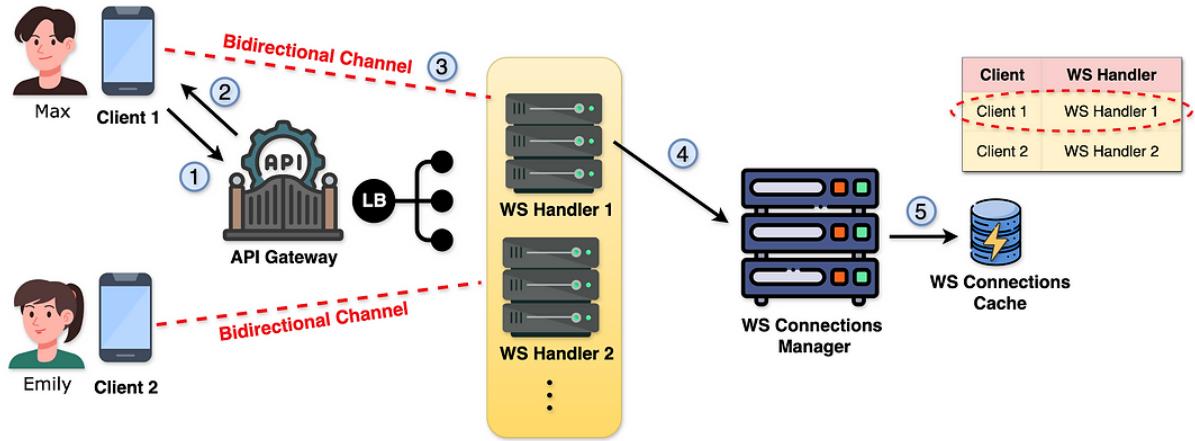
API Design: Group Messaging



Establishing WebSocket connection

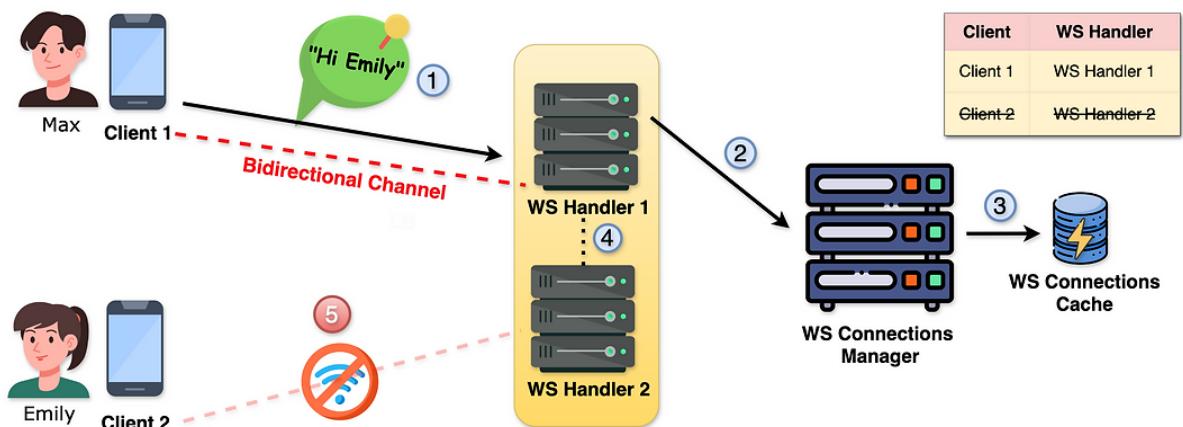
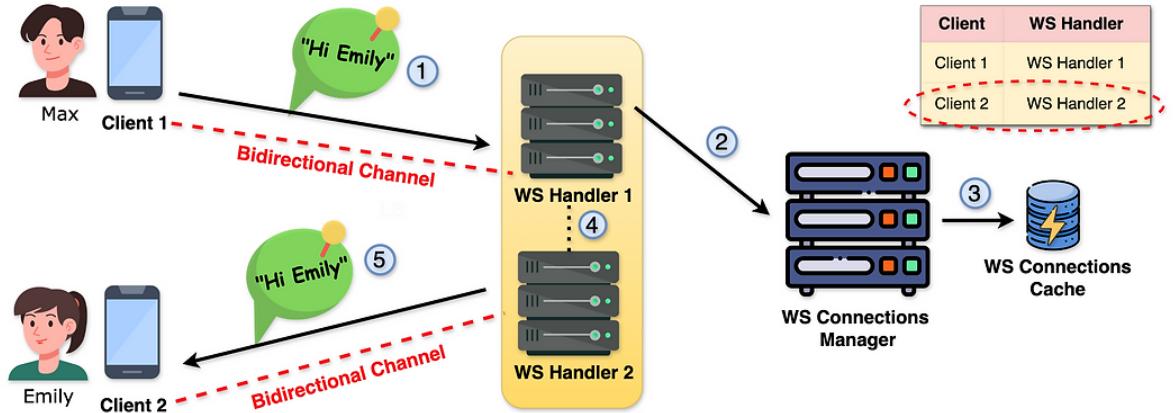




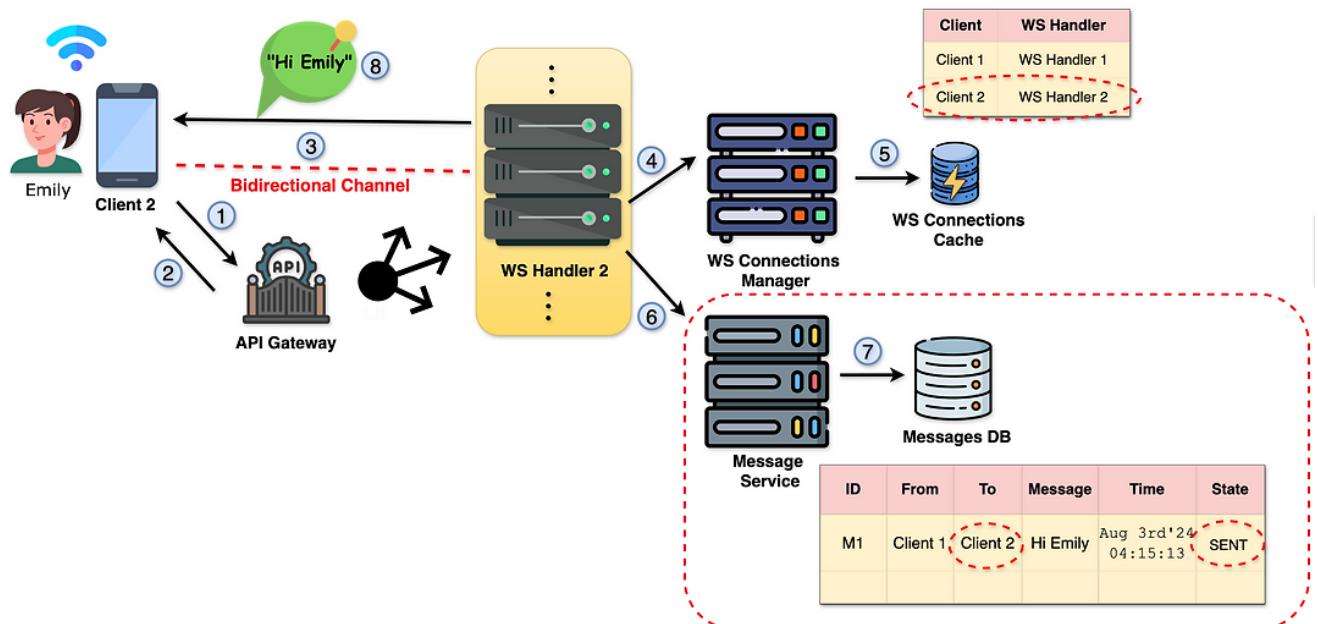
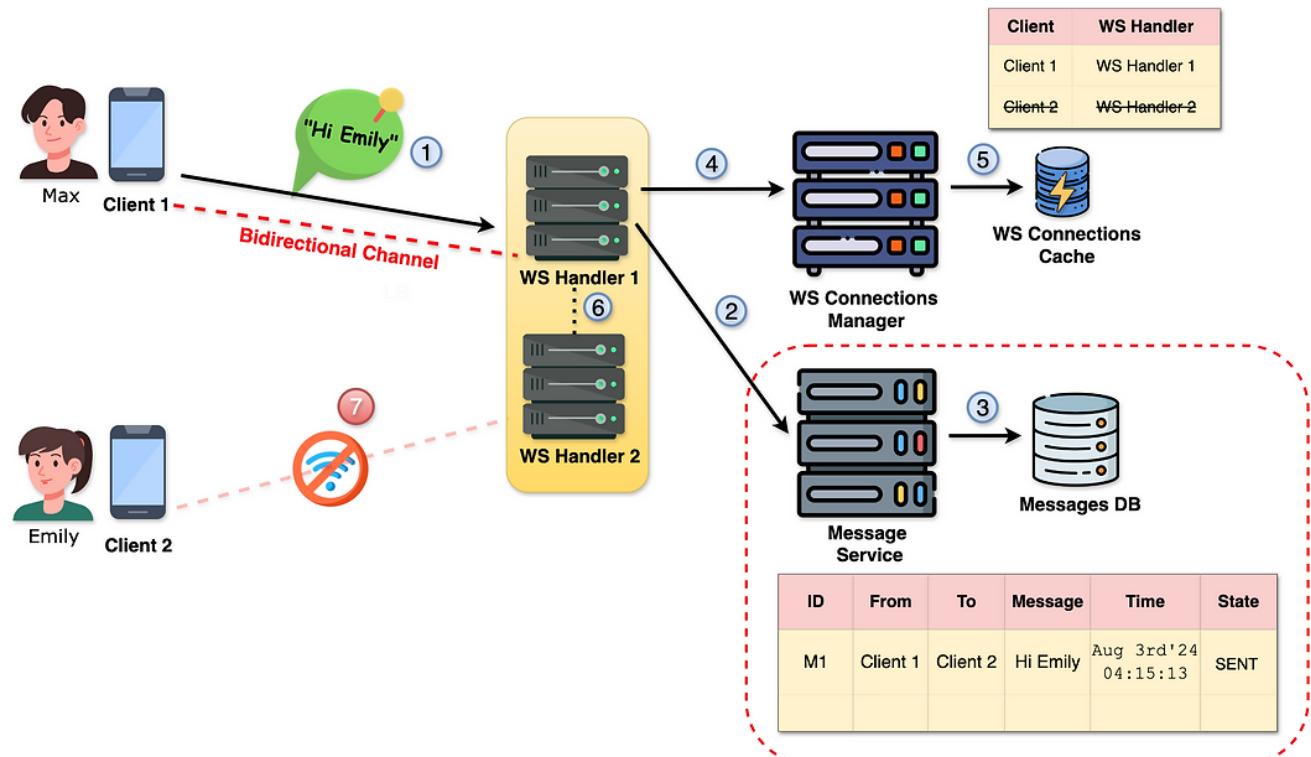


High Level Design: One-One Messaging

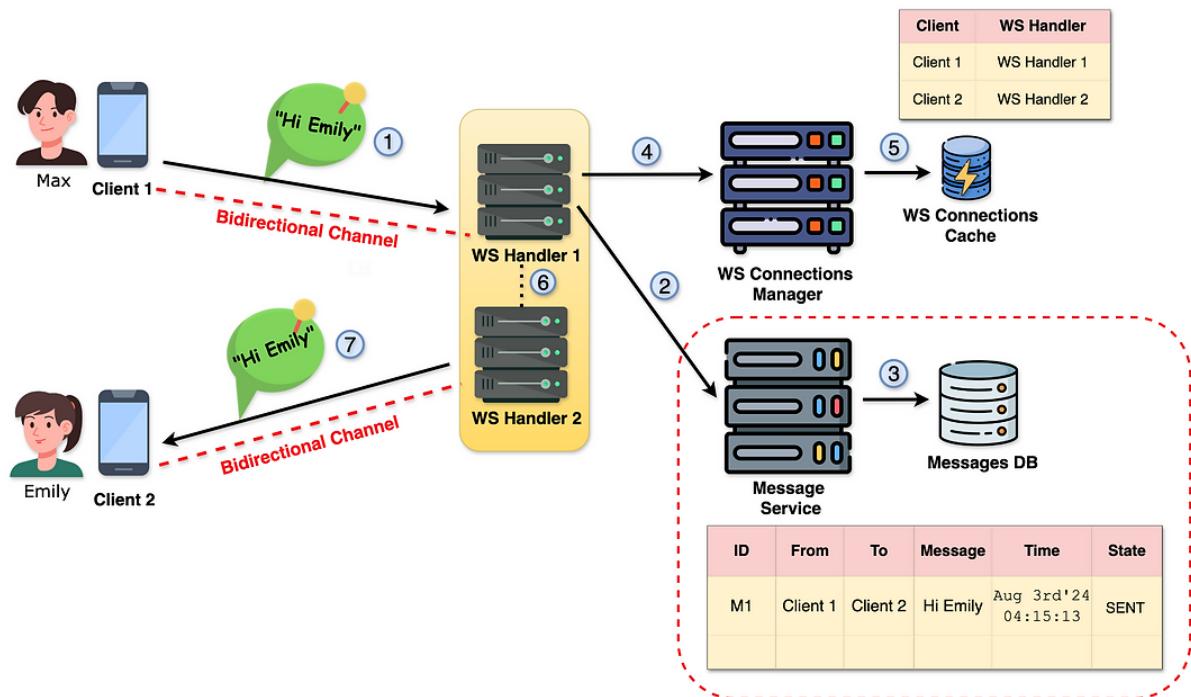
1:1 Messaging



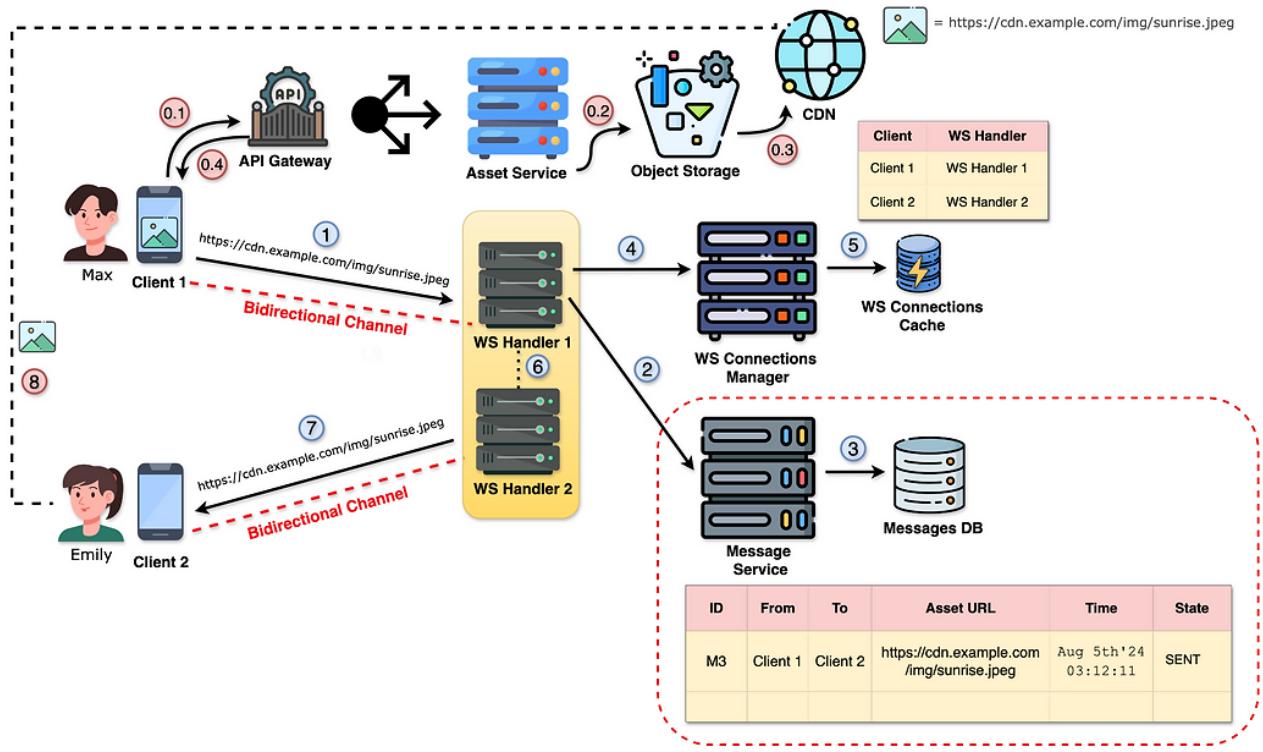
High Level Design: One-One Messaging (Offline Client)



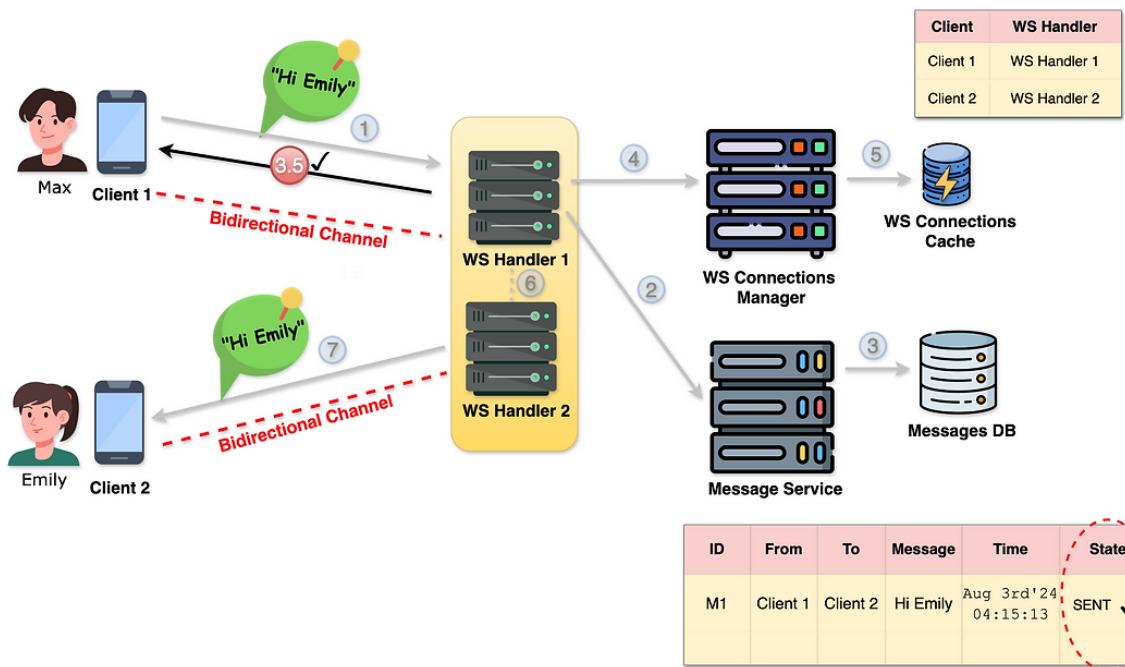
High Level Design: One-One Messaging Final Design



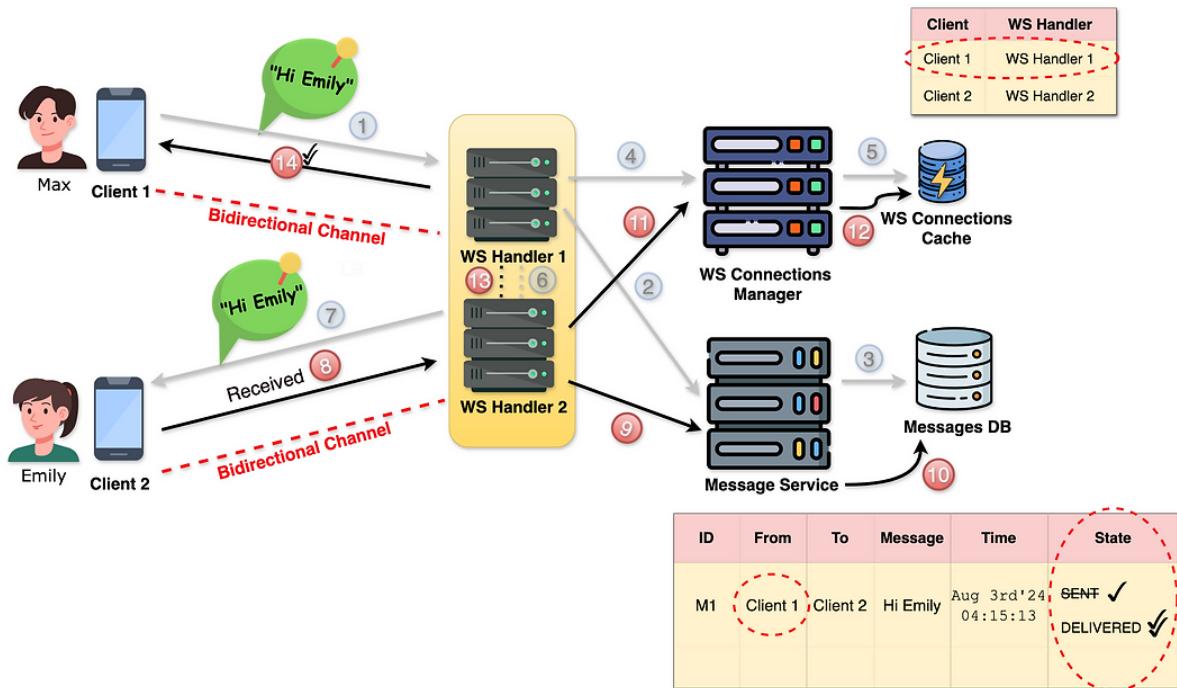
High Level Design: One-One Messaging (Image, Video, Document)



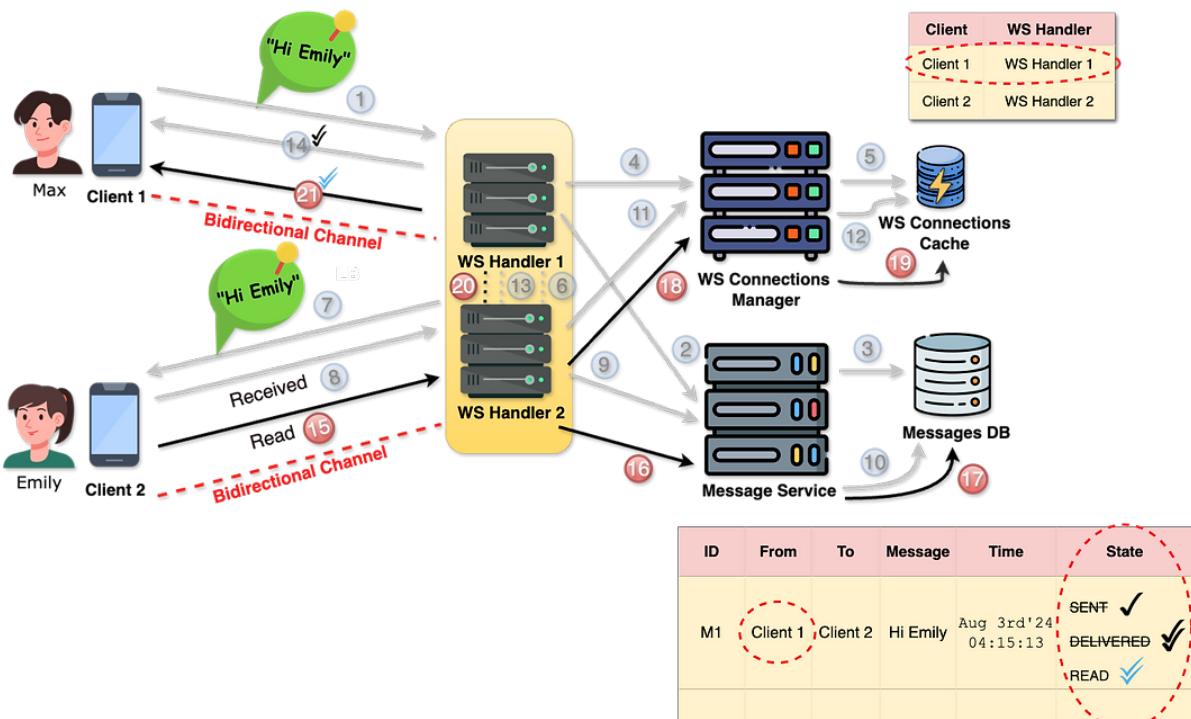
High Level Design: Message Status (Sent)



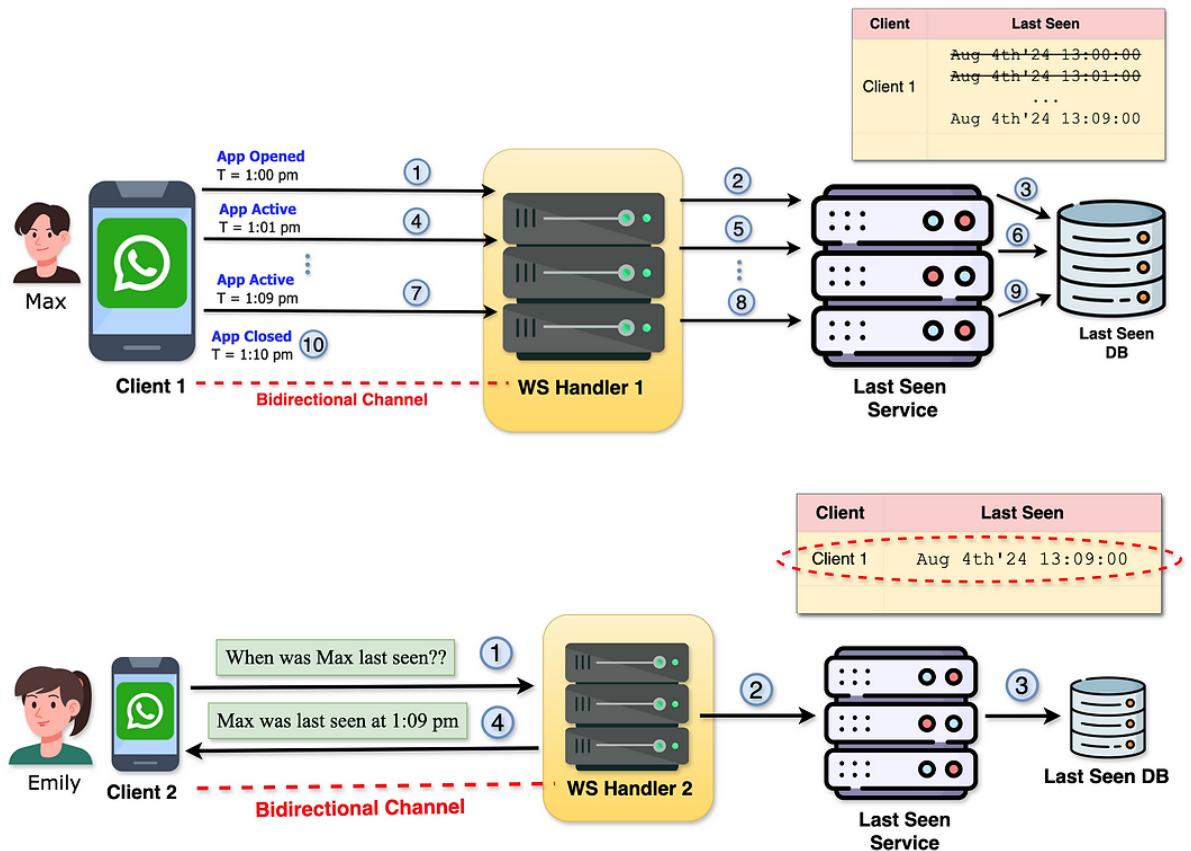
High Level Design: Message Status (Delivered)



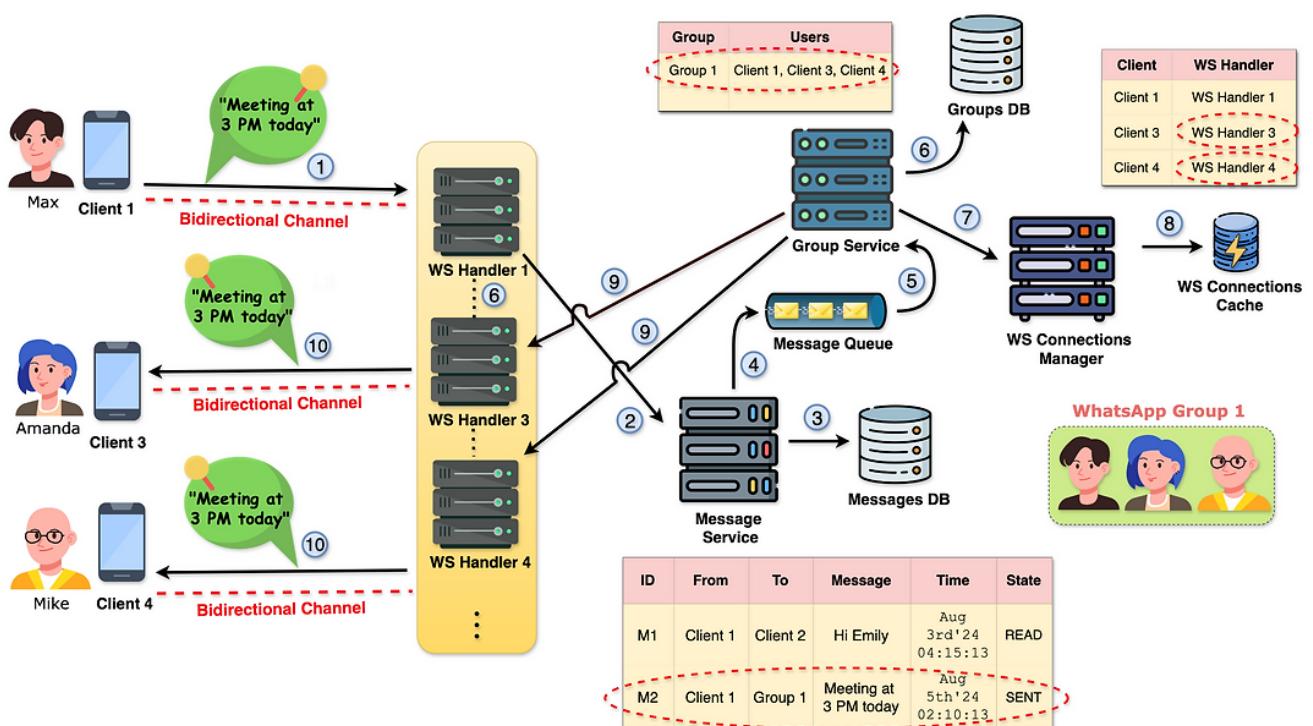
High Level Design: Message Status (Read)



High Level Design: Online Status



High Level Design: Group Messaging



Messages DB schema

{
}

MessageId: A unique ID for each message.
SenderId ("From"): ID of the user sending the message.
ReceiverId ("To"): ID of the user or group getting the message.
MessageType: Type of the message, such as text, image, or video.
Message: The actual text of the message.
AssetURL: Link to any images, videos, or files attached.
State: Current status of the message, like 'Sent', 'Delivered', or 'Read'.
Time: When the message was sent.

}
}

Groups DB schema

{
}

GroupId: A unique identifier for each group.
GroupName: The name of the group.
GroupDescription: A short description or purpose of the group.
CreatorId: The user ID of the person who created the group.
Members: A list of user IDs that are members of the group.
CreatedTime: Timestamp when the group was created.
LastUpdatedTime: Timestamp for the last update to the group such as adding or removing members.

}
}

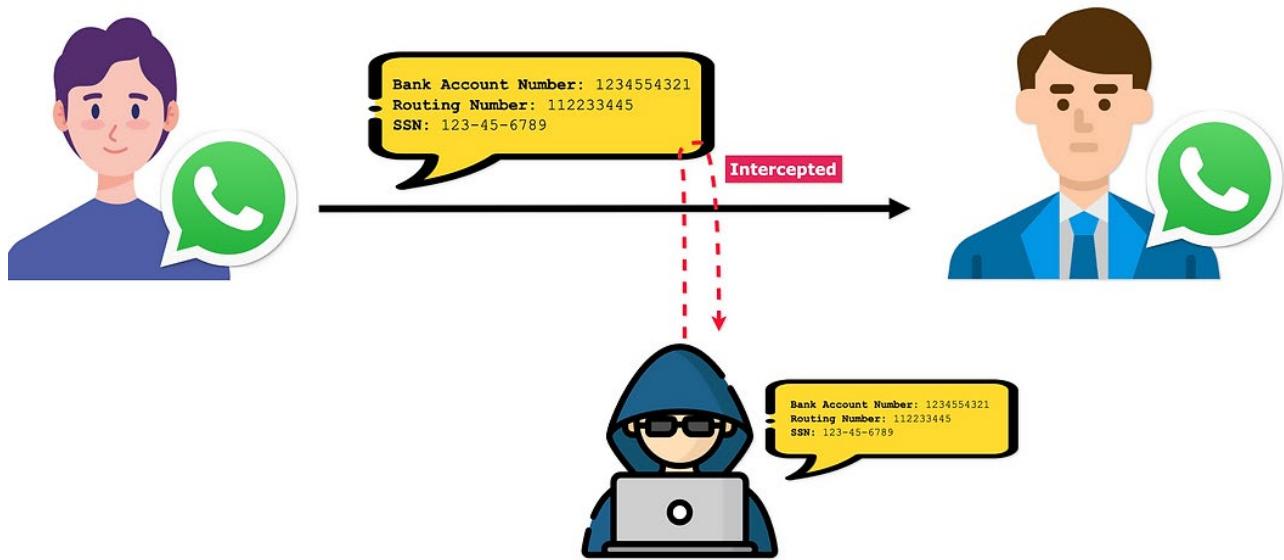
Last Seen DB schema

{
}

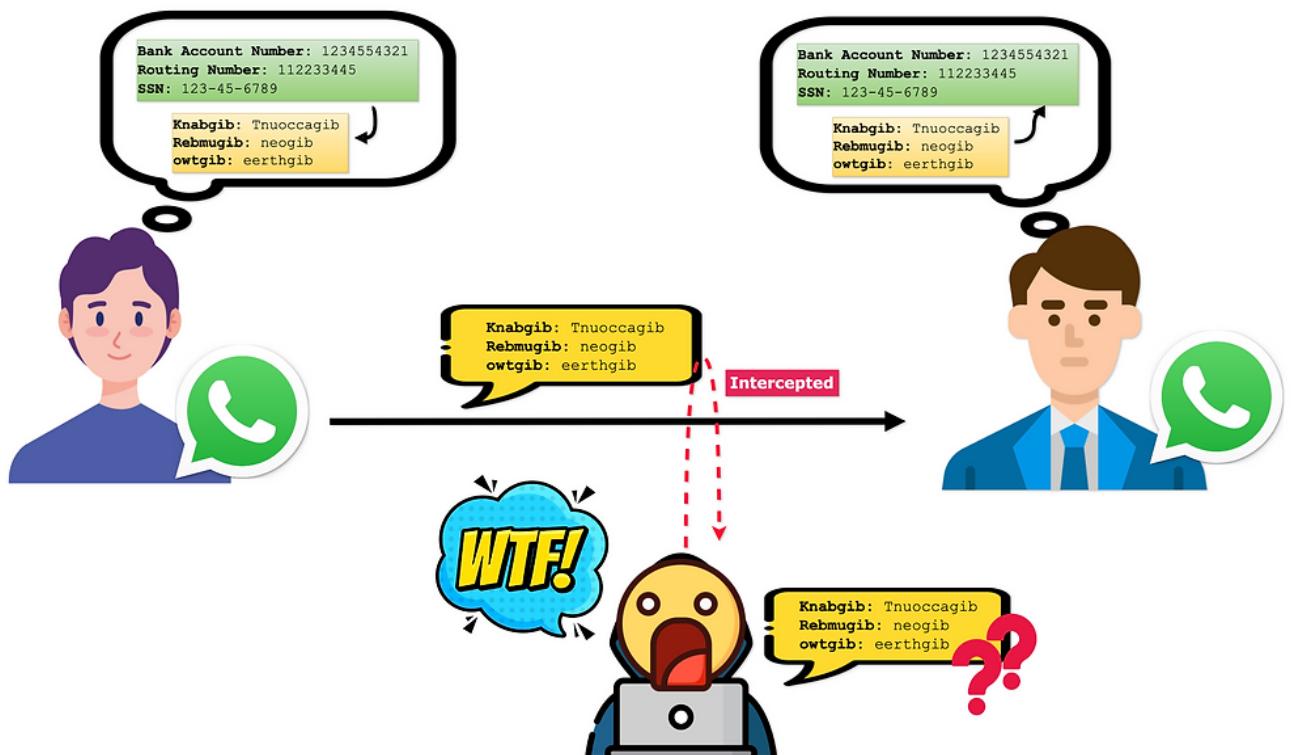
ClientId: A unique identifier for each user/client.
LastSeenTime: The timestamp when the user was last active on the app.

}
}

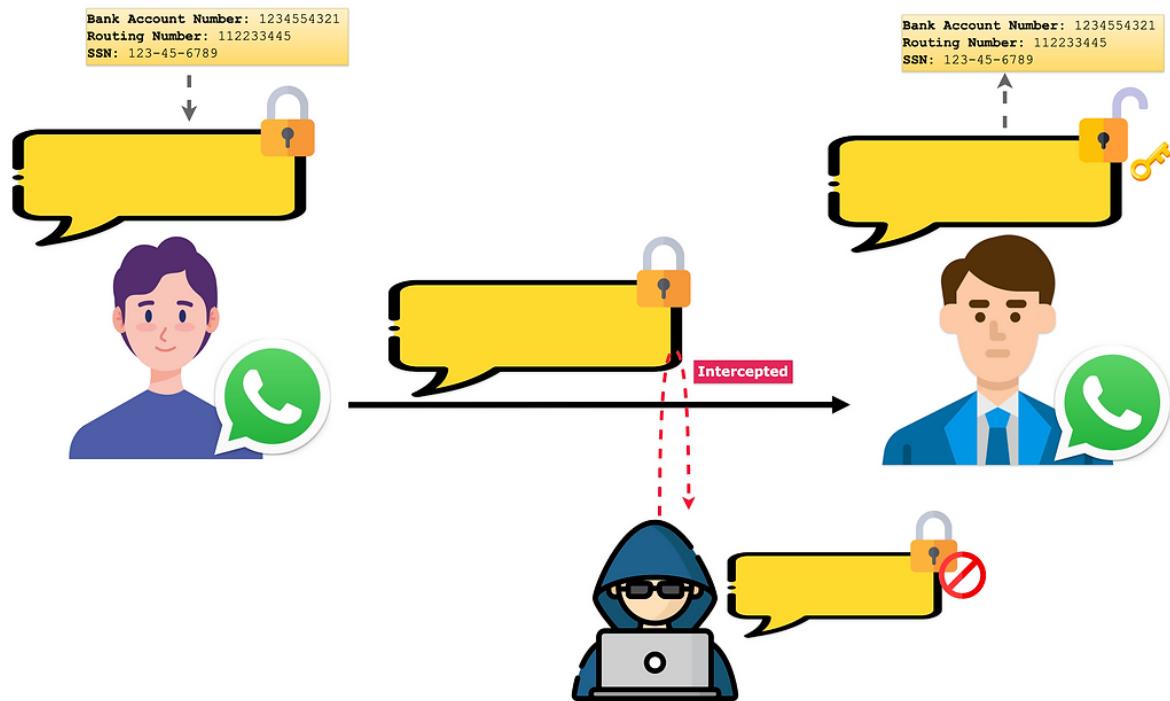
End-to-End Encryption (without Encryption)



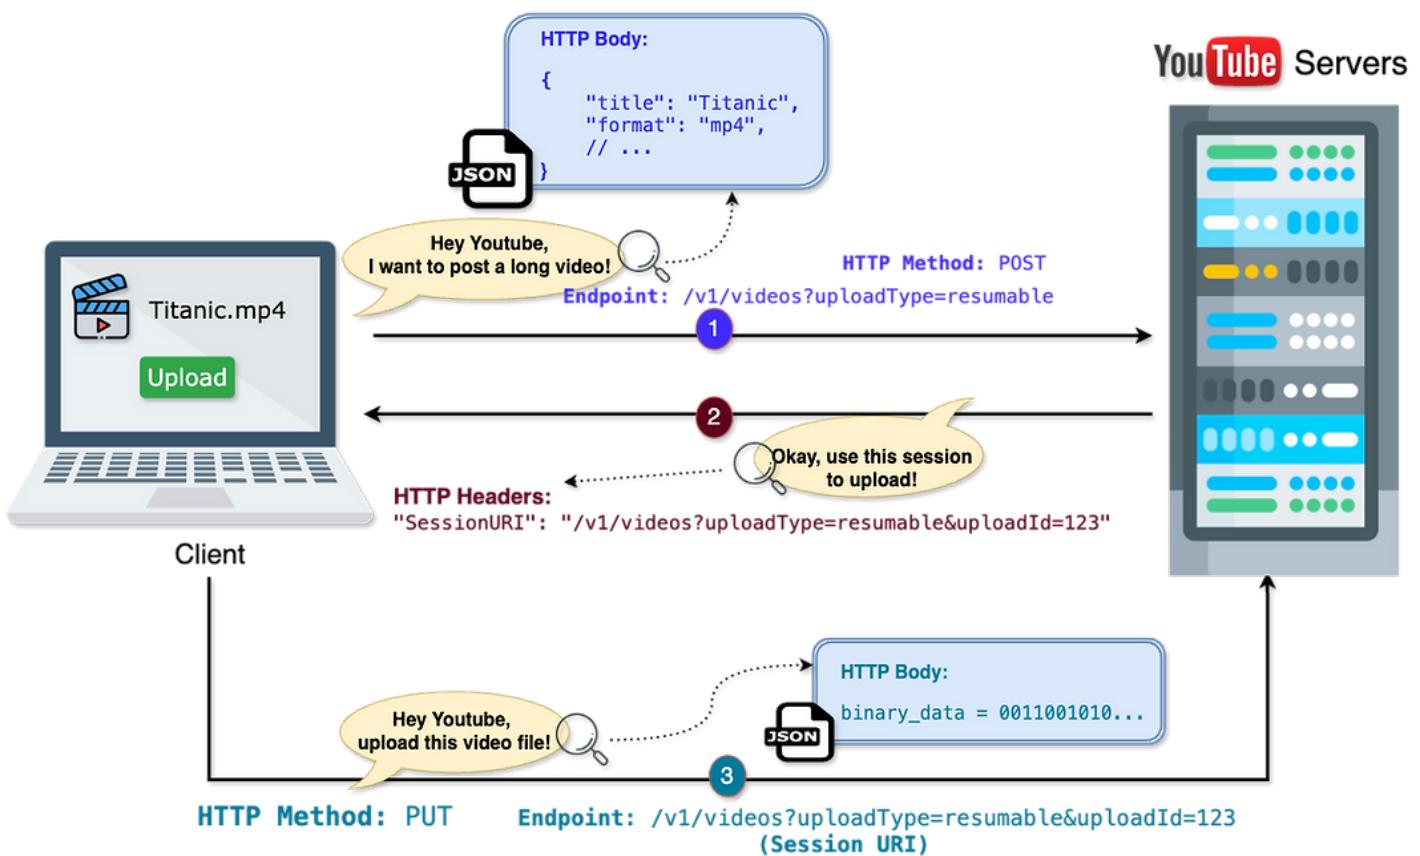
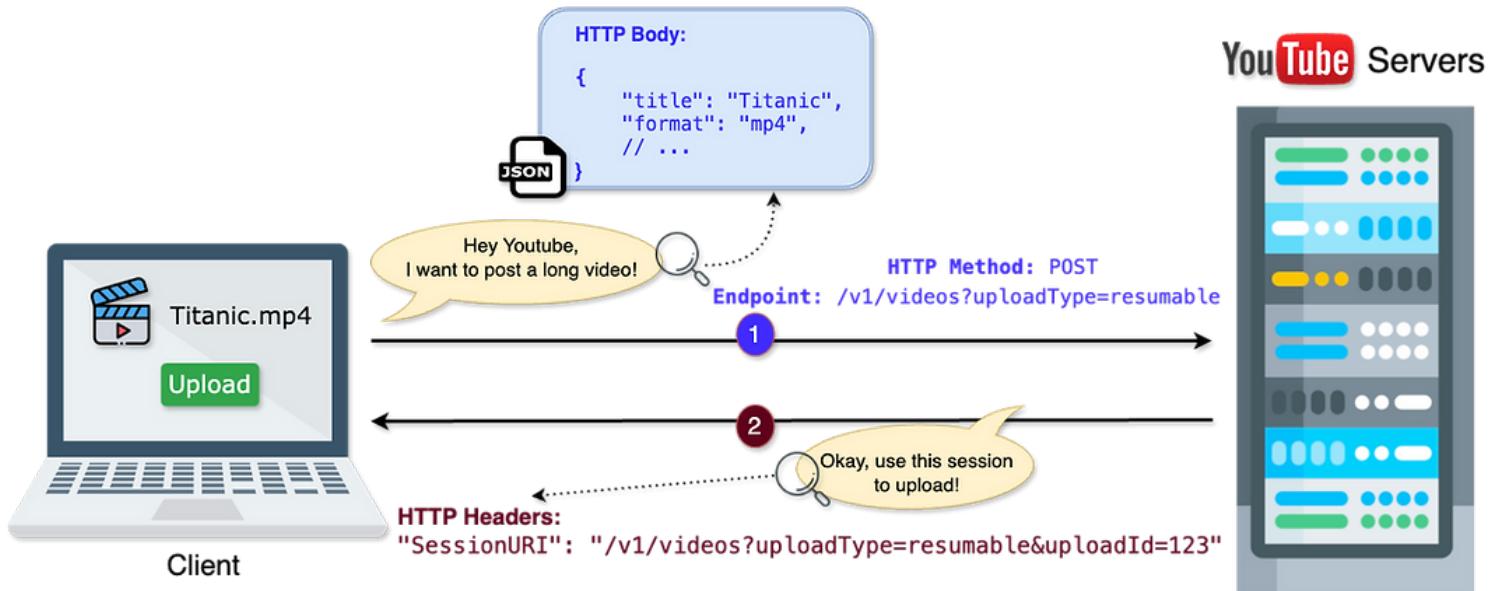
End-to-End Encryption (with Encryption)



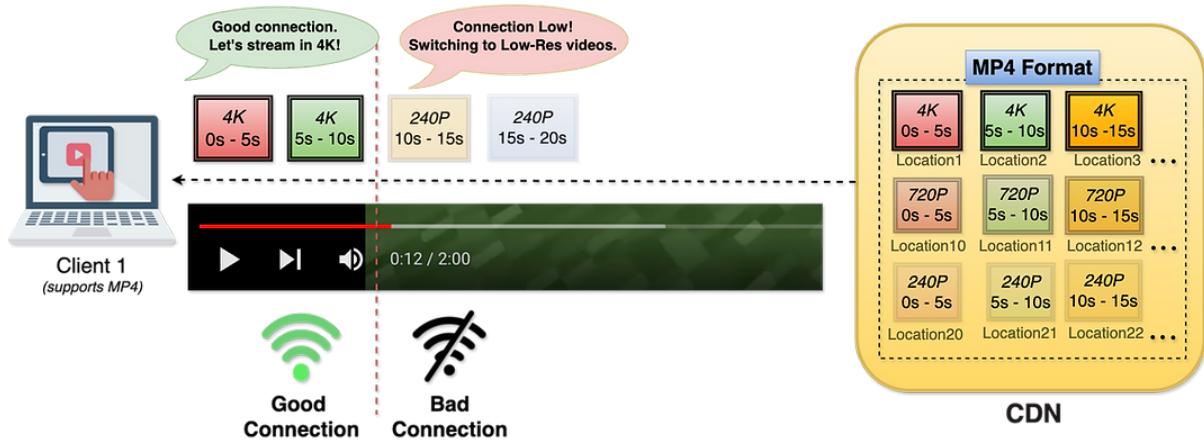
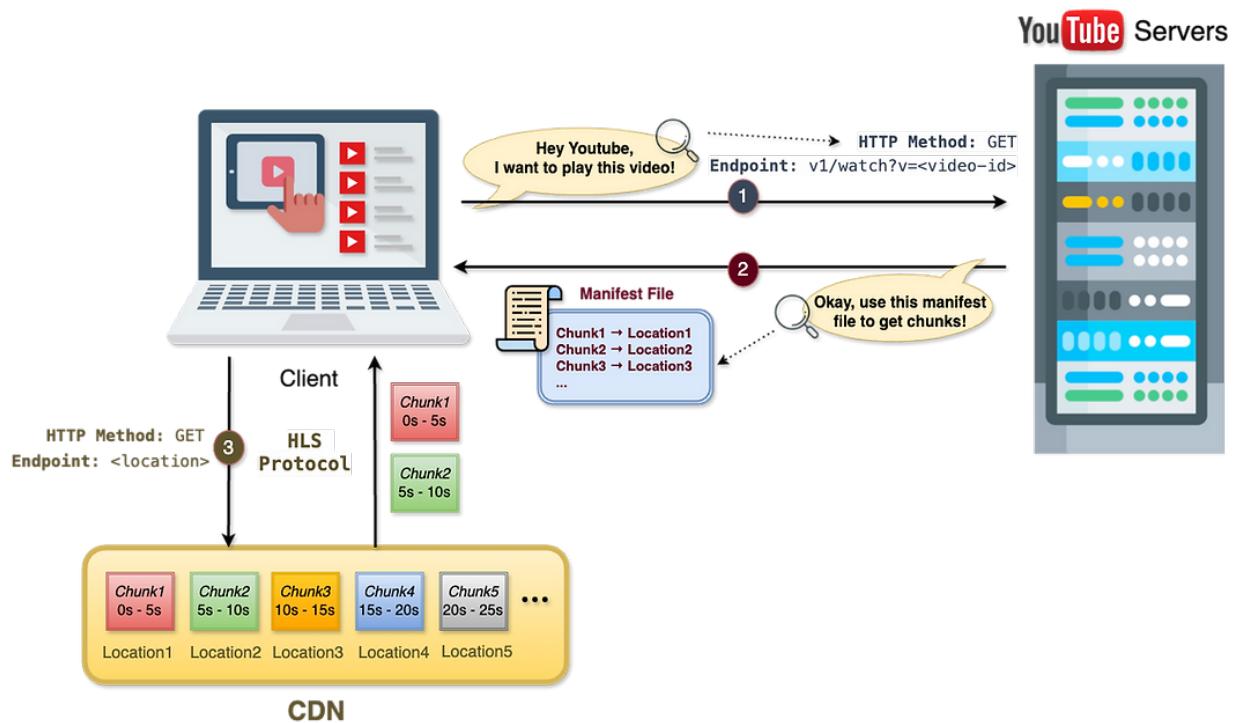
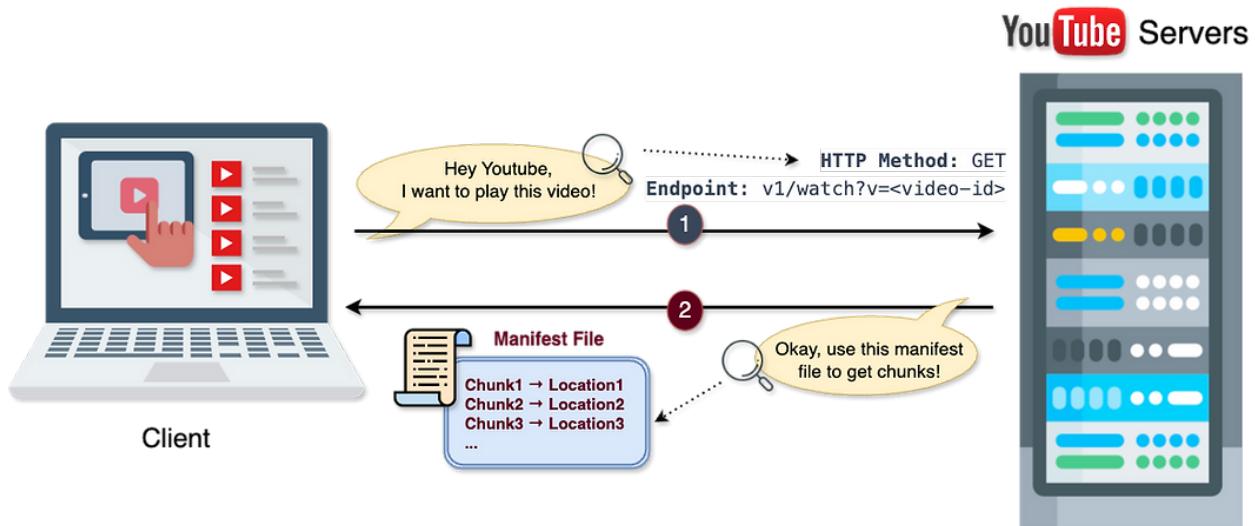
End-to-End Encryption (Encryption with lock & Key)



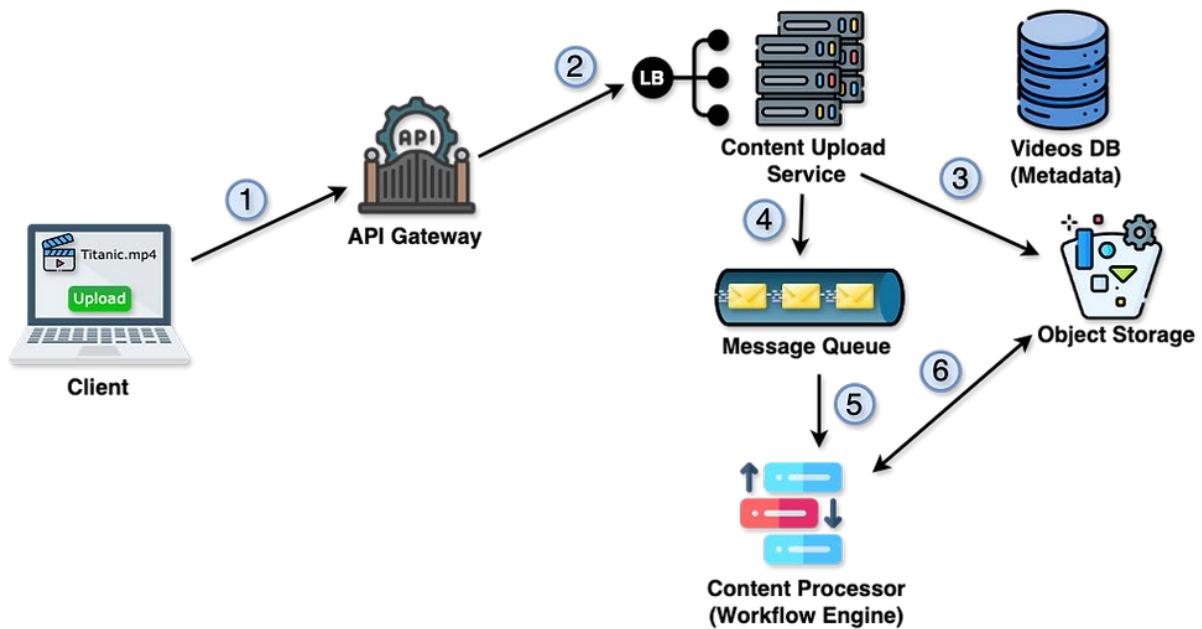
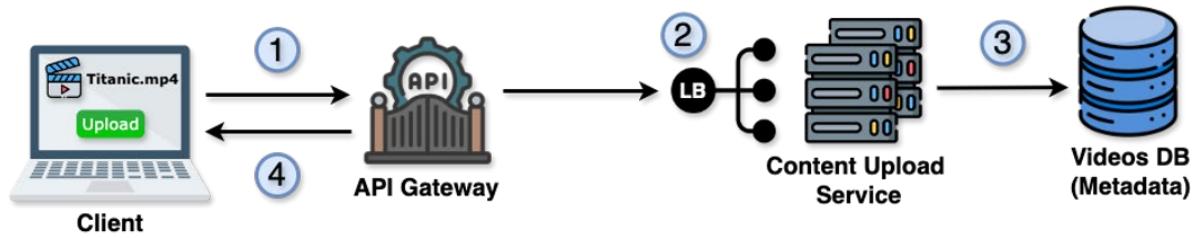
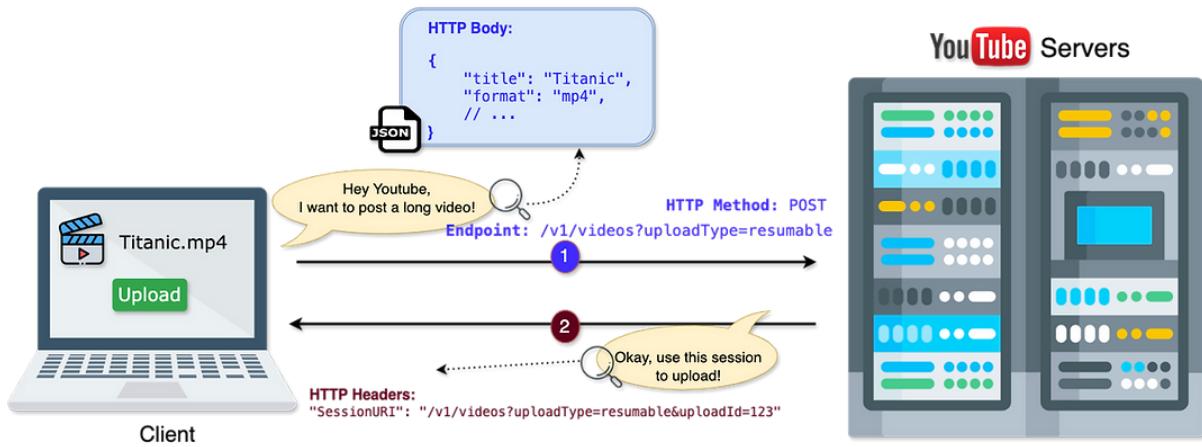
API DESIGN-Upload Content

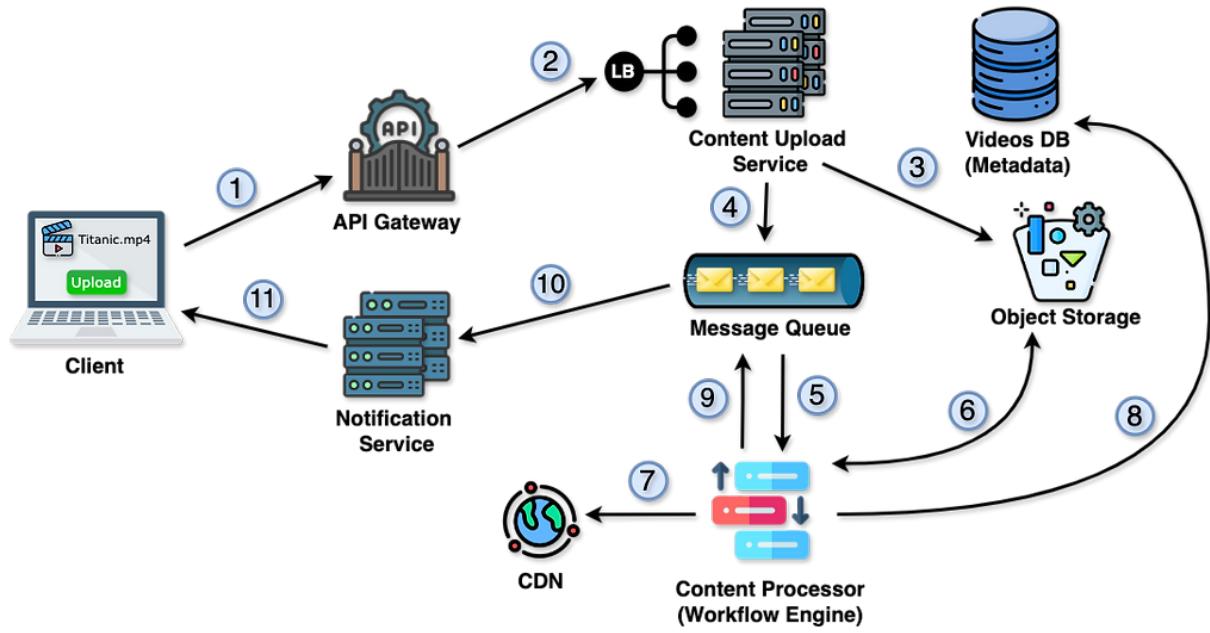


API DESIGN-Stream Content

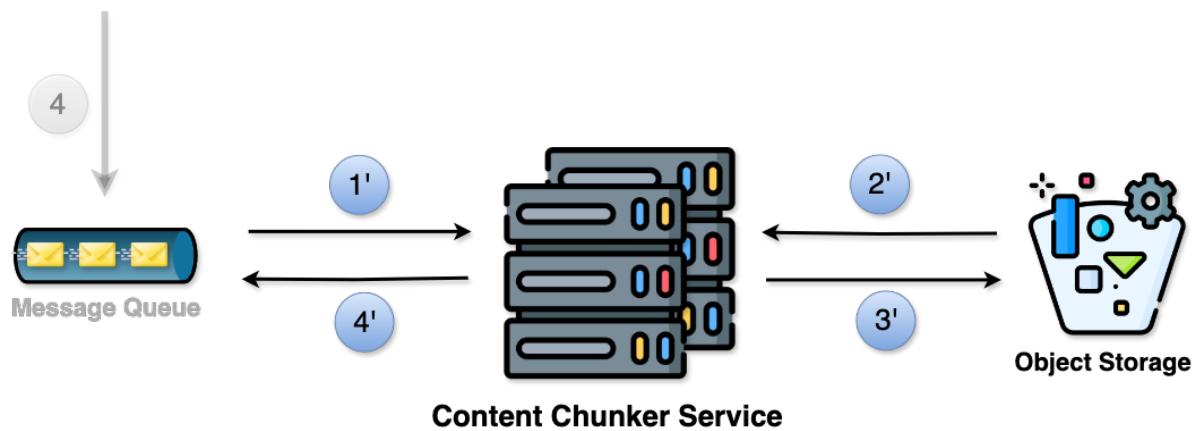


High Level Design-Upload Content

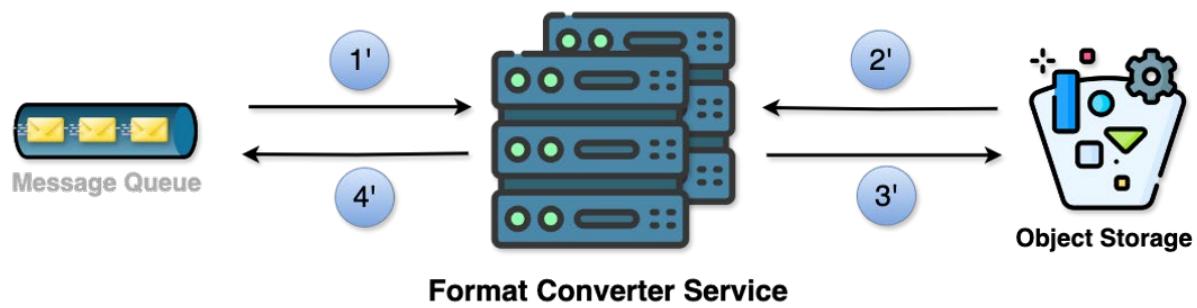




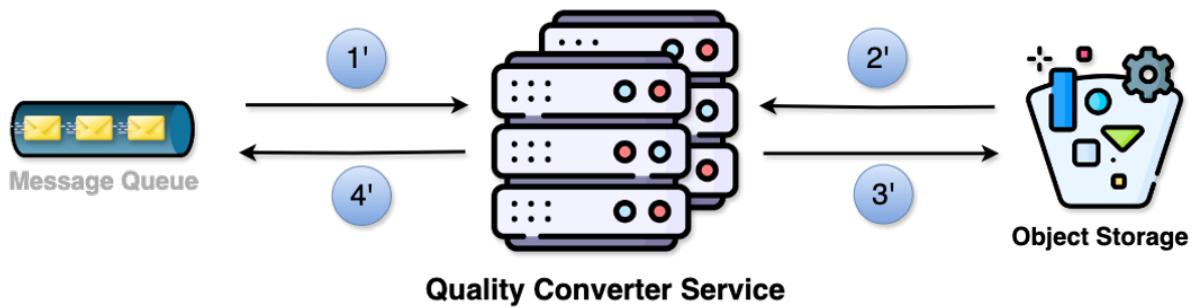
Content Chunker Service



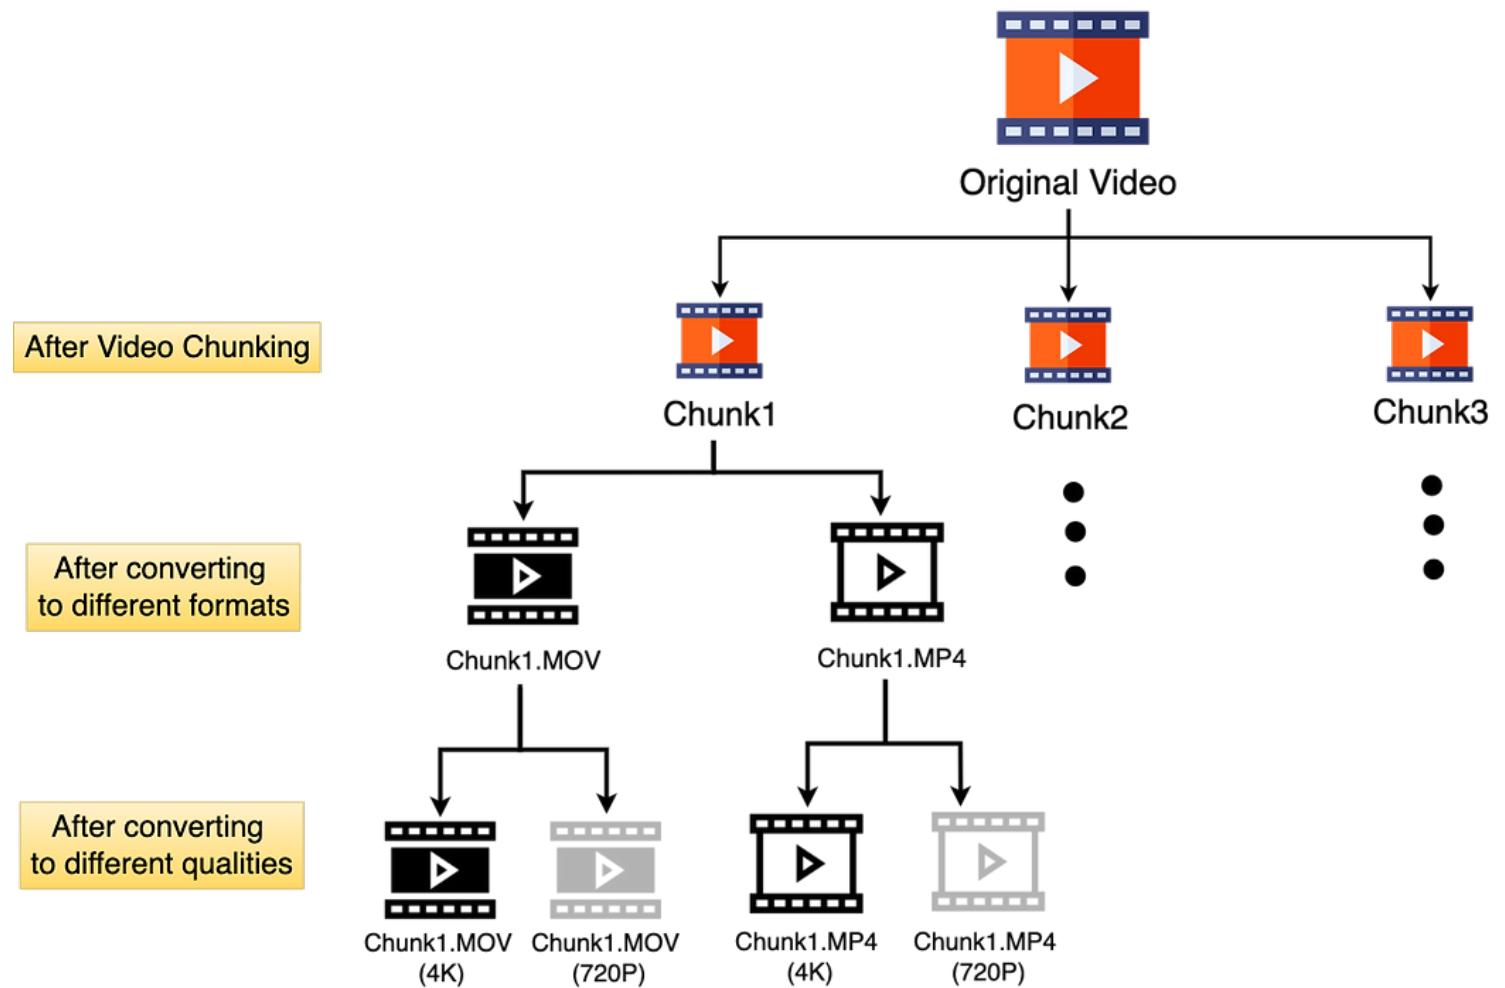
Format Converter Service



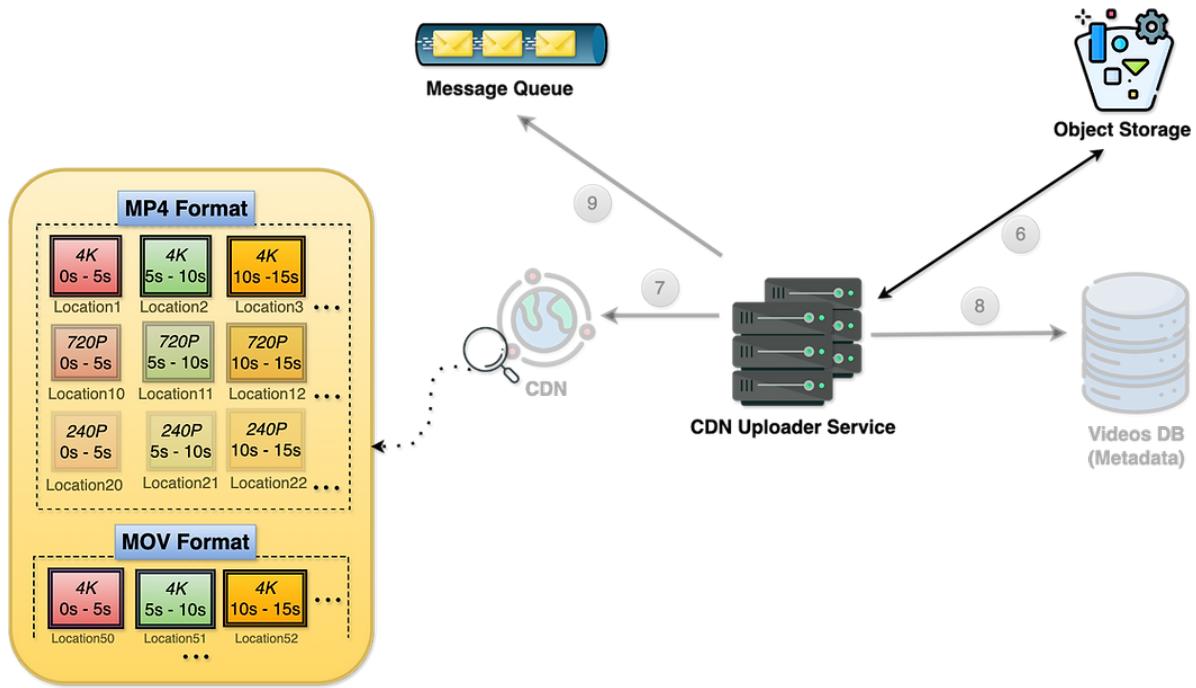
Quality Converter Service



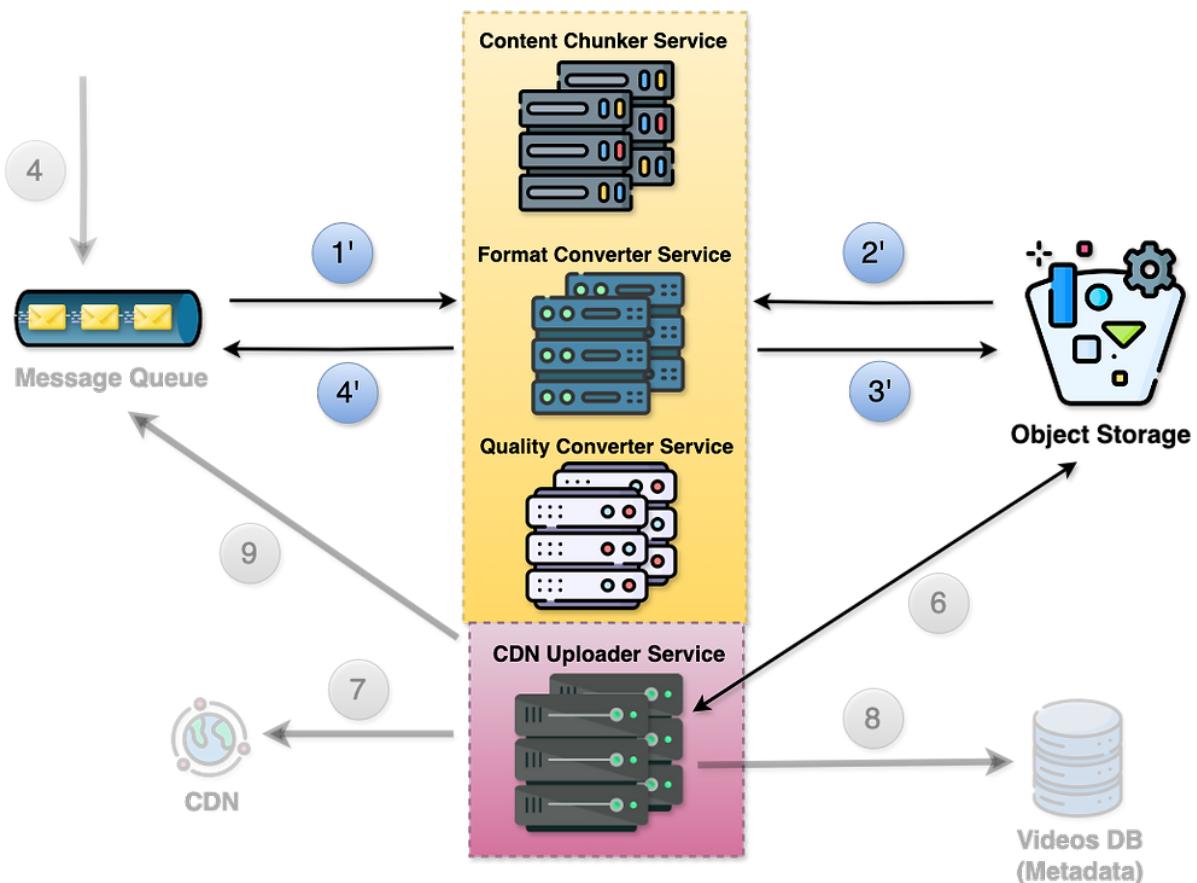
Chunking Videos



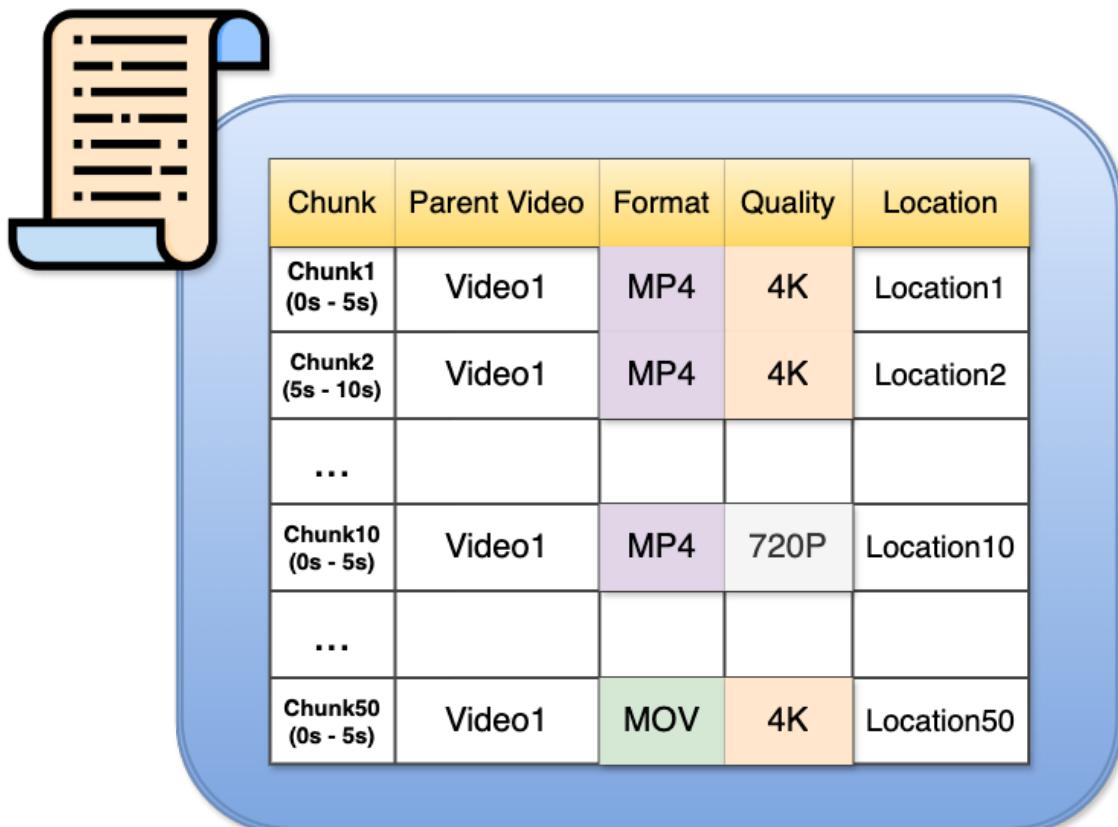
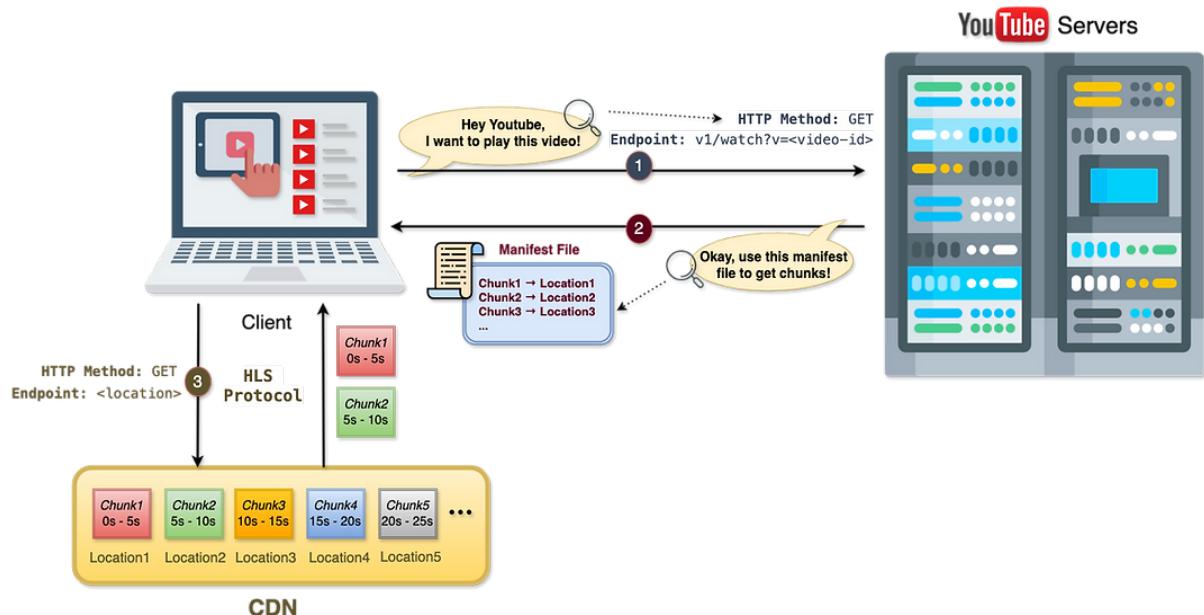
CDN Uploader Service

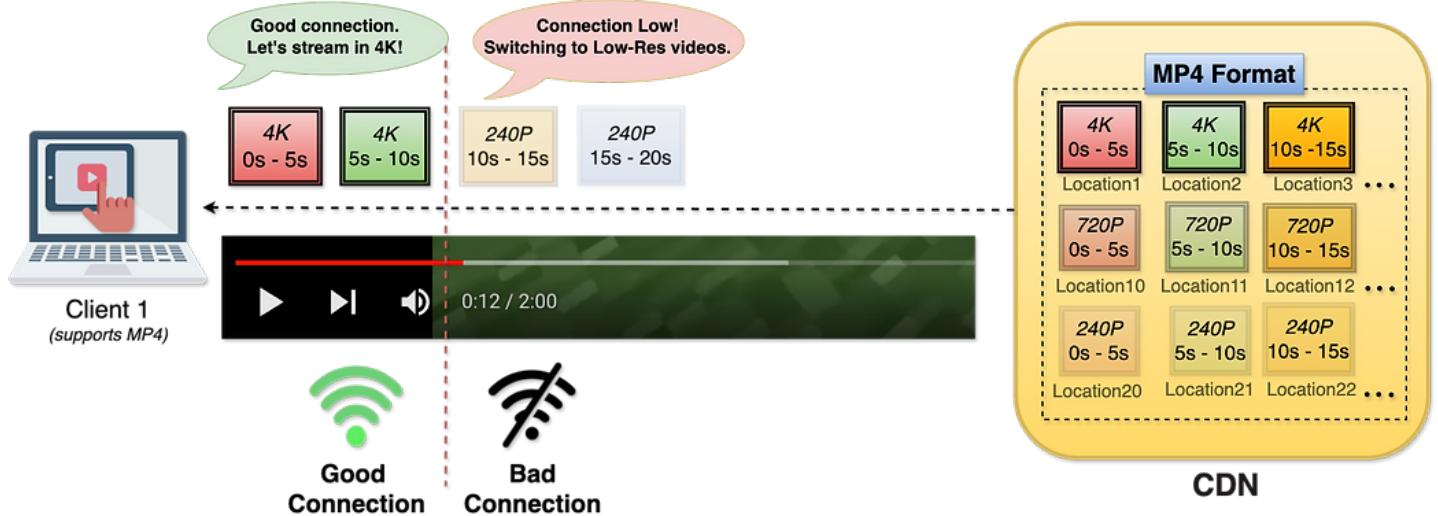
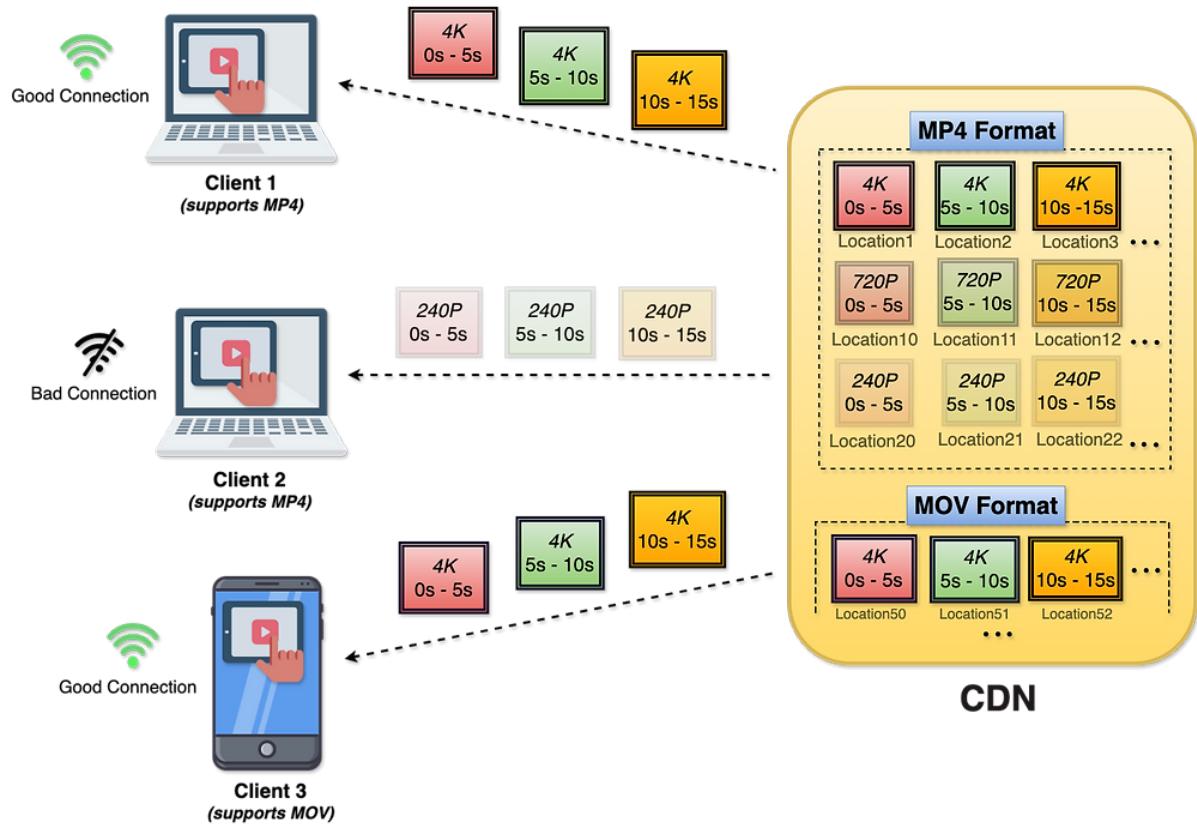


Overall Chunking Flow



High Level Design-Stream Content





Video's DB Schema



**Videos DB
(Metadata)**

```
{  
    "videoId": "unique_video_id",  
  
    // Technical Metadata  
    "format": "mp4",  
    "duration": "length_of_video_in_s",  
    ...  
  
    // General Metadata  
    "creatorId": "creator_id",  
    "title": "video_title",  
    "description": "video_description_text",  
    ...  
  
    // Content Delivery Metadata  
    "cdnUrls" : {  
        "MP4": {  
            "4K": ["chunkurl1", "chunkurl2", ...],  
            "720P": ["chunkurl1", "chunkurl2", ...],  
            ...  
        }  
        "MOV": {  
            "4K": ["chunkurl1", "chunkurl2", ...],  
            "720P": ["chunkurl1", "chunkurl2", ...],  
            ...  
        }  
        ...  
    }  
}
```

Chunking Workflow

